

# Detecting Automation of Twitter Accounts: Are You a Human, Bot, or Cyborg?

The Turtles Project: Design and  
Implementation of Nested Virtualization

**Deepak Kumar**



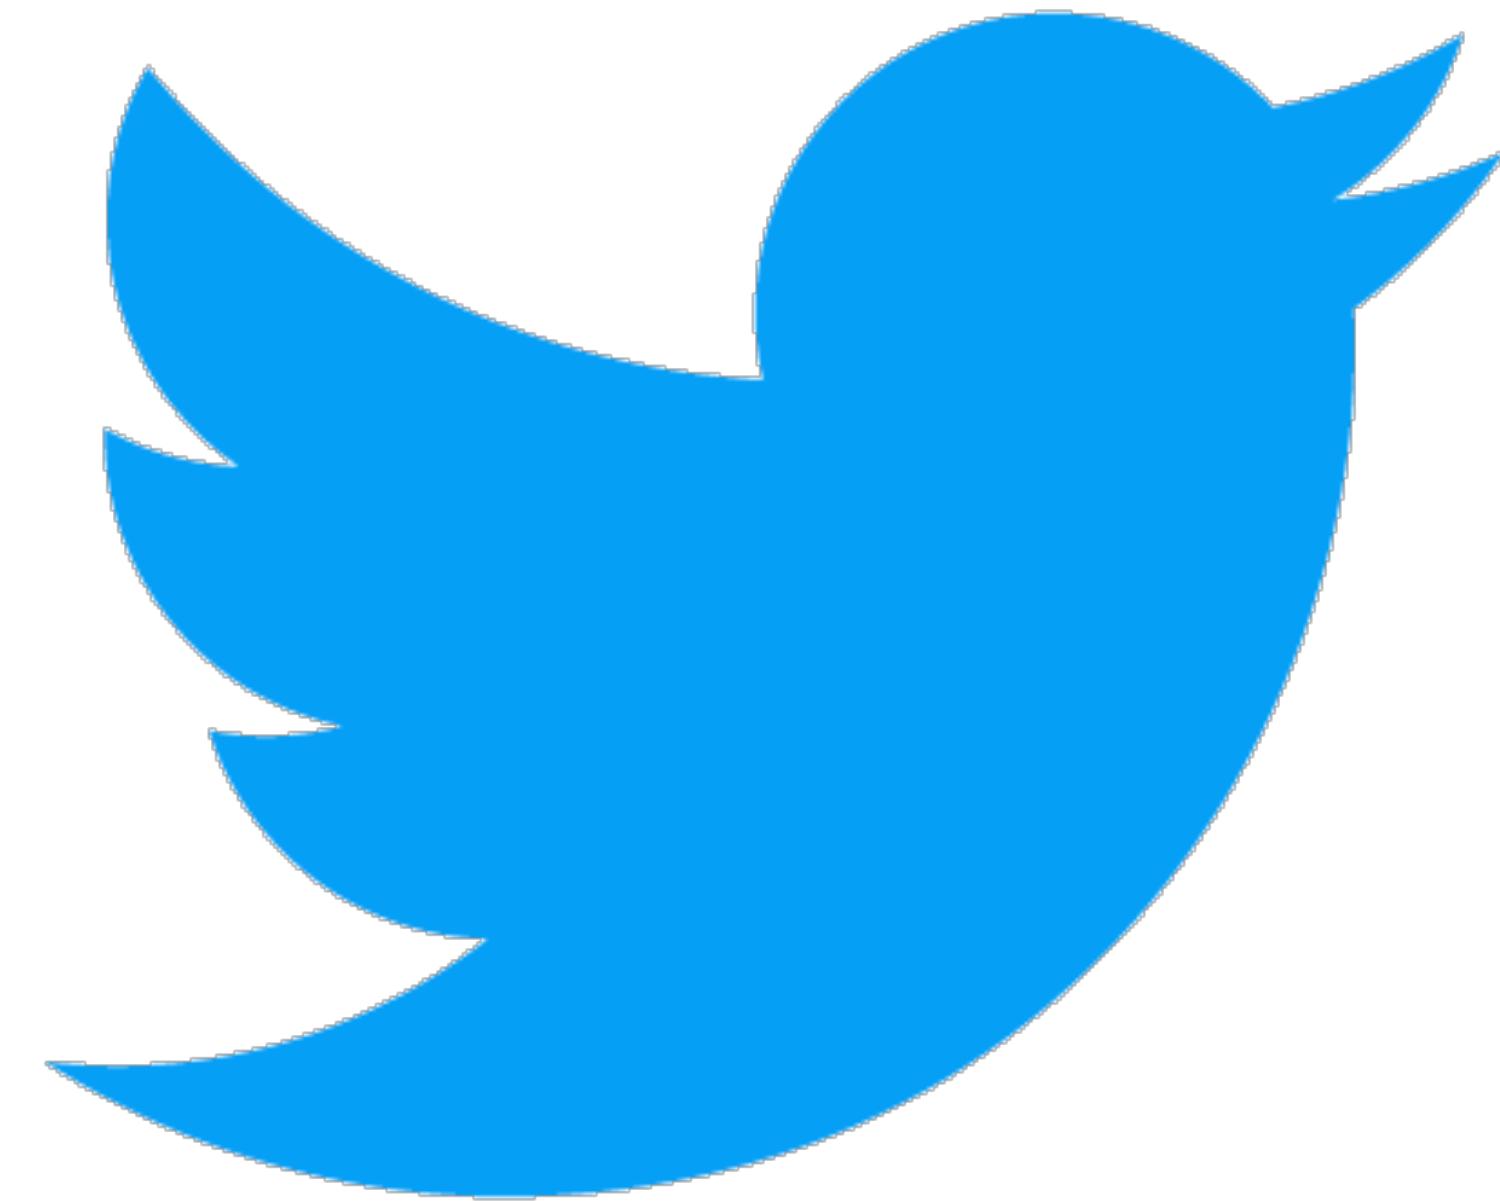
# Detecting Automation of Twitter Accounts: Are You a Human, Bot, or Cyborg?

Zi Chu, Steven Gianvecchio, Haining Wang, Sushil Jajodia  
IEEE Transactions in Dependable and Secure Computing



# Twitter

- Online social networking and microblogging tool



# Twitter

- Online social networking and microblogging tool
- Users post “tweets” - short, 280 character messages



# Twitter

- Online social networking and microblogging tool
- Users post “tweets” - short, 280 character messages
- Users can be *followers* or *friends* of other users



# Twitter

- Online social networking and microblogging tool
- Users post “tweets” - short, 280 character messages
- Users can be *followers* or *friends* of other users
- By end of Q4 2017, Twitter had ~330M monthly active users<sup>[1]</sup>



# Automation on Twitter

- Automated programs, or **bots**, are prevalent on Twitter
  - In fact, integral to Twitter's design
  - “News bots”



# Automation on Twitter

- Automated programs, or **bots**, are prevalent on Twitter
  - In fact, integral to Twitter's design
  - “News bots”
  - However, bots can also be used in *malicious* ways



# RUSSIAN BOTS ARE USING 2016 TACTICS TO HIJACK THE GUN DEBATE ON TWITTER

ERIN GRIFFITH SECURITY 02.15.18 02:00 PM

# PRO-GUN RUSSIAN BOTS FLOOD TWITTER AFTER PARKLAND SHOOTING

## *After Florida School Shooting, Russian ‘Bot’ Army Pounced*

By SHEERA FRENKEL and DAISUKE WAKABAYASHI FEB. 19, 2018



# How can we determine if a Twitter user is automated?



# Collecting Users + Tweets

Twitter Crawl

Timeline API



# Collecting Users + Tweets

Twitter Crawl

Timeline API



# Crawling Twitter

- Crawl Twitter with a Depth-First-Search



# Crawling Twitter

- Crawl Twitter with a Depth-First-Search
  - Select 5 users at random as a “seed”



# Crawling Twitter

- Crawl Twitter with a Depth-First-Search
  - Select 5 users at random as a “seed”
  - For each user, record their tweets and their followers



# Crawling Twitter

- Crawl Twitter with a Depth-First-Search
  - Select 5 users at random as a “seed”
  - For each user, record their tweets and their followers
  - Repeat for each follower...



# Crawling Twitter

- Crawl Twitter with a Depth-First-Search
  - Select 5 users at random as a “seed”
  - For each user, record their tweets and their followers
  - Repeat for each follower...
  - In total, crawled **429,423** users



# Collecting Users + Tweets

Twitter Crawl

Timeline API



# Timeline API

- Twitter publishes the 20 most recent Tweets in global scope
  - Collected tweets and users during same timeframe as Twitter crawl
  - Added an additional **89,984** users



# Twitter Dataset

Data Volume

Metadata Collected?

Users

512,407

username, followers,  
following, registration  
date, verified

Tweets

41,995,545

user, time, device



# Types of Users on Twitter



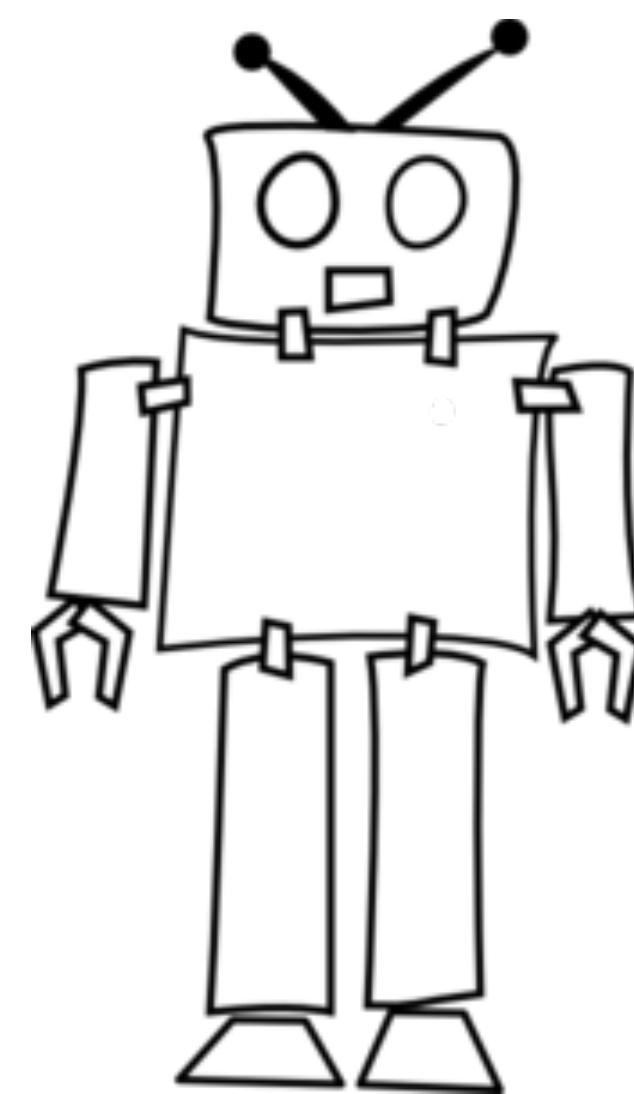
Human



# Types of Users on Twitter



Human

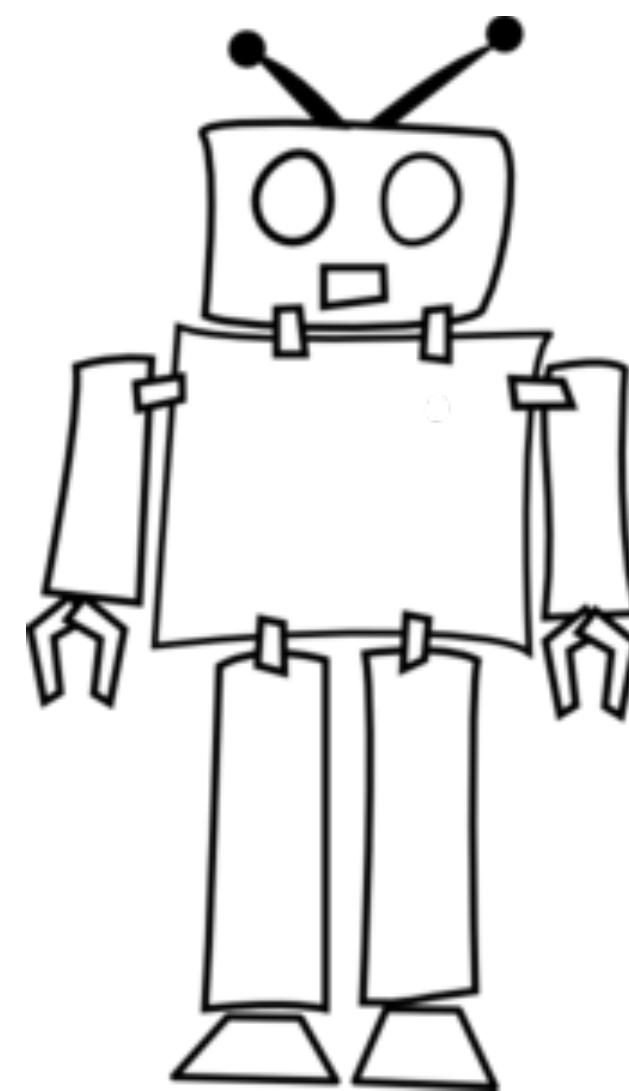


Bot

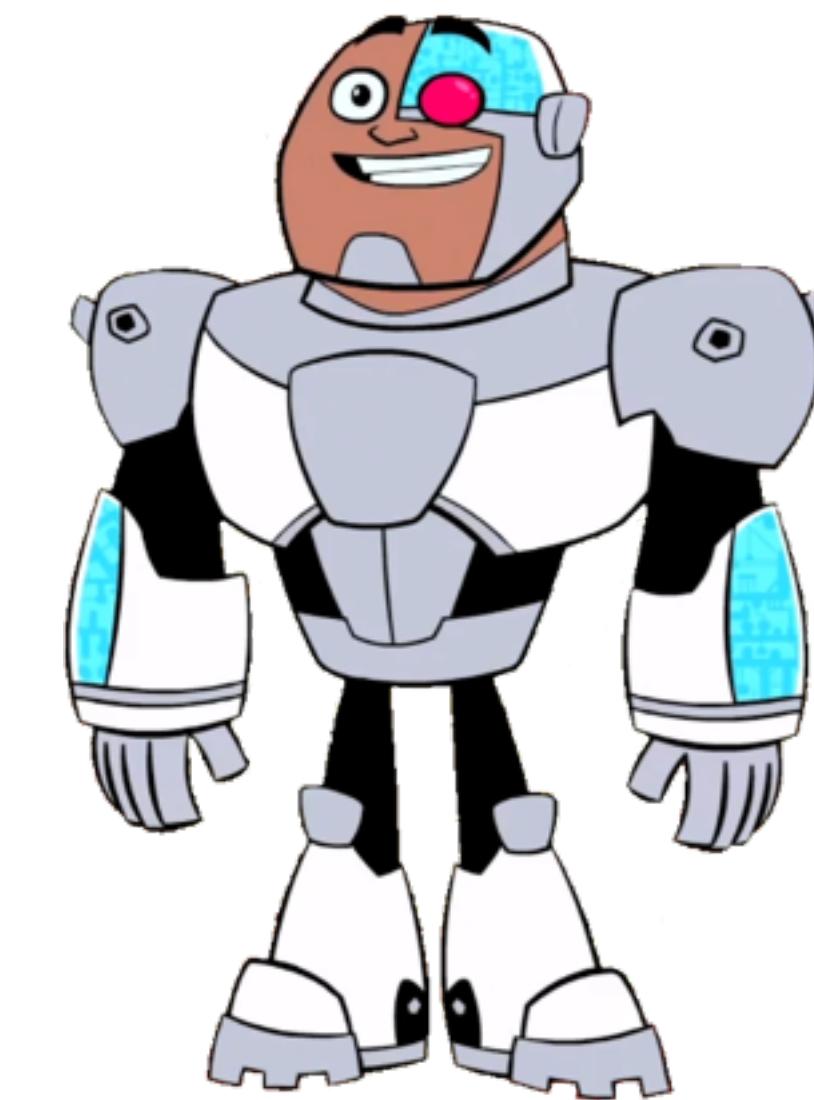
# Types of Users on Twitter



Human



Bot



Cyborg

# Classifying Users on Twitter: Ground Truth

- Select random sample of Twitter users from dataset
- Manually classify each user based on inspecting their profile
  - Follows the principle of “Turing Test” - 5 minutes to determine what class of user



# Classifying Users on Twitter: Ground Truth

- Collected **2,000** unique users in each class
- Contributed **8,350,095** tweets



# Classifying Users on Twitter: Ground Truth

- Collected **2,000** unique users in each class
- Contributed **8,350,095** tweets
- We can now investigate the *differences in behavior between user-classes*



Hypothesis: Each class of users exhibits unique behavior



# Behavioral Patterns

Social Relationships

Tweeting Habits

Tweet Properties



# Behavioral Patterns

## **Social Relationships**

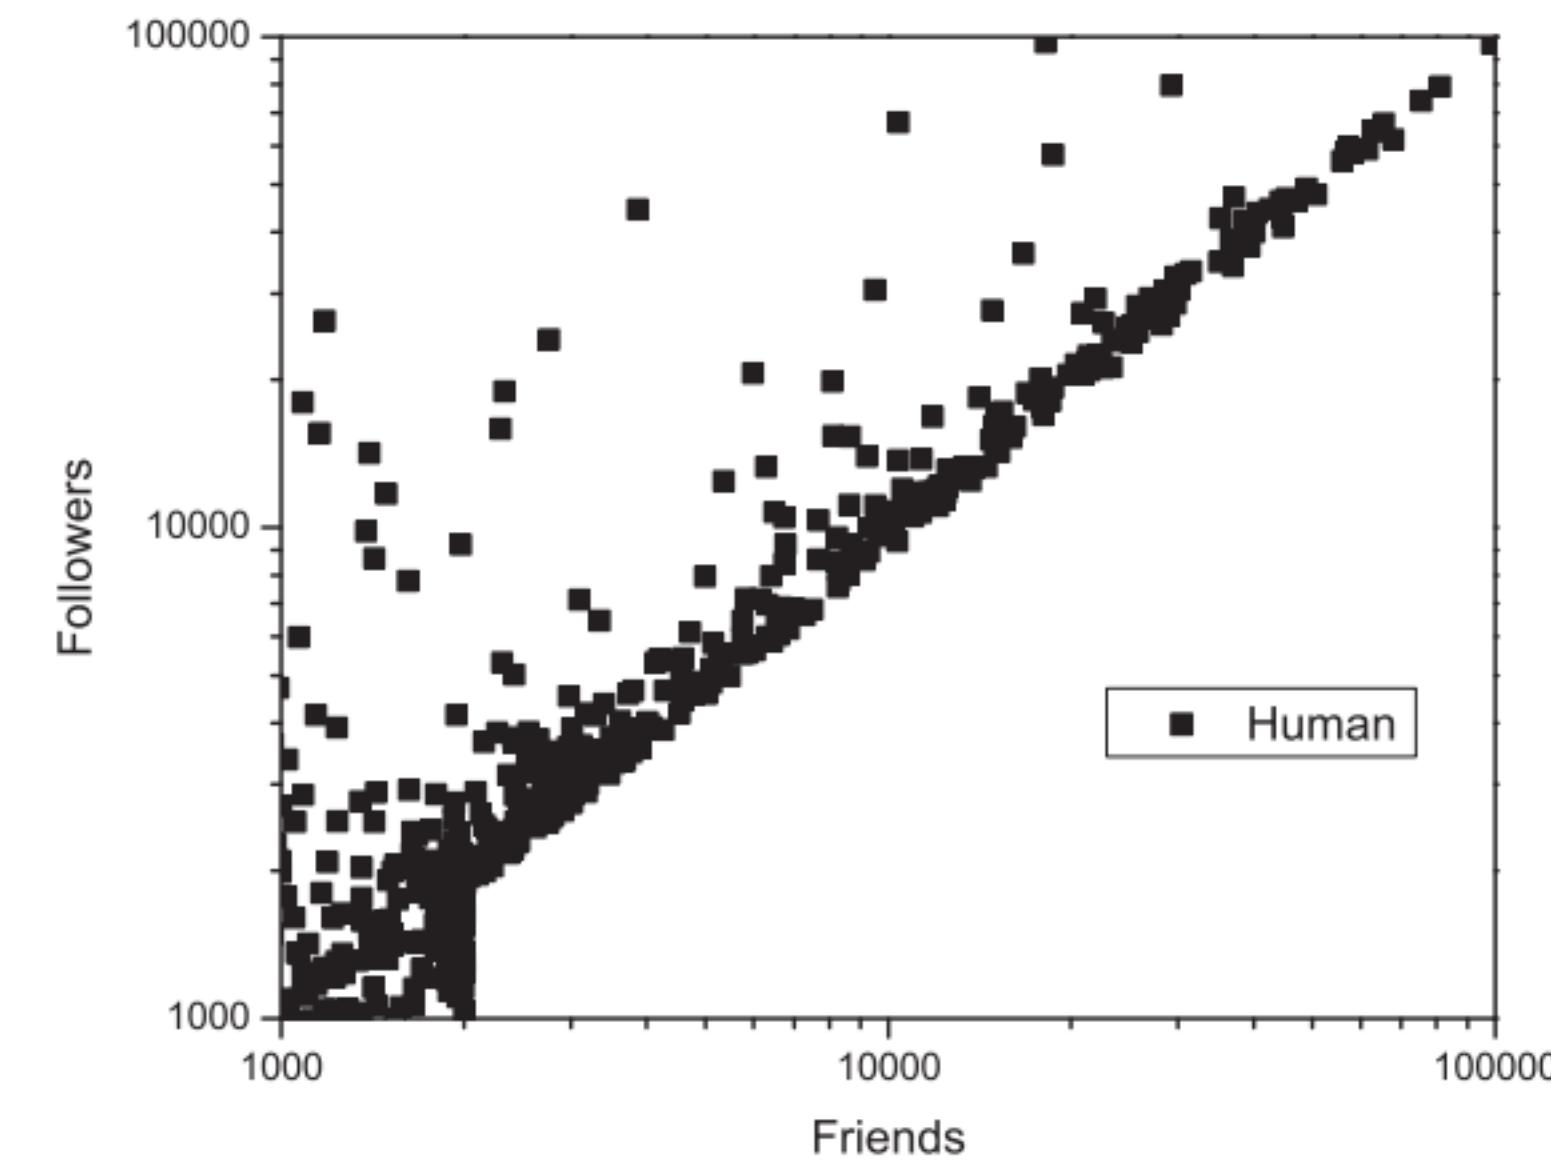
Tweeting Habits

Tweet Properties



# Social Relationships: Humans

- Typically *reciprocal* relationships

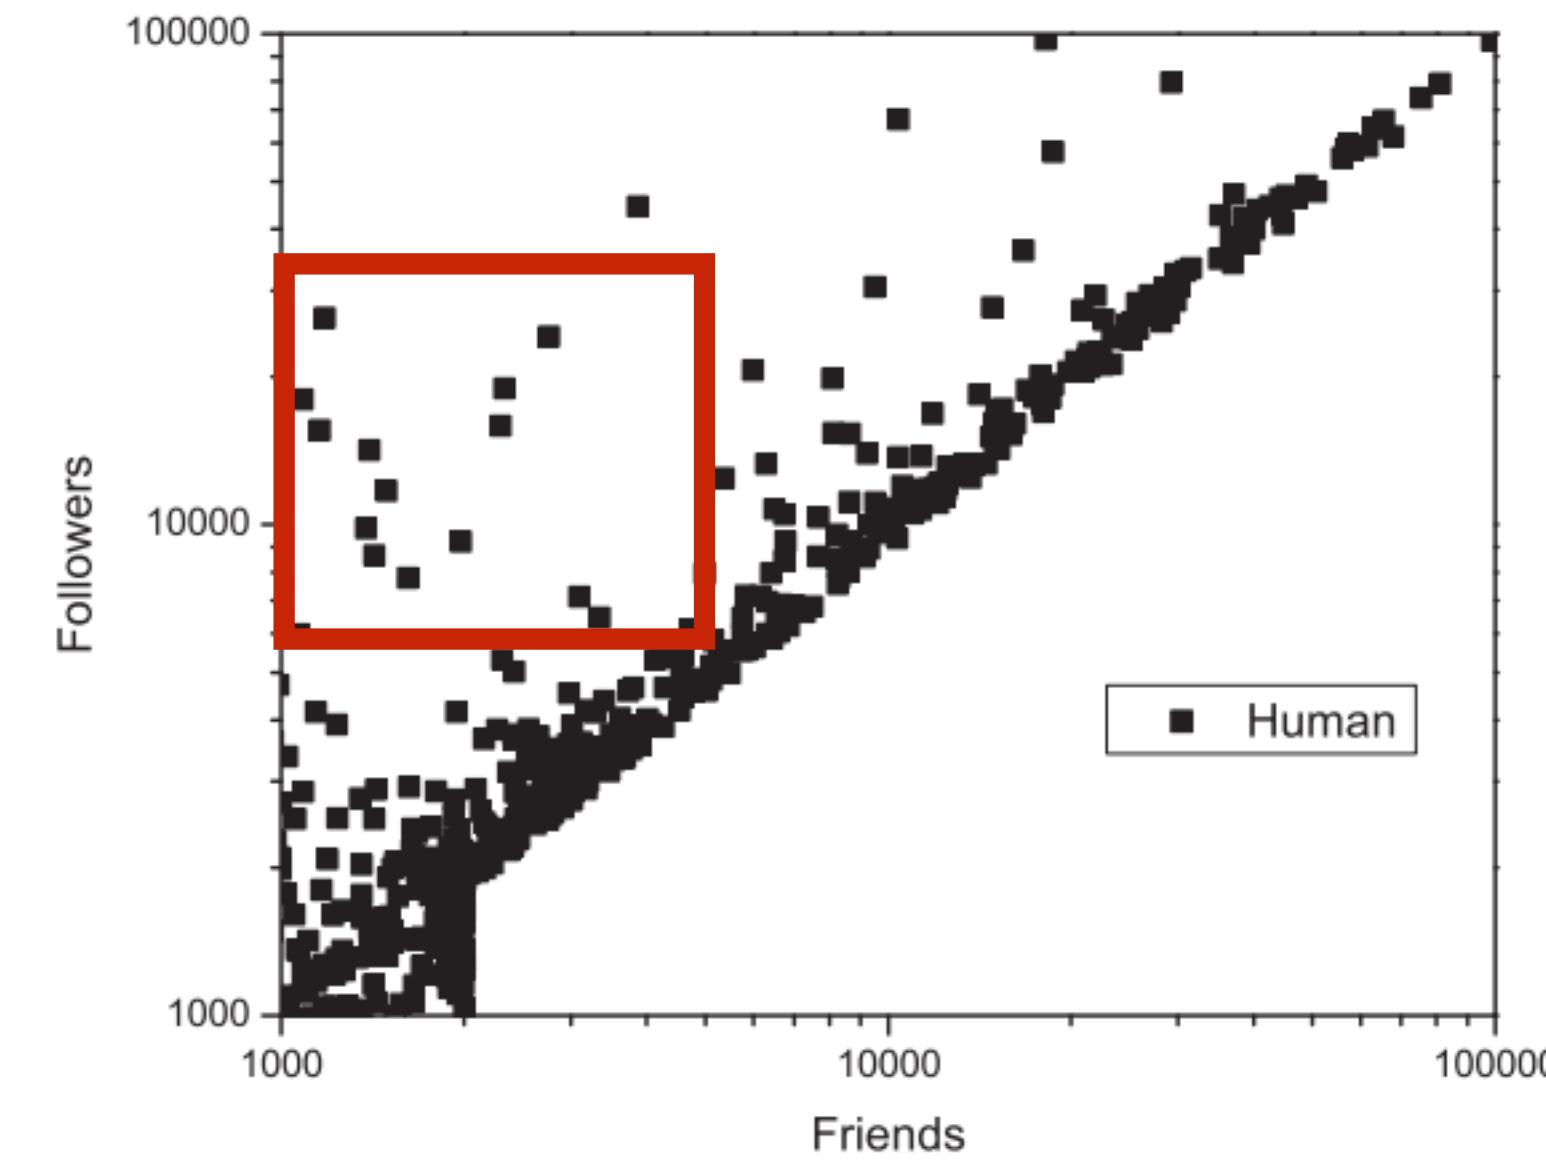


(a) Human



# Social Relationships: Humans

- Typically *reciprocal* relationships
- Handful of high follower/low friends, these are mainly celebrities and important human figures

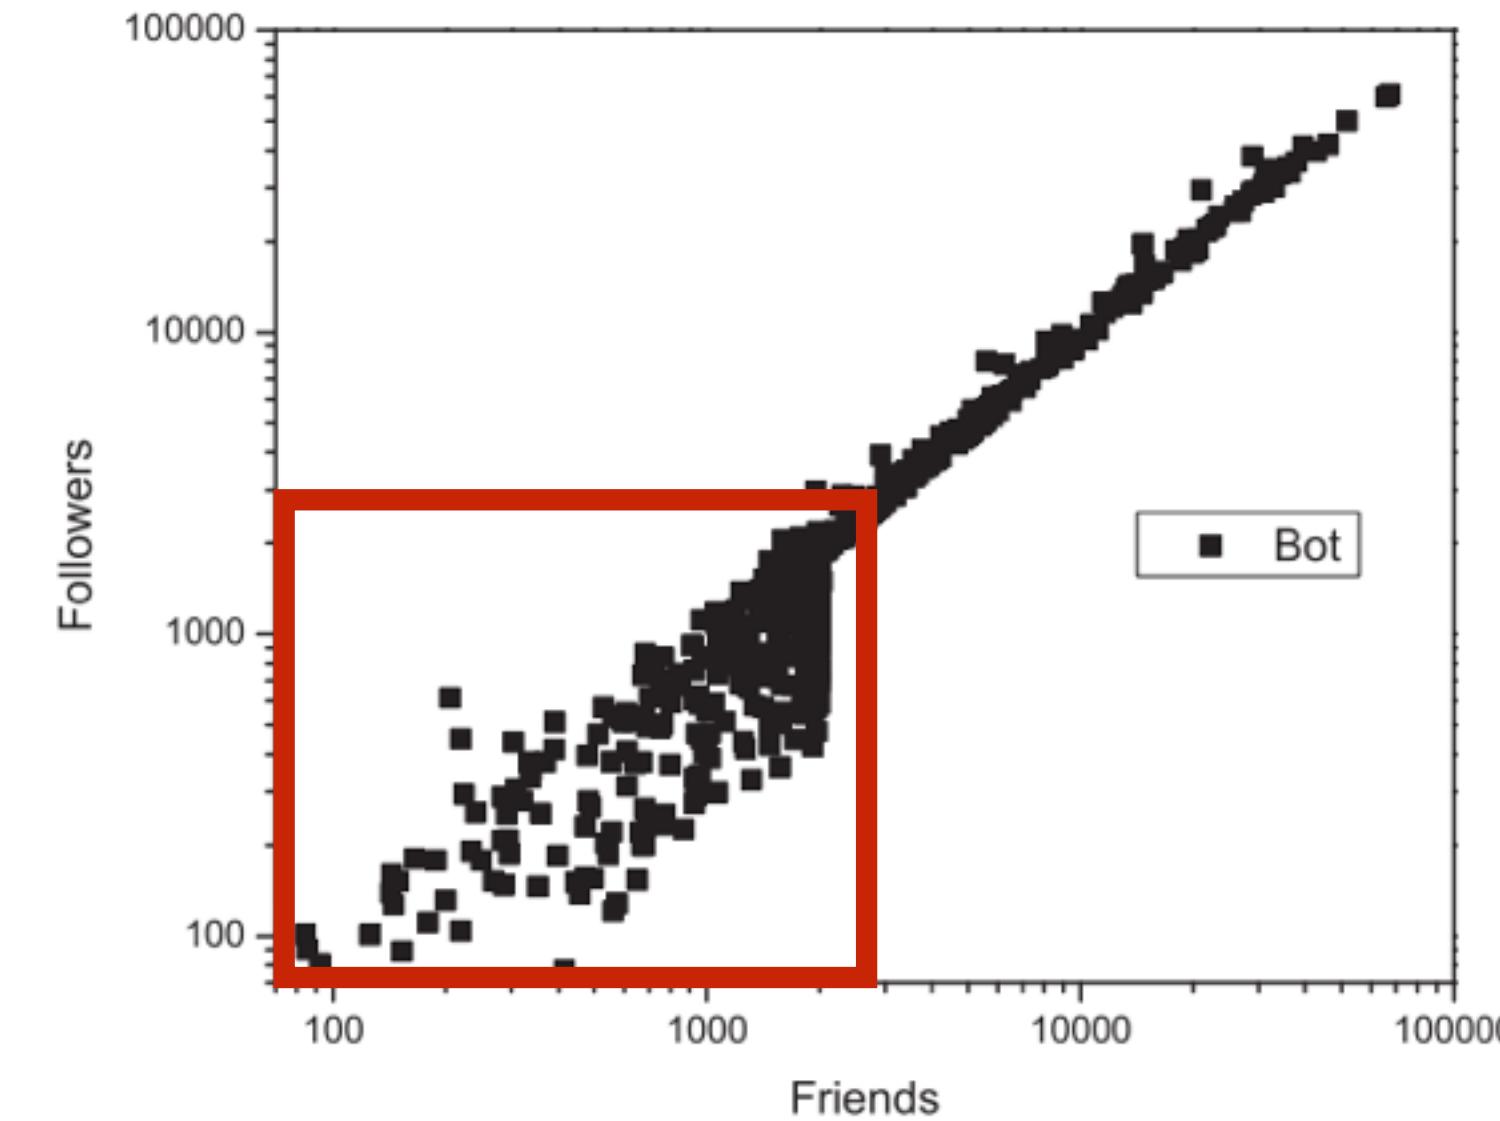


(a) Human



# Social Relationships: Bots

- Most bots have lots of friends (i.e., they follow many other accounts), but not followers

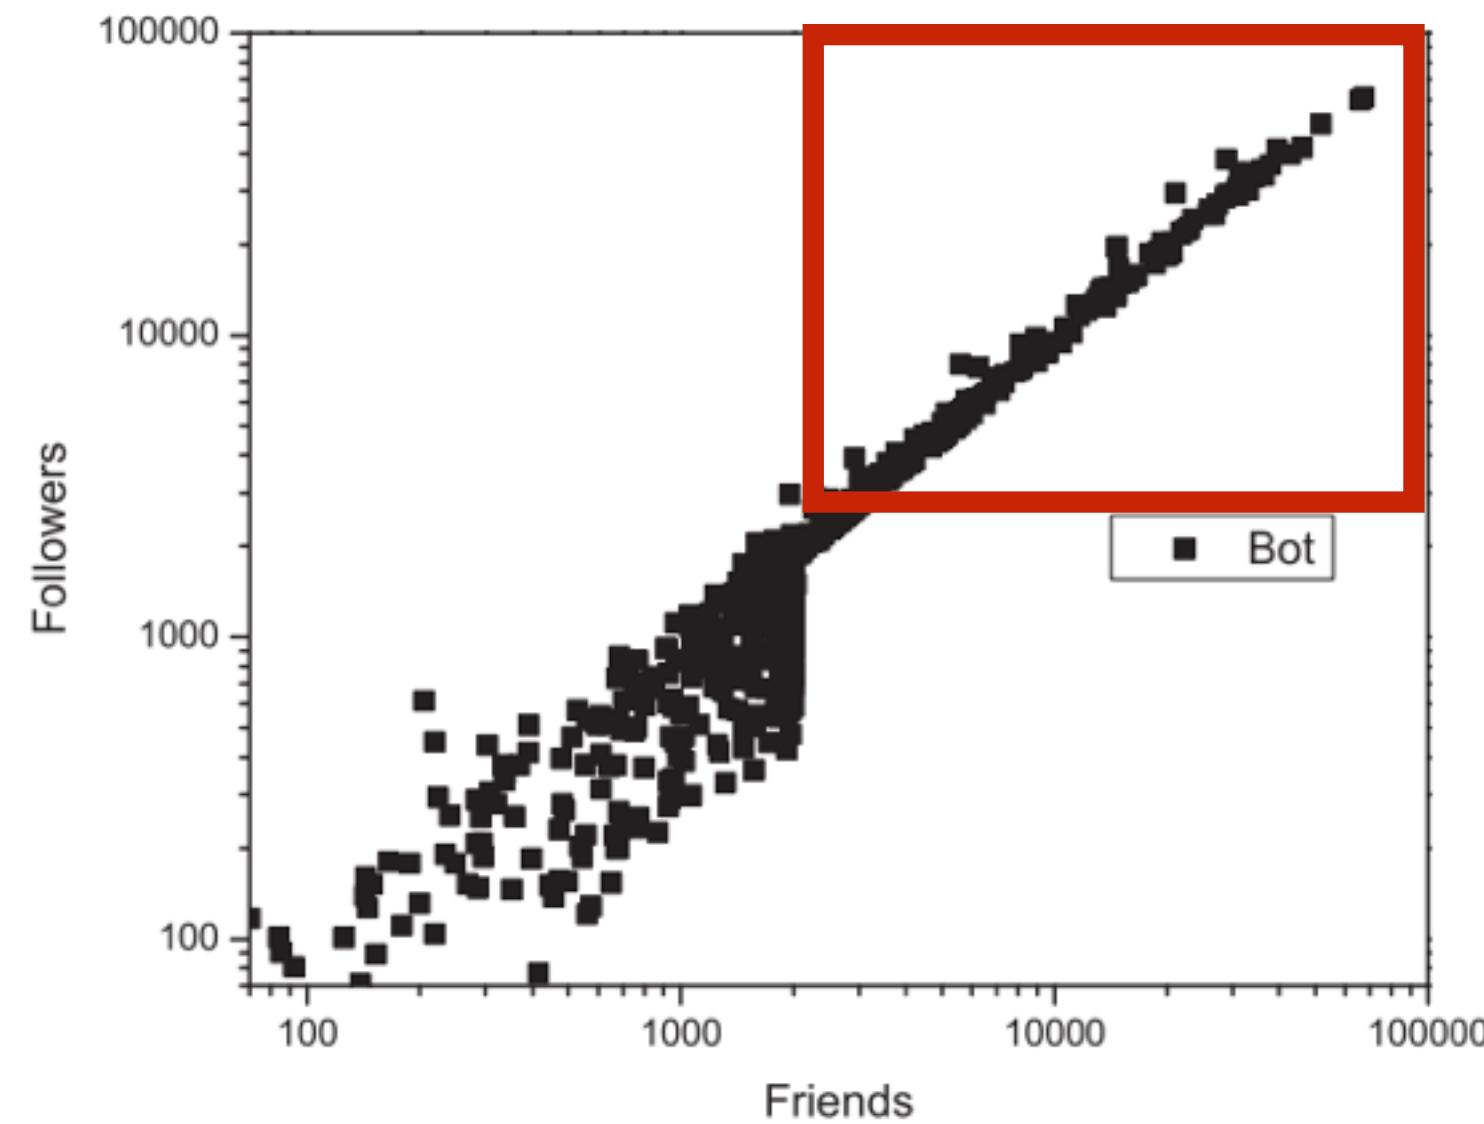


(b) Bot



# Social Relationships: Bots

- Most bots have lots of friends (i.e., they follow many other accounts), but not followers
- Some bots are “smart” - they continually unfollow and follow to maintain a 1:1 ratio

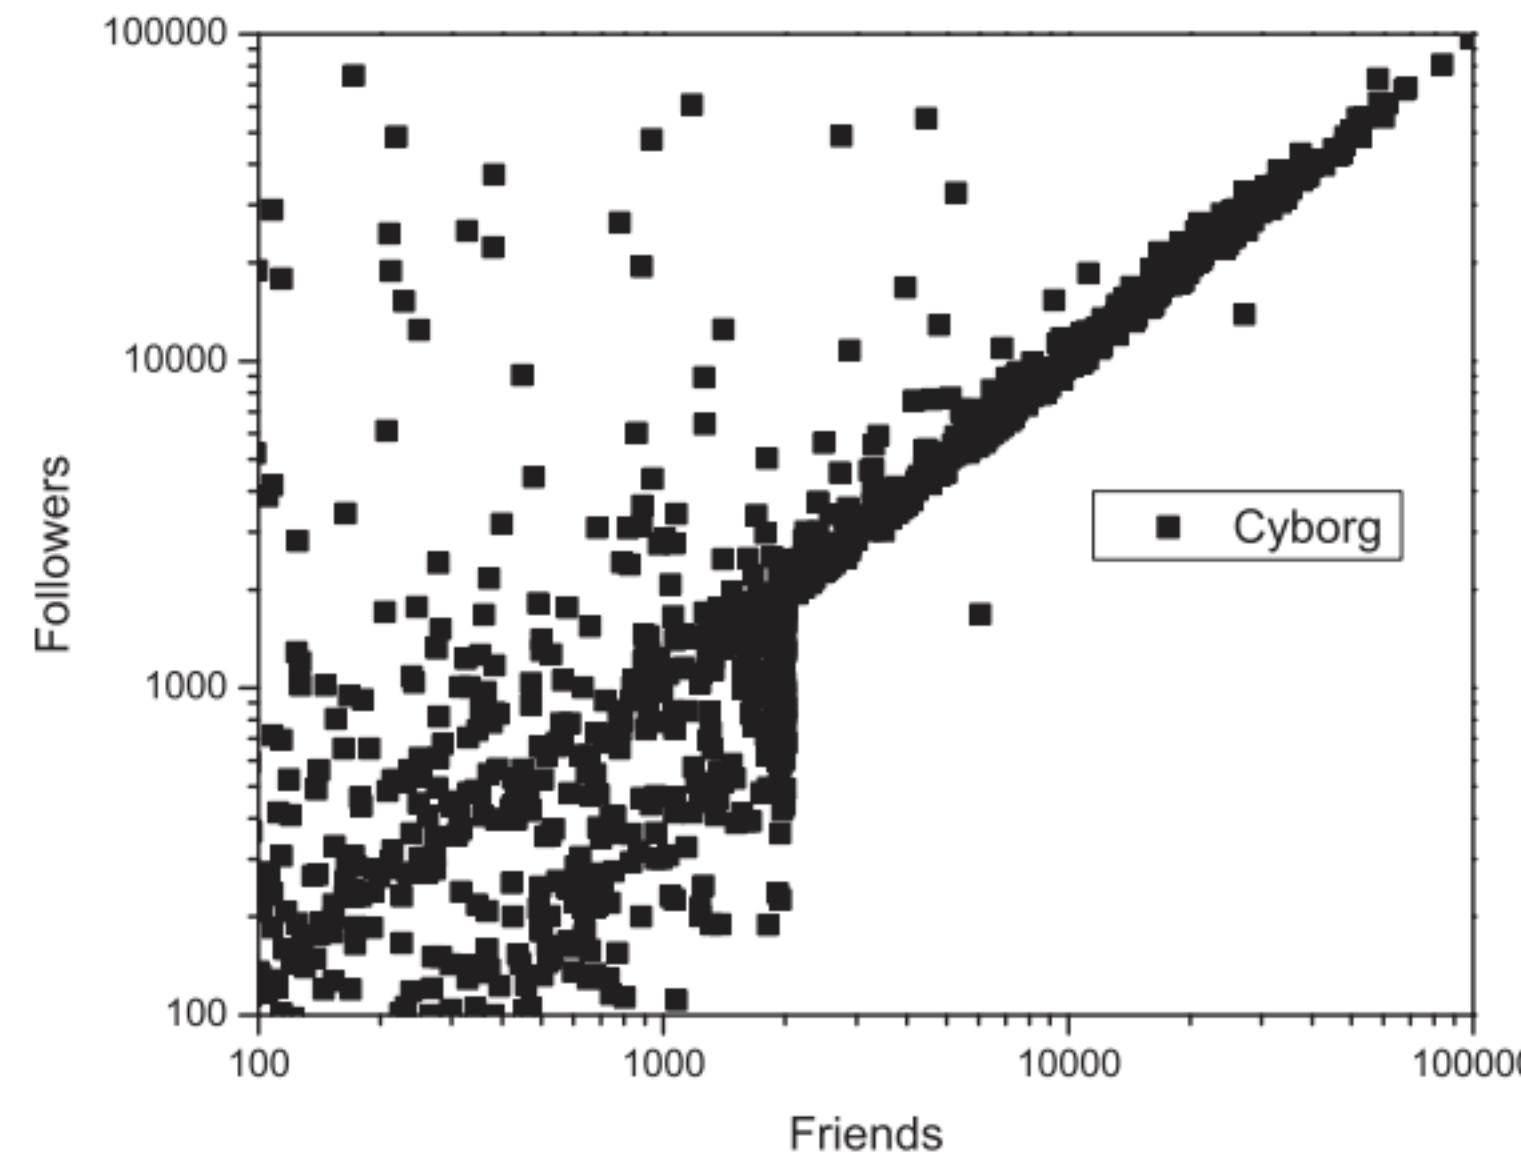


(b) Bot



# Social Relationships: Cyborgs

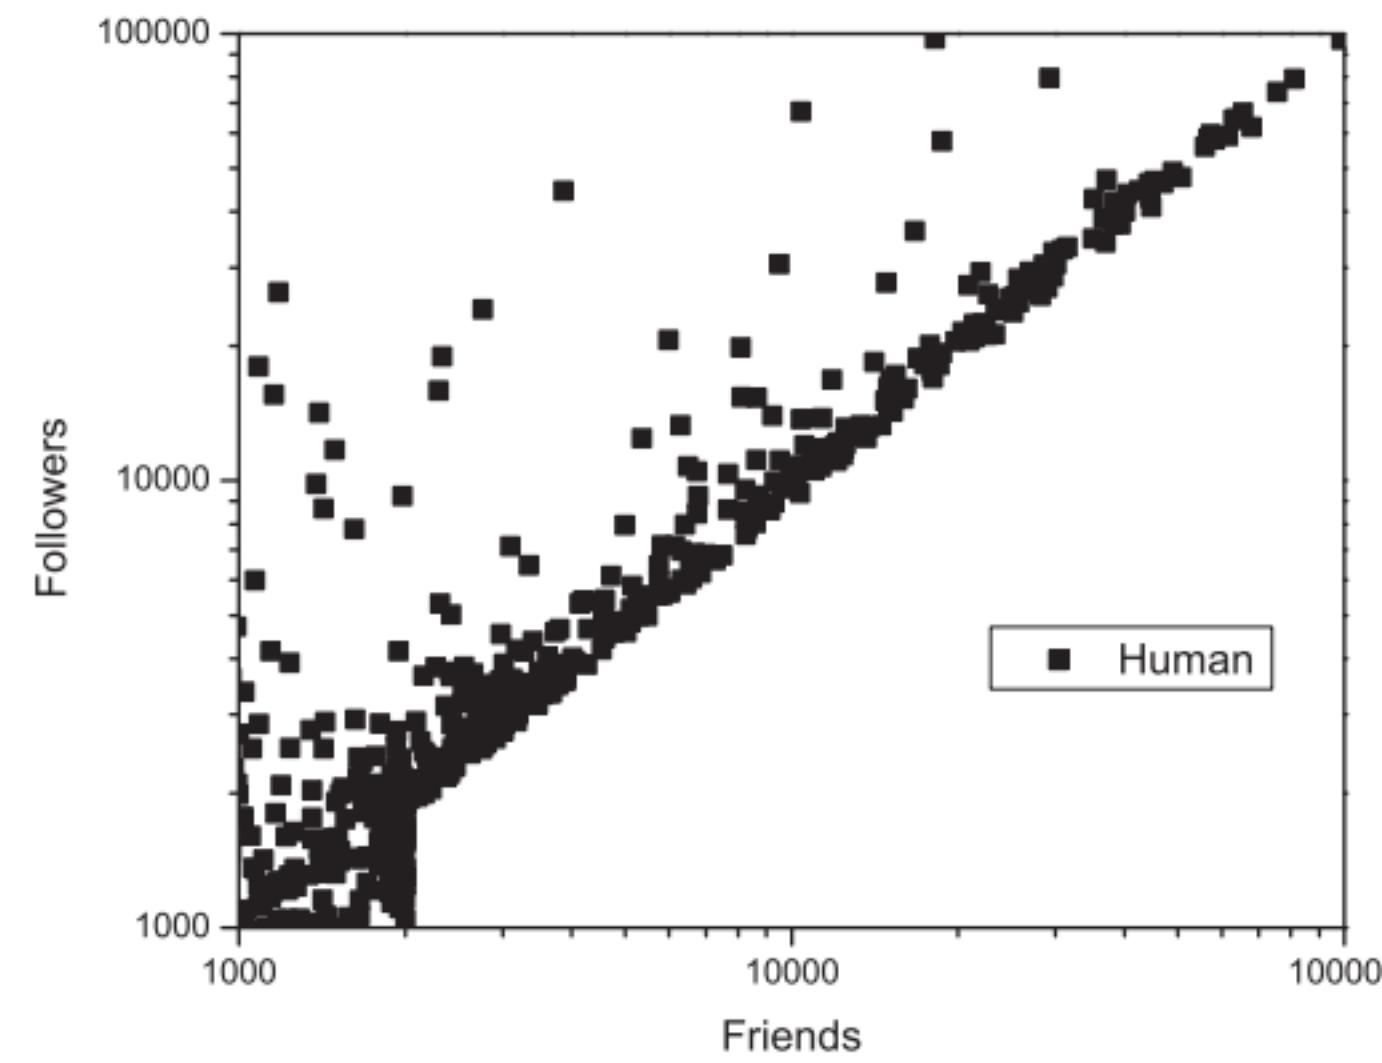
- Cyborgs exhibit a mix of human and bot social relationships



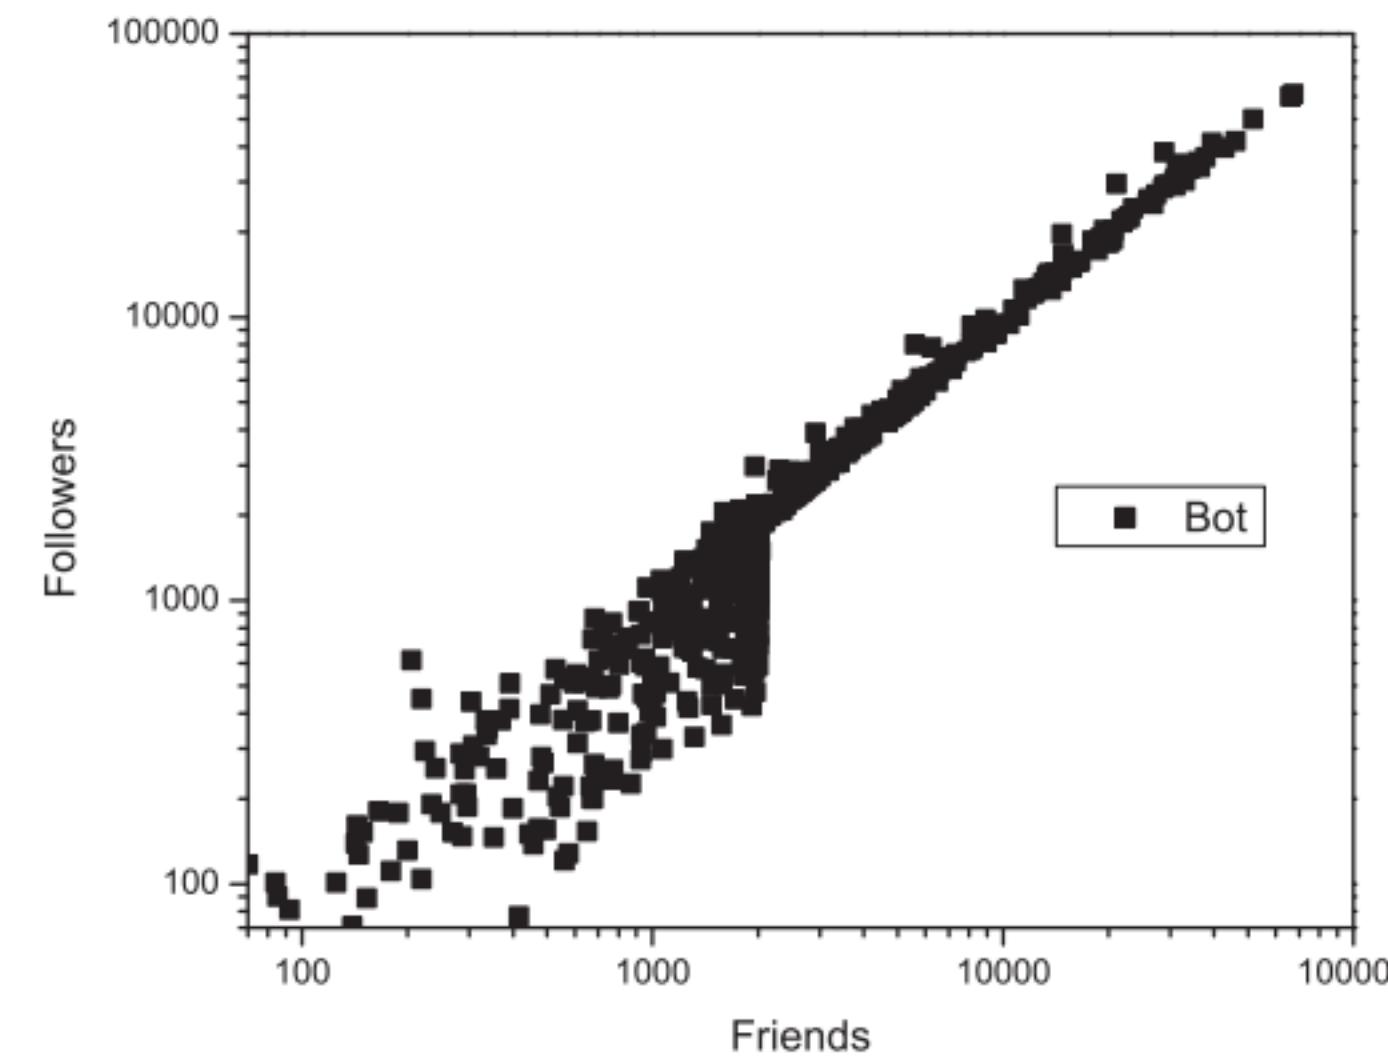
(c) Cyborg



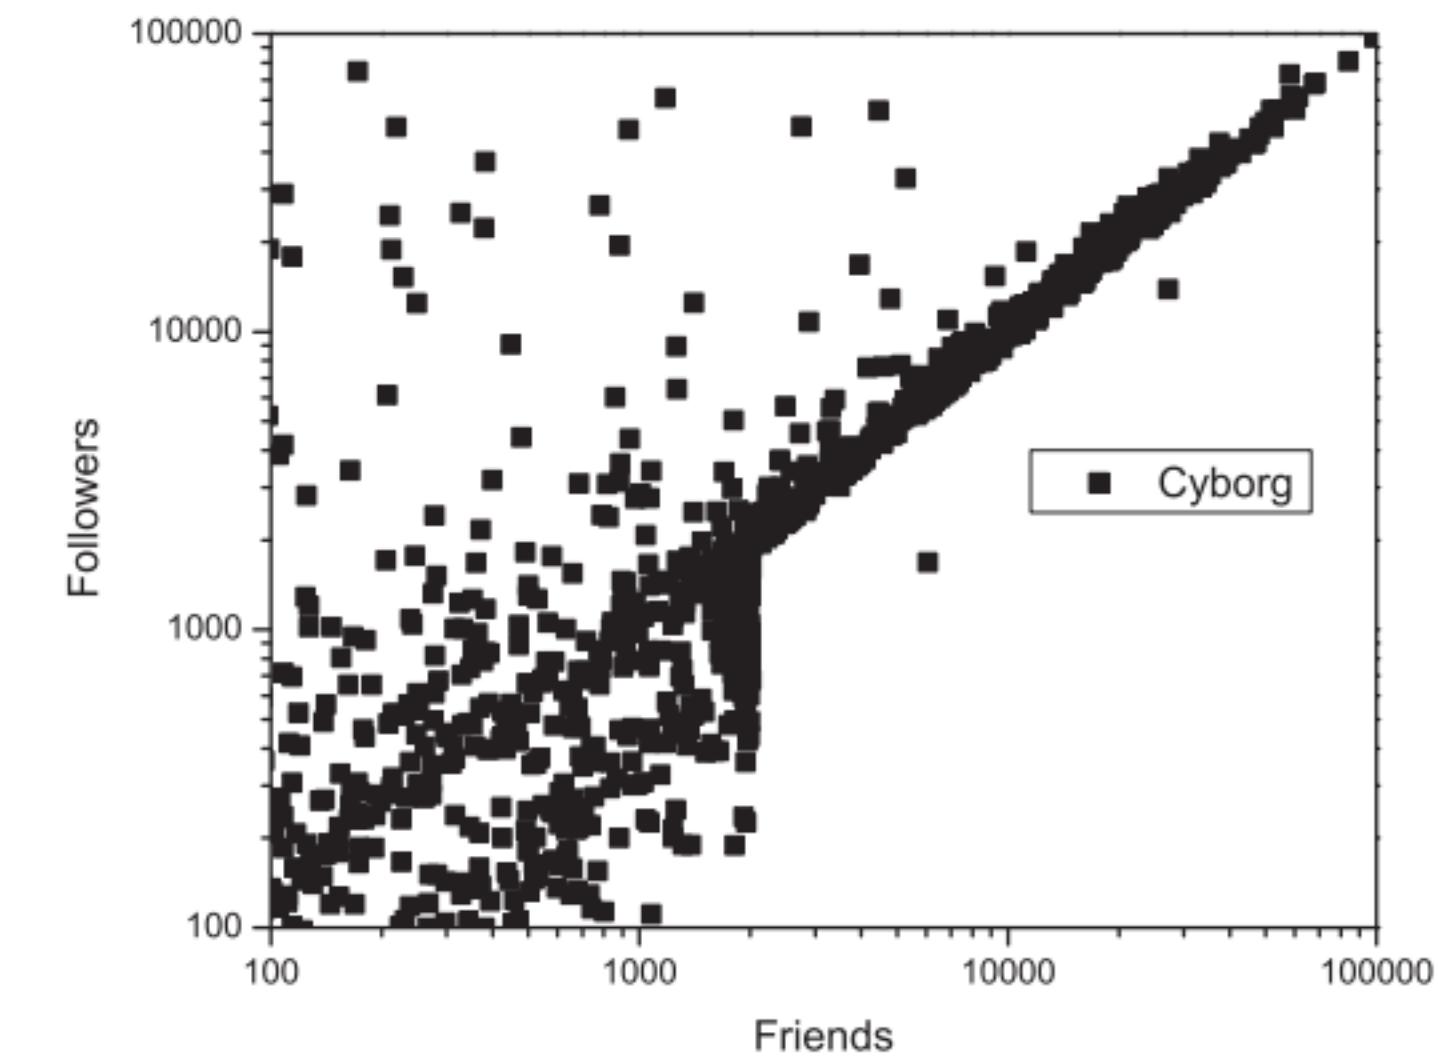
# Social Relationships



(a) Human



(b) Bot



(c) Cyborg

**Takeaway: Automated users follow more users than those that follow them**



# Behavioral Patterns

Social Relationships

**Tweeting Habits**

Tweet Properties



# Tweeting Habits: Entropy of Tweet Timing

- Entropy is the measure of *complexity* or *chaos* of a process
- In this context, the “process” is the *timing* of tweets



# Tweeting Habits: Entropy of Tweet Timing

- Entropy is the measure of *complexity* or *chaos* of a process
- In this context, the “process” is the *timing* of tweets

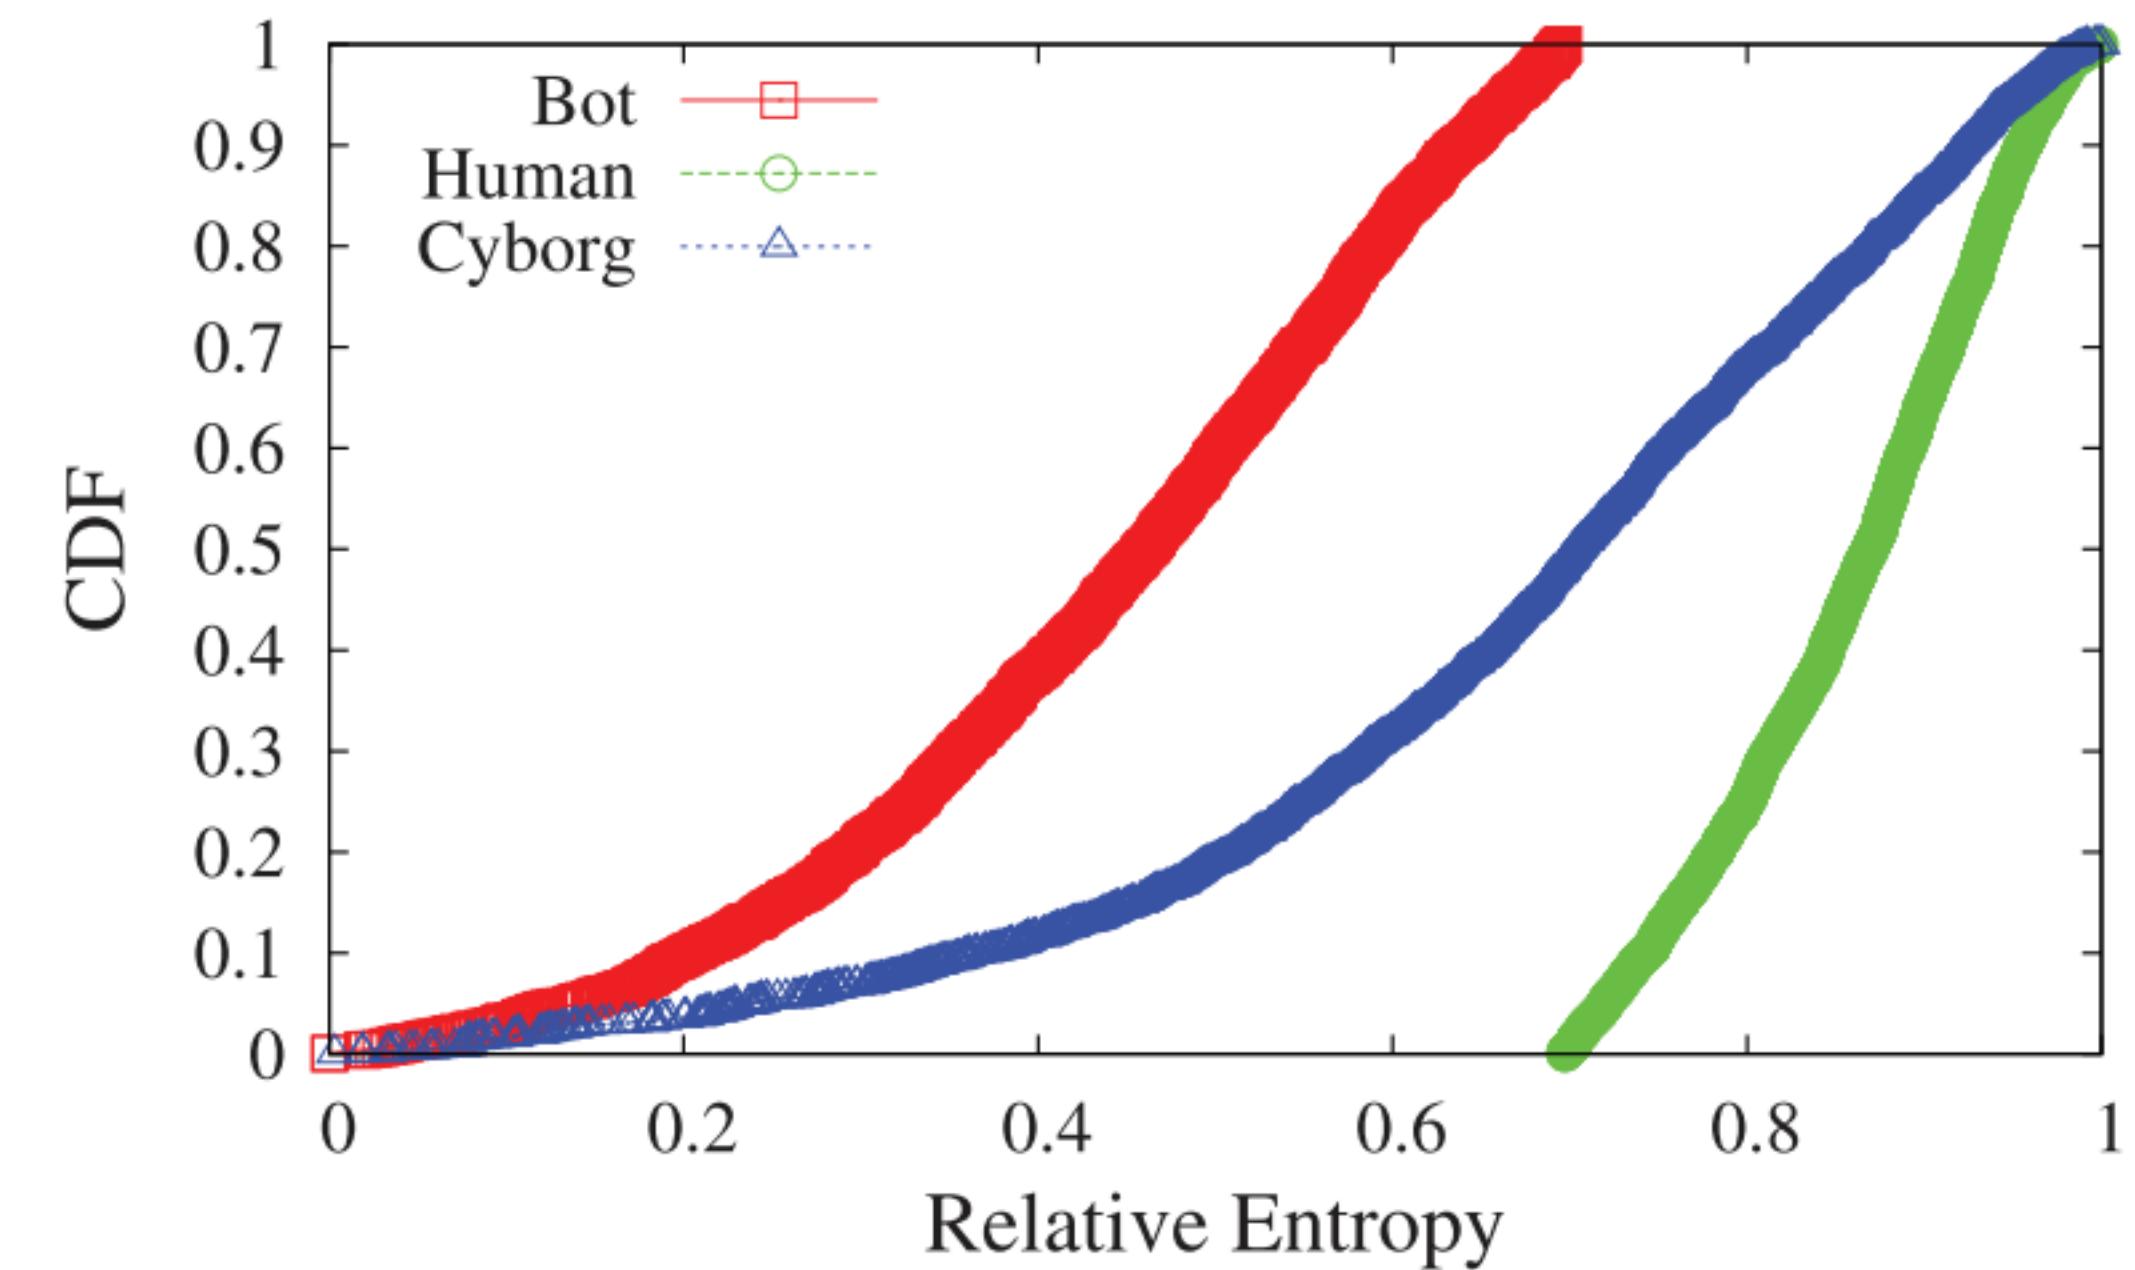
**Key intuition: Automated accounts tweet at regular intervals, while humans do not**



# Tweeting Habits: Entropy of Tweet Timing

$$\frac{entropy(u_i)}{max(entropy(u_1), entropy(u_2), \dots, entropy(u_n))}$$

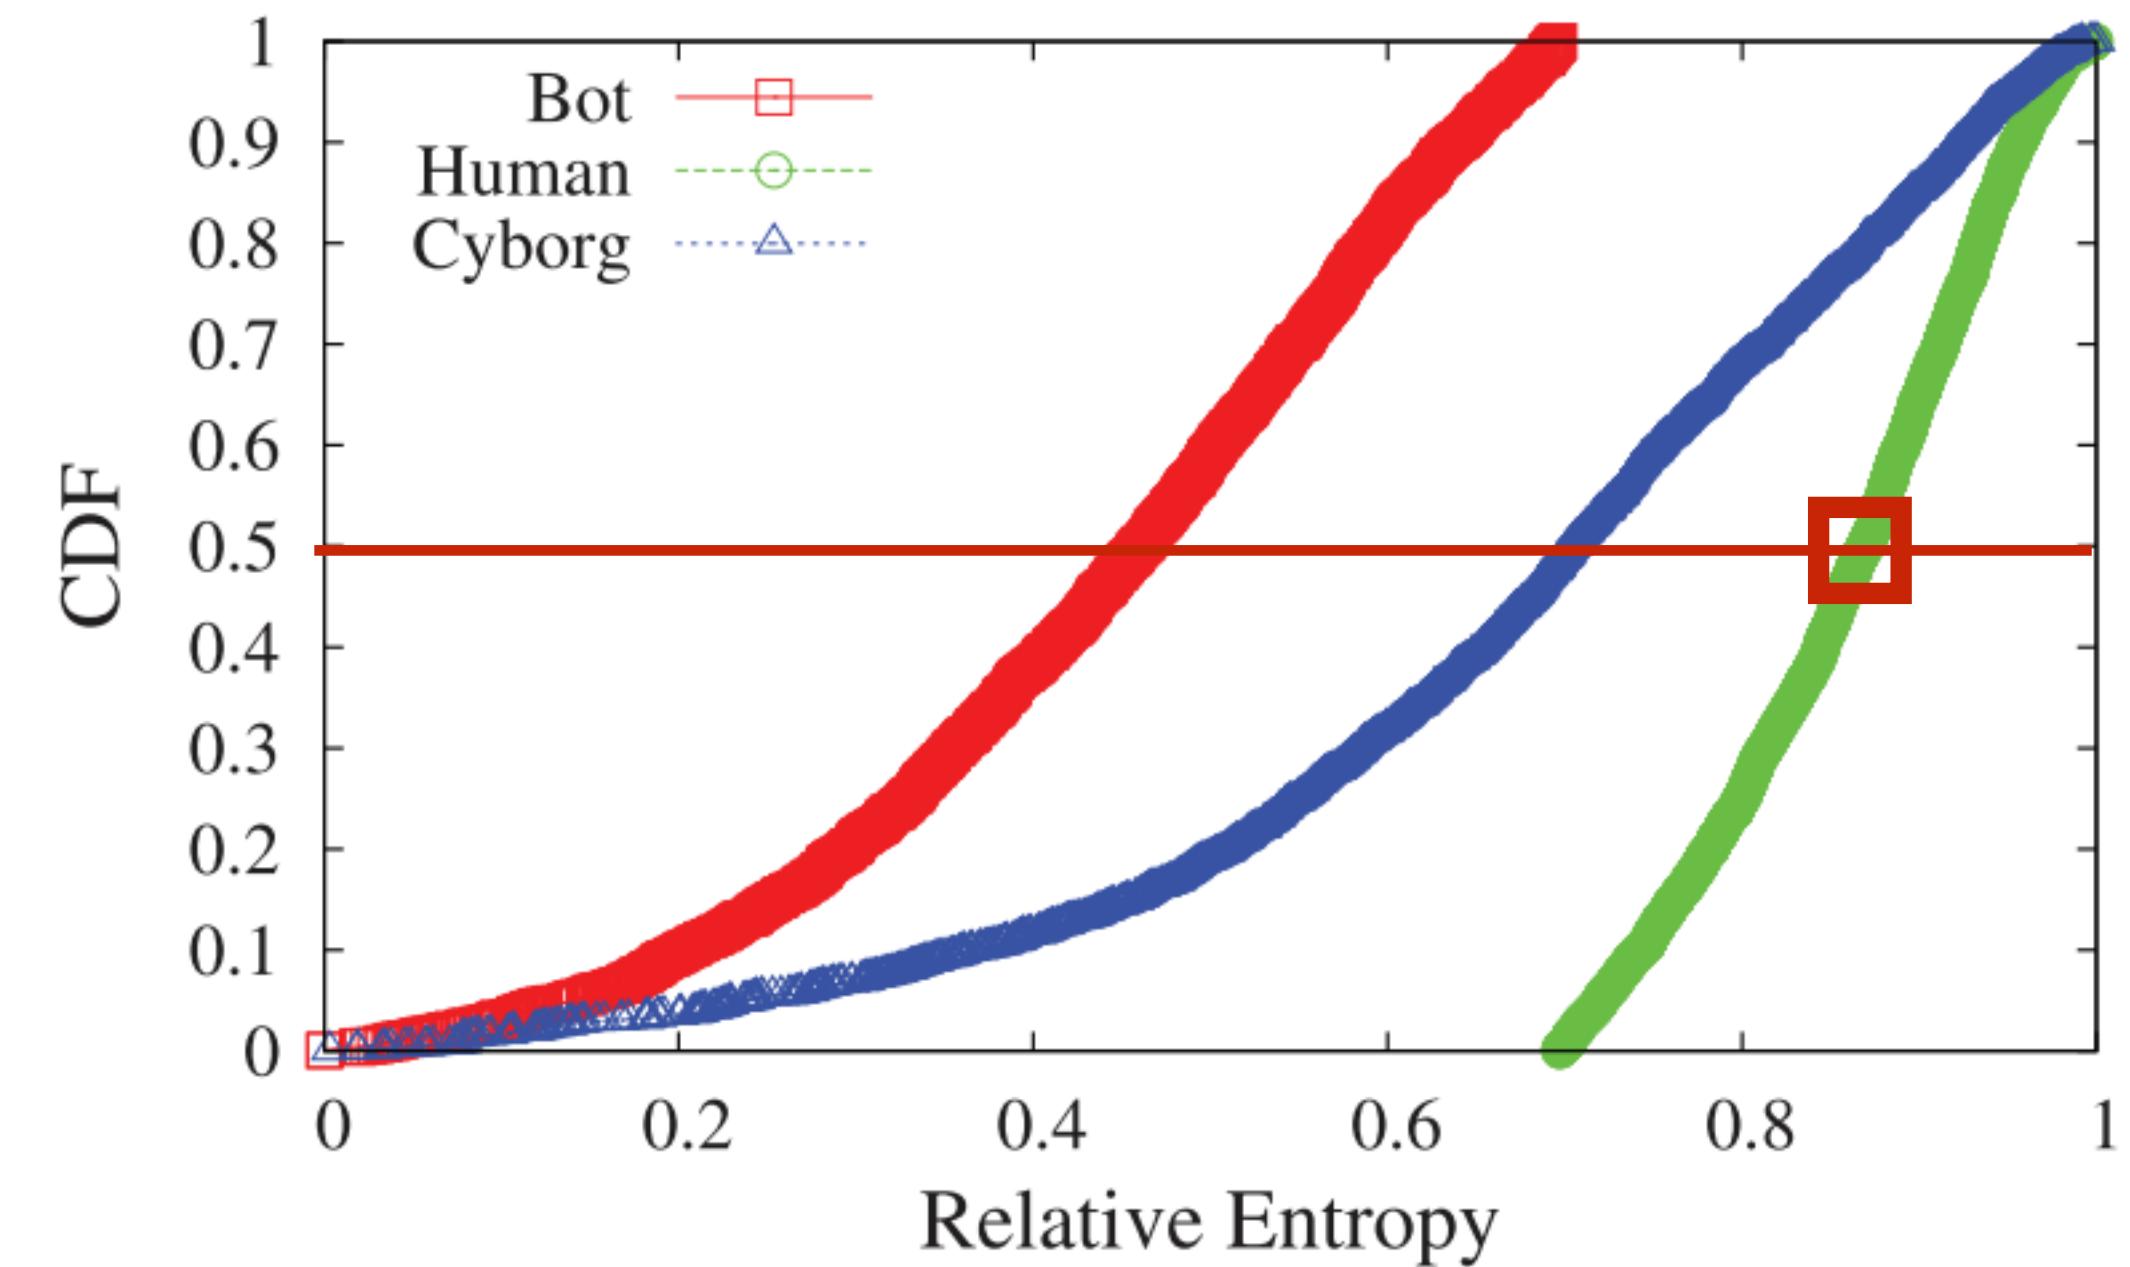
- Normalize entropy of each account by largest entropy in the dataset



# Tweeting Habits: Entropy of Tweet Timing

$$\frac{entropy(u_i)}{\max(entropy(u_1), entropy(u_2), \dots, entropy(u_n))}$$

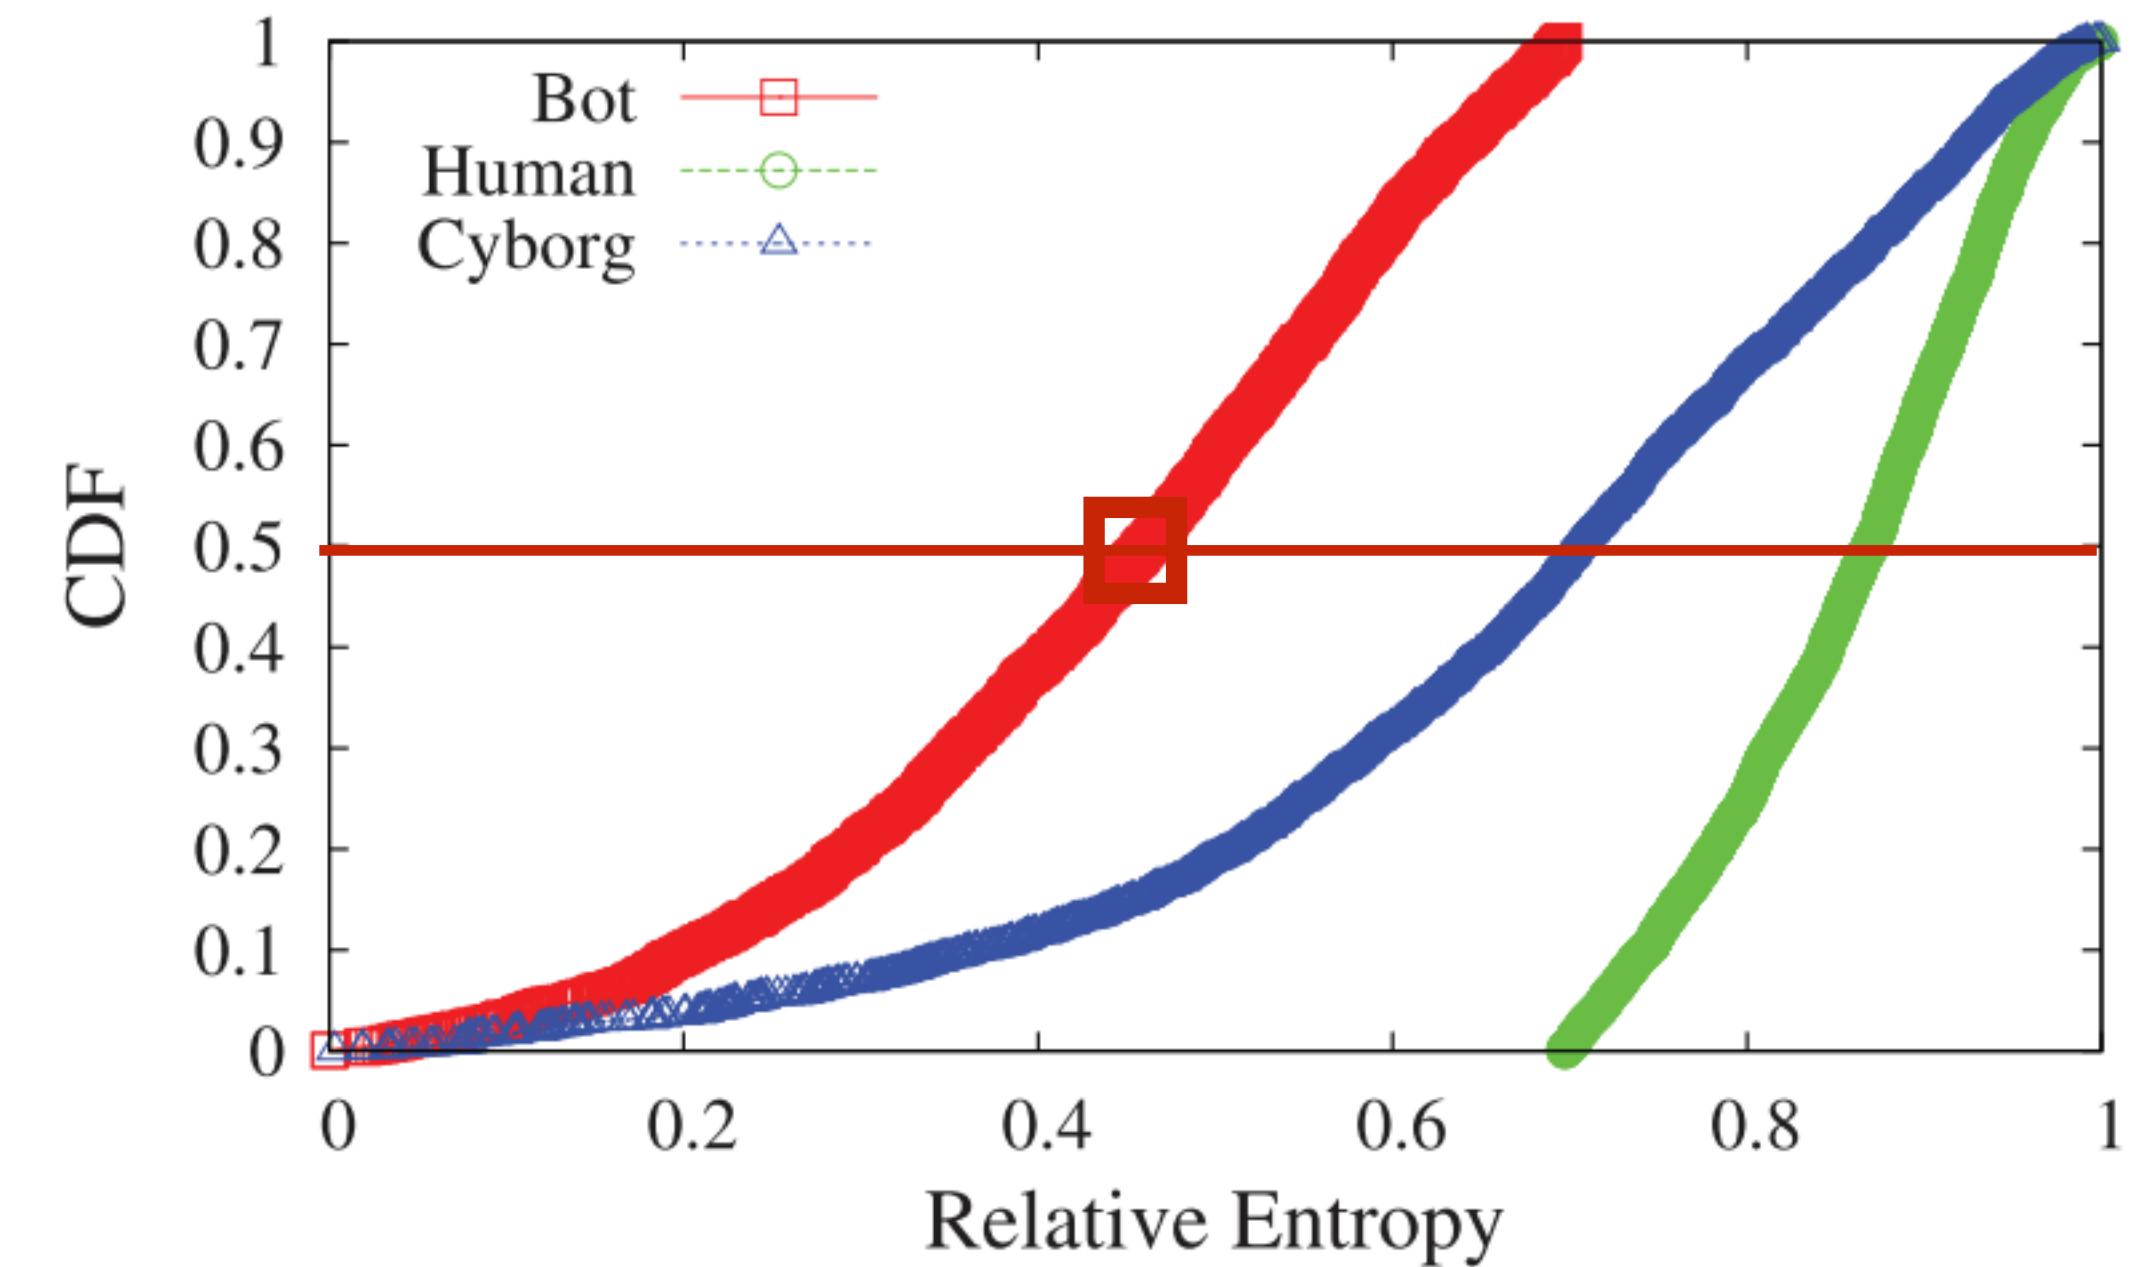
- Normalize entropy of each account by largest entropy in the dataset
- Humans exhibit highest entropy - median rel. entropy is 0.9



# Tweeting Habits: Entropy of Tweet Timing

$$\frac{entropy(u_i)}{\max(entropy(u_1), entropy(u_2), \dots, entropy(u_n))}$$

- Normalize entropy of each account by largest entropy in the dataset
- Humans exhibit highest entropy - median rel. entropy is 0.9
- Bots exhibit lowest entropy - median rel. entropy is 0.4

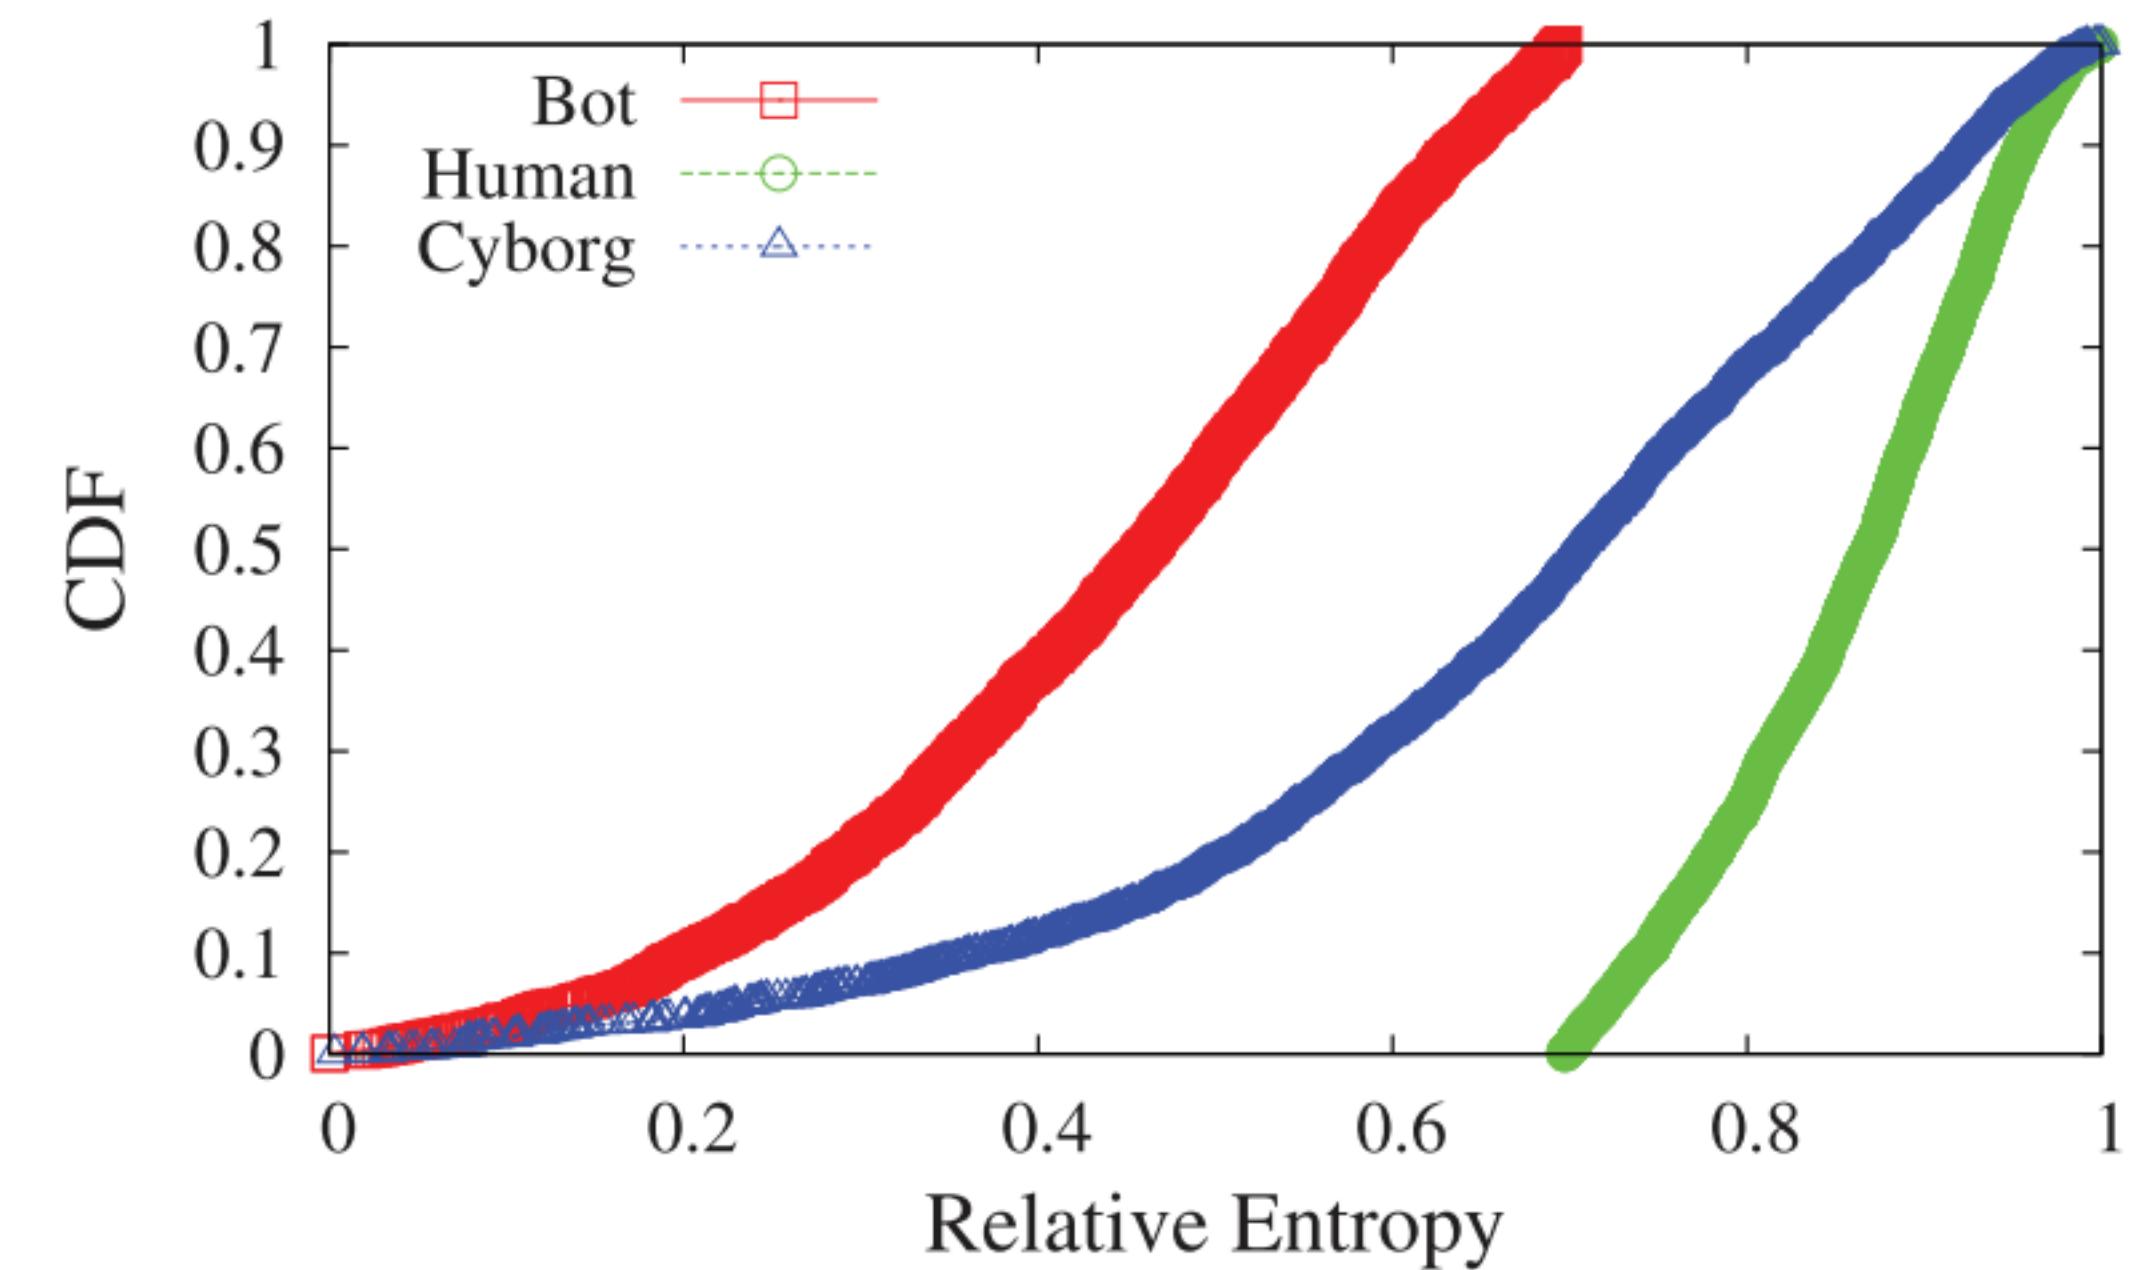


# Tweeting Habits: Entropy of Tweet Timing

$$\frac{entropy(u_i)}{\max(entropy(u_1), entropy(u_2), \dots, entropy(u_n))}$$

- Normalize entropy of each account by largest entropy in the dataset
- Humans exhibit highest entropy - median rel. entropy is 0.9
- Bots exhibit lowest entropy - median rel. entropy is 0.4

**Takeaway: Human tweeting behavior is more random than bot tweeting behavior**



# Tweeting Habits: What tools are used to Tweet?

- Users can Tweet via a wide-array of channels
  - Web
  - Mobile
  - Third party applications
  - APIs

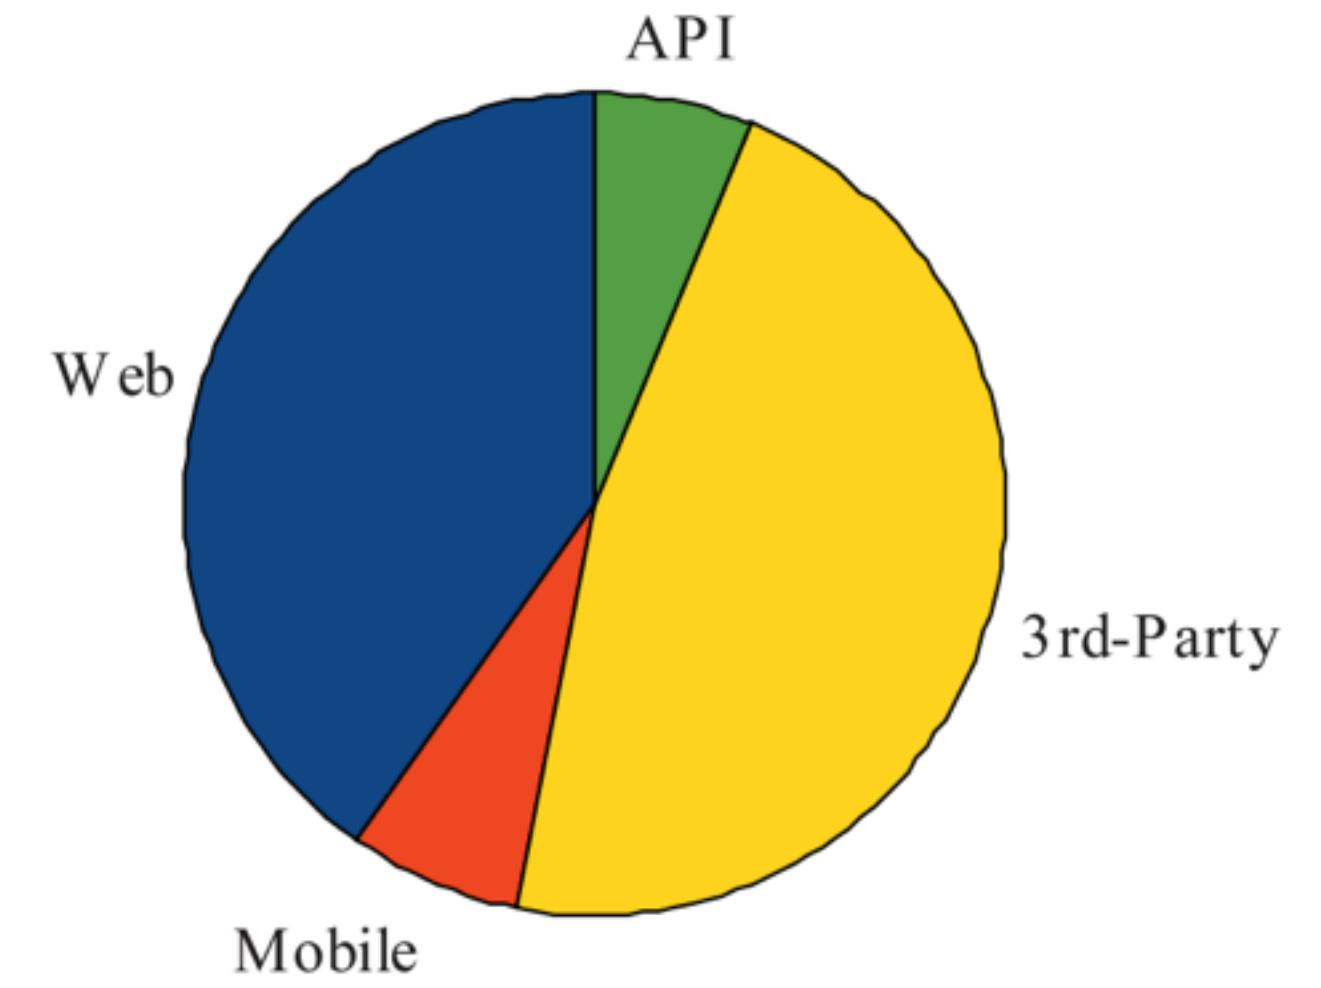
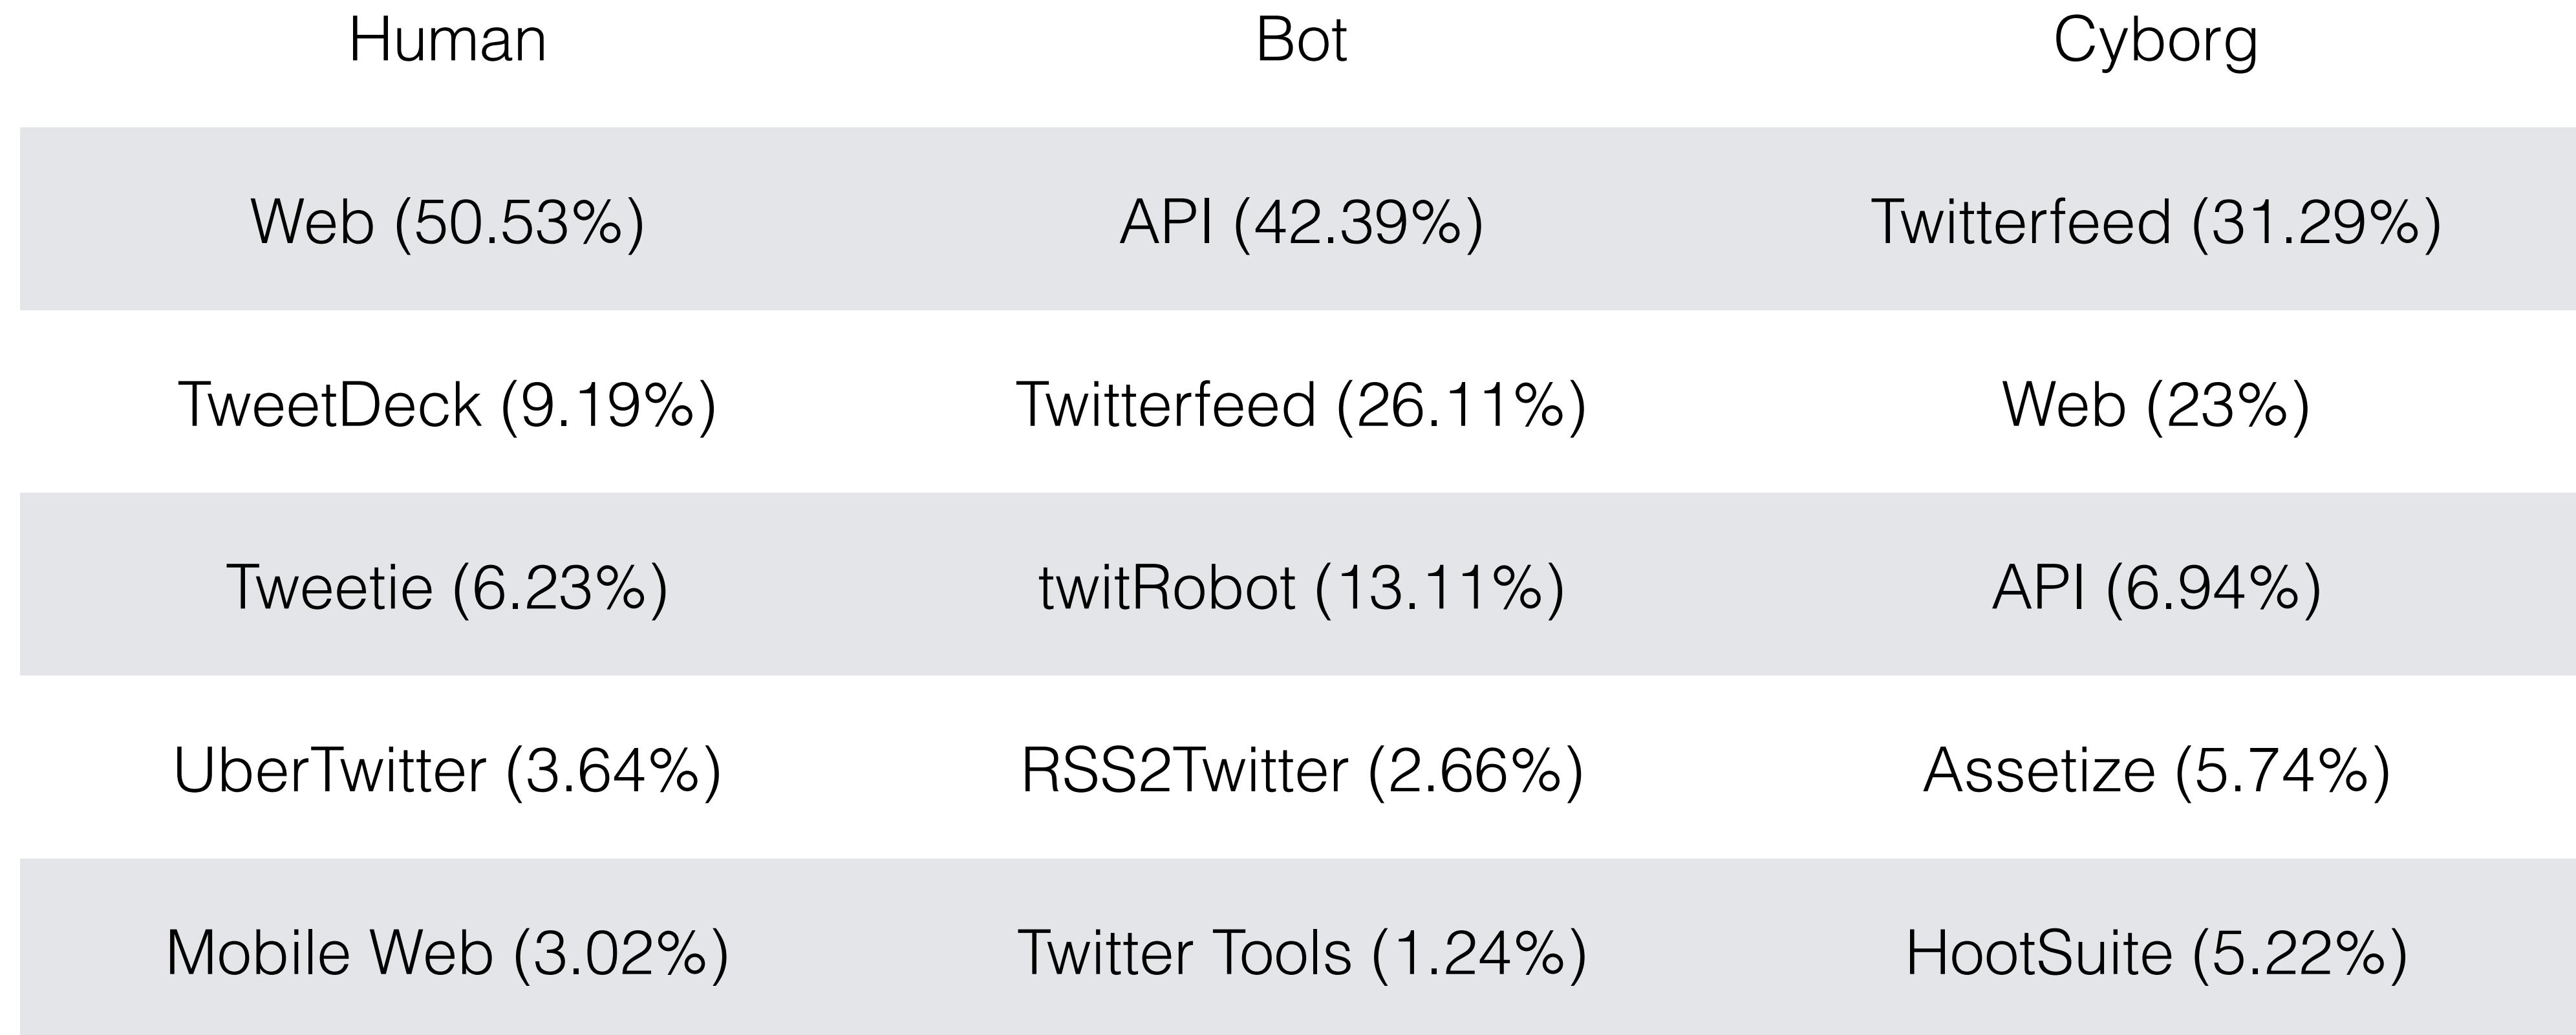


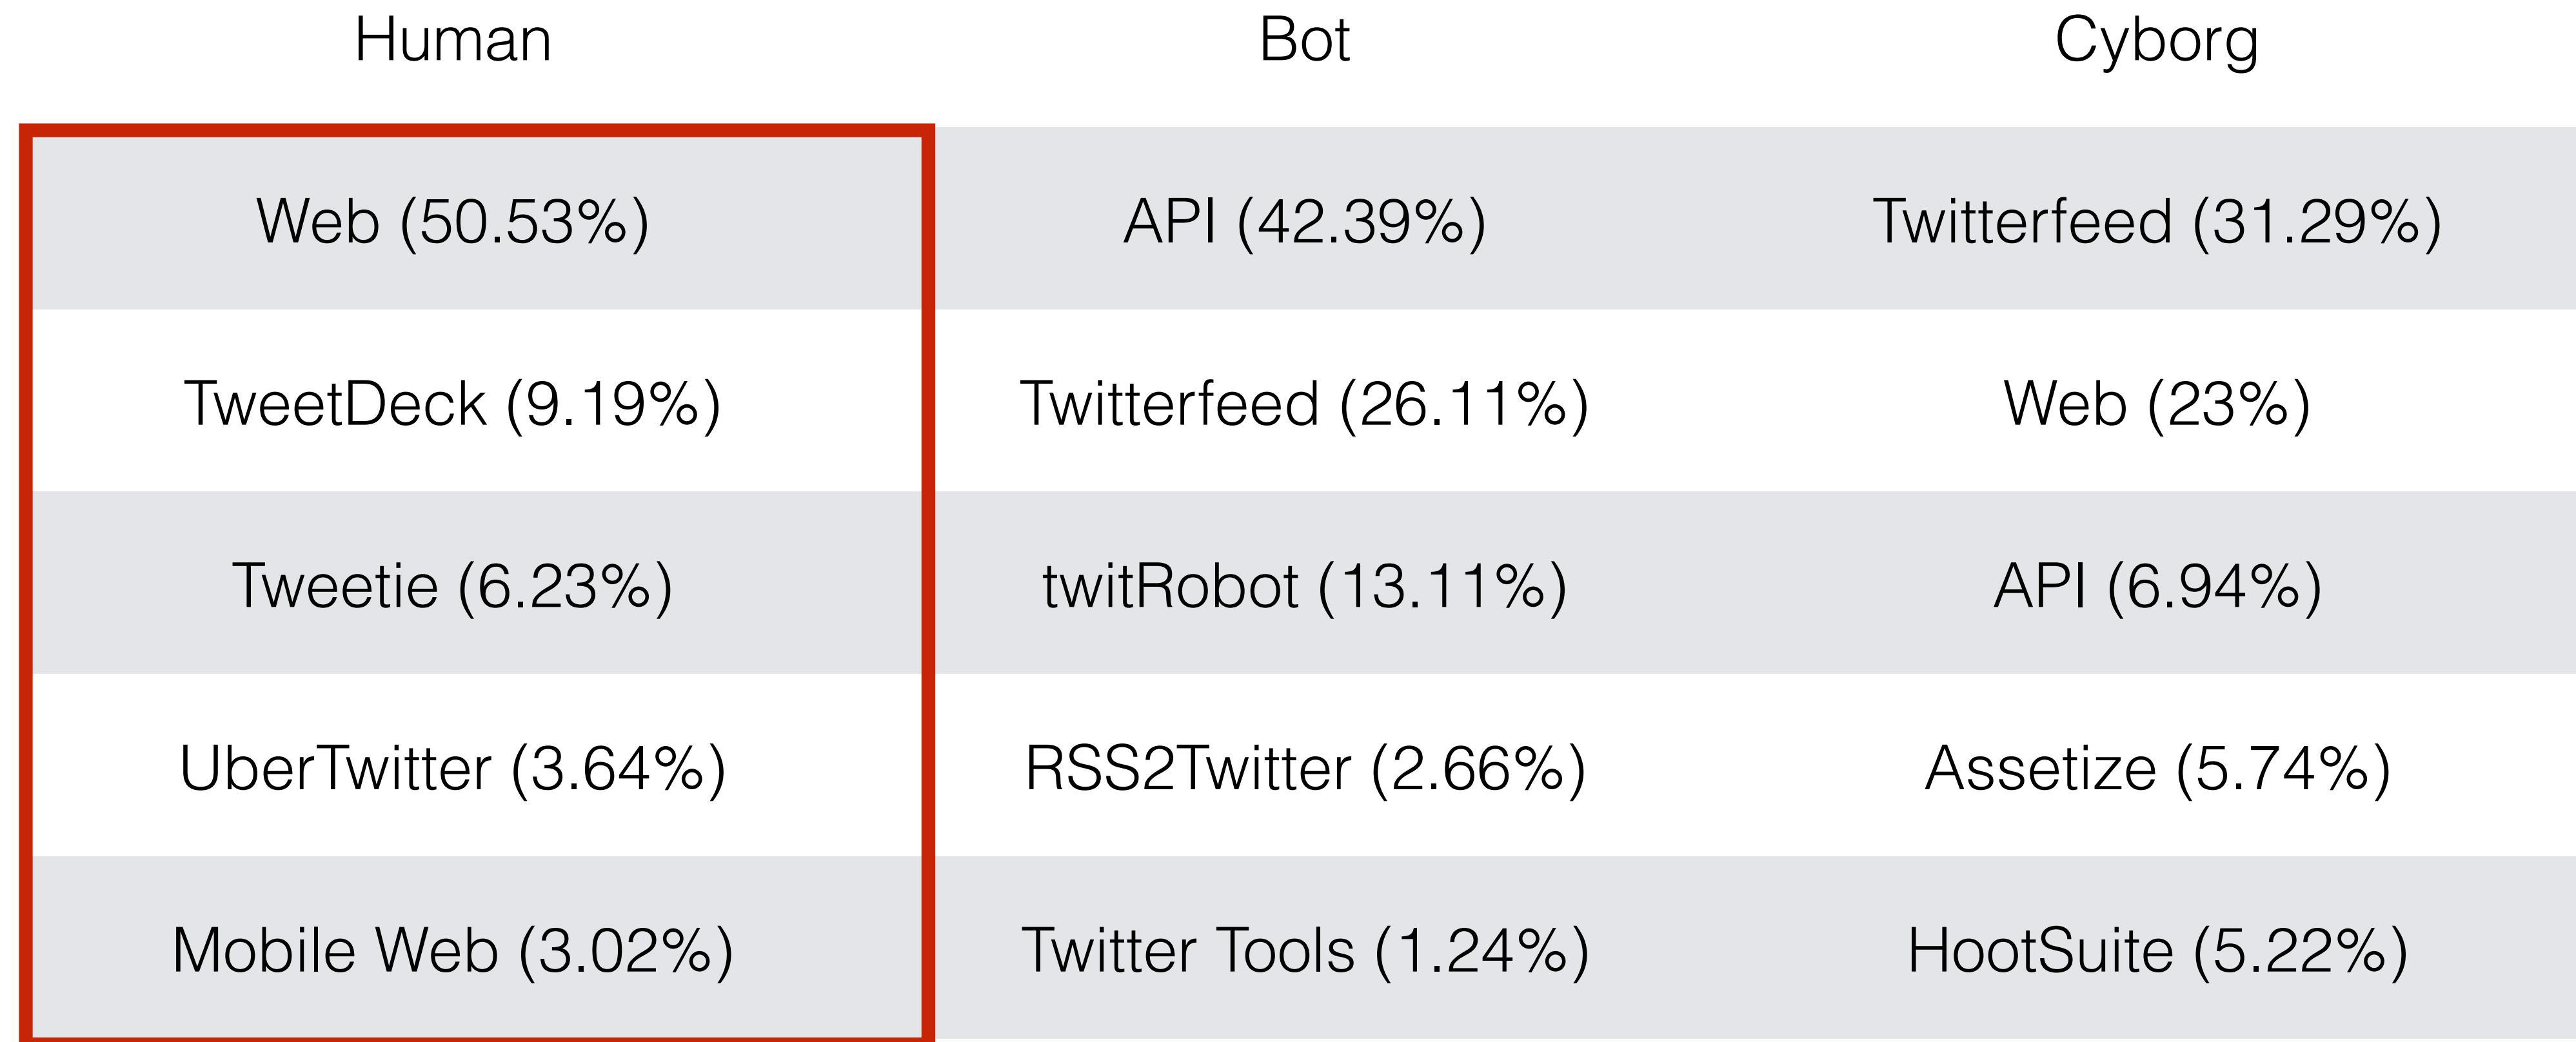
Fig. 6. Tweeting device makeup.



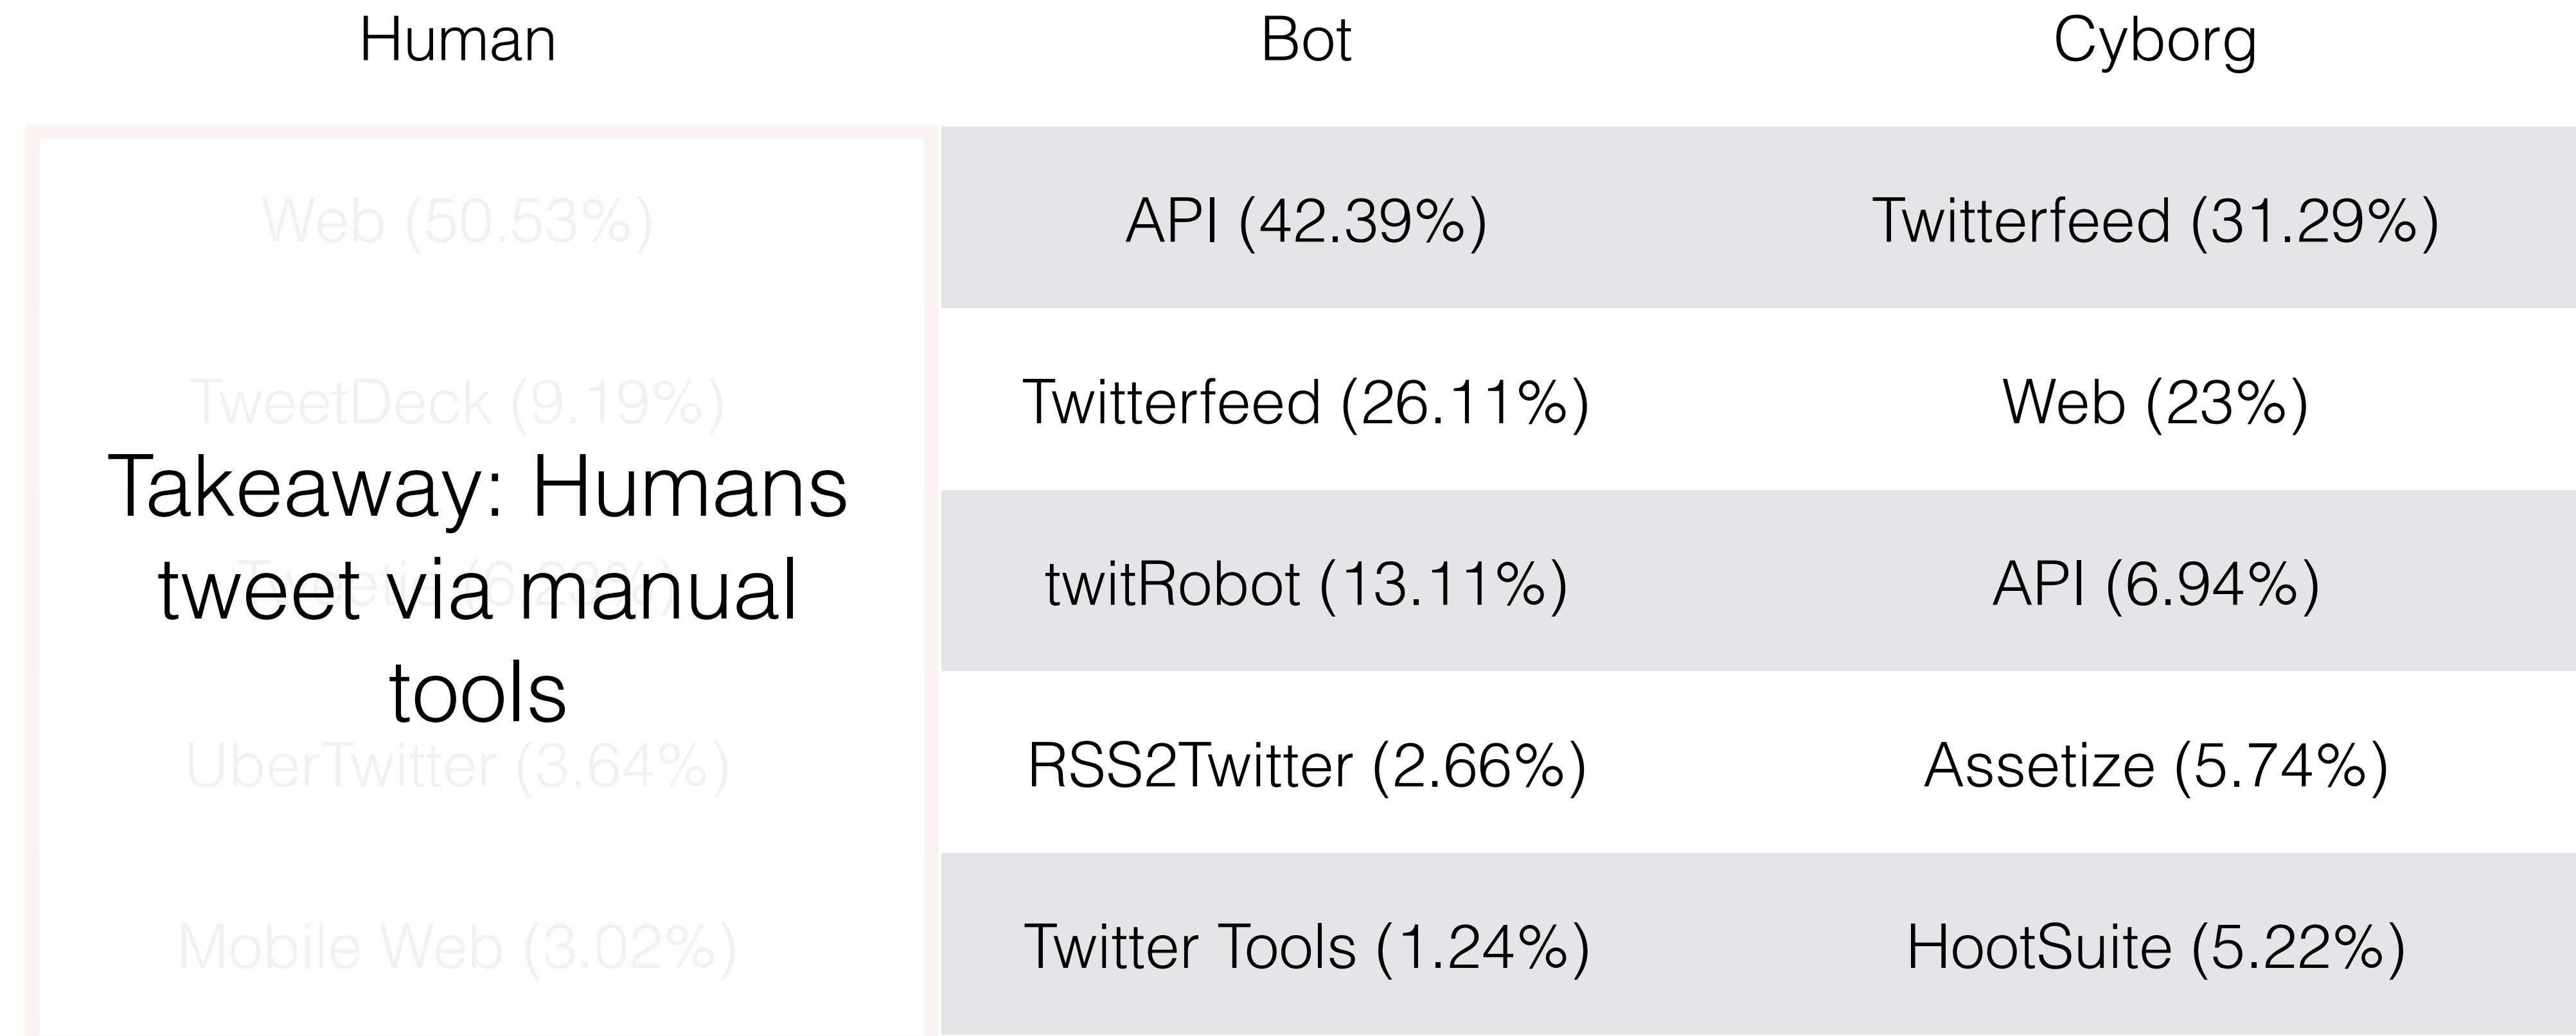
# Tweeting Habits: What tools are used to Tweet?



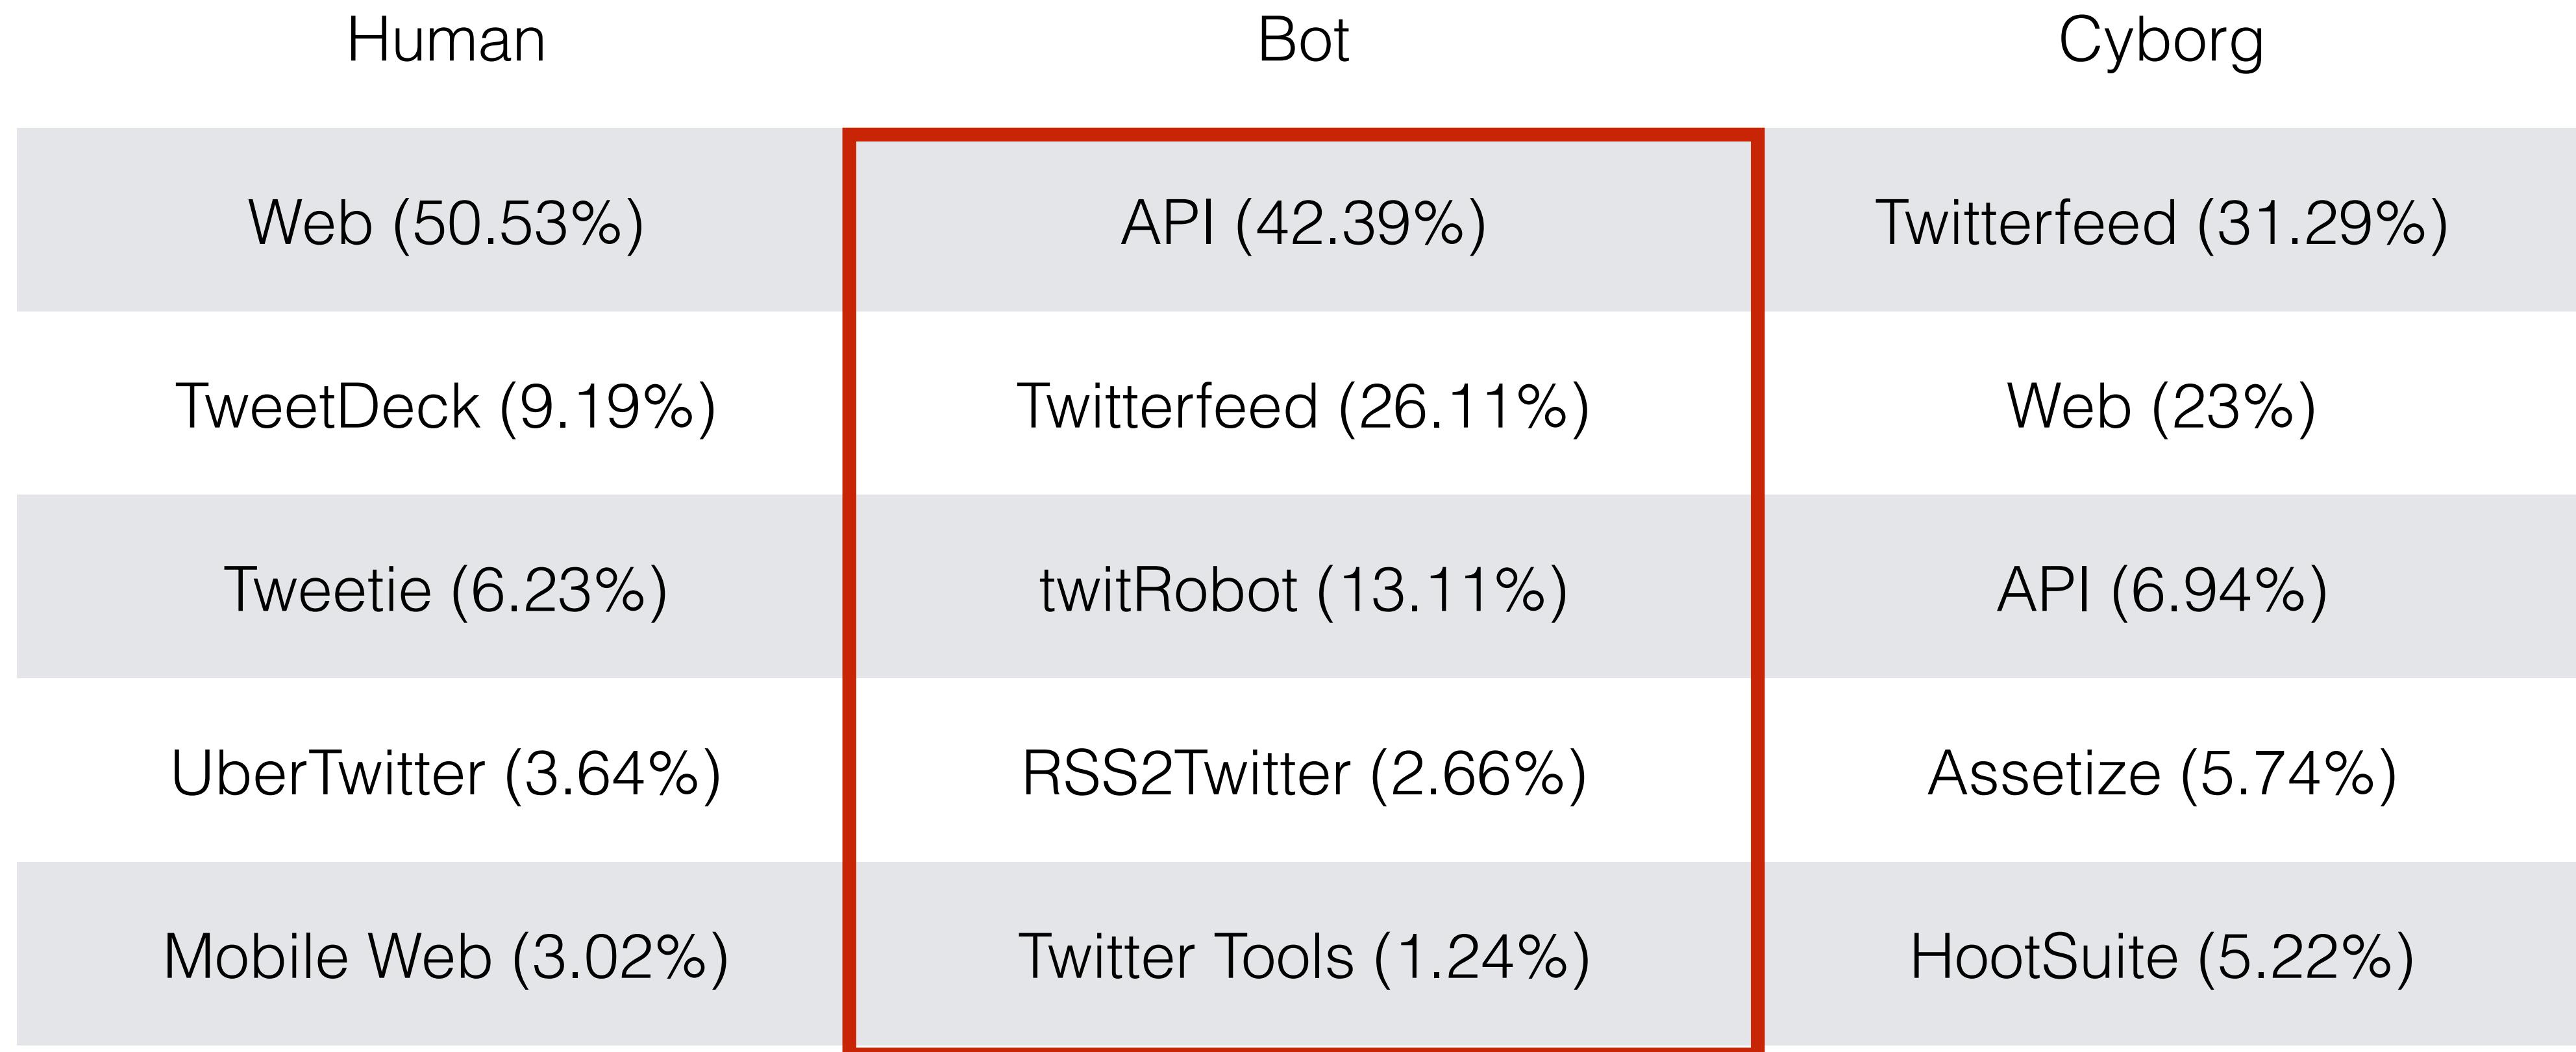
# Tweeting Habits: What tools are used to Tweet?



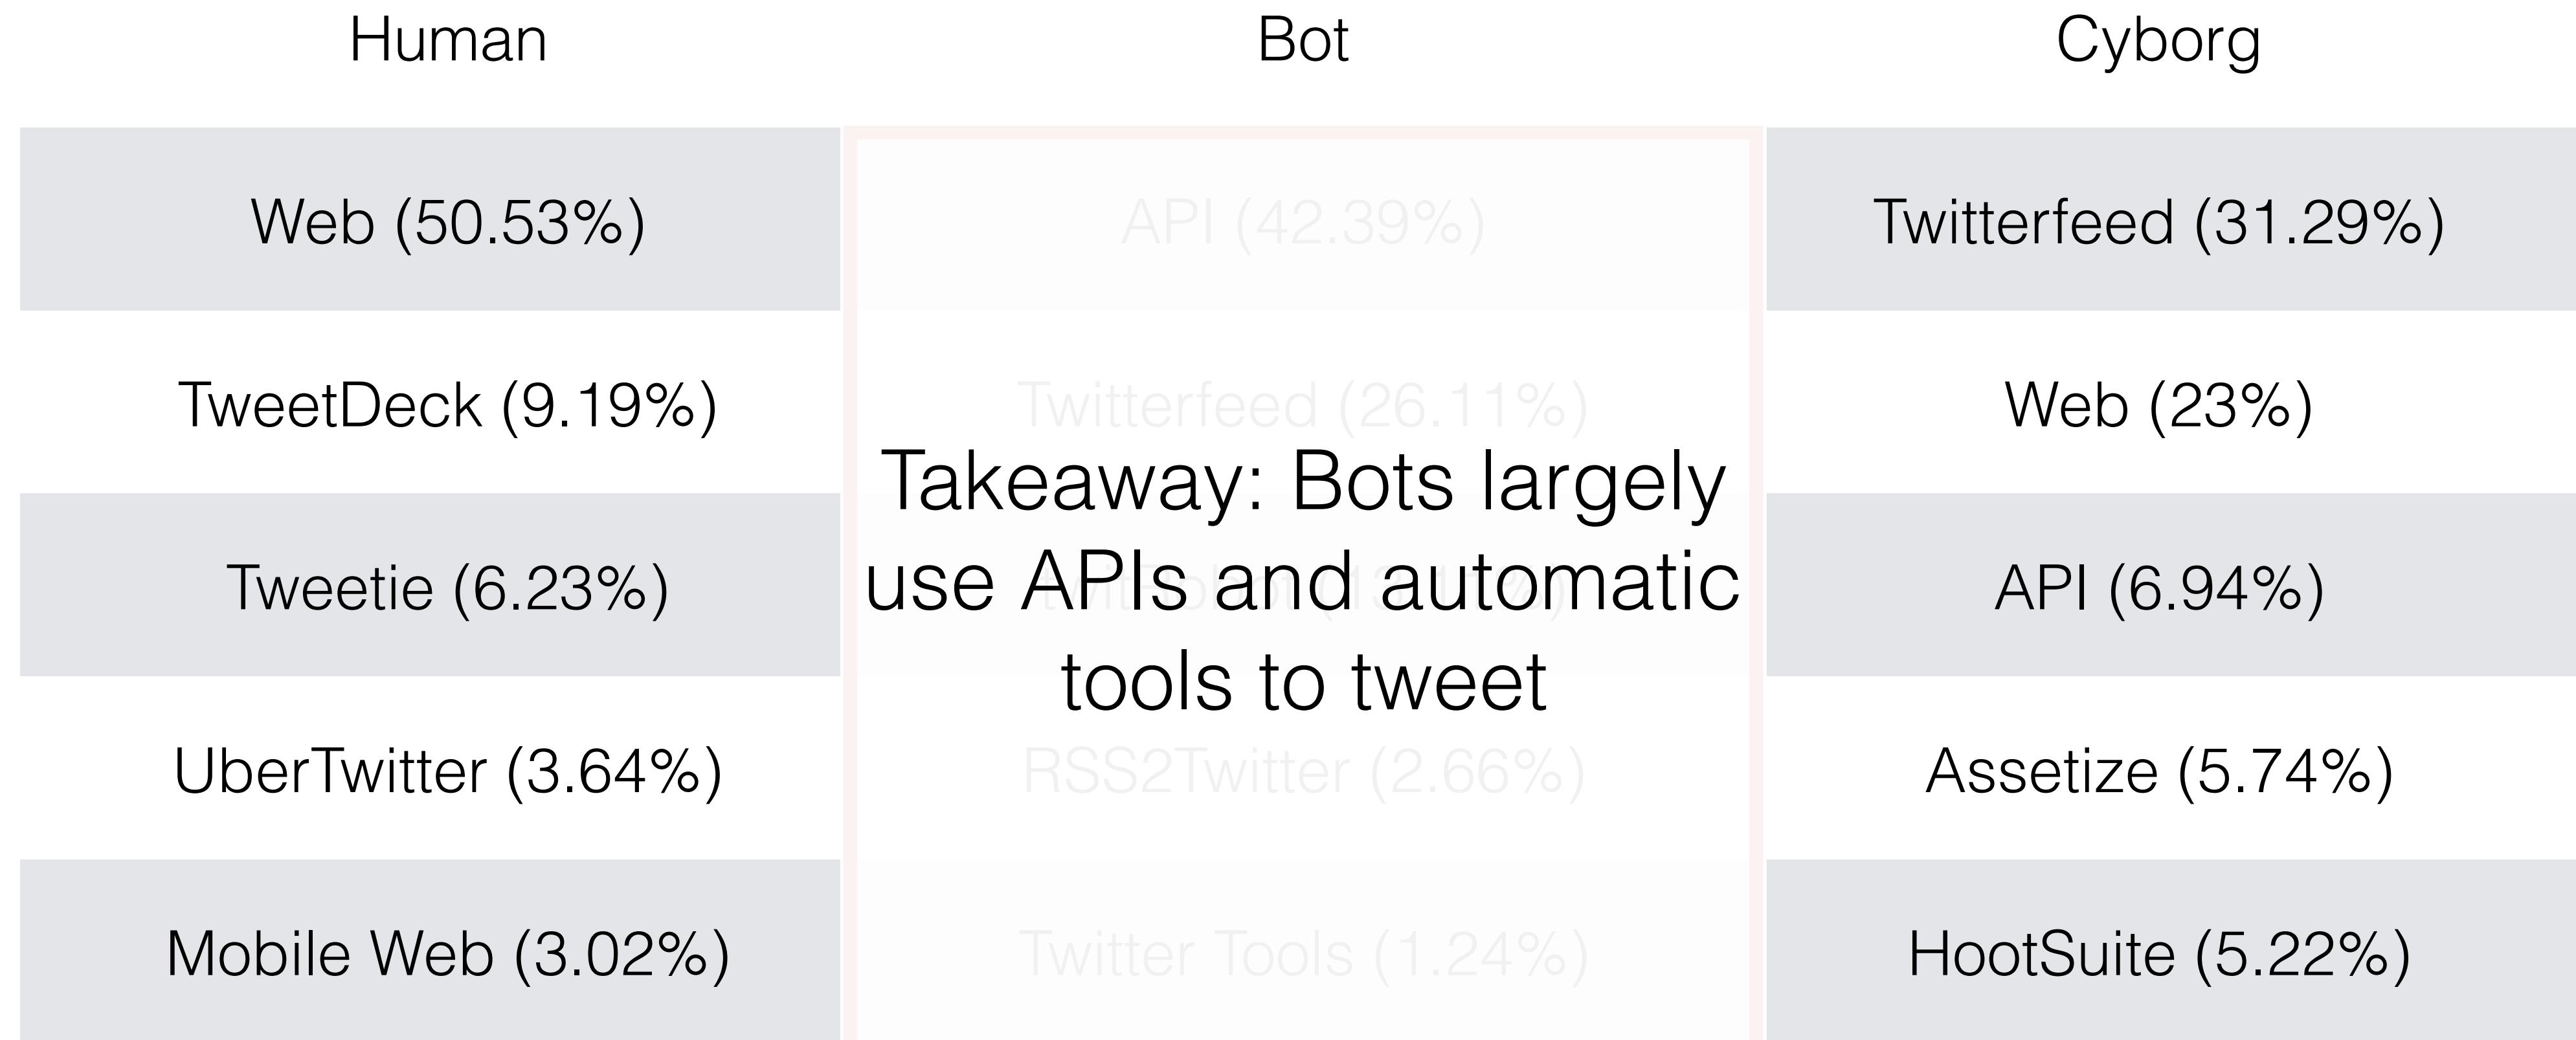
# Tweeting Habits: What tools are used to Tweet?



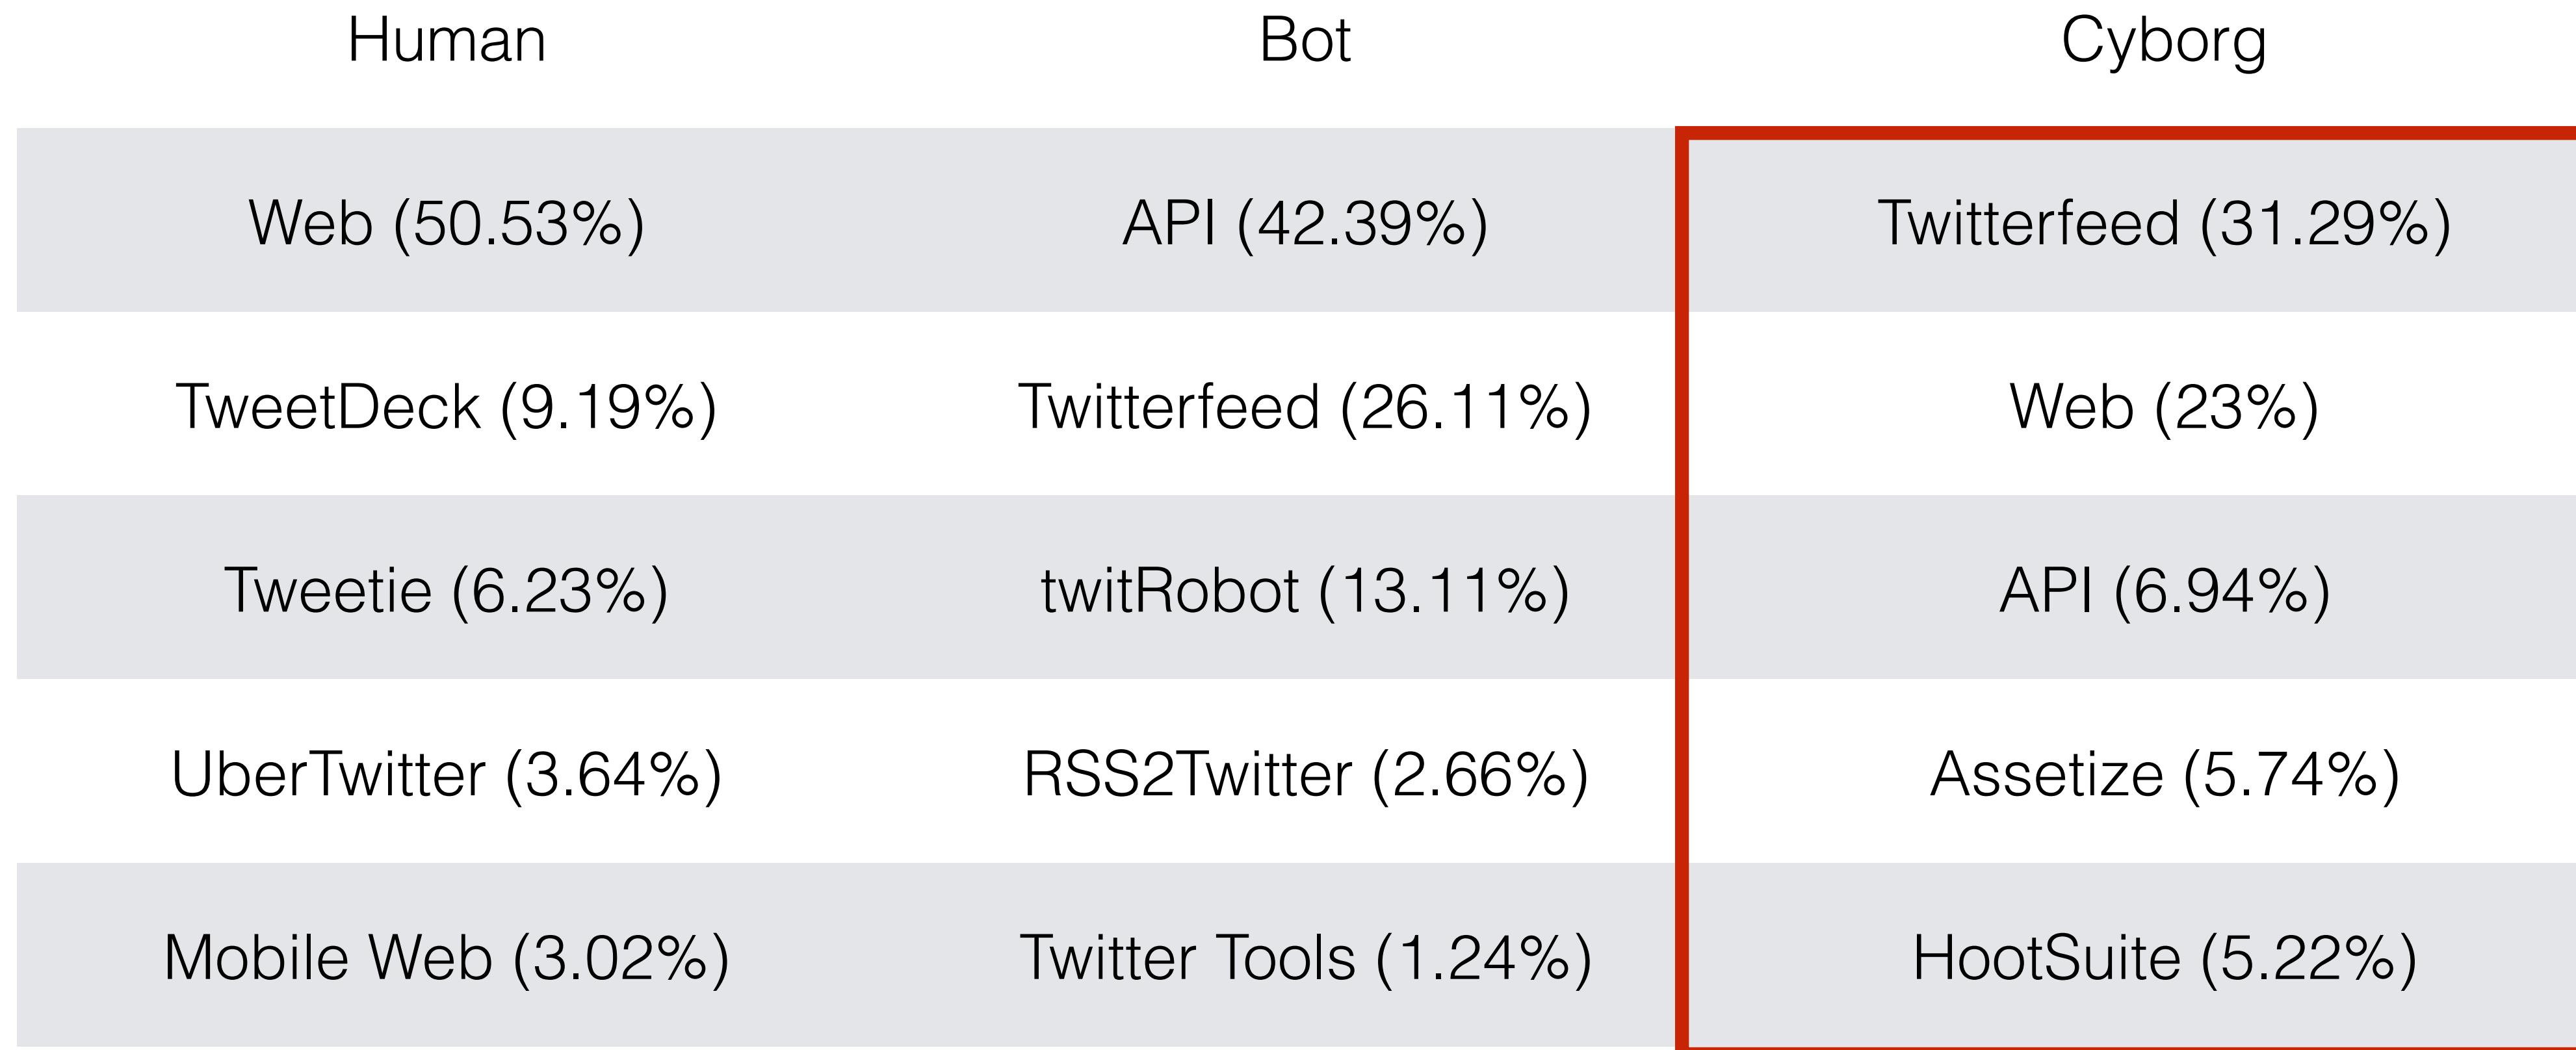
# Tweeting Habits: What tools are used to Tweet?



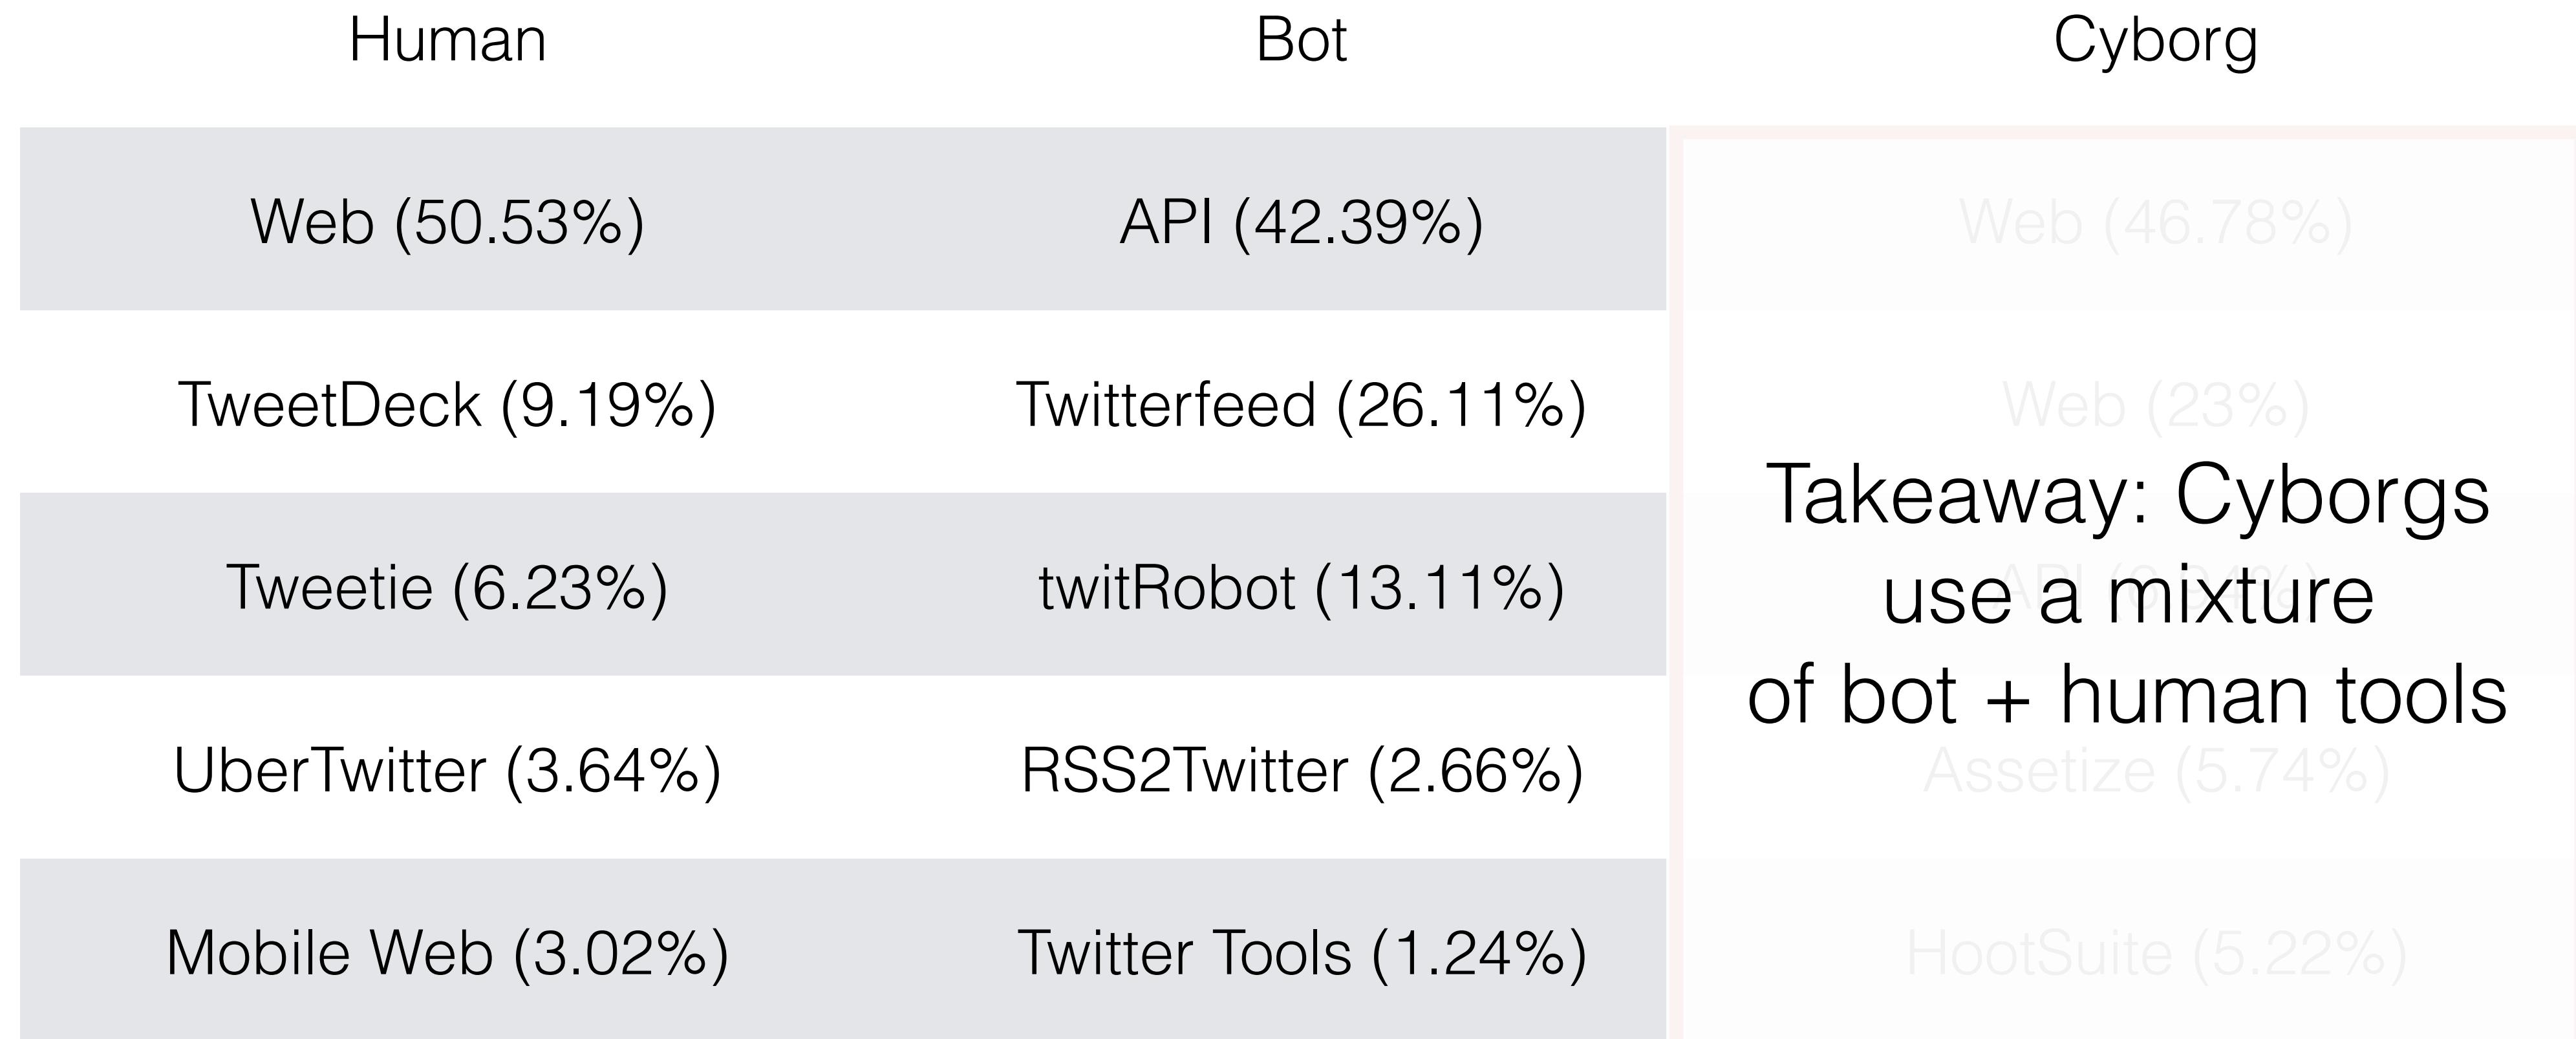
# Tweeting Habits: What tools are used to Tweet?



# Tweeting Habits: What tools are used to Tweet?



# Tweeting Habits: What tools are used to Tweet?



# Behavioral Patterns

Social Relationships

Tweeting Habits

**Tweet Properties**



# Tweet Properties: URLs

- The # of URLs in bot tweets is higher than those of human + cyborg
- Some bots tweet more than one URL per tweet!

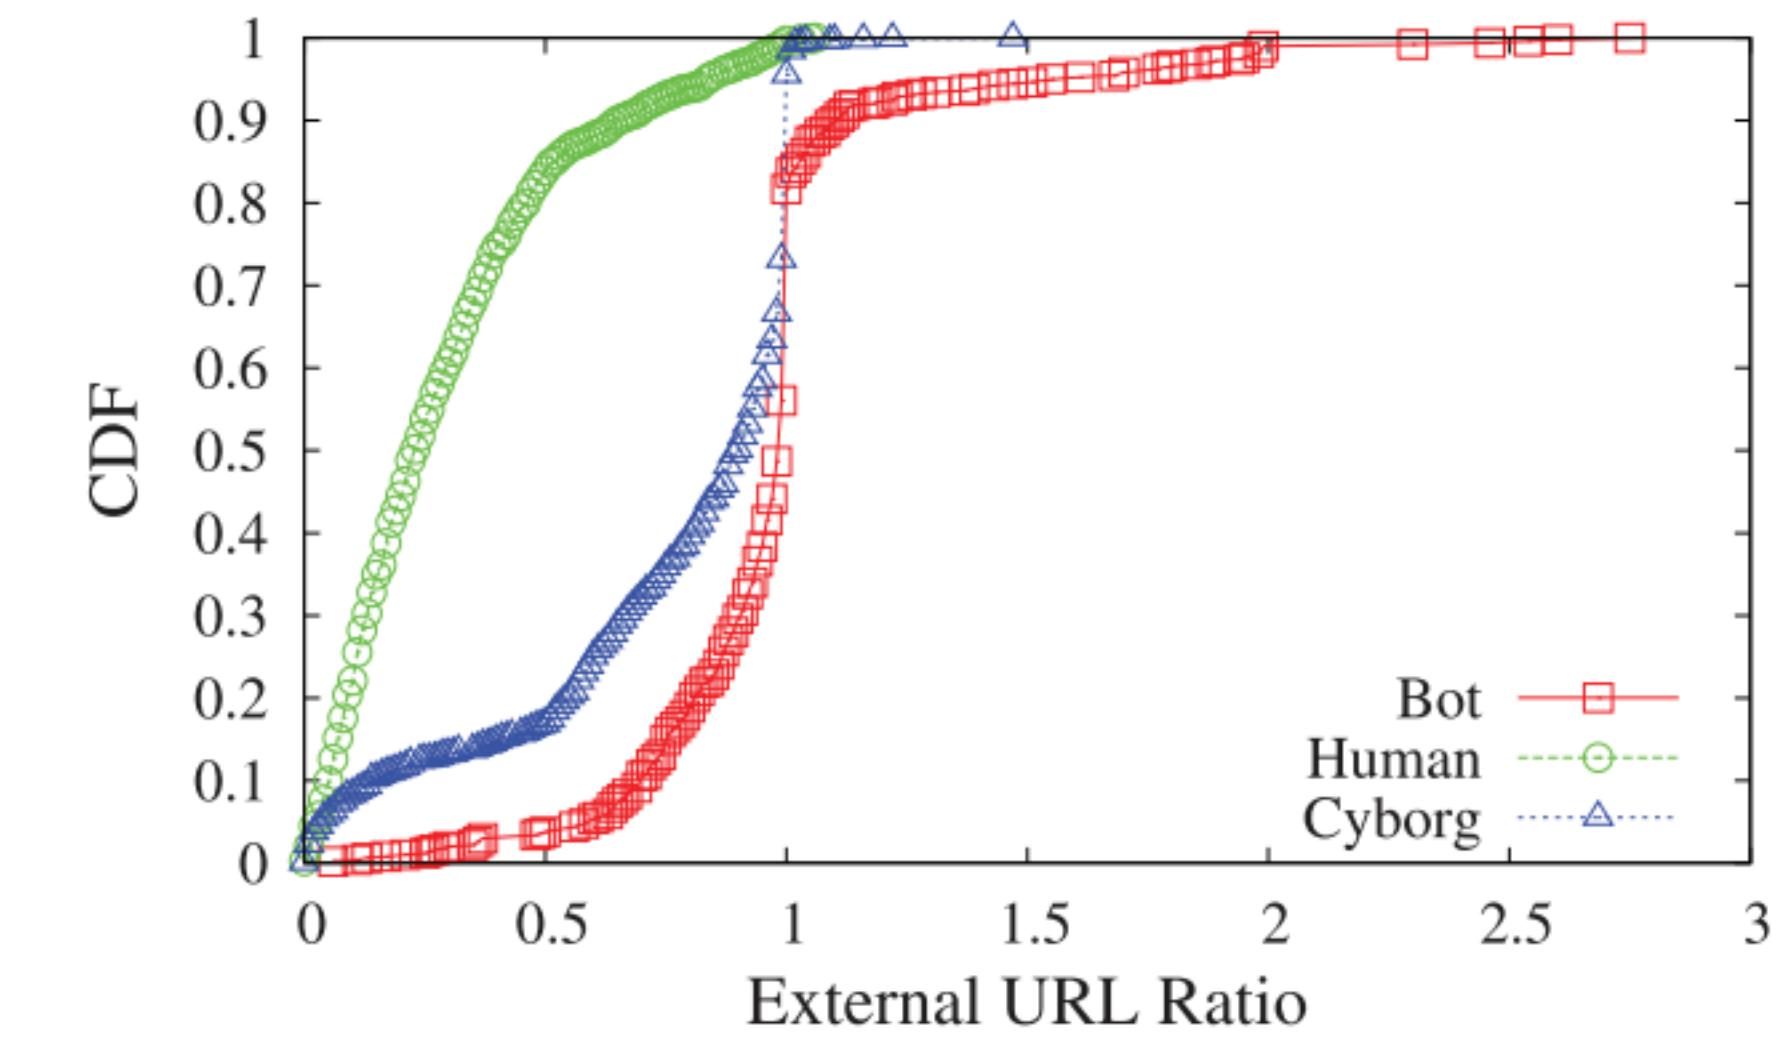


Fig. 7. External URL ratio in tweets.



# Tweet Properties: URLs

- The # of URLs in bot tweets is higher than those of human + cyborg
- Some bots tweet more than one URL per tweet!
- Another useful property of URLs: appearing in **blacklists**

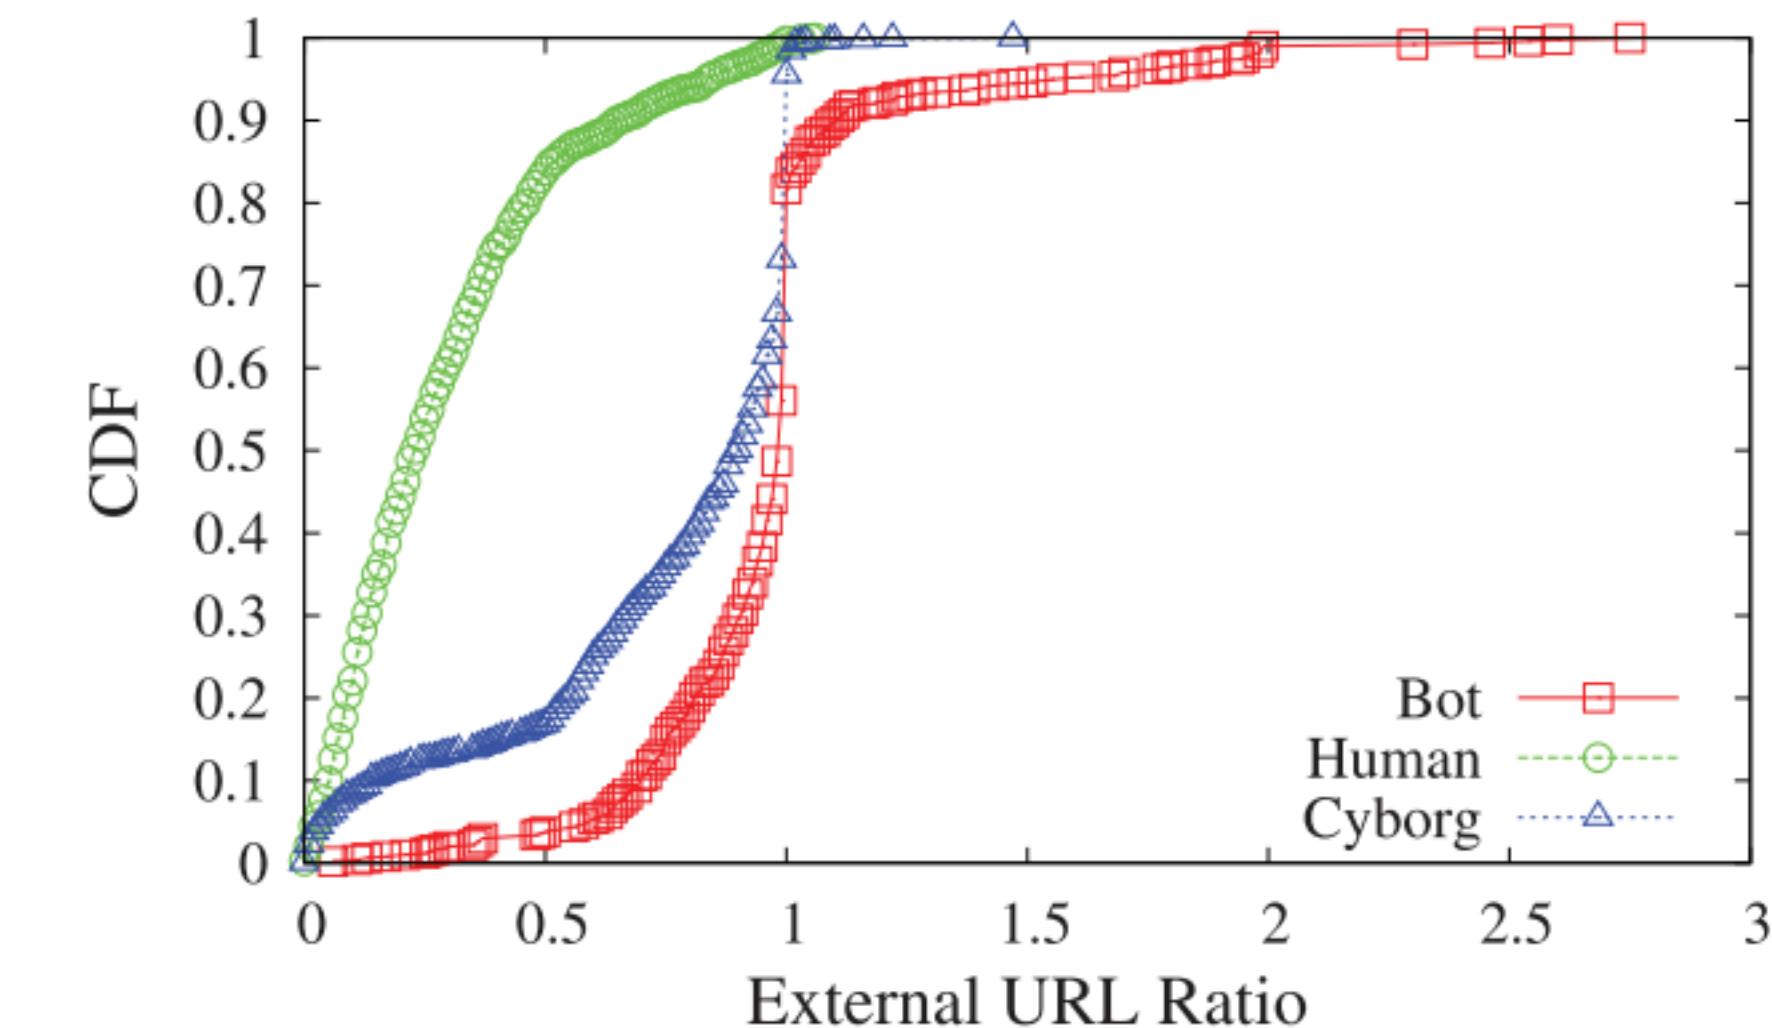


Fig. 7. External URL ratio in tweets.



# Tweet Properties: Spam

- Intuition: Use **spam** as a mechanism for identifying automation



# Tweet Properties: Spam

- Intuition: Use **spam** as a mechanism for identifying automation
- How? Naive Bayes for Spam Classification



# Identifying Spam via Naive Bayes

- Simple classification model built off of conditional probabilities
  - “What is the likelihood a message is spam given previous spam messages?”
  - Treat each word (or set of words!) as *independent events*
  - Researchers labeled tweets as “spam” in ground-truth collection

$$\Pr(S|W) = \frac{\Pr(W|S) \cdot \Pr(S)}{\Pr(W|S) \cdot \Pr(S) + \Pr(W|H) \cdot \Pr(H)}$$



**Takeaway: Humans, Bots,  
and Cyborgs exhibit  
different behavioral patterns**



Can a standard machine  
learning classifier use these  
differences to detect bots?



# Classifying Automation: Features

## Features

Timing Entropy

URL Ratio

Automated Device %

Bayesian Spam Detection

Manual Device %

Registration Date

Mention Ratio

Link Safety

Hashtag Ratio

Followers to Friends Ratio

Account Verification



# Classifying Automation: Features

## **Features**

Timing Entropy

URL Ratio

Automated Device %

Bayesian Spam Detection

Manual Device %

## **Registration Date**

## **Mention Ratio**

Link Safety

## **Hashtag Ratio**

Followers to Friends Ratio

## **Account Verification**



# Classifying Automation: Training and Testing

- Use a *random forest*, popular ML algorithm for labeled classification
- Use 10-fold cross validation to determine accuracy of the model
- Split ground-truth data randomly into 10 partitions
  - Train on 9, test on 1
  - Repeat and average results



# Classifying Automation: Training and Testing

		Classified			
		Human	Cyborg	Bot	True Positive
Actual	Human	1972	27	1	98.6%
	Cyborg	65	1833	102	91.7%
	Bot	2	46	1952	97.6%
<b>Total</b>					96.0%



# Classifying Automation: Training and Testing

		Classified			
		Human	Cyborg	Bot	True Positive
Actual	Human	1972	27	1	98.6%
	Cyborg	65	1833	102	91.7%
	Bot	2	46	1952	97.6%
Total					96.0%



# Classifying Automation: Training and Testing

		Classified			
		Human	Cyborg	Bot	True Positive
Actual	Human	1972	27	1	98.6%
	Cyborg	65	1833	102	91.7%
	Bot	2	46	1952	97.6%
Total					96.0%



# Classifying Automation: Feature Importance

Features	Accuracy
Timing Entropy	82.8%
URL Ratio	74.9%
Automated Device %	71.0%
Bayesian Spam	69.5%
Manual Device %	69.2%
Registration Date	62.9%
Mention Ratio	56.2%
Link Safety	49.3%
Hashtag Ratio	47.0%
Followers to Friends	45.3%
Account Verification	35.0%



# Classifying Automation: Feature Importance

Features      Accuracy

Takeaway: No one  
feature tells the whole

Manual Device %	69.2%
Registration Late	52.9%
Mention Ratio	56.2%
Link Safety	49.3%
Hashtag Ratio	47.0%
Followers to Friends	45.3%
Account Verification	35.0%



# Twitter Composition

- Apply classification algorithm to set of 500,000 users
- 53.2% Human
- 36.2% Cyborg
- 10.5% Bot
- 5:4:1 Ratio on Twitter in 2012



# My Takeaways

- Good first step towards automation detection (96% TP!), but shelf life of technique is limited
  - Behavioral properties are the easiest to change
  - In 2018, bots behave just like humans



# My Takeaways

- Good first step towards automation detection (96% TP!), but shelf life of technique is limited
  - Behavioral properties are the easiest to change
  - In 2018, bots behave just like humans
- We still don't understand how effective bots are at their tasks



# My Takeaways

- Good first step towards automation detection (96% TP!), but shelf life of technique is limited
  - Behavioral properties are the easiest to change
  - In 2018, bots behave just like humans
- We still don't understand how effective bots are at their tasks
- **Paper: Understanding the network infrastructure behind automation**
- **Paper: Attribution of bot effectiveness through network analysis**



# Classifying Users on Twitter: Ground Truth

	Human	Bot	Cyborg
Criteria	<ul style="list-style-type: none"><li>• Original, intelligent, specific content</li></ul>	<ul style="list-style-type: none"><li>• Unoriginal content</li><li>• <i>Clearly</i> automated (i.e., RSS feed)</li><li>• Spammy, filled with shady URLs</li></ul>	<ul style="list-style-type: none"><li>• Evidence of both bot and human participation</li></ul>



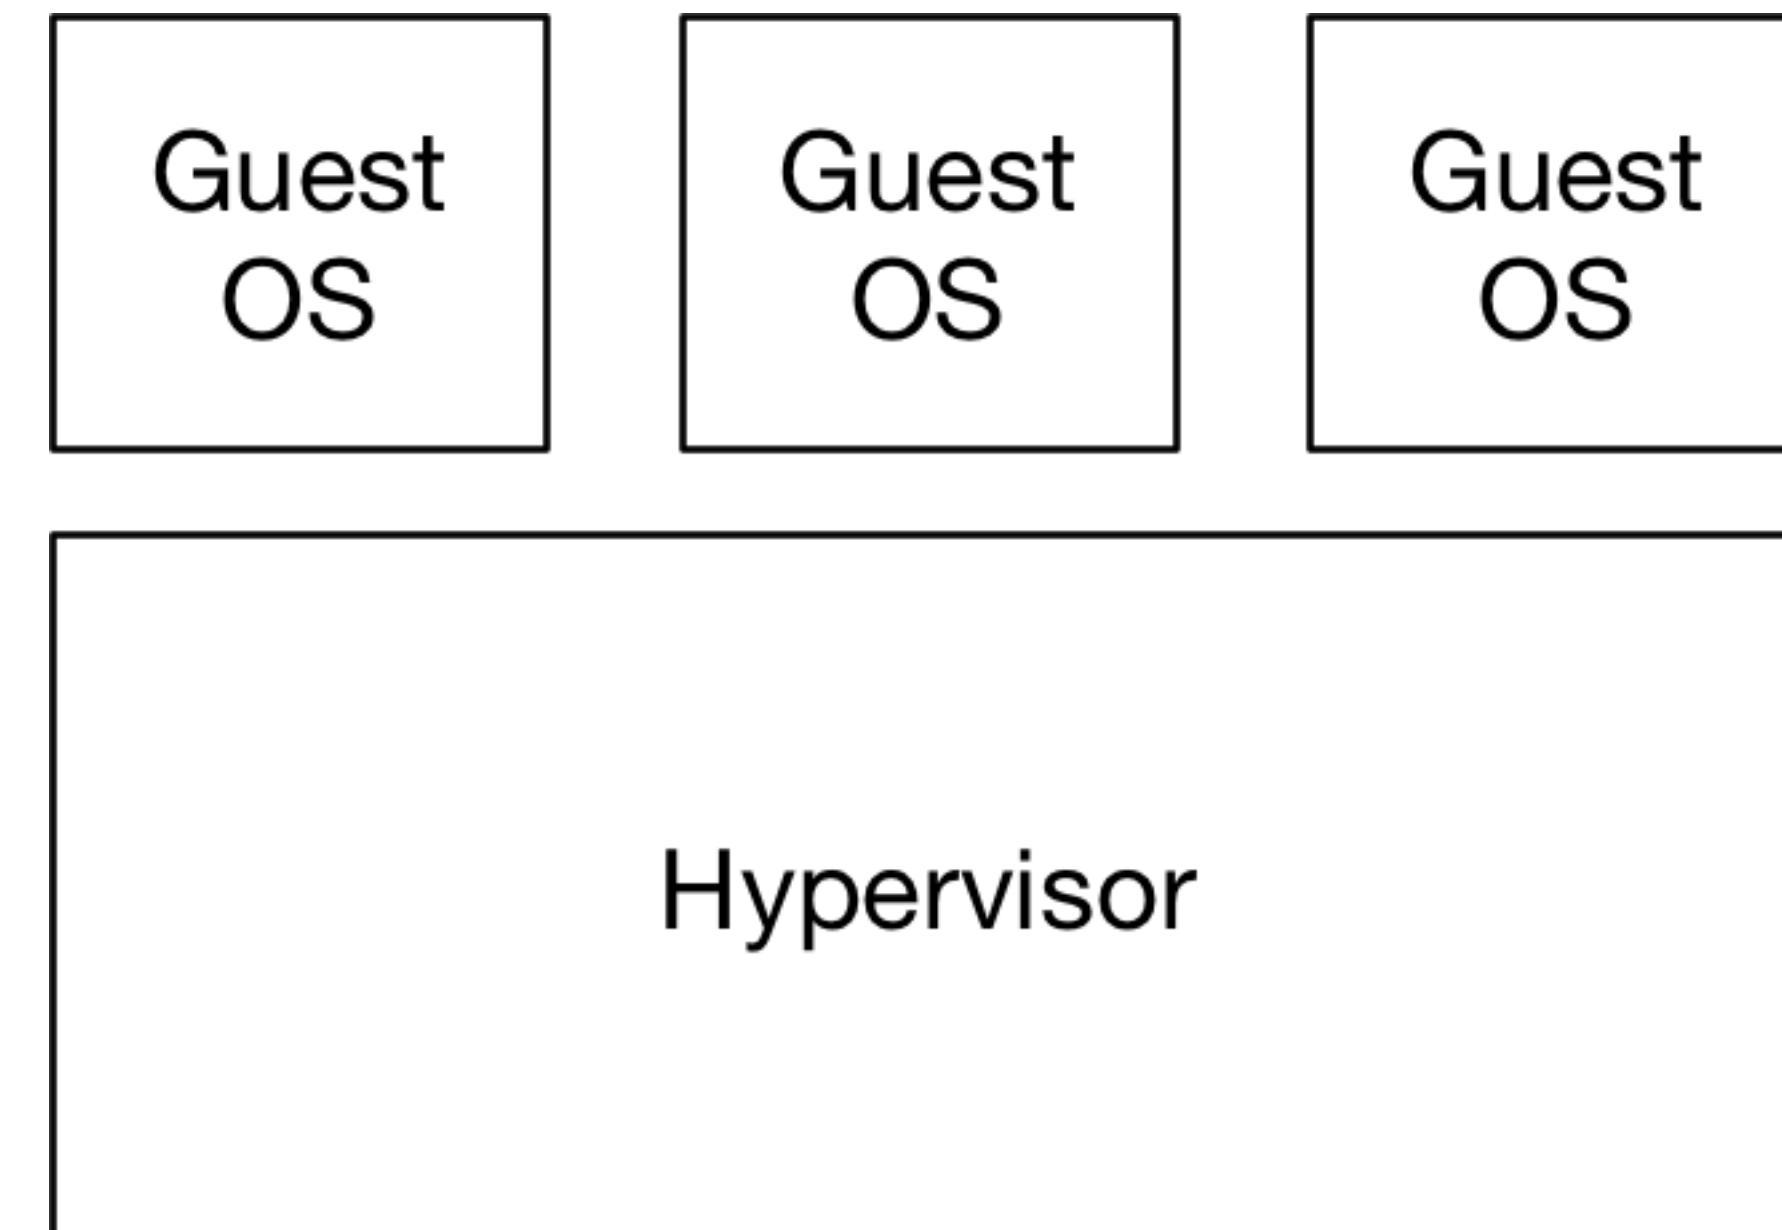
# The Turtles Project: Design and Implementation of Nested Virtualization

Muli Ben-Yehuda, Michael D. Day, Zvi Dubitzky, Michael Factor, Nadav Har'El, Abel Gordon, Anthony Liguori, Orit Wasserman, Ben-Ami Yassour



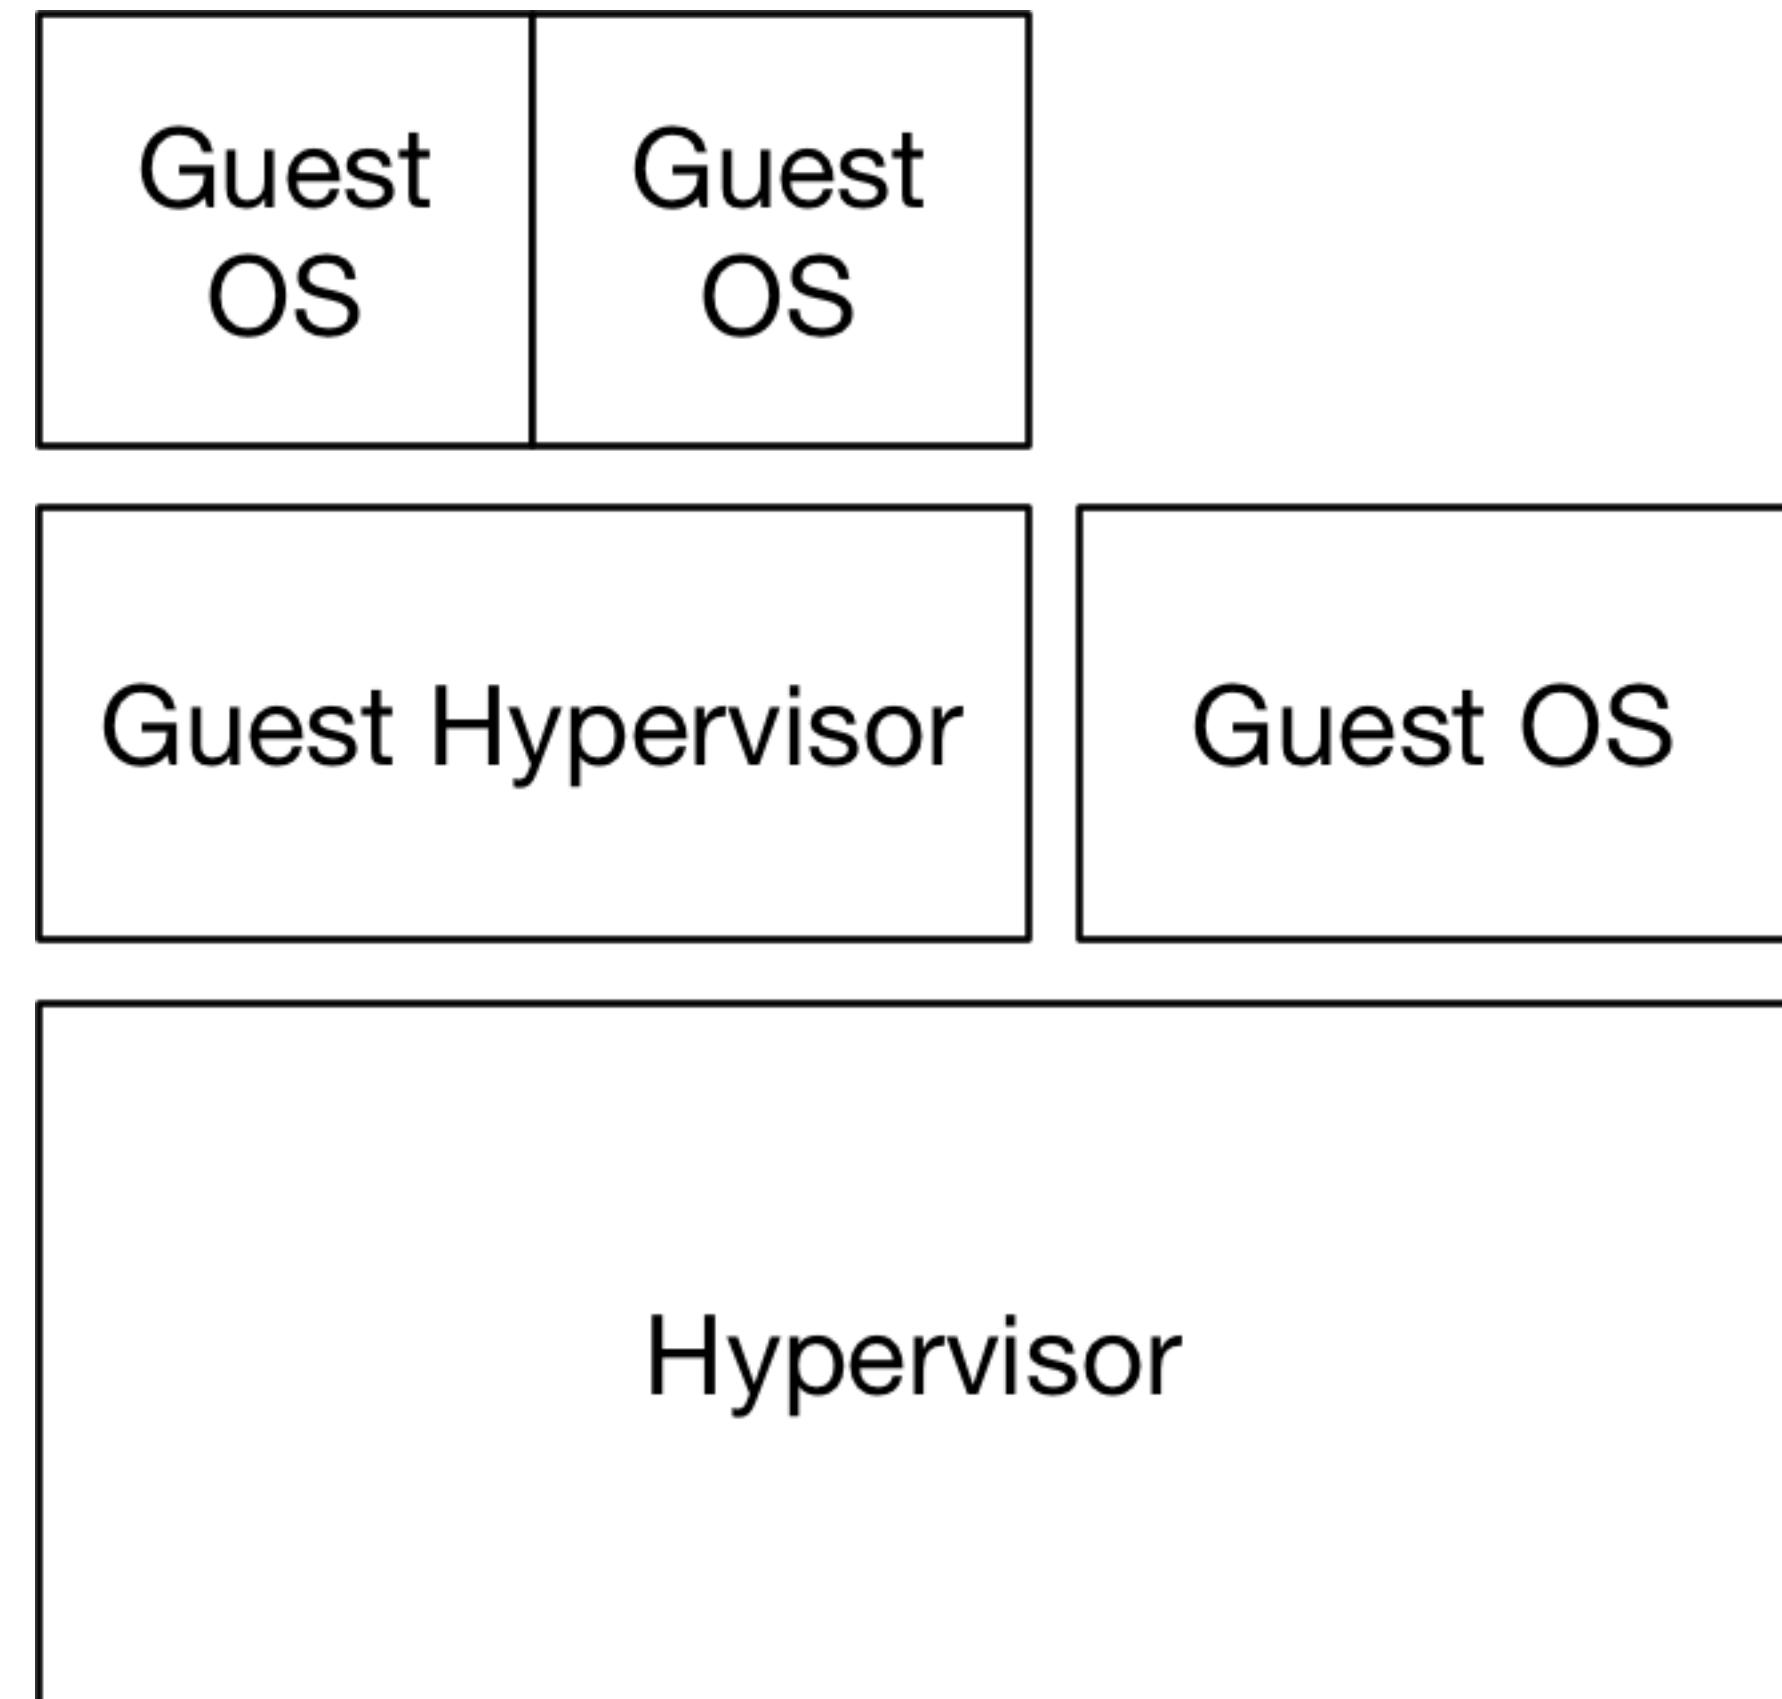
# Classic Virtualization

- One hypervisor
  - Runs multiple guest operating systems
  - Hardware support from Intel (Intel VMX), AMD (AMD SVM) for x86



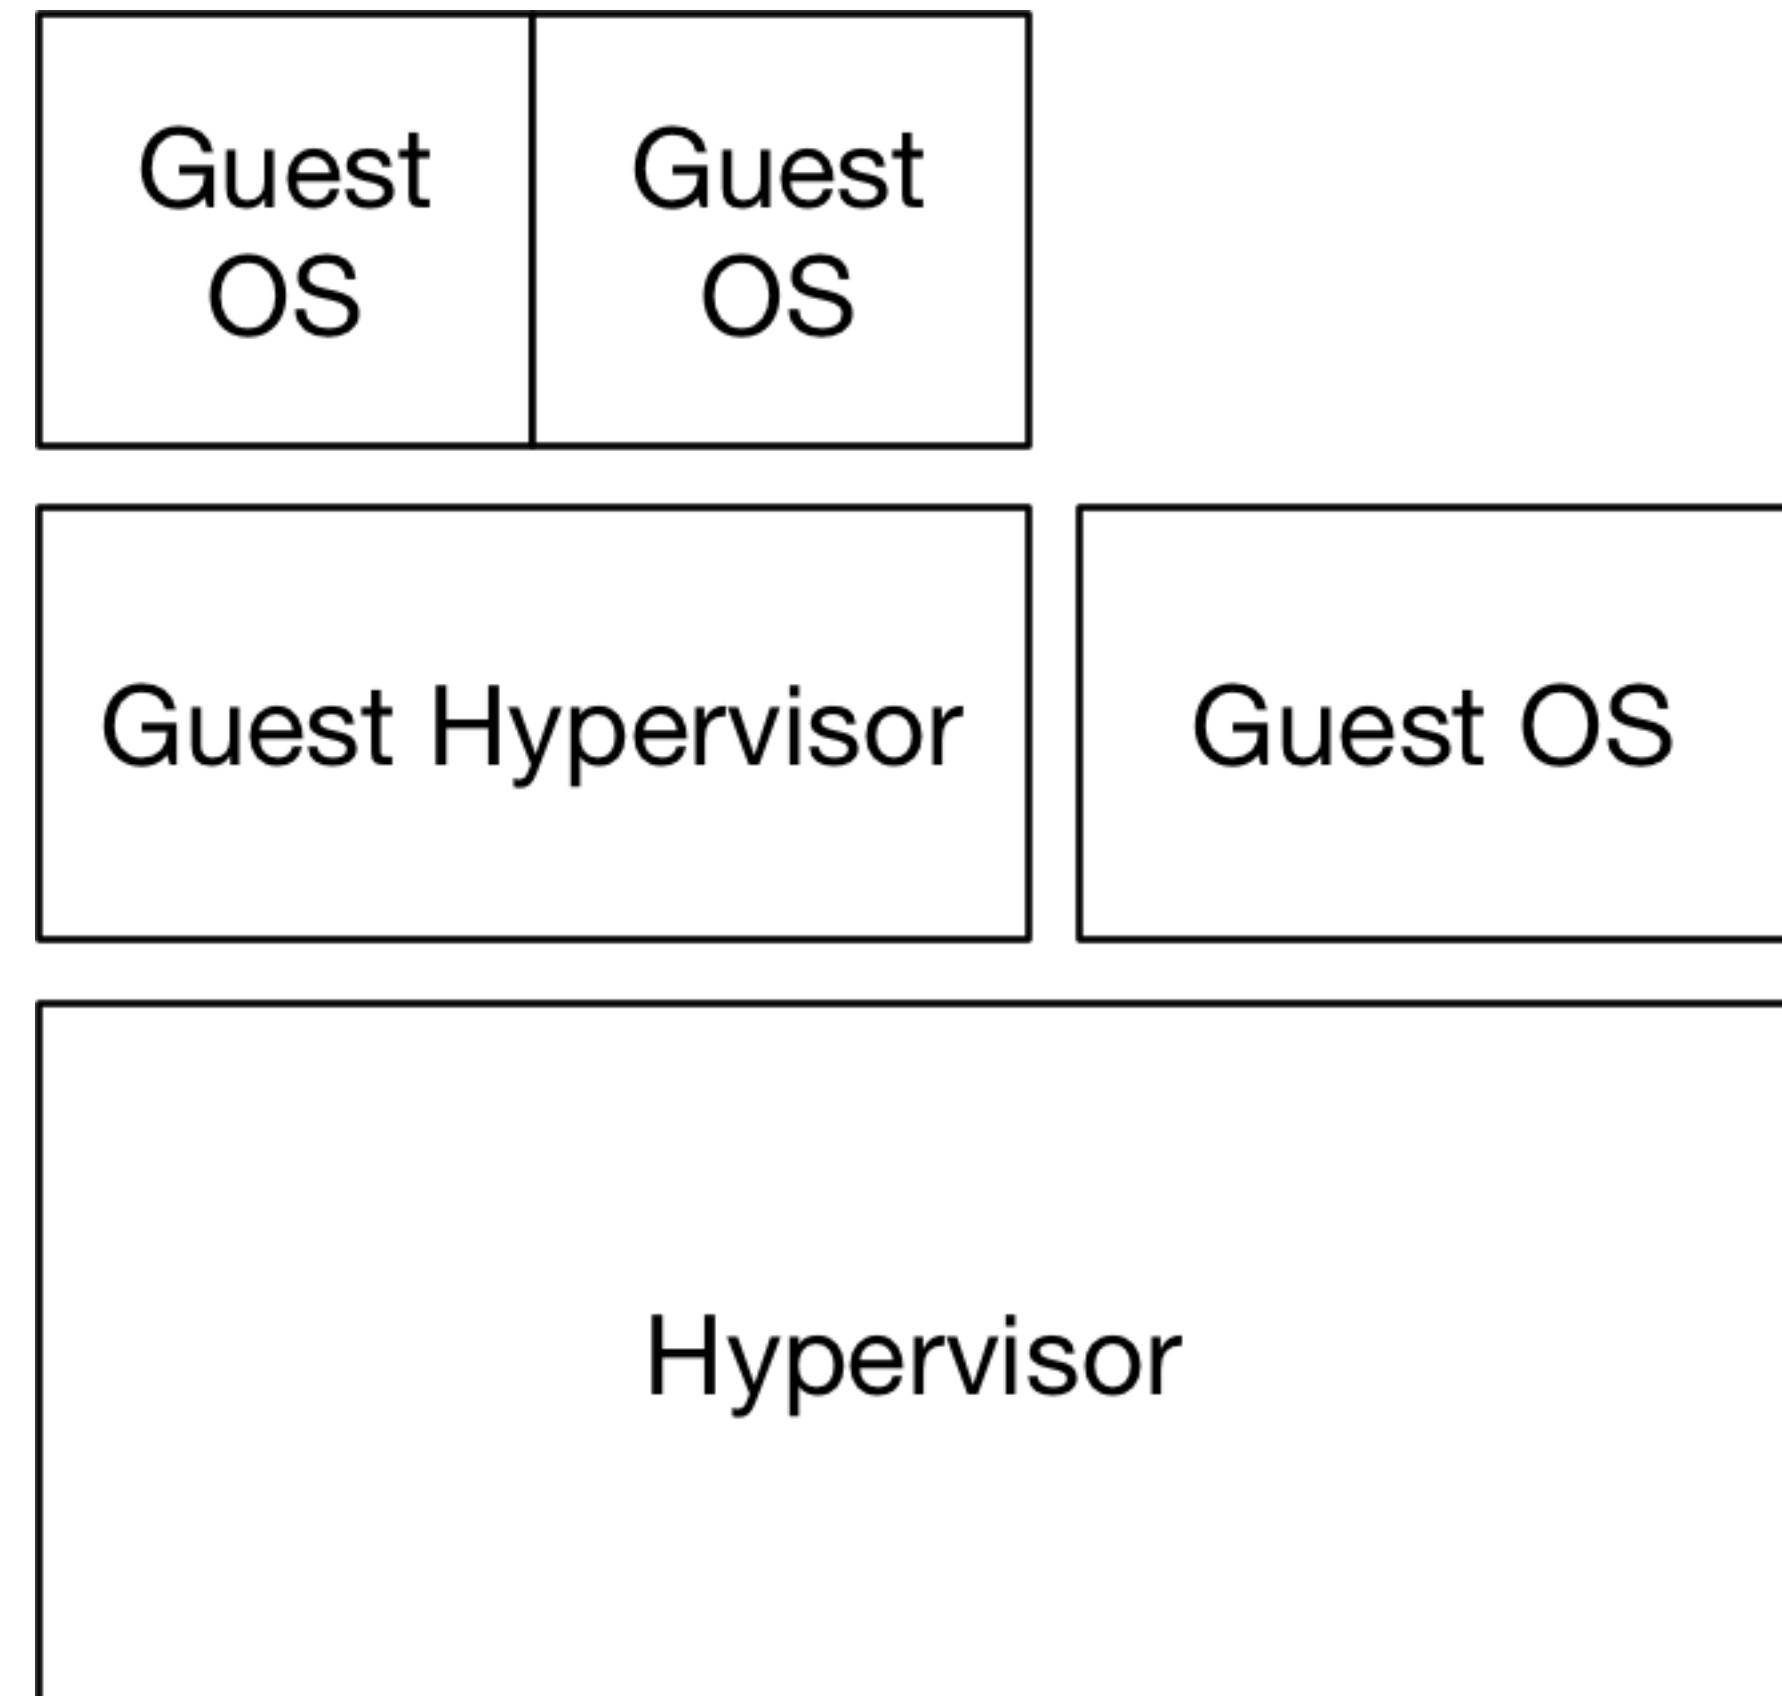
# Nested Virtualization

- Run multiple *unmodified* hypervisors
  - Unmodified guest operating systems



# Nested Virtualization

- Run multiple *unmodified* hypervisors
  - Unmodified guest operating systems
  - Nested virtualization increasing in importance
  - Many commodity OSes make use of virtualization already



# Why Nested Virtualization?

- Cloud
  - Expose user-controlled hypervisor as a virtual machine



# Why Nested Virtualization?

- Cloud
  - Expose user-controlled hypervisor as a virtual machine
- Security
  - Hypervisor-level rootkit protection, hypervisor-level intrusion detection



# Why Nested Virtualization?

- Cloud
  - Expose user-controlled hypervisor as a virtual machine

## **Problem: No x86 support for nested virtualization!**

- Security
  - Hypervisor-level rootkit protection, hypervisor-level intrusion detection



# Why Nested Virtualization?

- Cloud
  - Expose user machine
  - Scale infrastructure on host
- Security
  - Hypervisor-level detection



# The Turtles Project

- *Single-level architectural support for nested virtualization*



# The Turtles Project

- *Single-level architectural support for nested virtualization*
  - Single hypervisor mode



# The Turtles Project

- *Single-level architectural support for nested virtualization*
  - Single hypervisor mode
  - Any trap at a nested level is handled by the “base” hypervisor, called  $L_0$



# The Turtles Project

- *Single-level architectural support for nested virtualization*
  - Single hypervisor mode
  - Any trap at a nested level is handled by the “base” hypervisor, called  $L_0$
  - Traps are then *forwarded* by  $L_0$  to the appropriate level

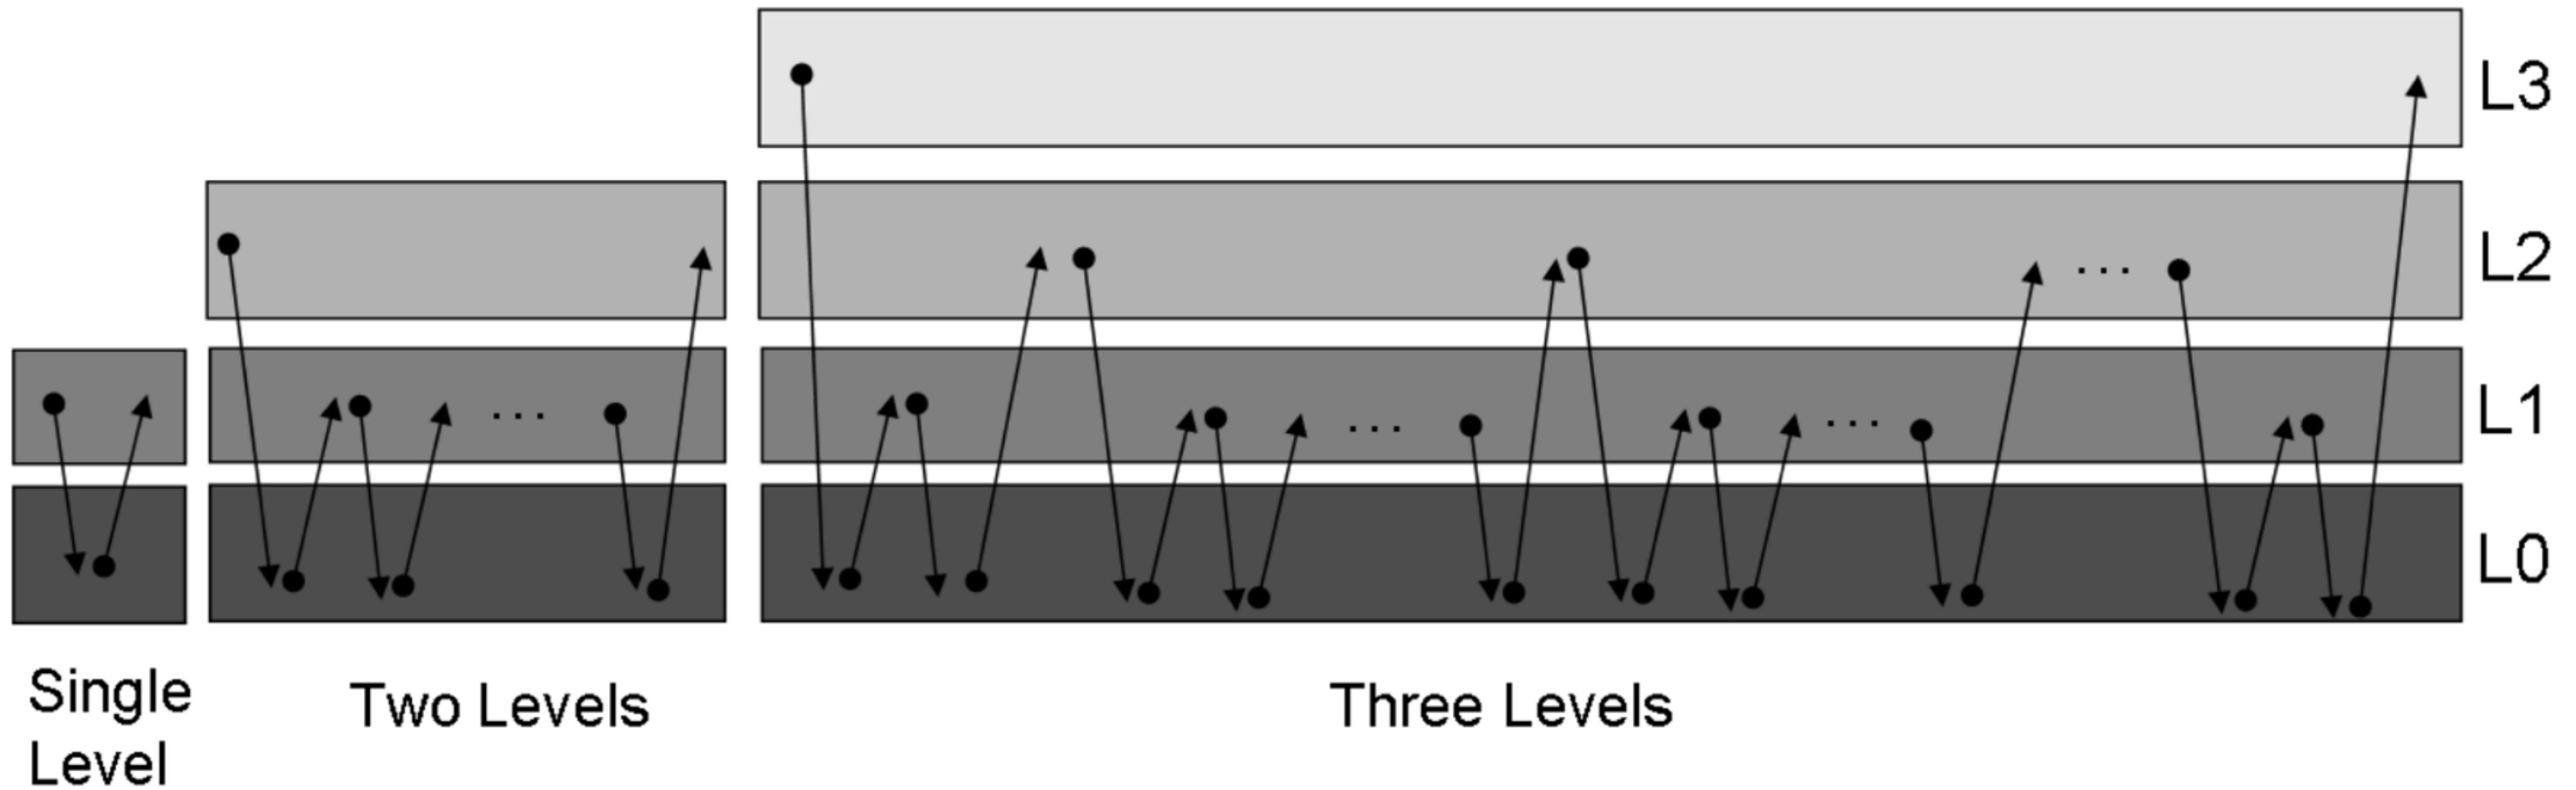


# The Turtles Project

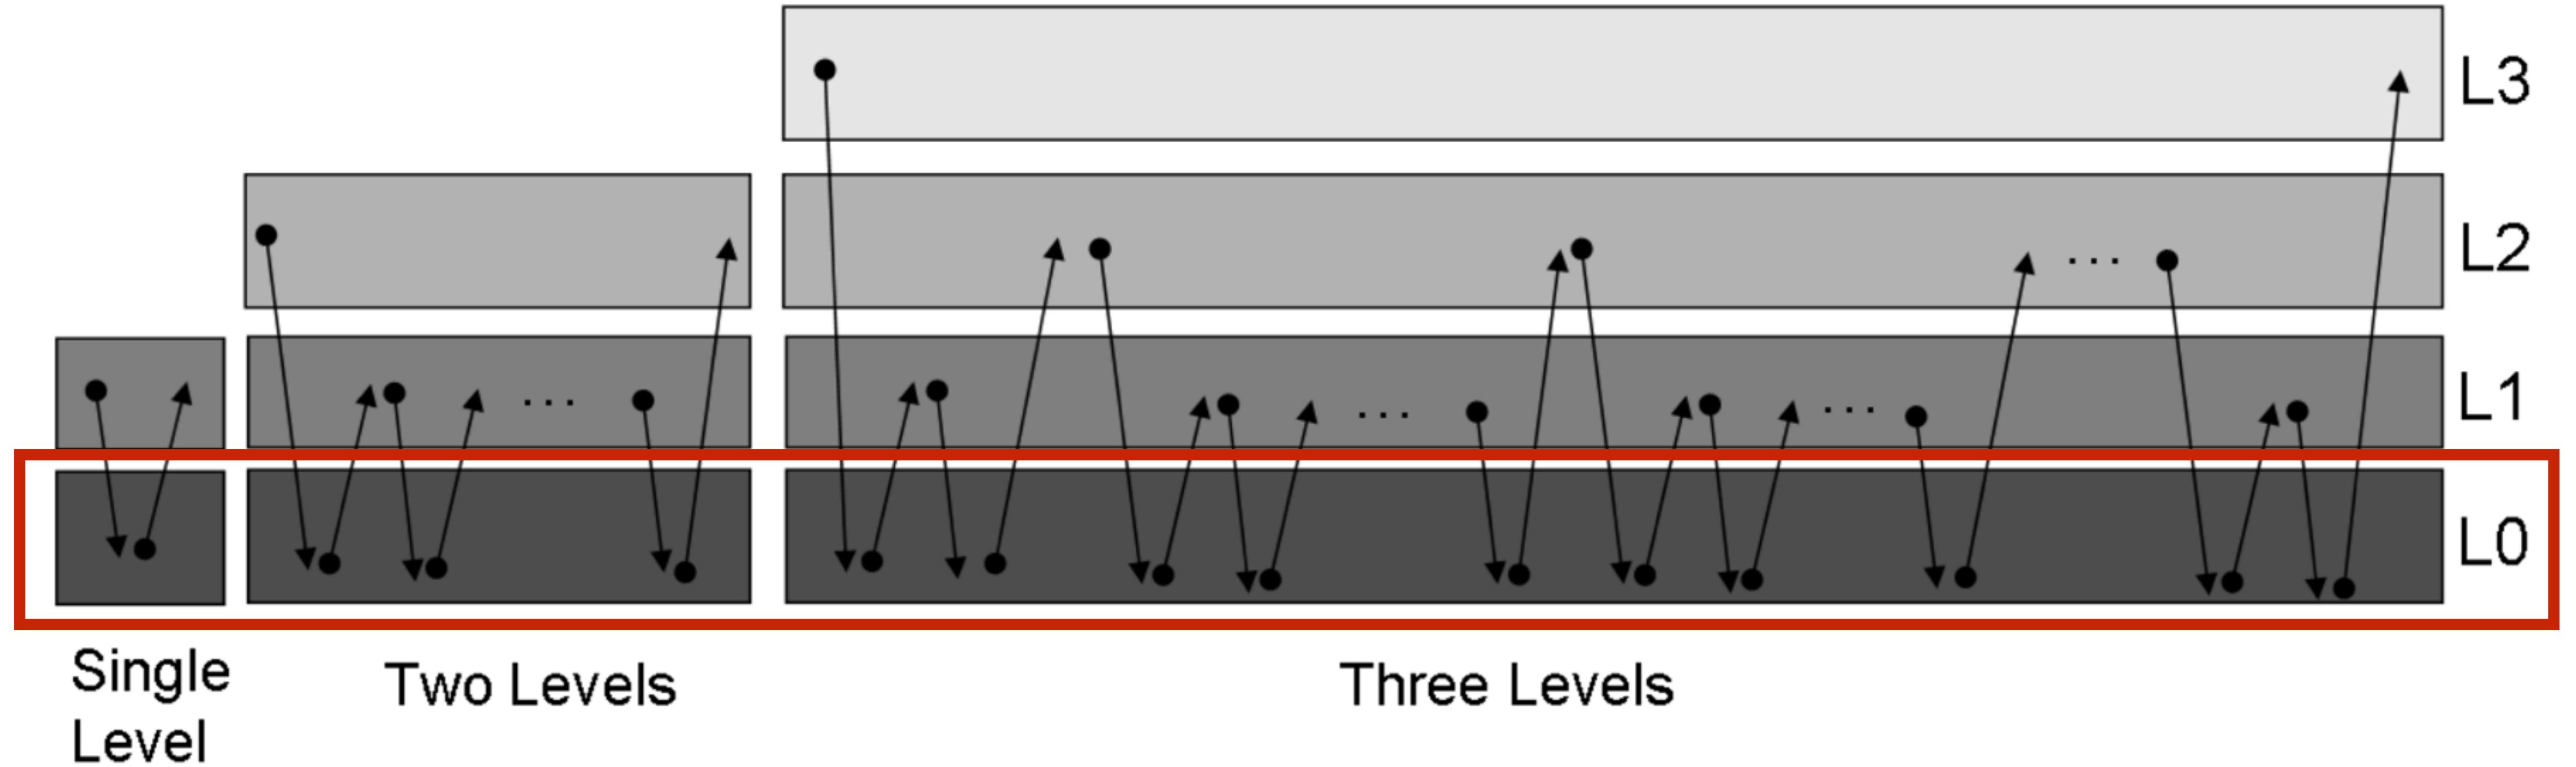
- *Single-level architectural support for nested virtualization*
  - Single hypervisor mode
  - Any trap at a nested level is handled by the “base” hypervisor, called  $L_0$
  - Traps are then *forwarded* by  $L_0$  to the appropriate level
  - $L_0$  thus *multiplexes* hardware between higher level hypervisors and OSes



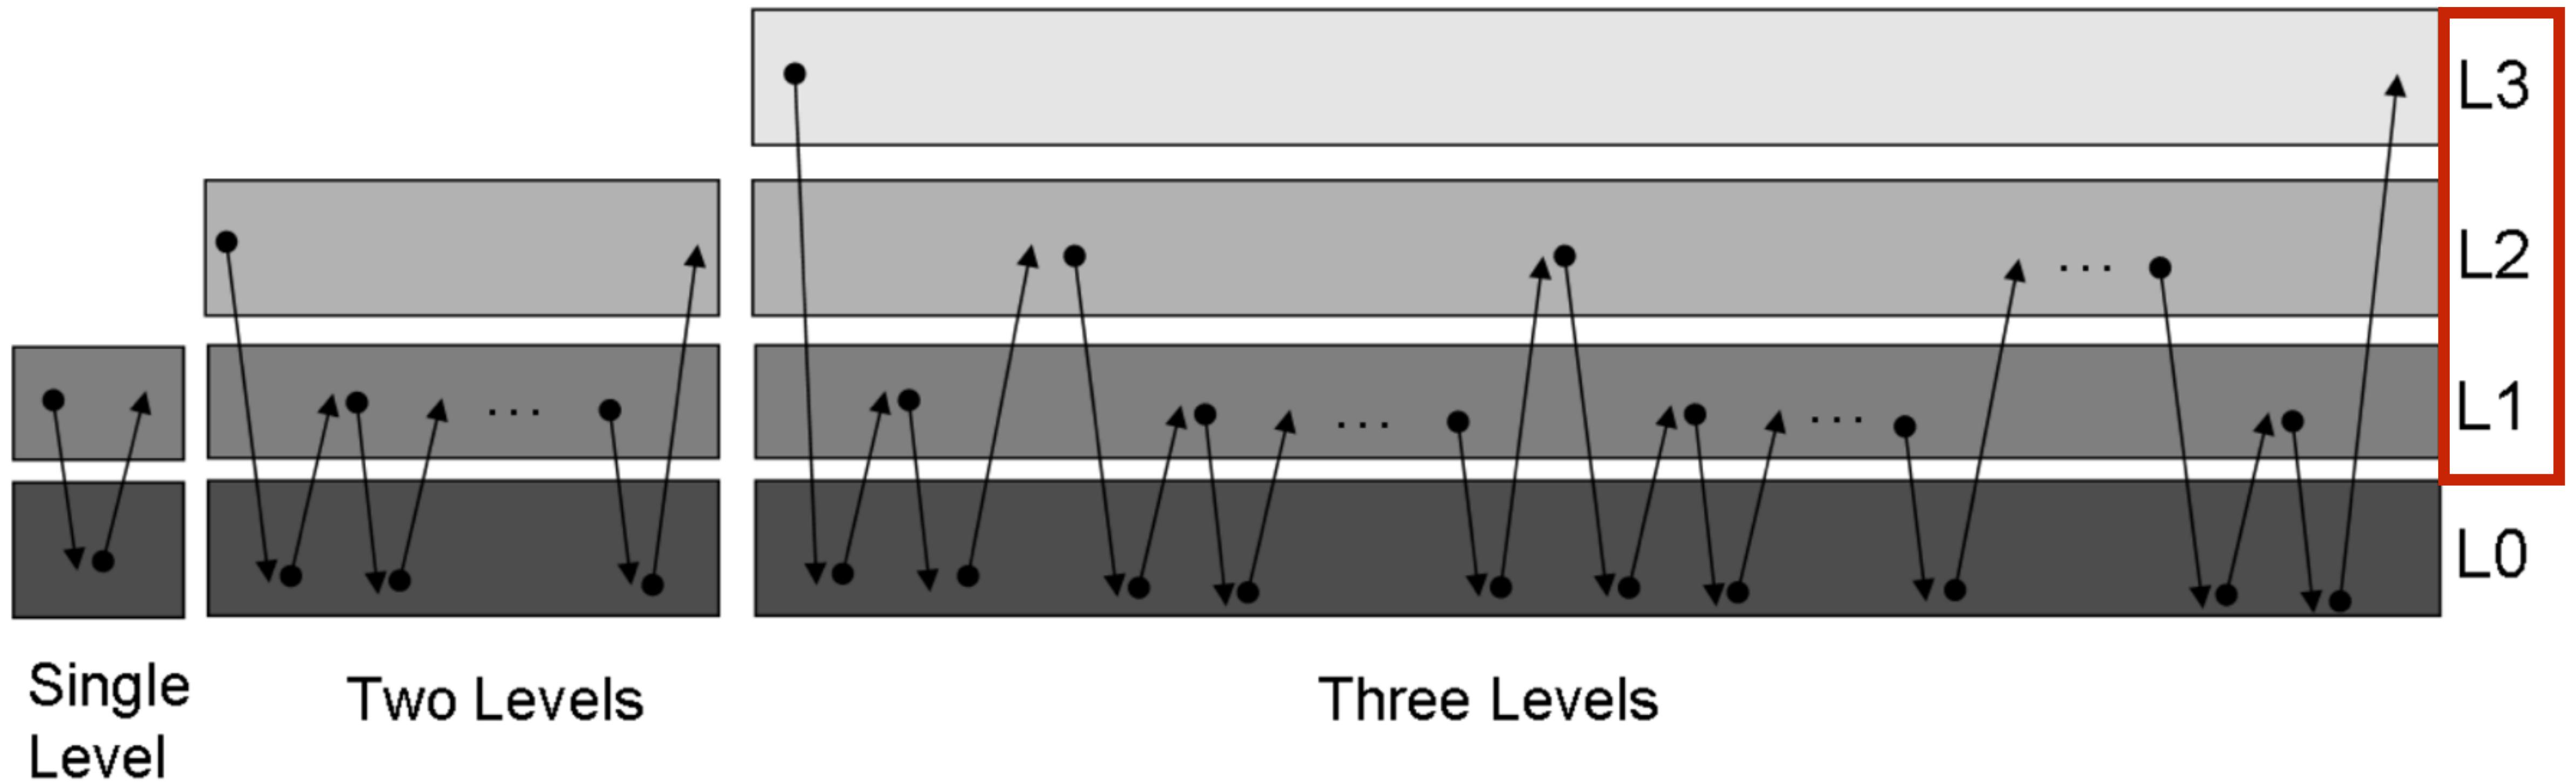
# Turtles Theory



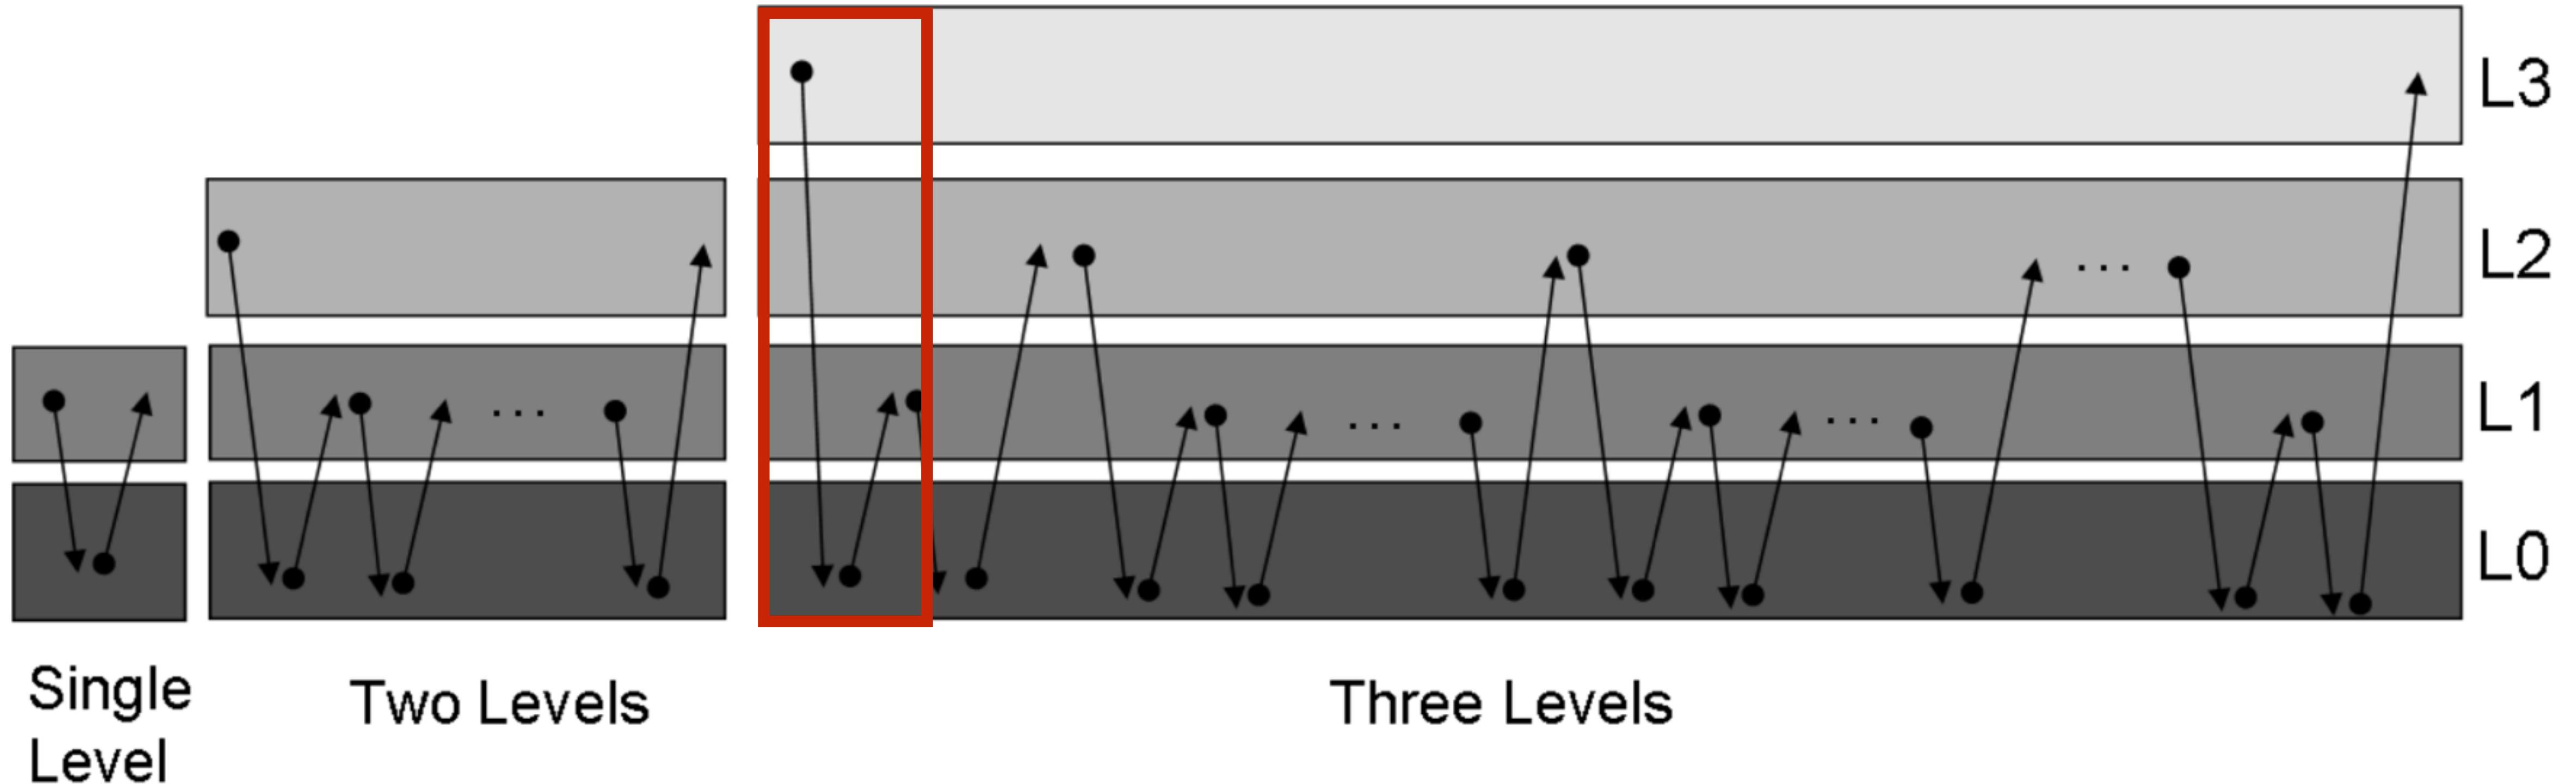
# Turtles Theory



# Turtles Theory



# Turtles Theory



# Implementing Turtles

- Nested VMX Virtualization for CPUs
- Multidimensional Paging for MMUs
- Multi-level Device Assignment for I/O



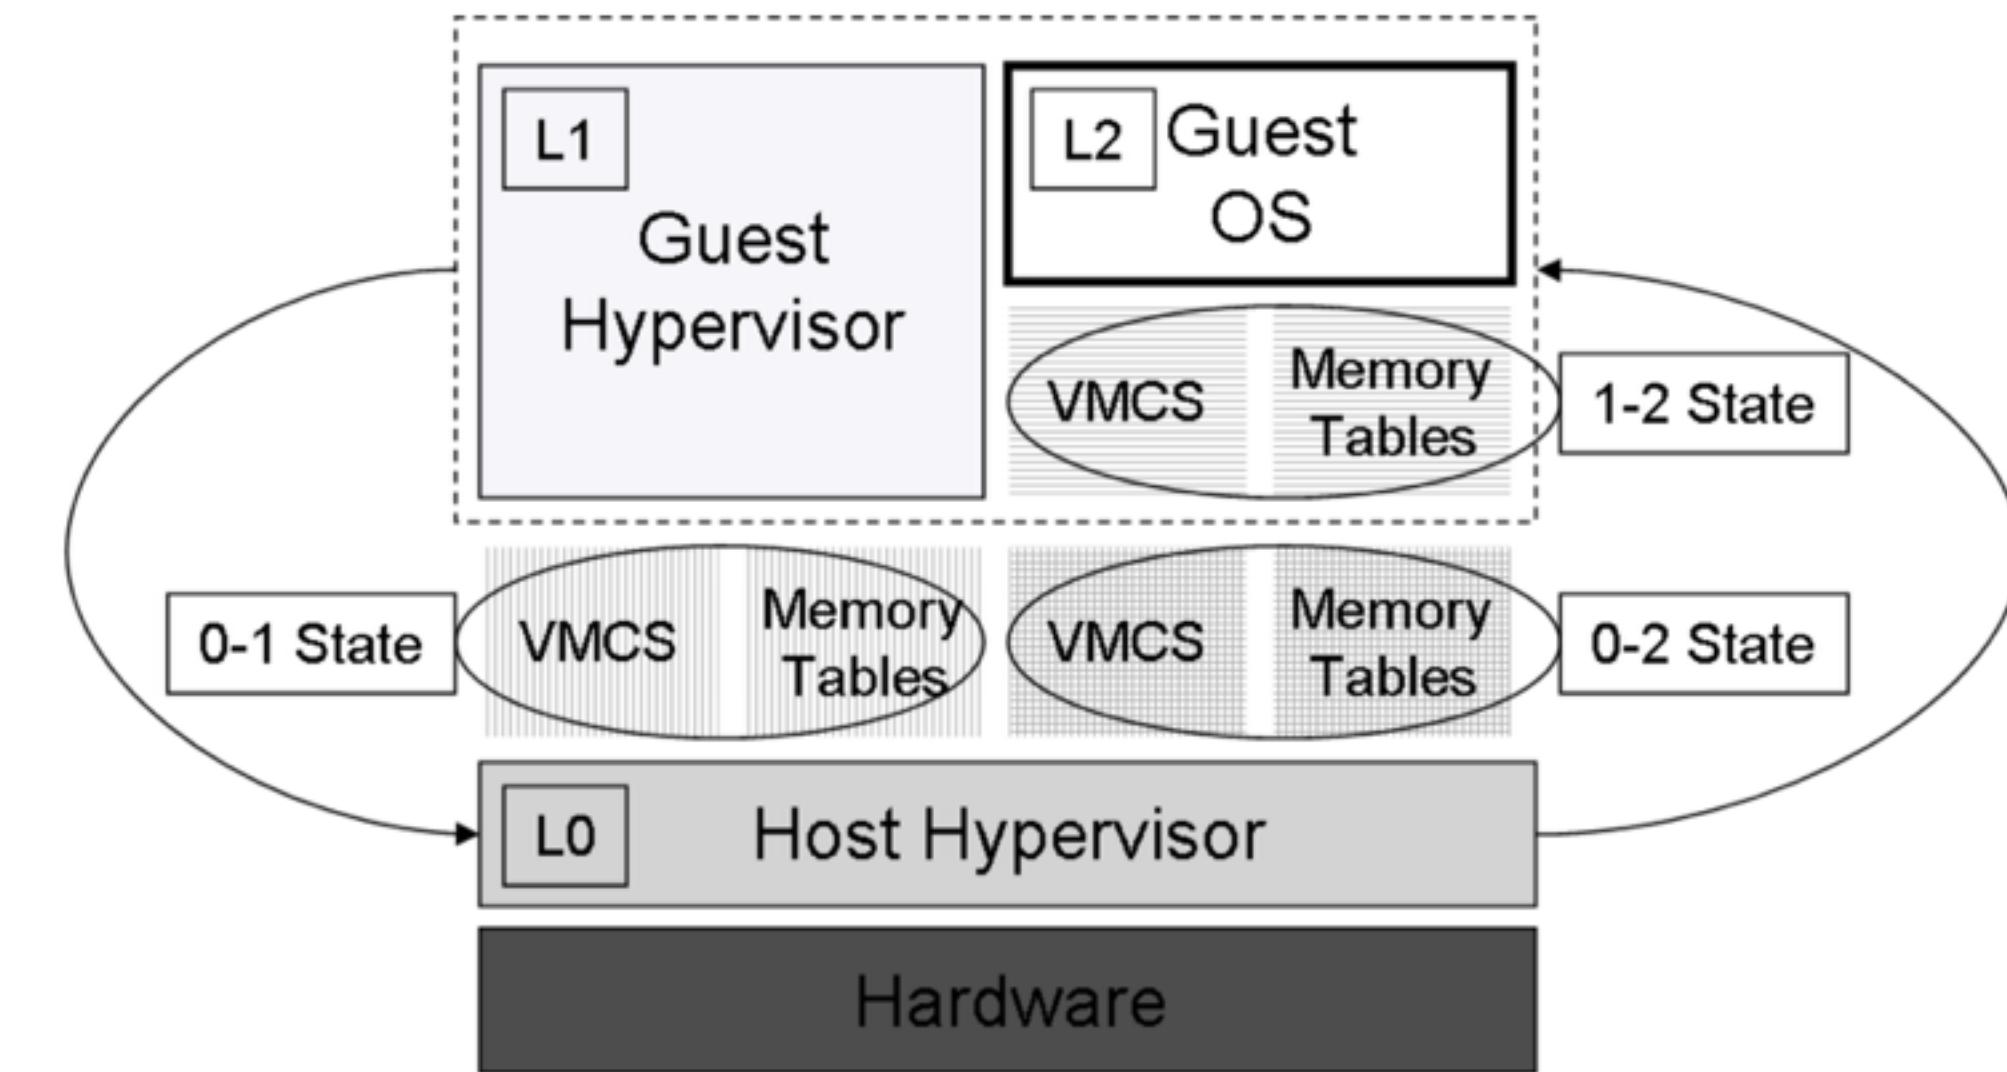
# Implementing Turtles

- **Nested VMX Virtualization for CPUs**
- Multidimensional Paging for MMUs
- Multi-level Device Assignment for I/O



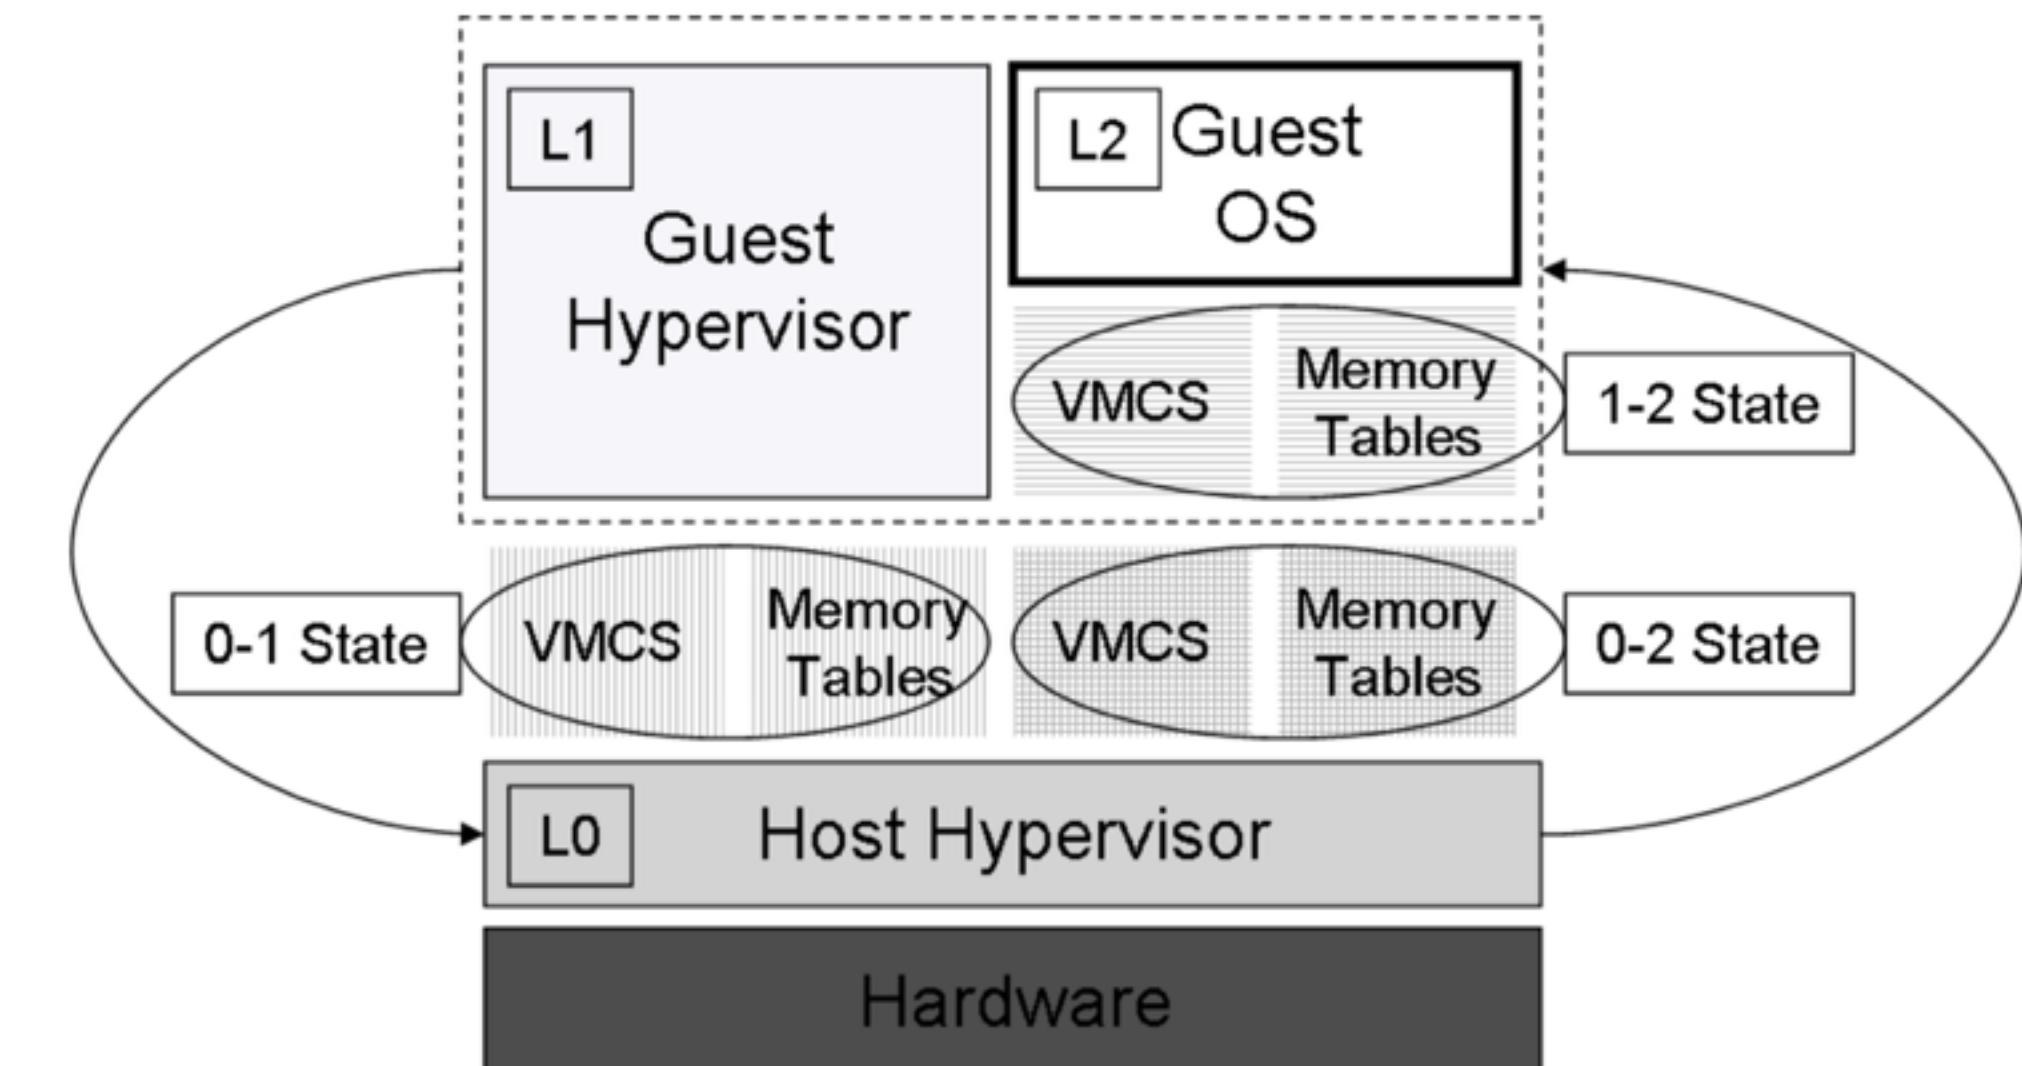
# Nested CPU Virtualization: Design

- L<sub>0</sub> can leverage virtualization support from hardware (VMX)
  - Supports *guest mode* and *host mode* for CPUs



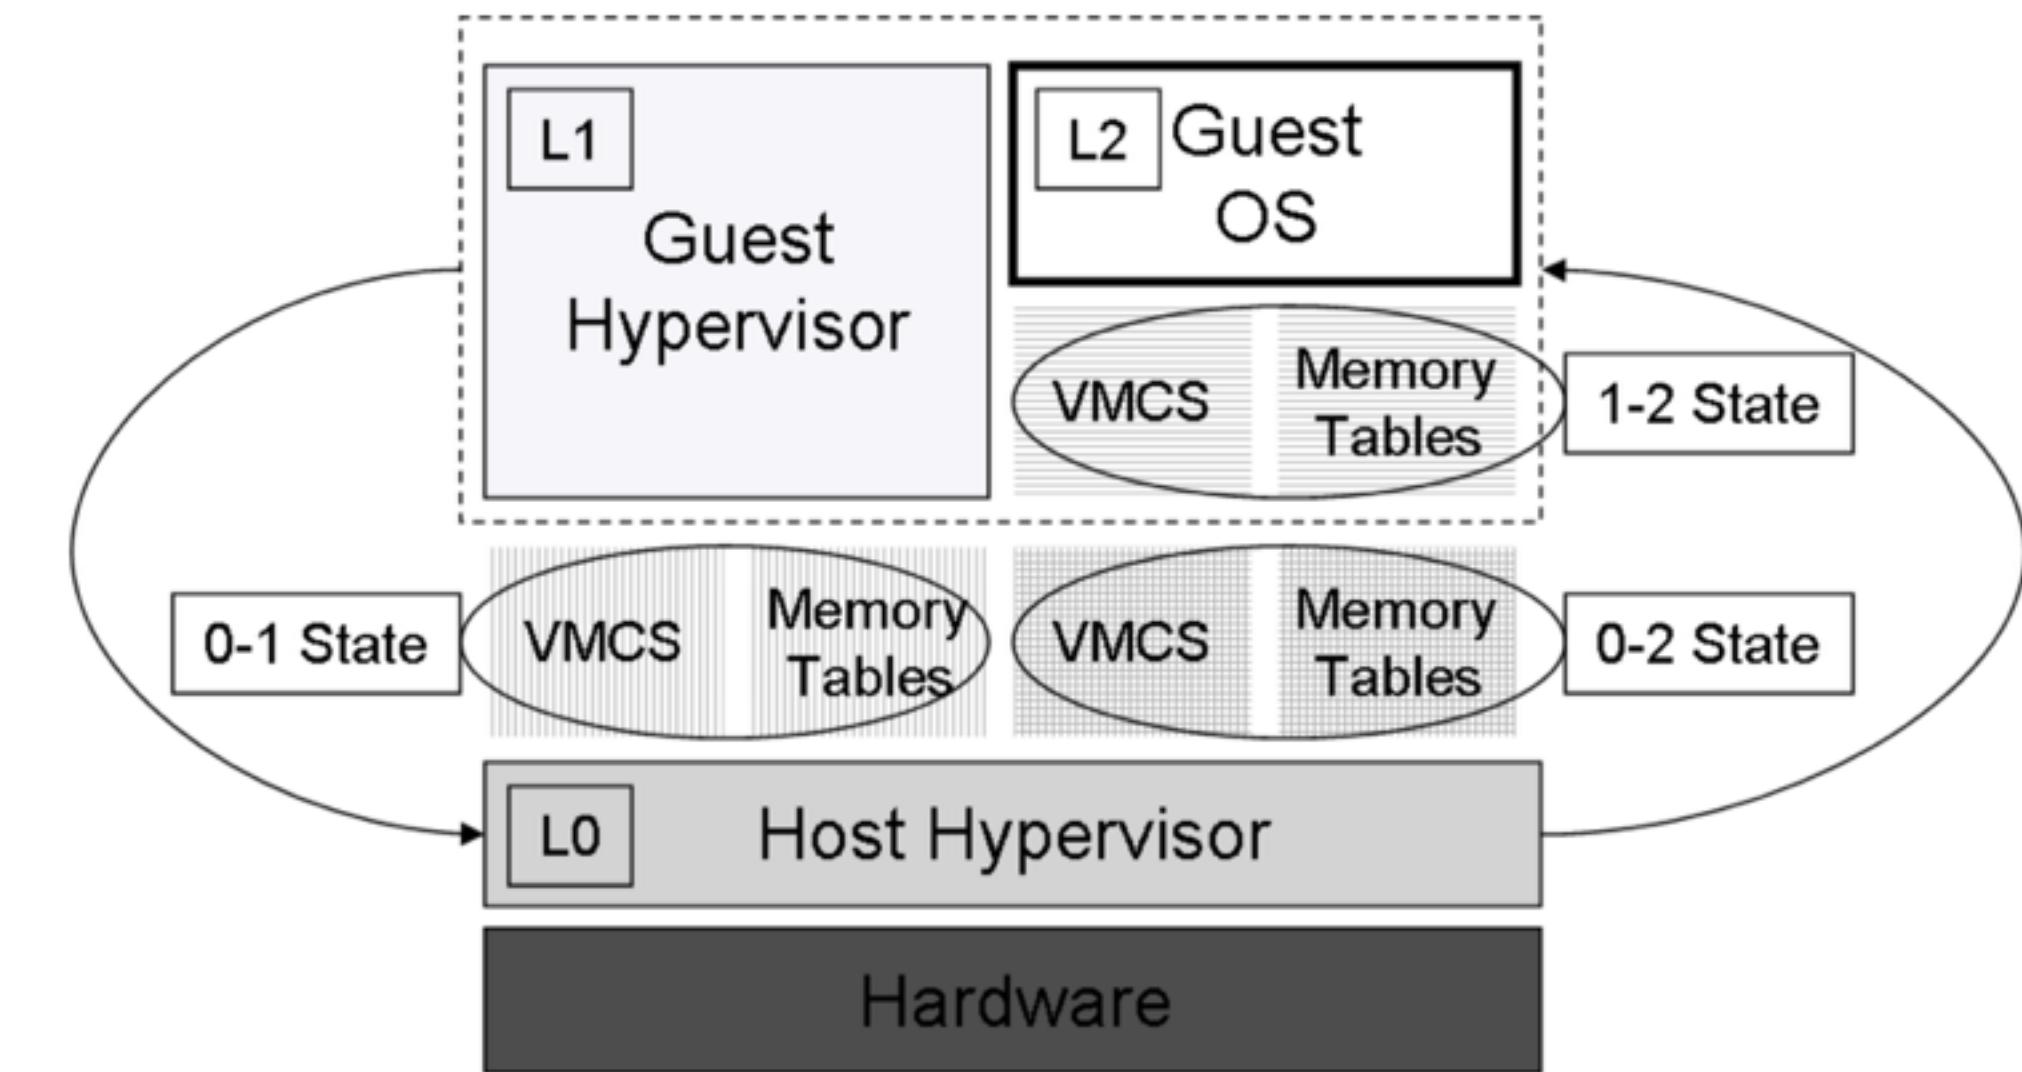
# Nested CPU Virtualization: Design

- L<sub>0</sub> can leverage virtualization support from hardware (VMX)
  - Supports *guest mode* and *host mode* for CPUs
  - VMCS: Control structures for defining virtual machines environments
    - Defines traps, CPU registers, host state, guest state



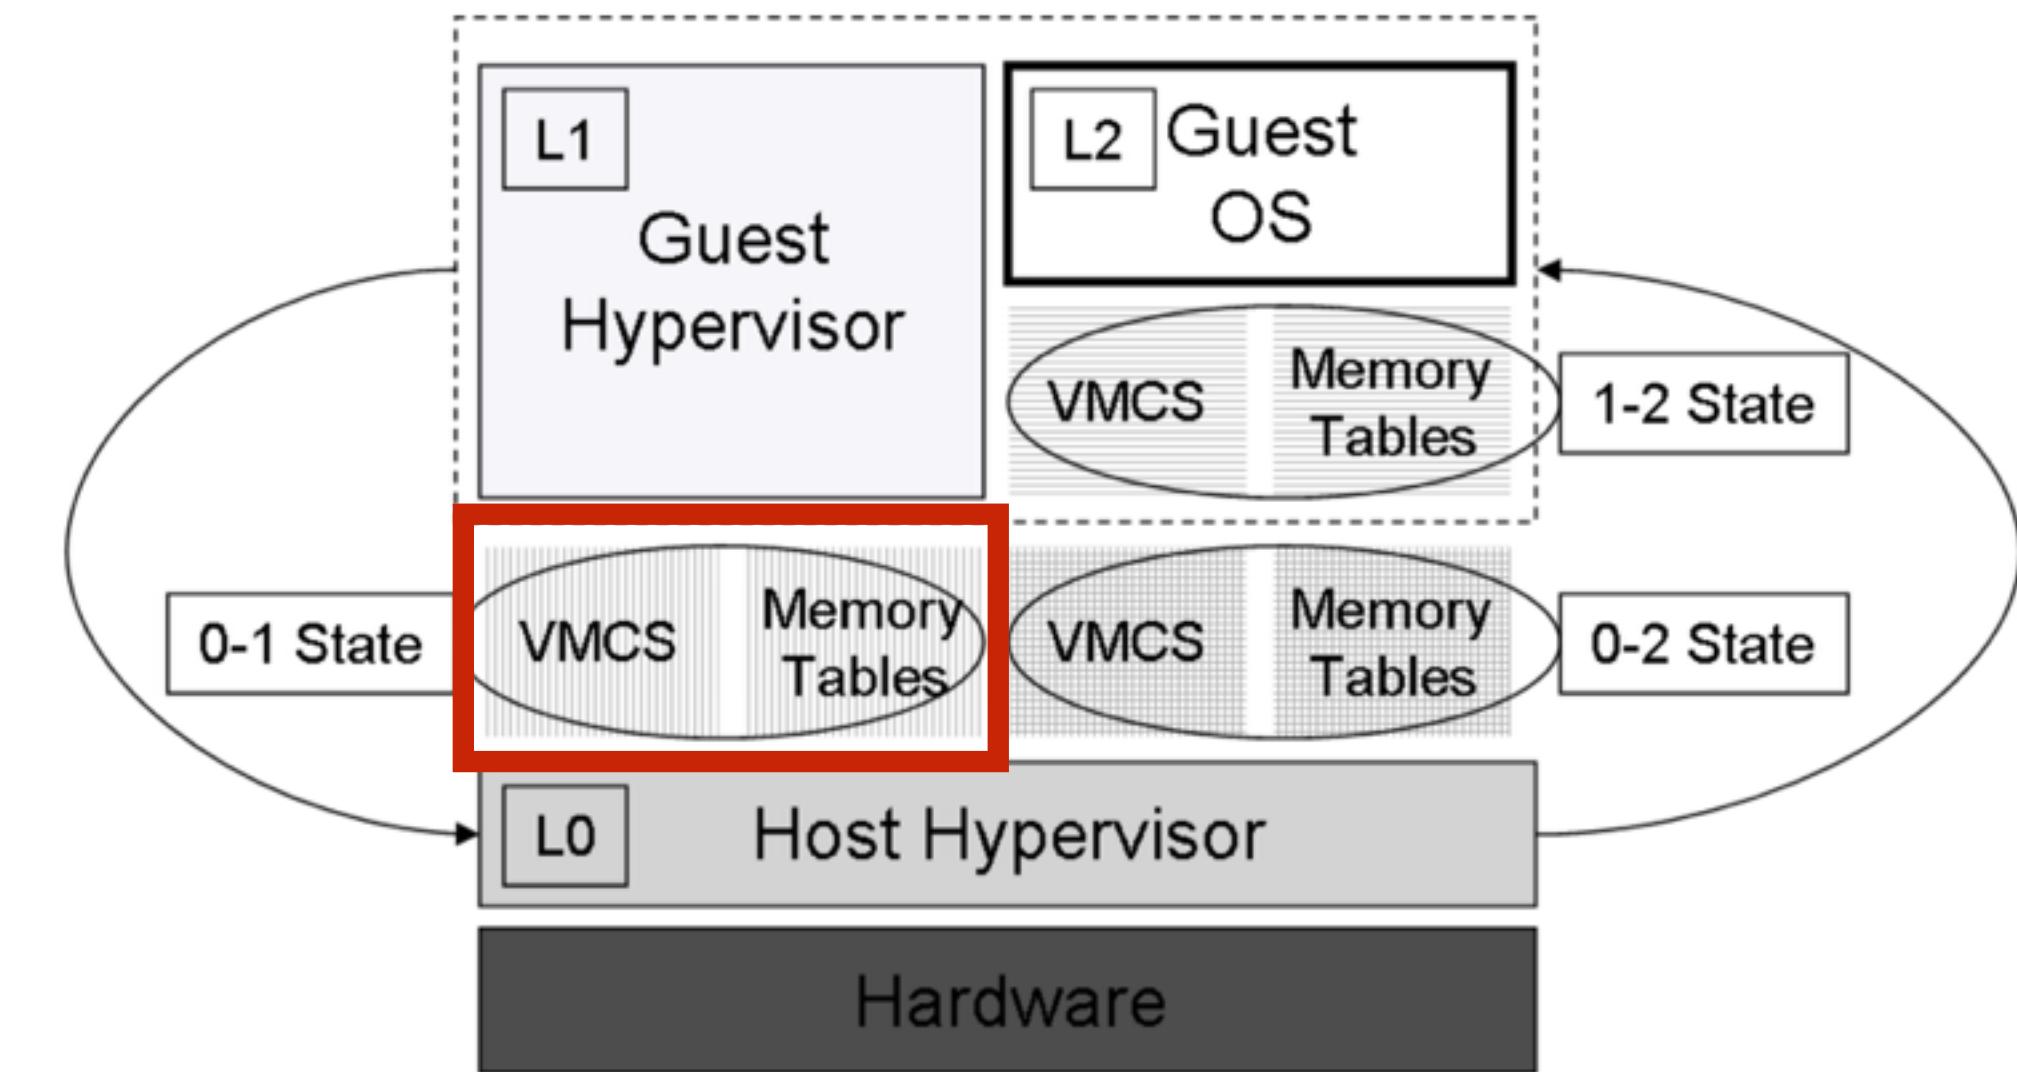
# Nested CPU Virtualization: Design

- Recursive VMX emulation



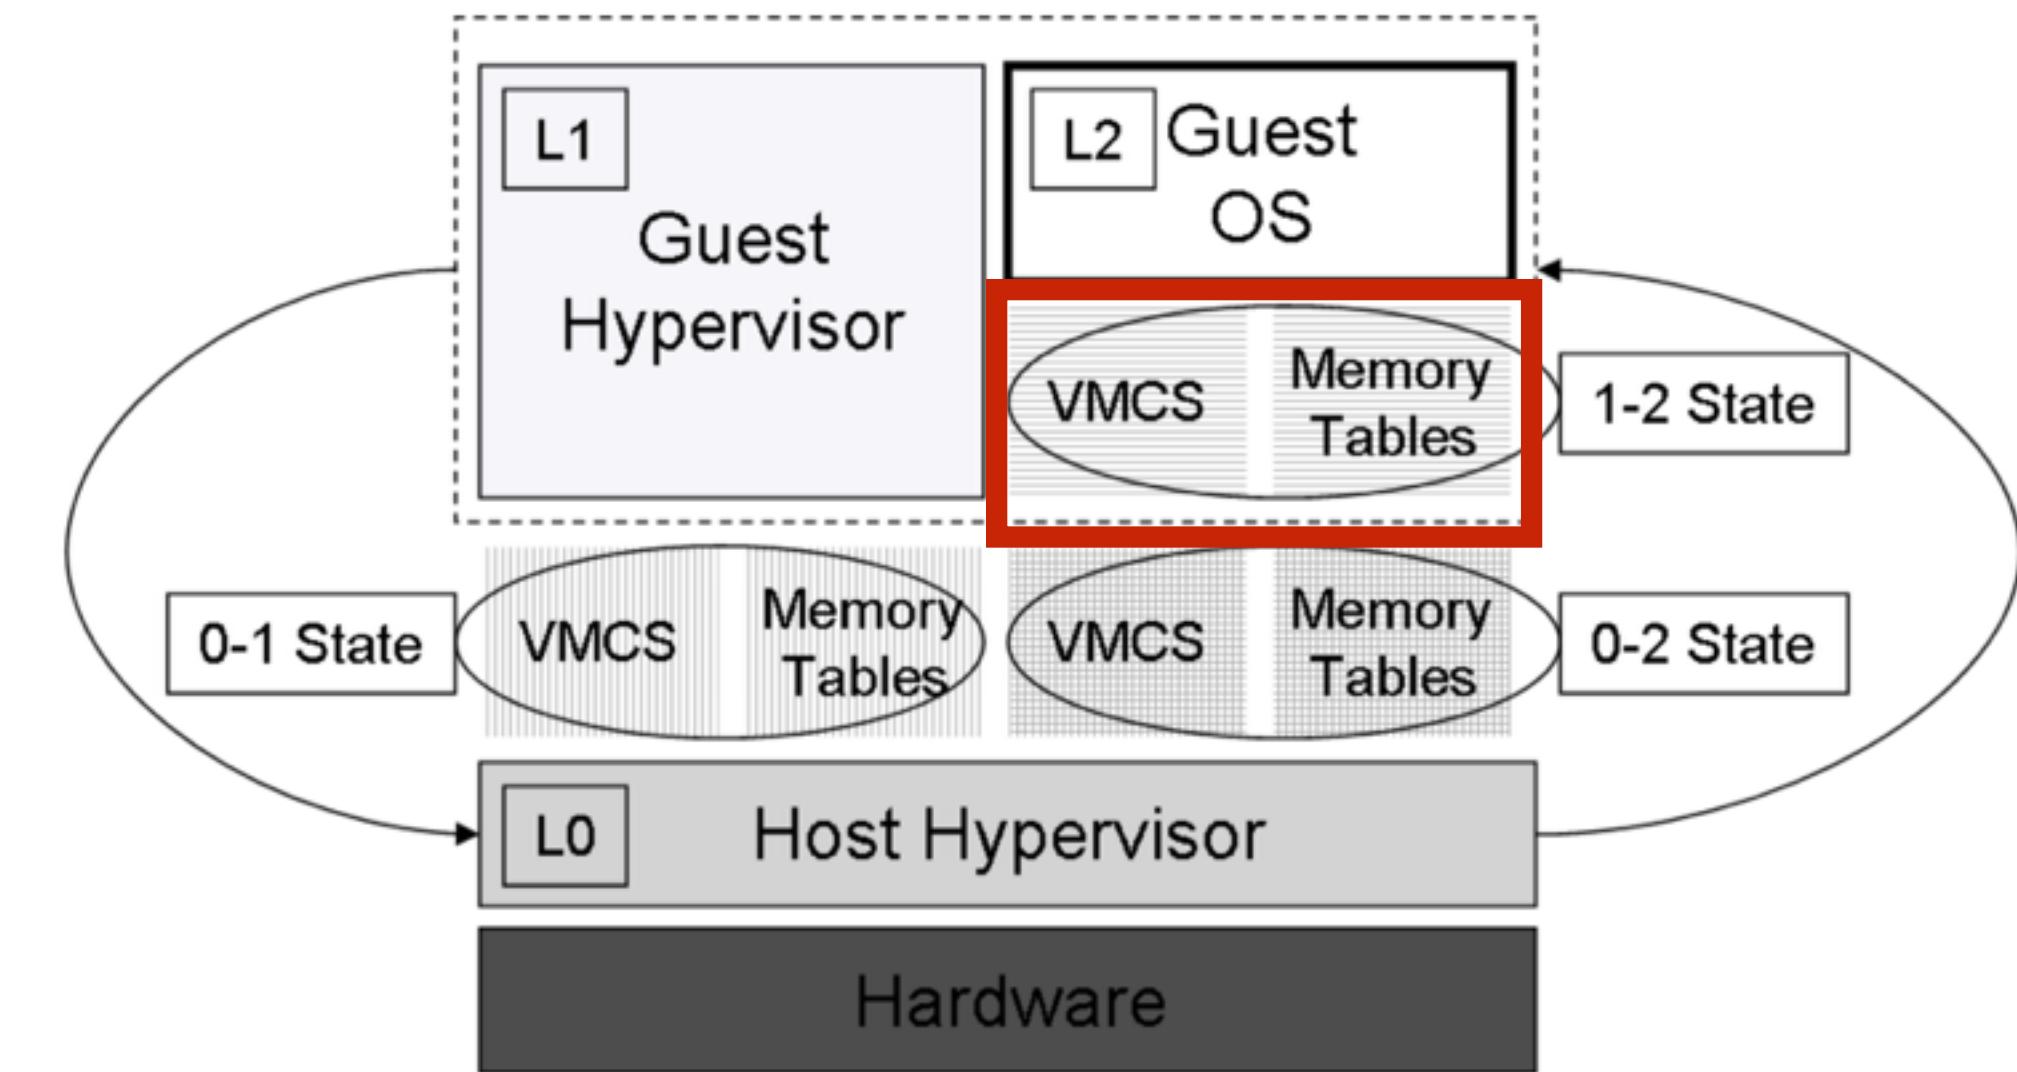
# Nested CPU Virtualization: Design

- Recursive VMX emulation
  - $L_0$  needs to *emulate* VMX to  $L_1$



# Nested CPU Virtualization: Design

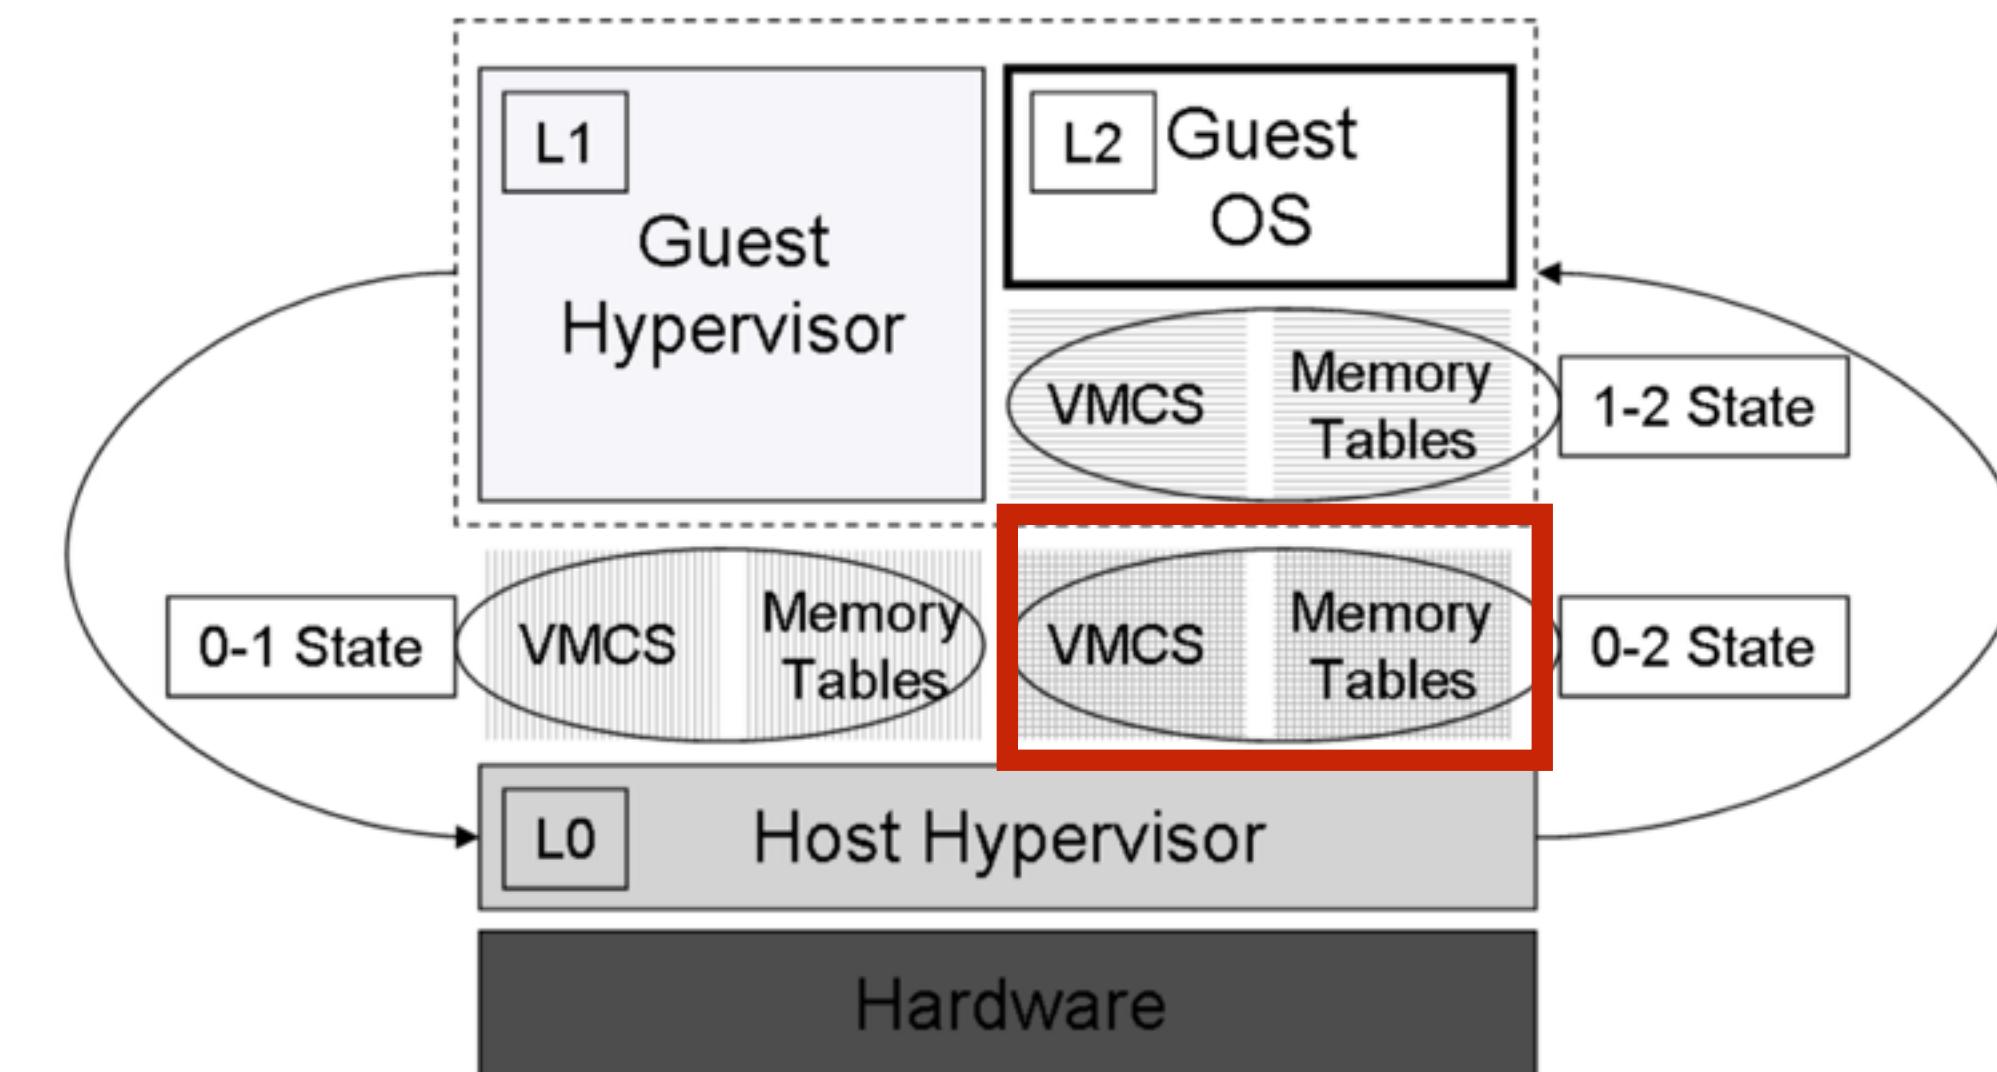
- Recursive VMX emulation
  - $L_0$  needs to *emulate* VMX to  $L_1$
  - $L_n$  will do the same for  $L_{n+1}$



# Nested CPU Virtualization: Design

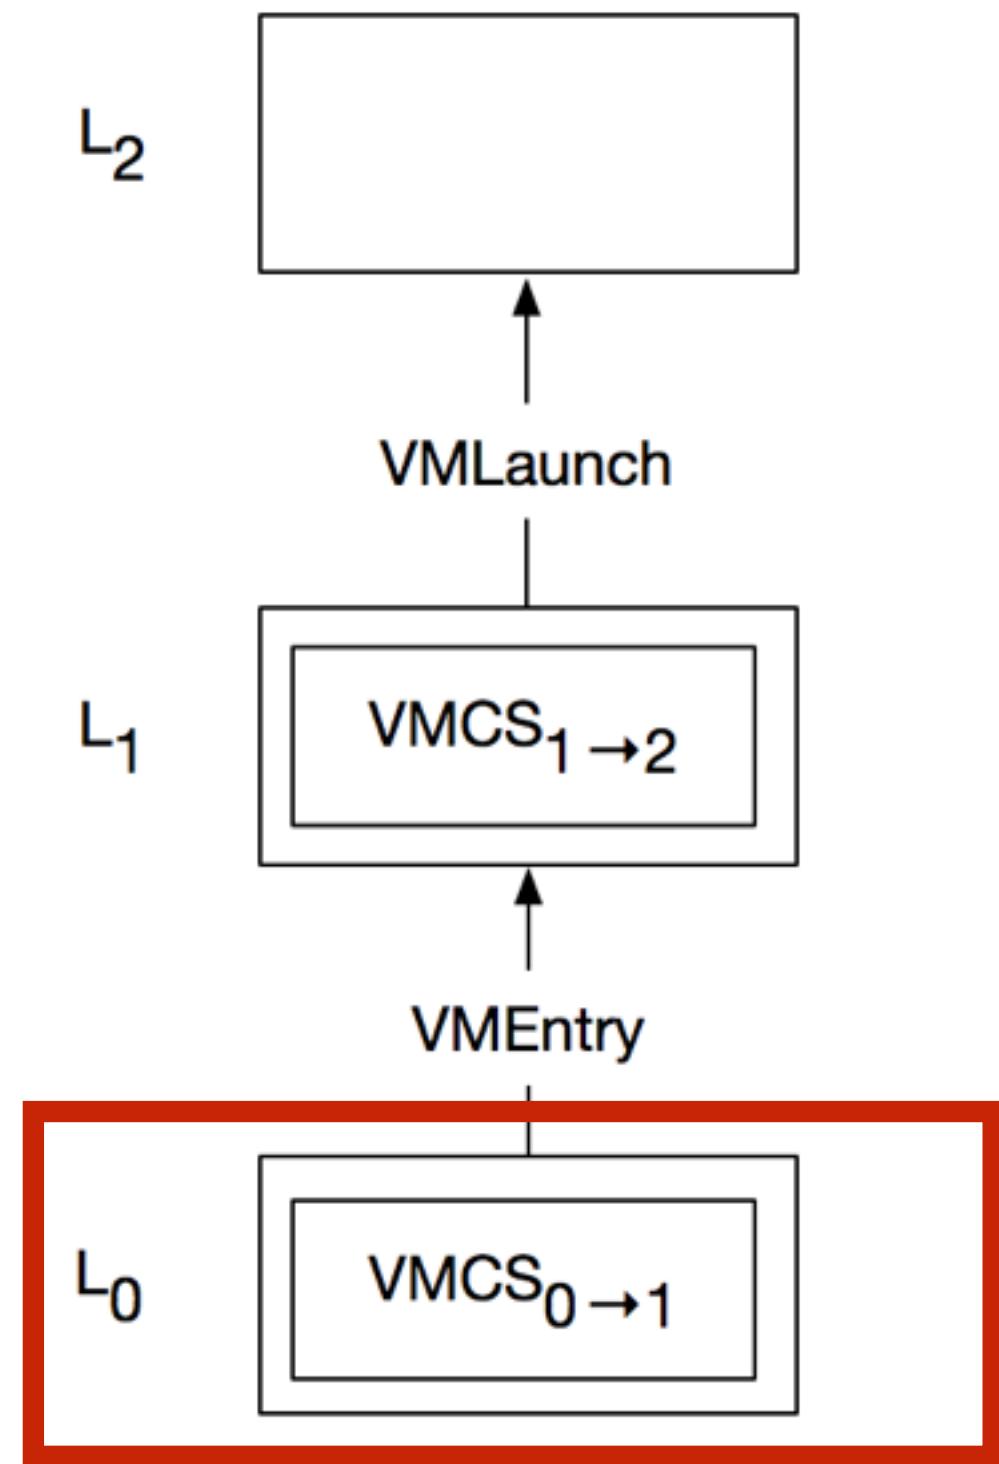
**Key point: L<sub>0</sub> must maintain state (VMCS) about L<sub>2</sub> to actually run L<sub>2</sub> on the hardware**

- run needs to map L<sub>1</sub> VMX to L<sub>1</sub>
- L<sub>n</sub> will do the same for L<sub>n+1</sub>



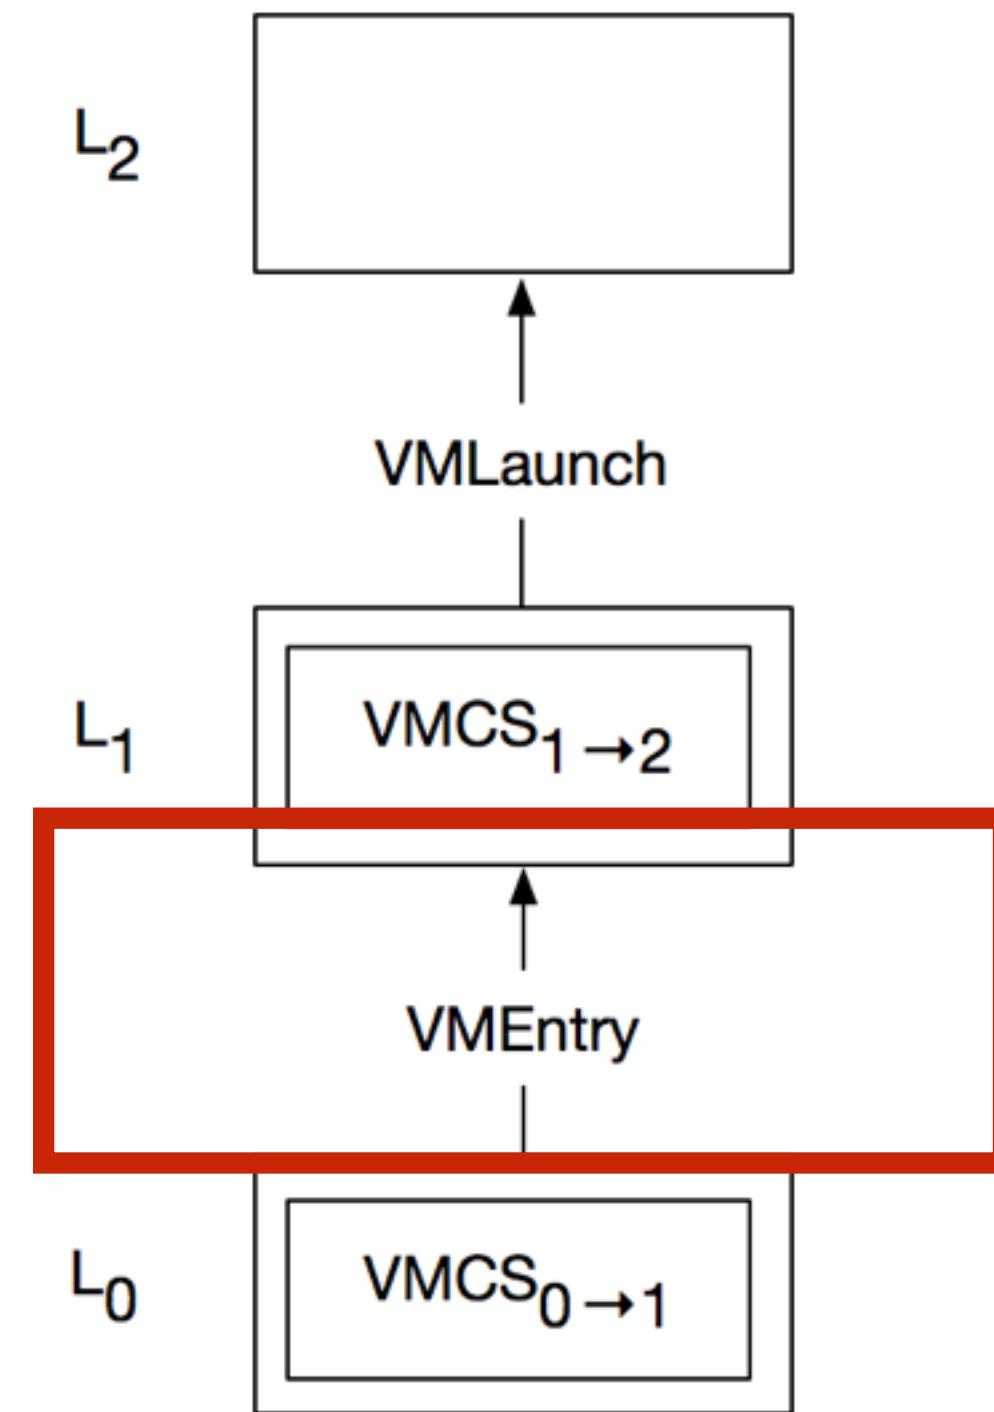
# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$



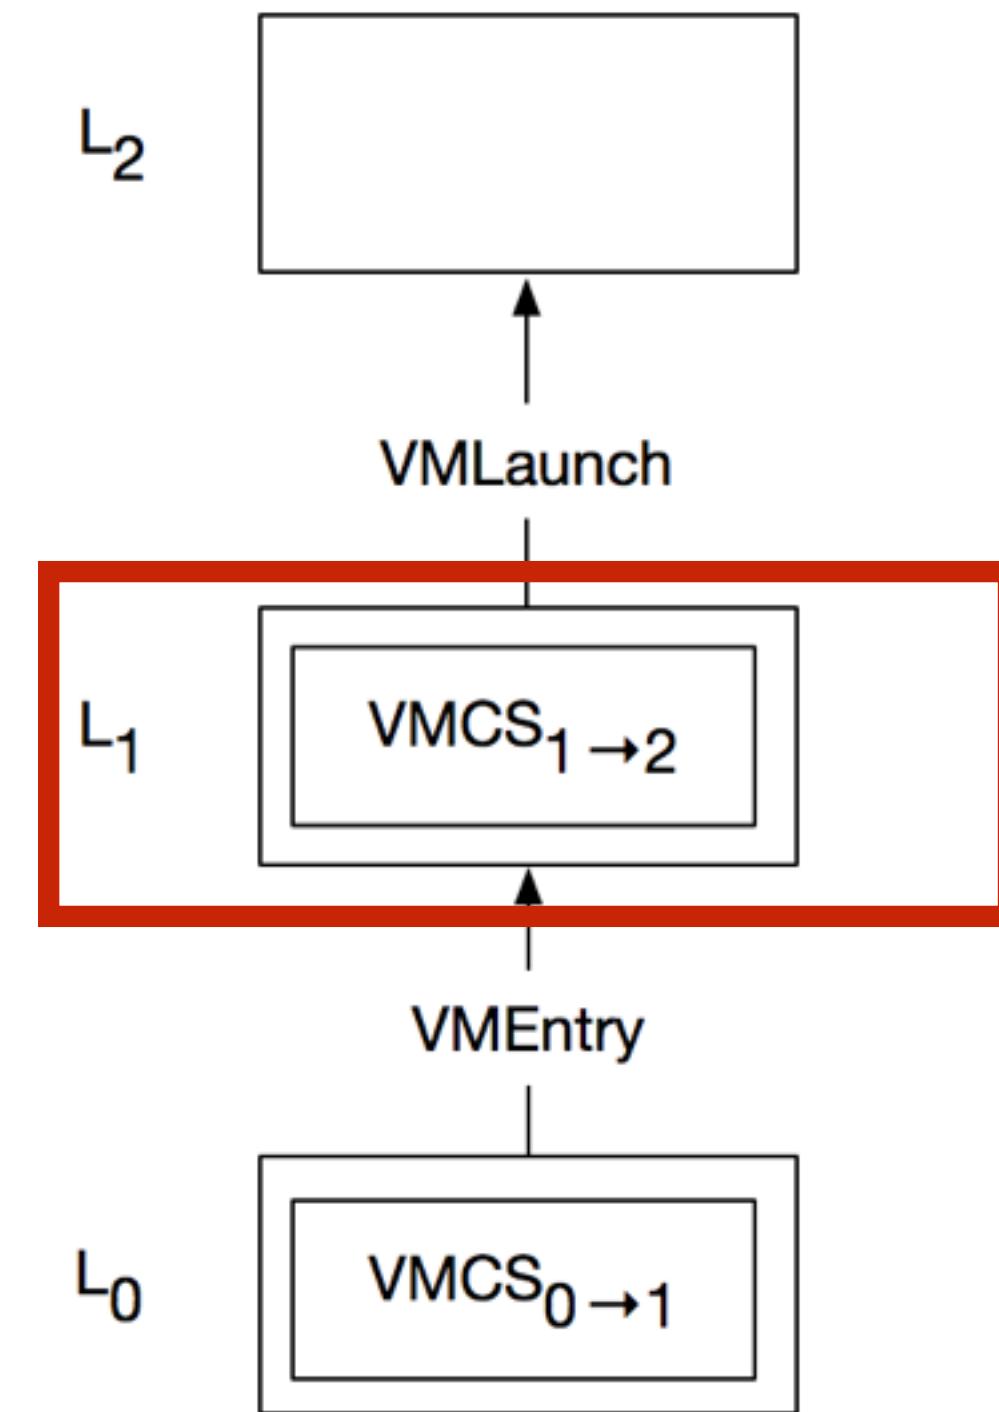
# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$



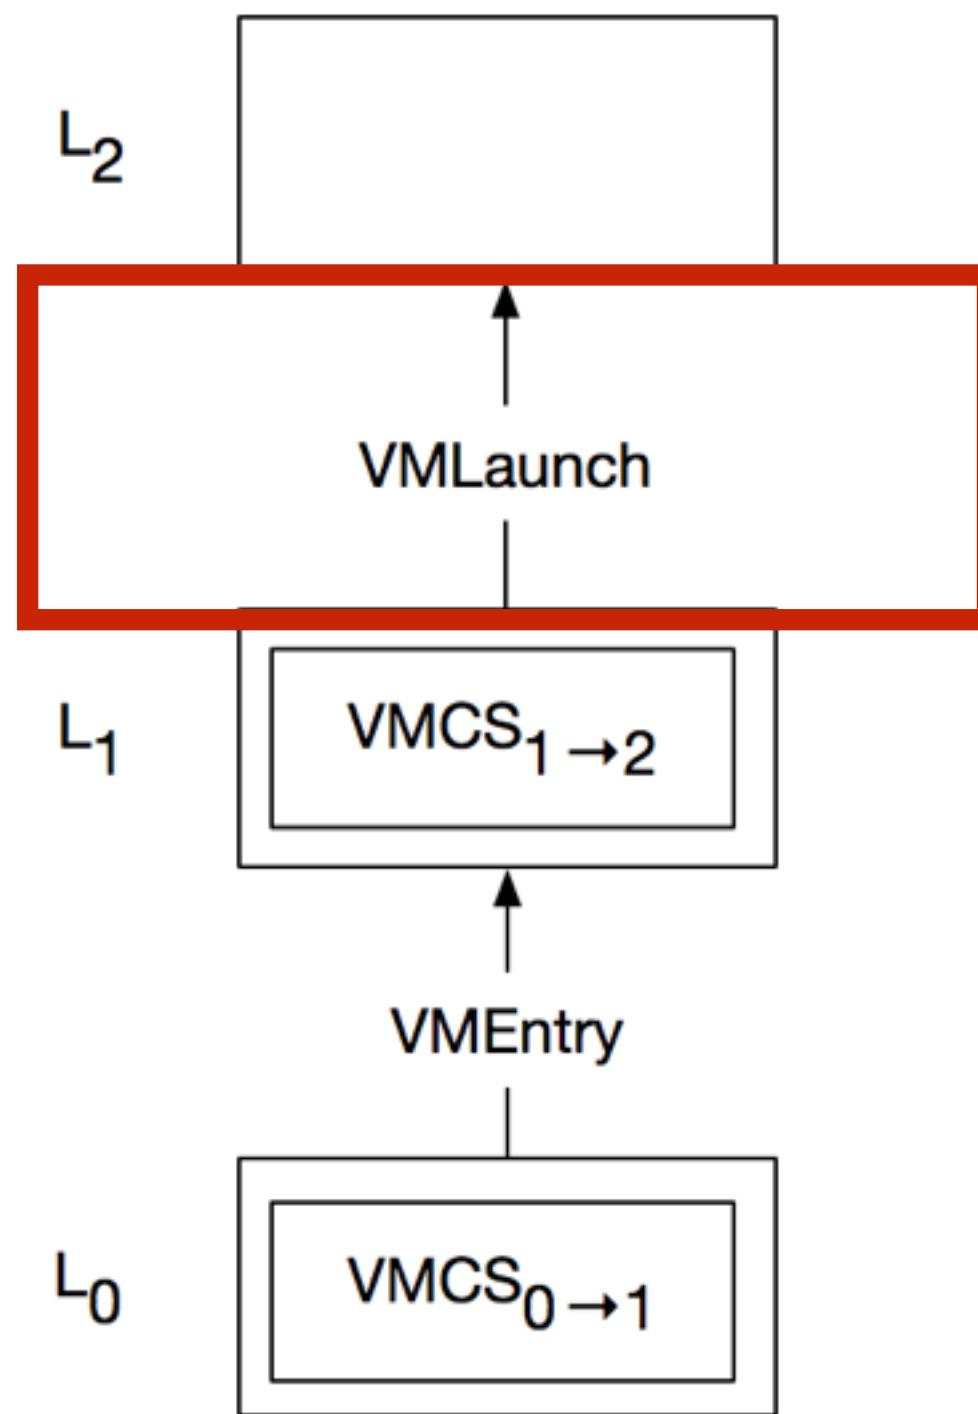
# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$
- $L_1$  runs  $L_2$  by creating  $\text{VMCS}_{1 \rightarrow 2}$ , then calls `vmlaunch`



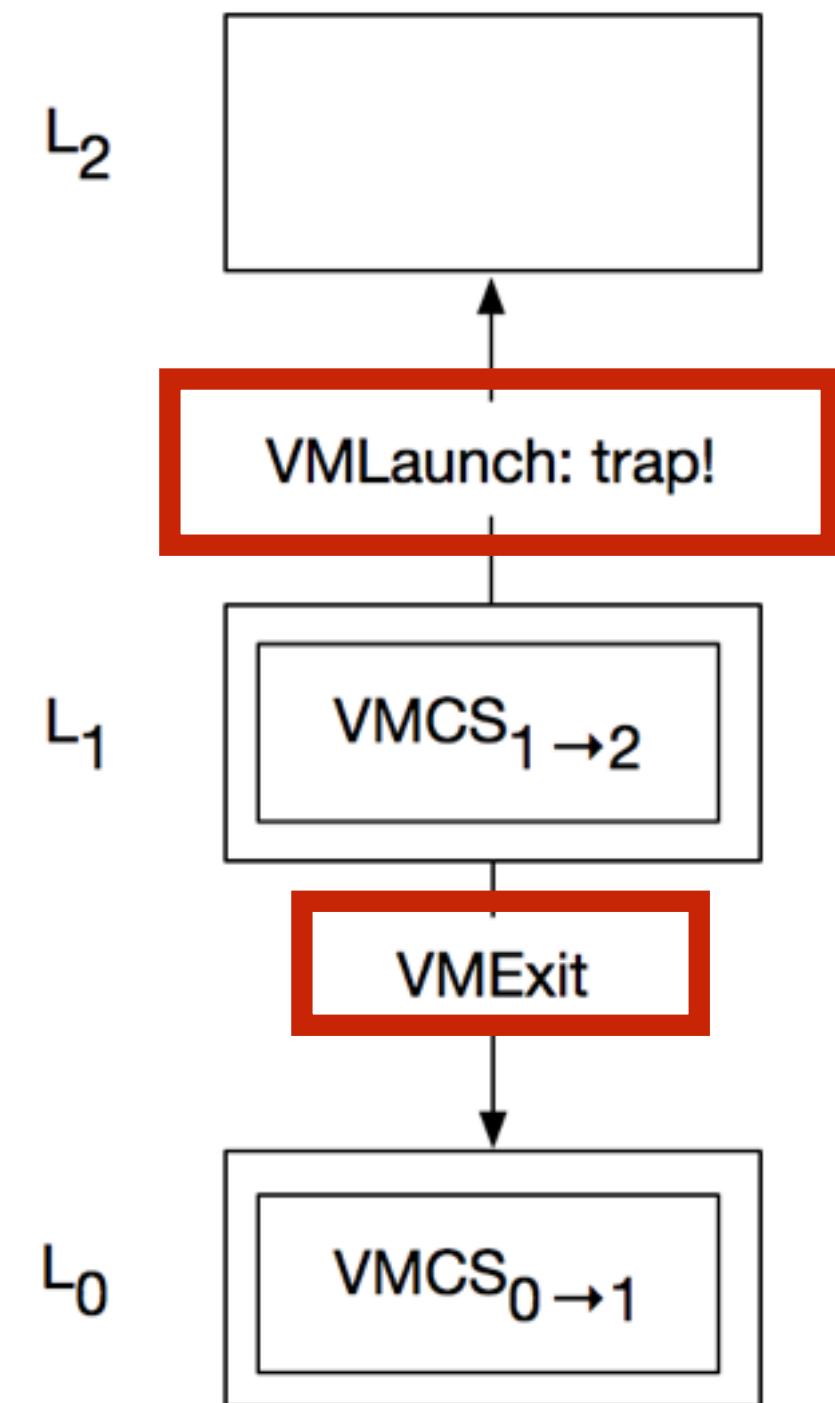
# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$
- $L_1$  runs  $L_2$  by creating  $\text{VMCS}_{1 \rightarrow 2}$  , then calls `vmlaunch`



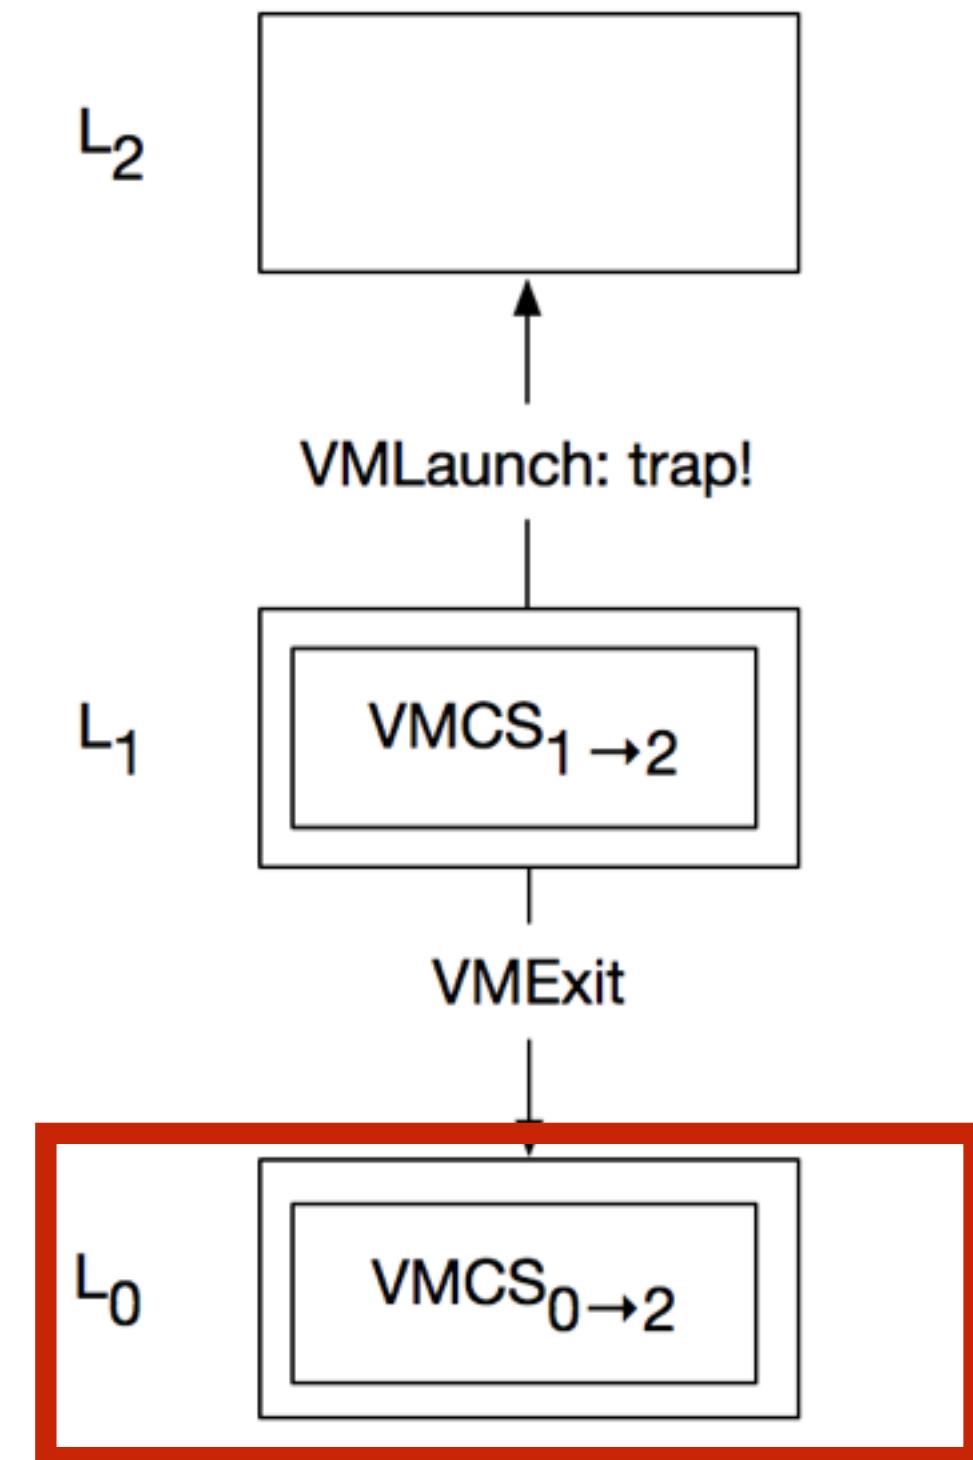
# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$
- $L_1$  runs  $L_2$  by creating  $\text{VMCS}_{1 \rightarrow 2}$ , then calls `vmlaunch`
- `vmlaunch` traps to  $L_0$



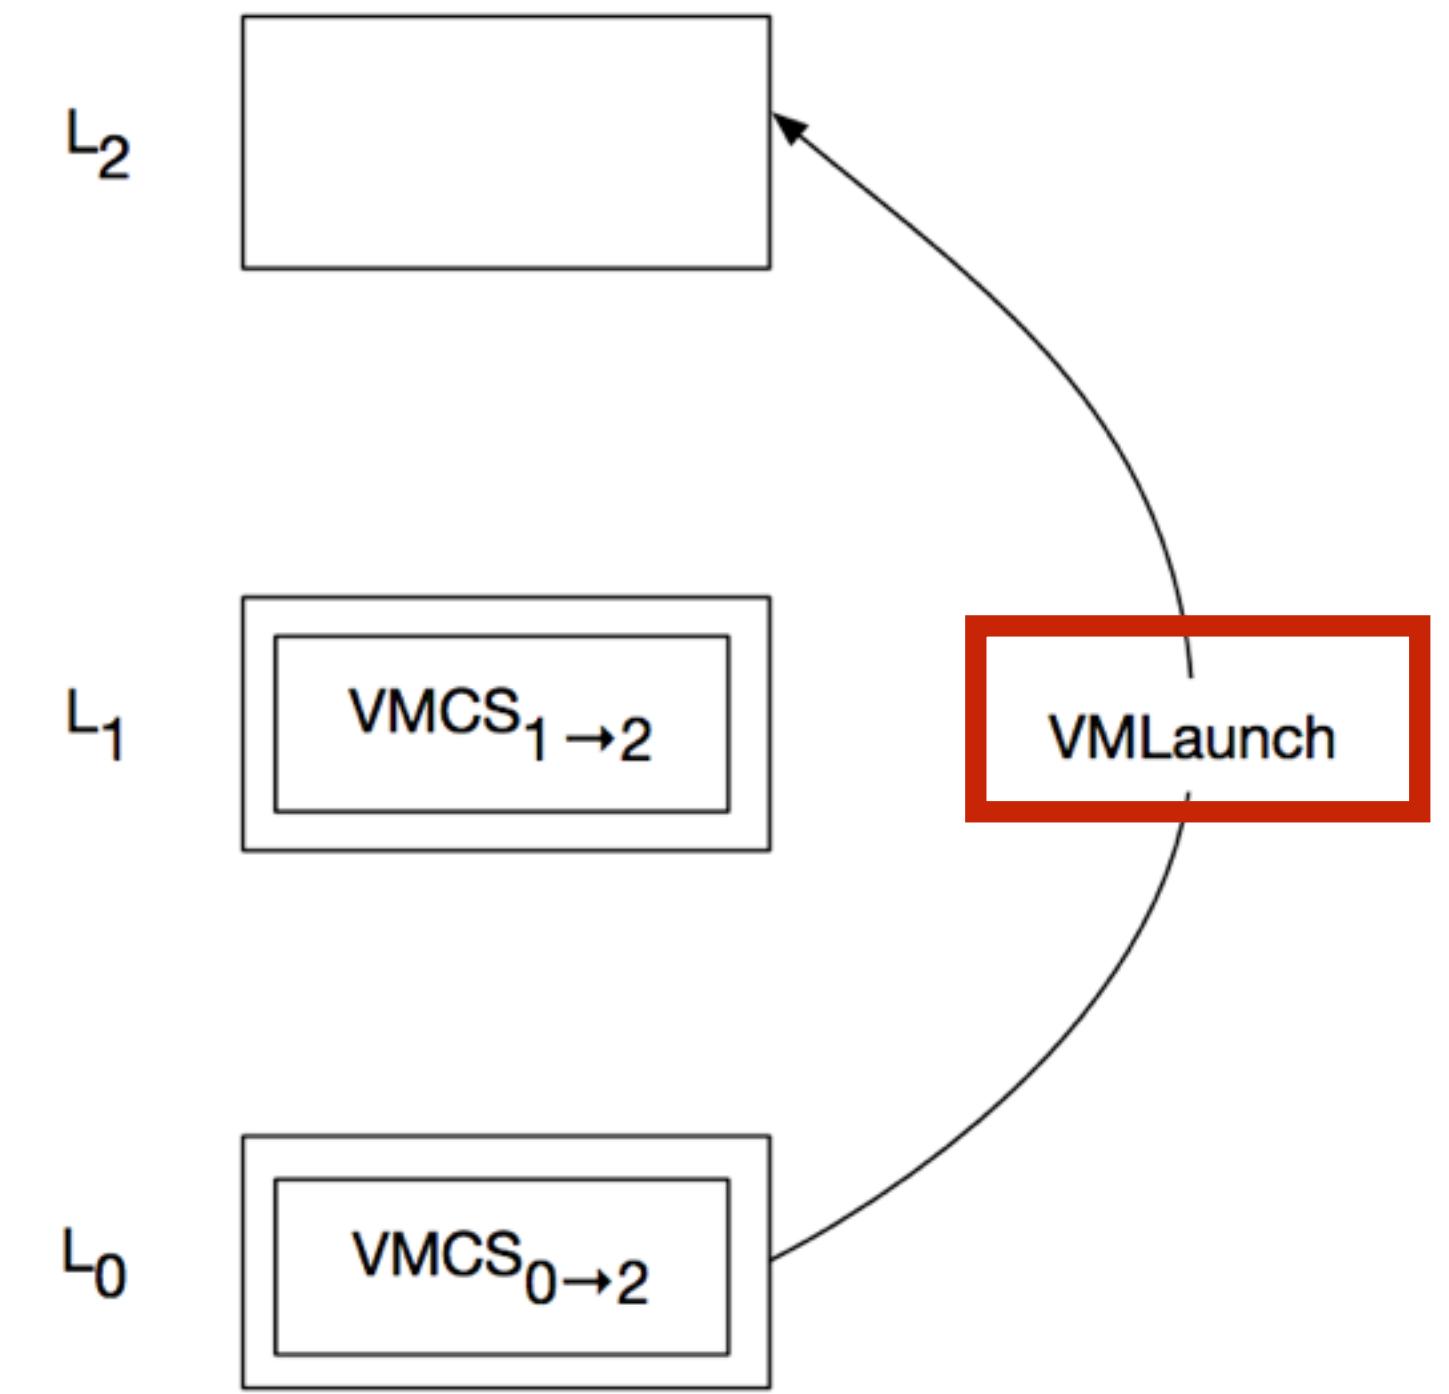
# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$
- $L_1$  runs  $L_2$  by creating  $\text{VMCS}_{1 \rightarrow 2}$ , then calls `vmlaunch`
- `vmlaunch` traps to  $L_0$
- $L_0$  merges  $\text{VMCS}_{0 \rightarrow 1}$  and  $\text{VMCS}_{1 \rightarrow 2}$  to create  $\text{VMCS}_{0 \rightarrow 2}$



# Nested CPU Virtualization: Launching a new VM

- $L_0$  runs  $L_1$  using  $\text{VMCS}_{0 \rightarrow 1}$
- $L_1$  runs  $L_2$  by creating  $\text{VMCS}_{1 \rightarrow 2}$ , then calls `vmlaunch`
- `vmlaunch` traps to  $L_0$
- $L_0$  merges  $\text{VMCS}_{0 \rightarrow 1}$  and  $\text{VMCS}_{1 \rightarrow 2}$  to create  $\text{VMCS}_{0 \rightarrow 2}$
- $L_0$  launches  $L_2$



# Implementing Turtles

- Nested VMX Virtualization for CPUs
- **Multidimensional Paging for MMUs**
- Multi-level Device Assignment for I/O



# MMU Virtualization

- With nested virtualization, we have to make n translations for n levels  
 $L_2^{virtual} \rightarrow L_2^{physical}, L_2 \rightarrow L_1, L_1 \rightarrow L_0$



# MMU Virtualization

- With nested virtualization, we have to make n translations for n levels  
 $L_2^{virtual} \rightarrow L_2^{physical}, L_2 \rightarrow L_1, L_1 \rightarrow L_0$
- Unfortunately, only **two dimensions** of hardware support



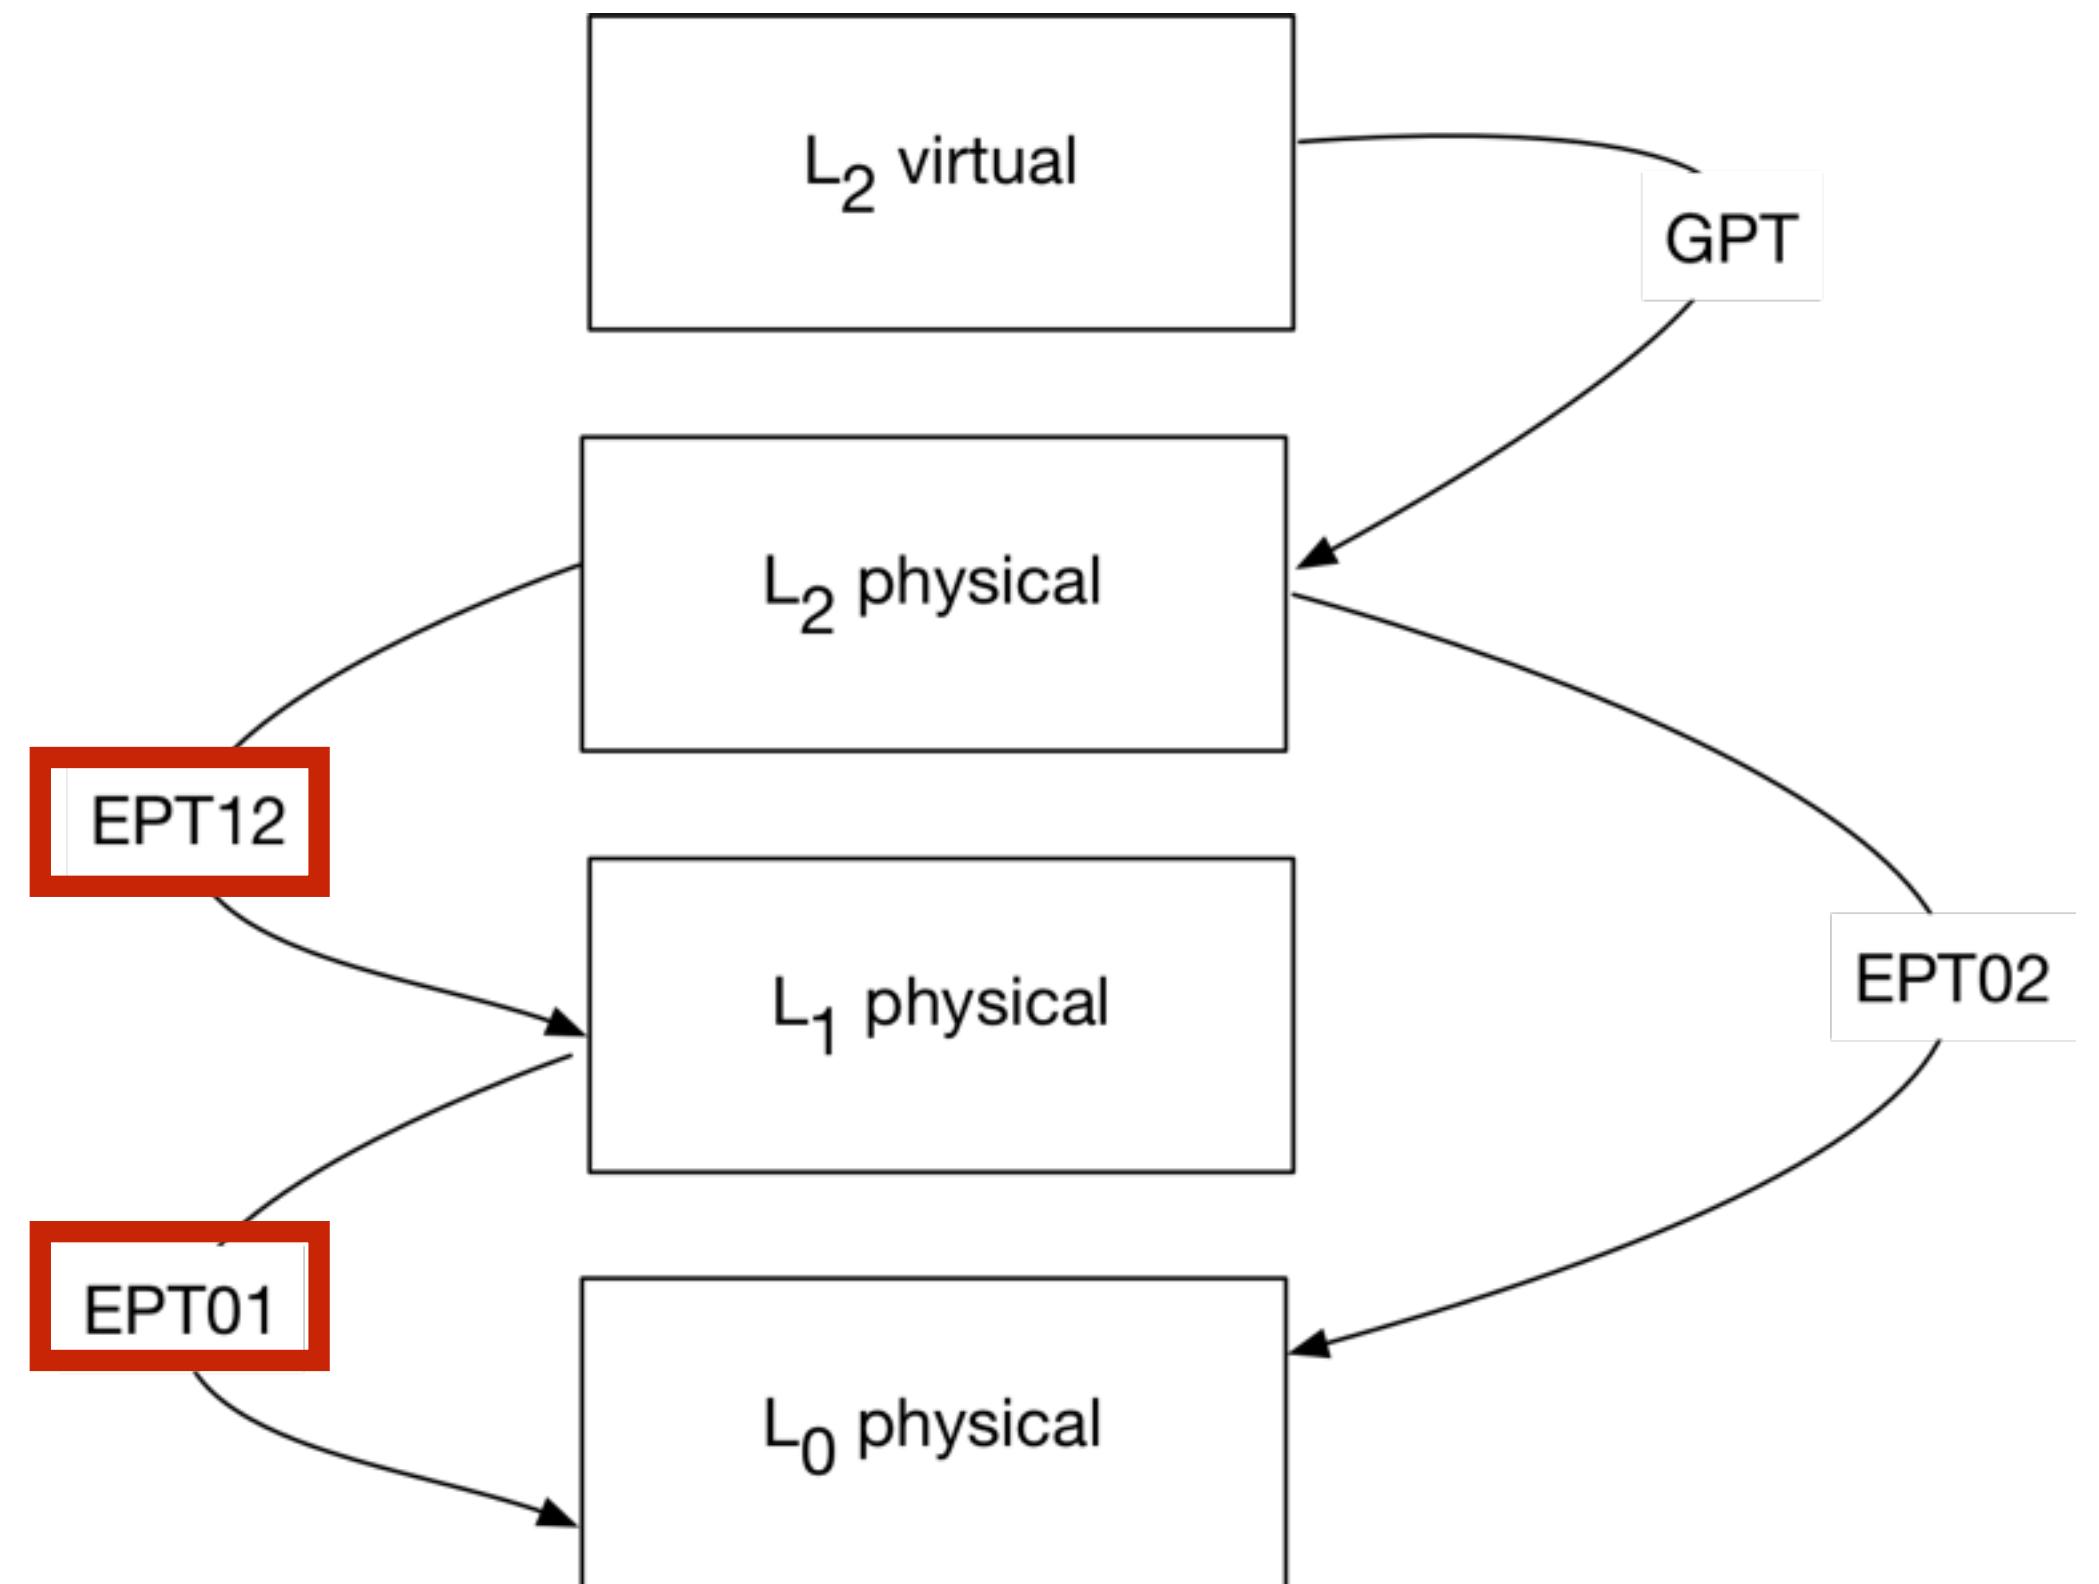
# MMU Virtualization

- With nested virtualization, we have to make n translations for n levels  
 $L_2^{virtual} \rightarrow L_2^{physical}, L_2 \rightarrow L_1, L_1 \rightarrow L_0$
- Unfortunately, only **two dimensions** of hardware support
  - EPT (extended page table) provides translation from guest physical to host physical



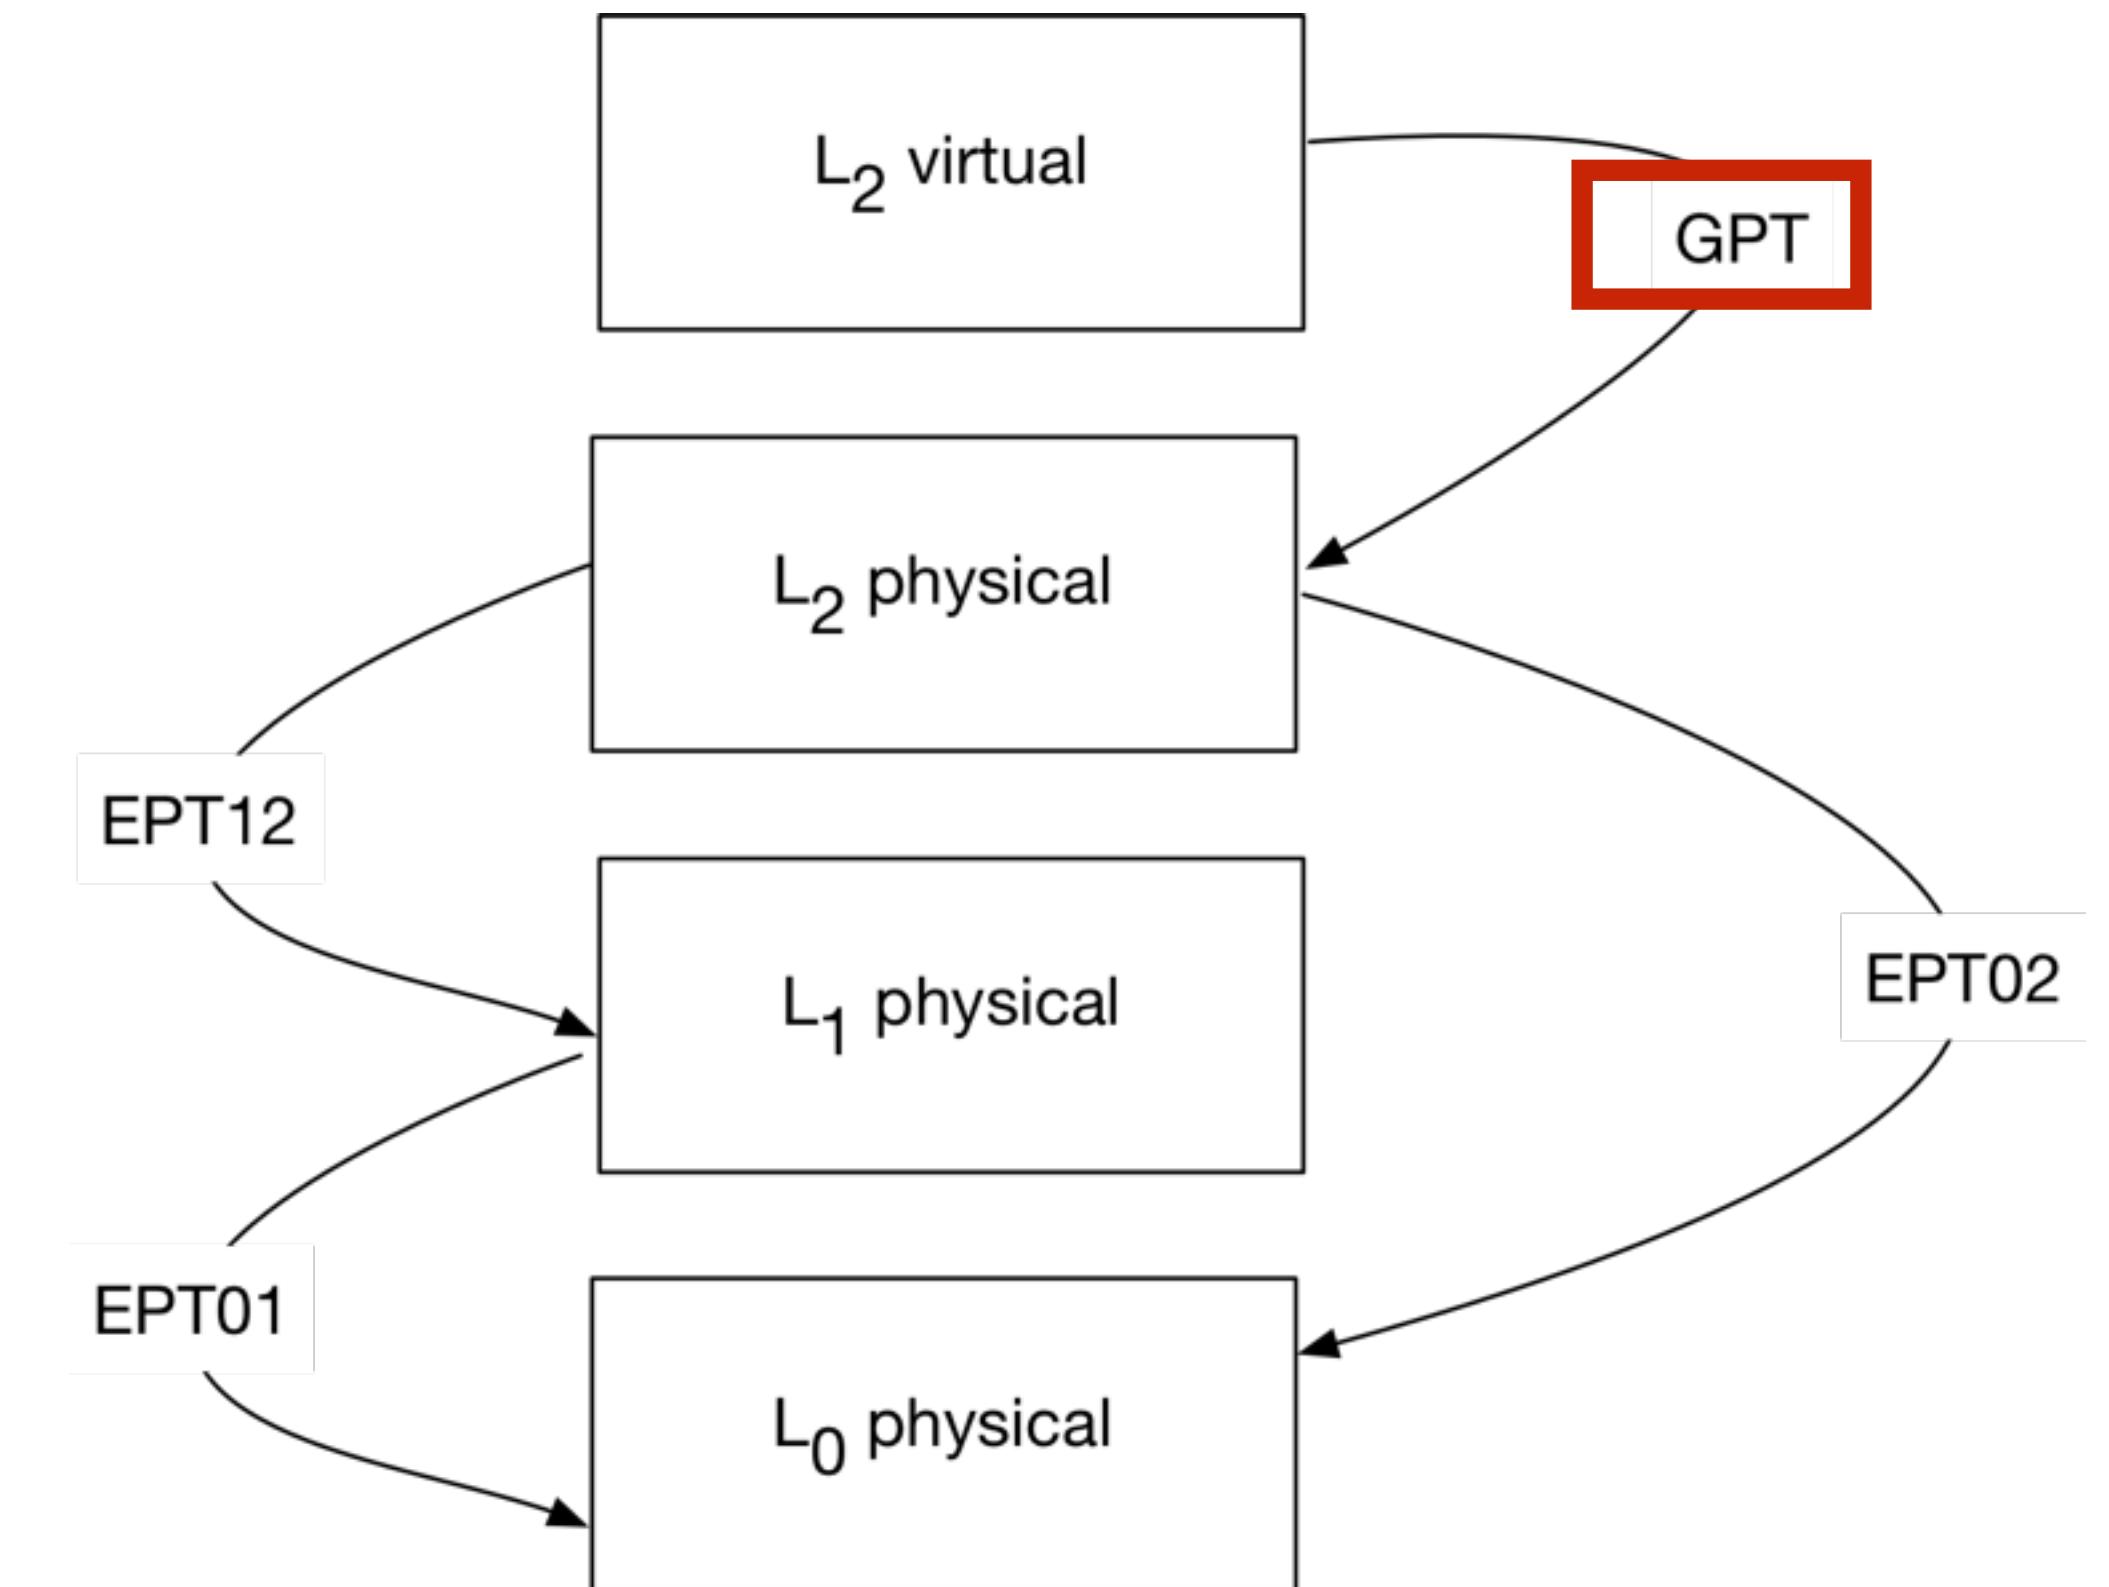
# Multi-dimensional Paging

- Each level creates a virtual EPT that it exposes to higher levels



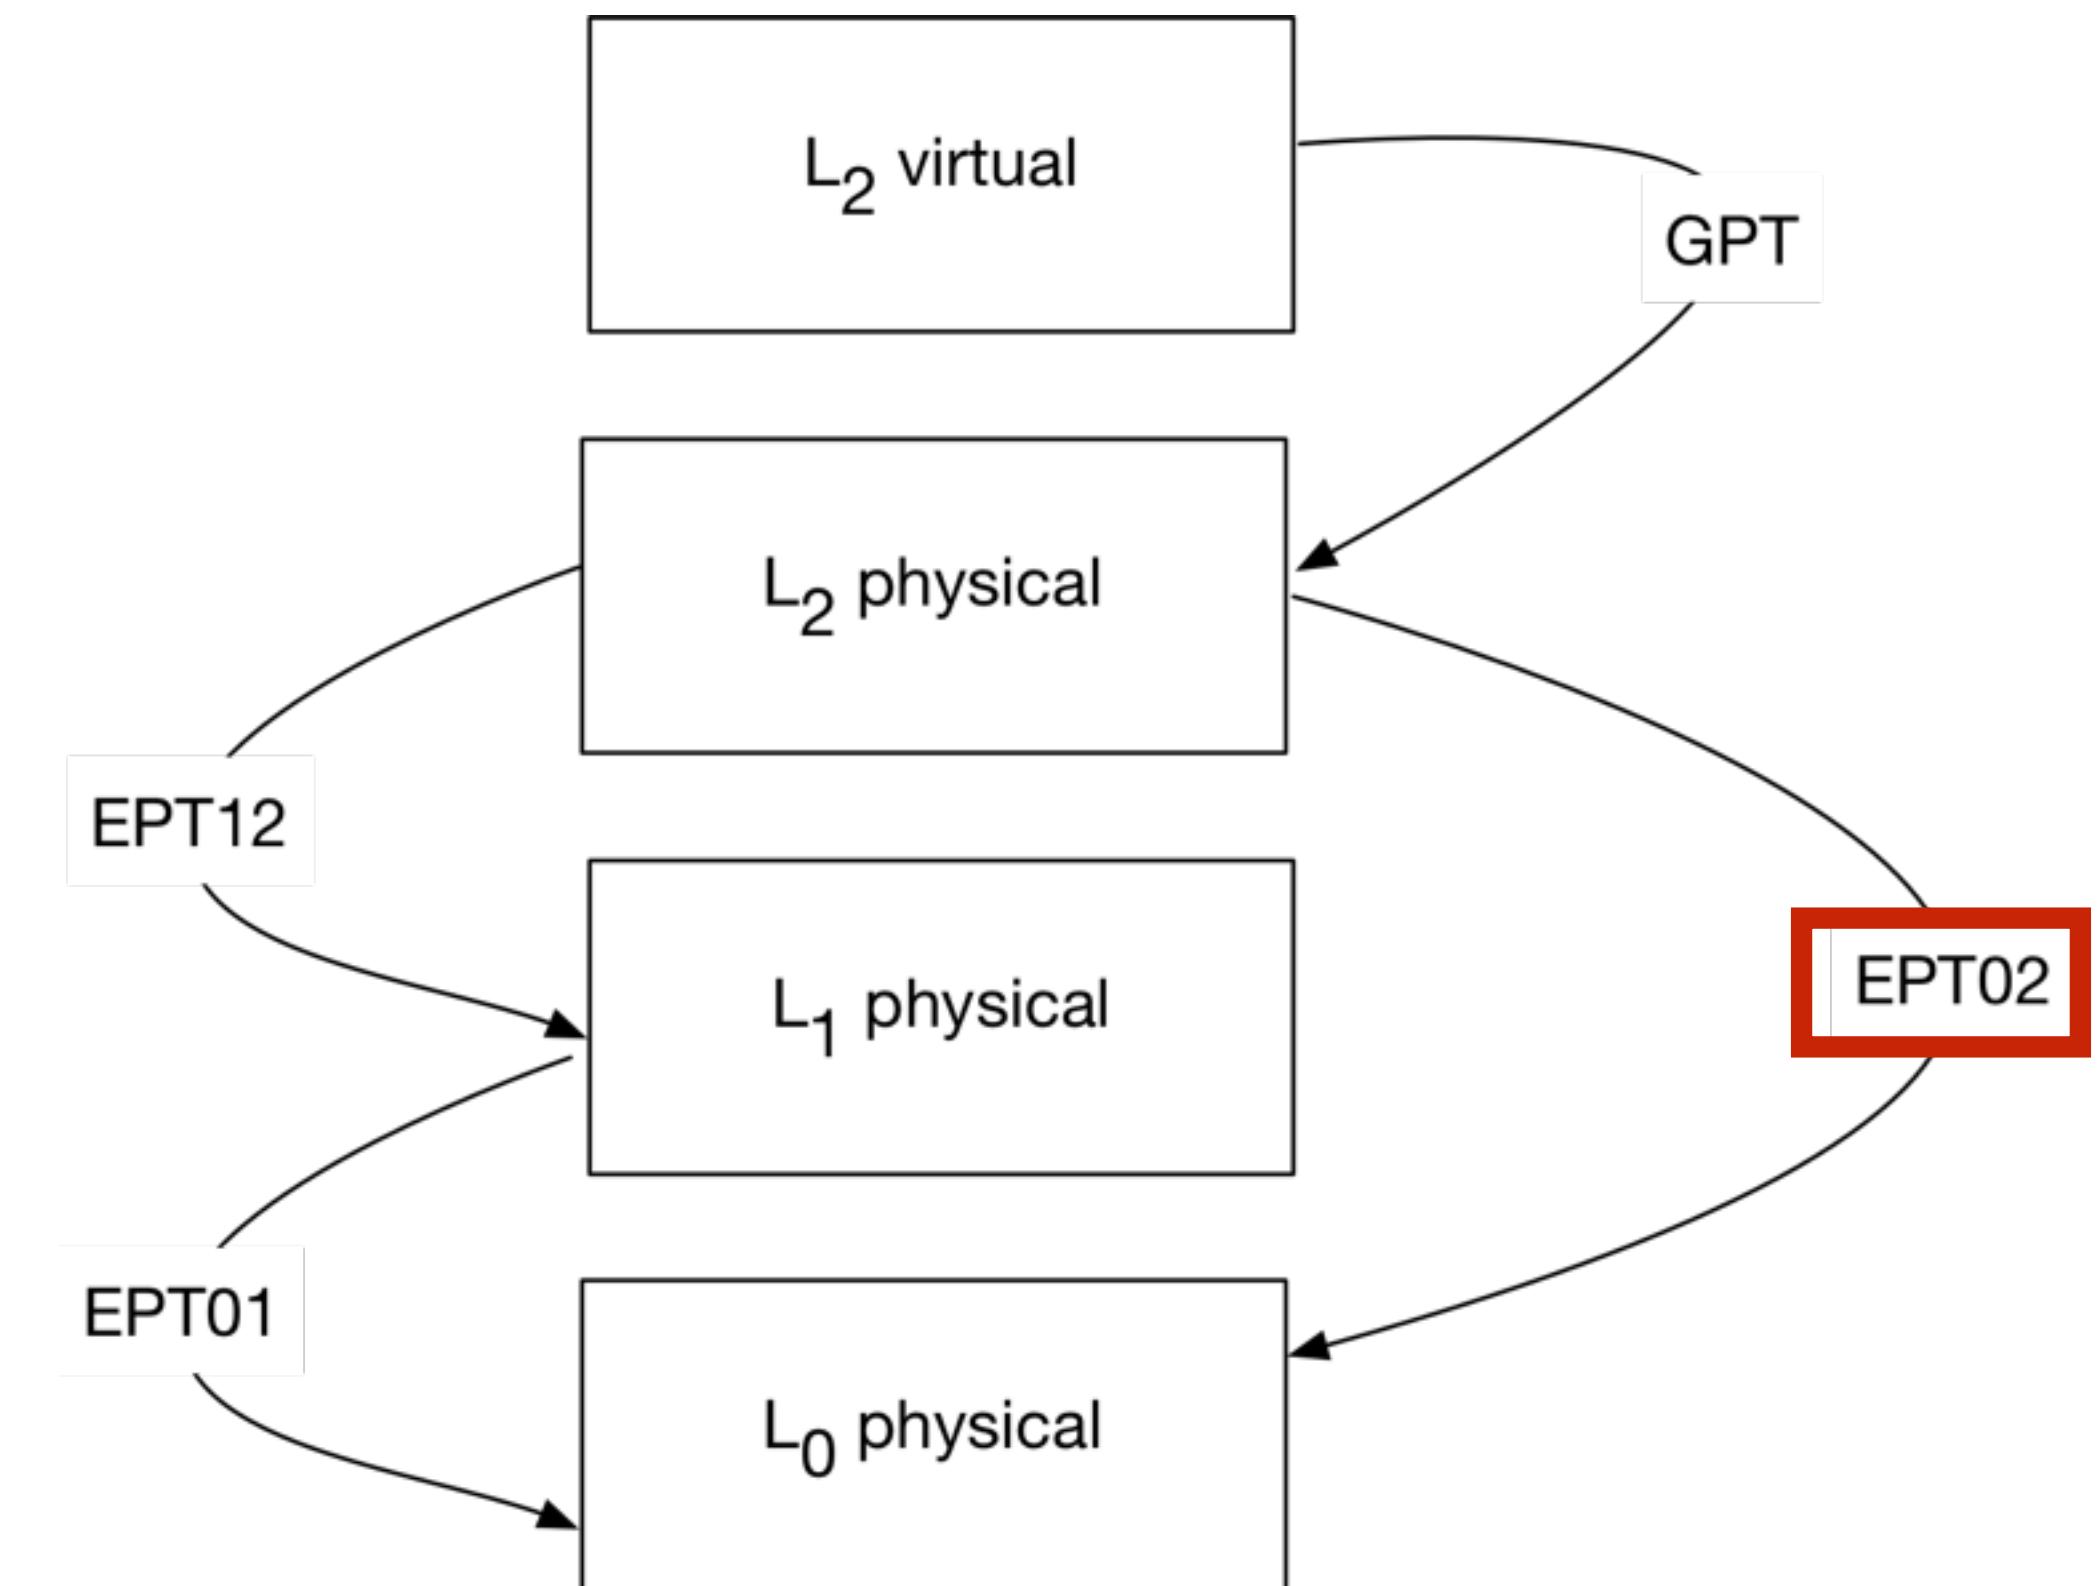
# Multi-dimensional Paging

- Each level creates a virtual EPT that it exposes to higher levels
- Guest handles their own page table



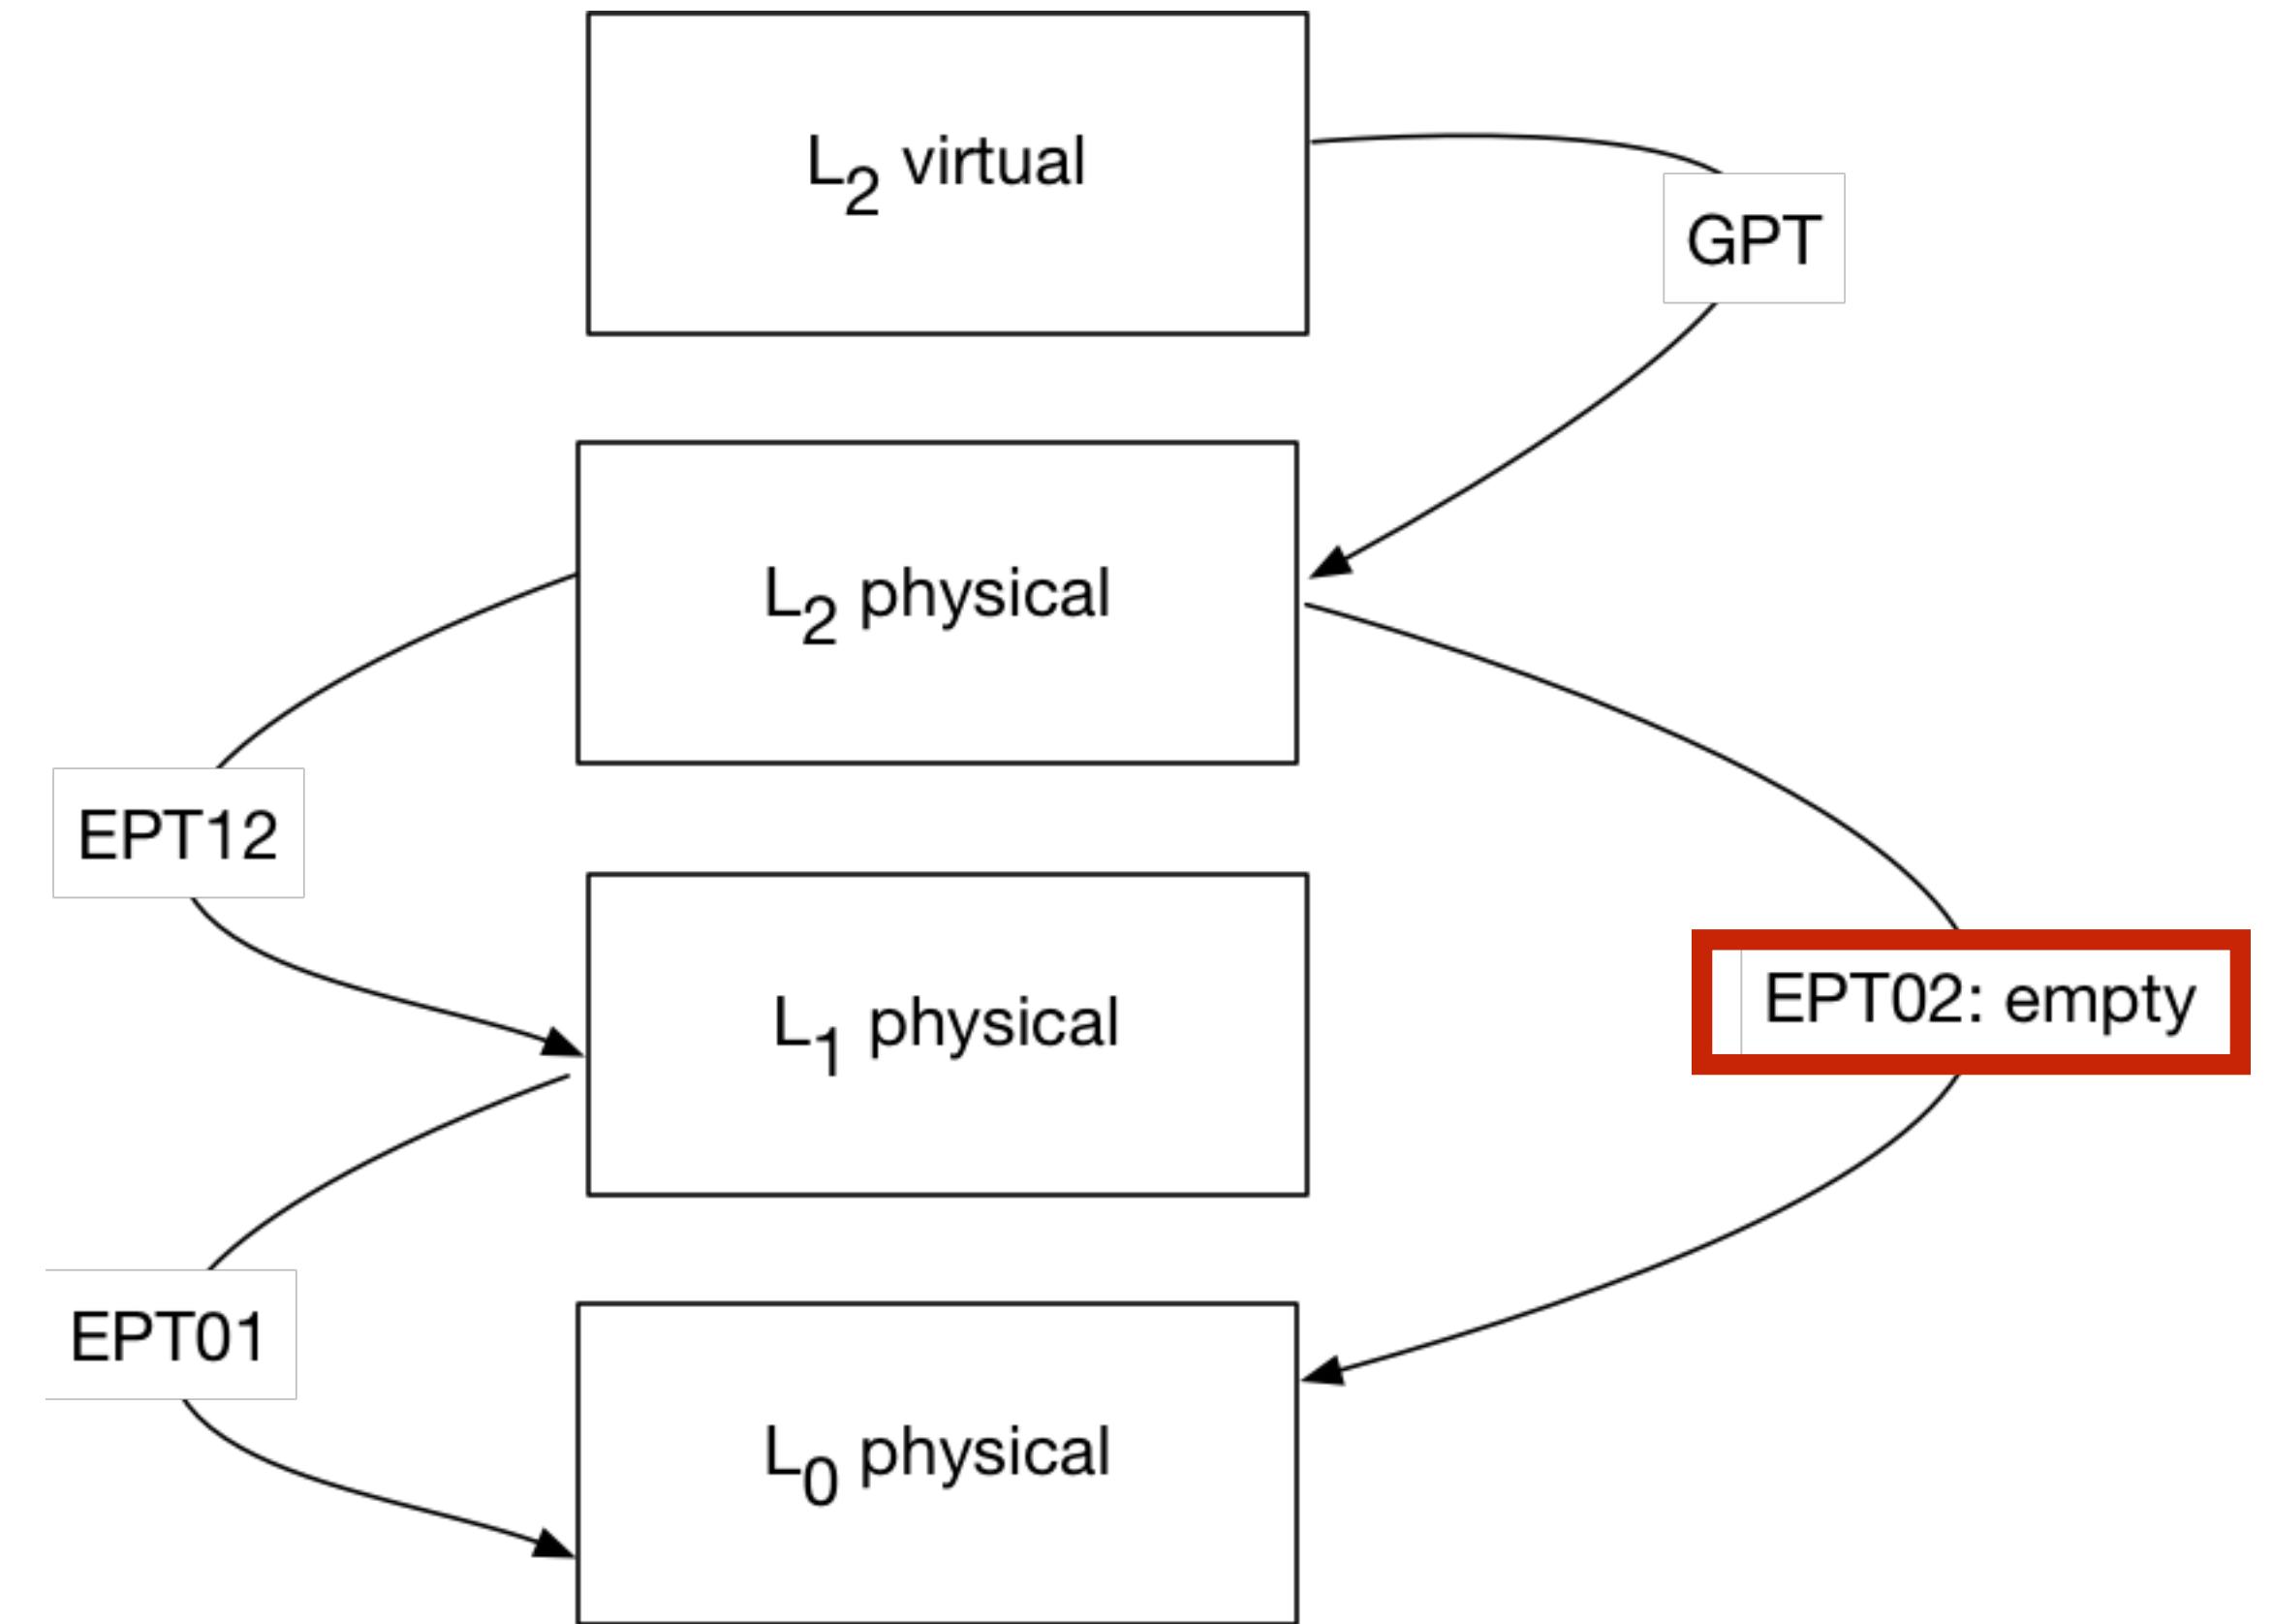
# Multi-dimensional Paging

- Each level creates a virtual EPT that it exposes to higher levels
- Guest handles their own page table
- “Base” hypervisor translates address using *compressed* EPT



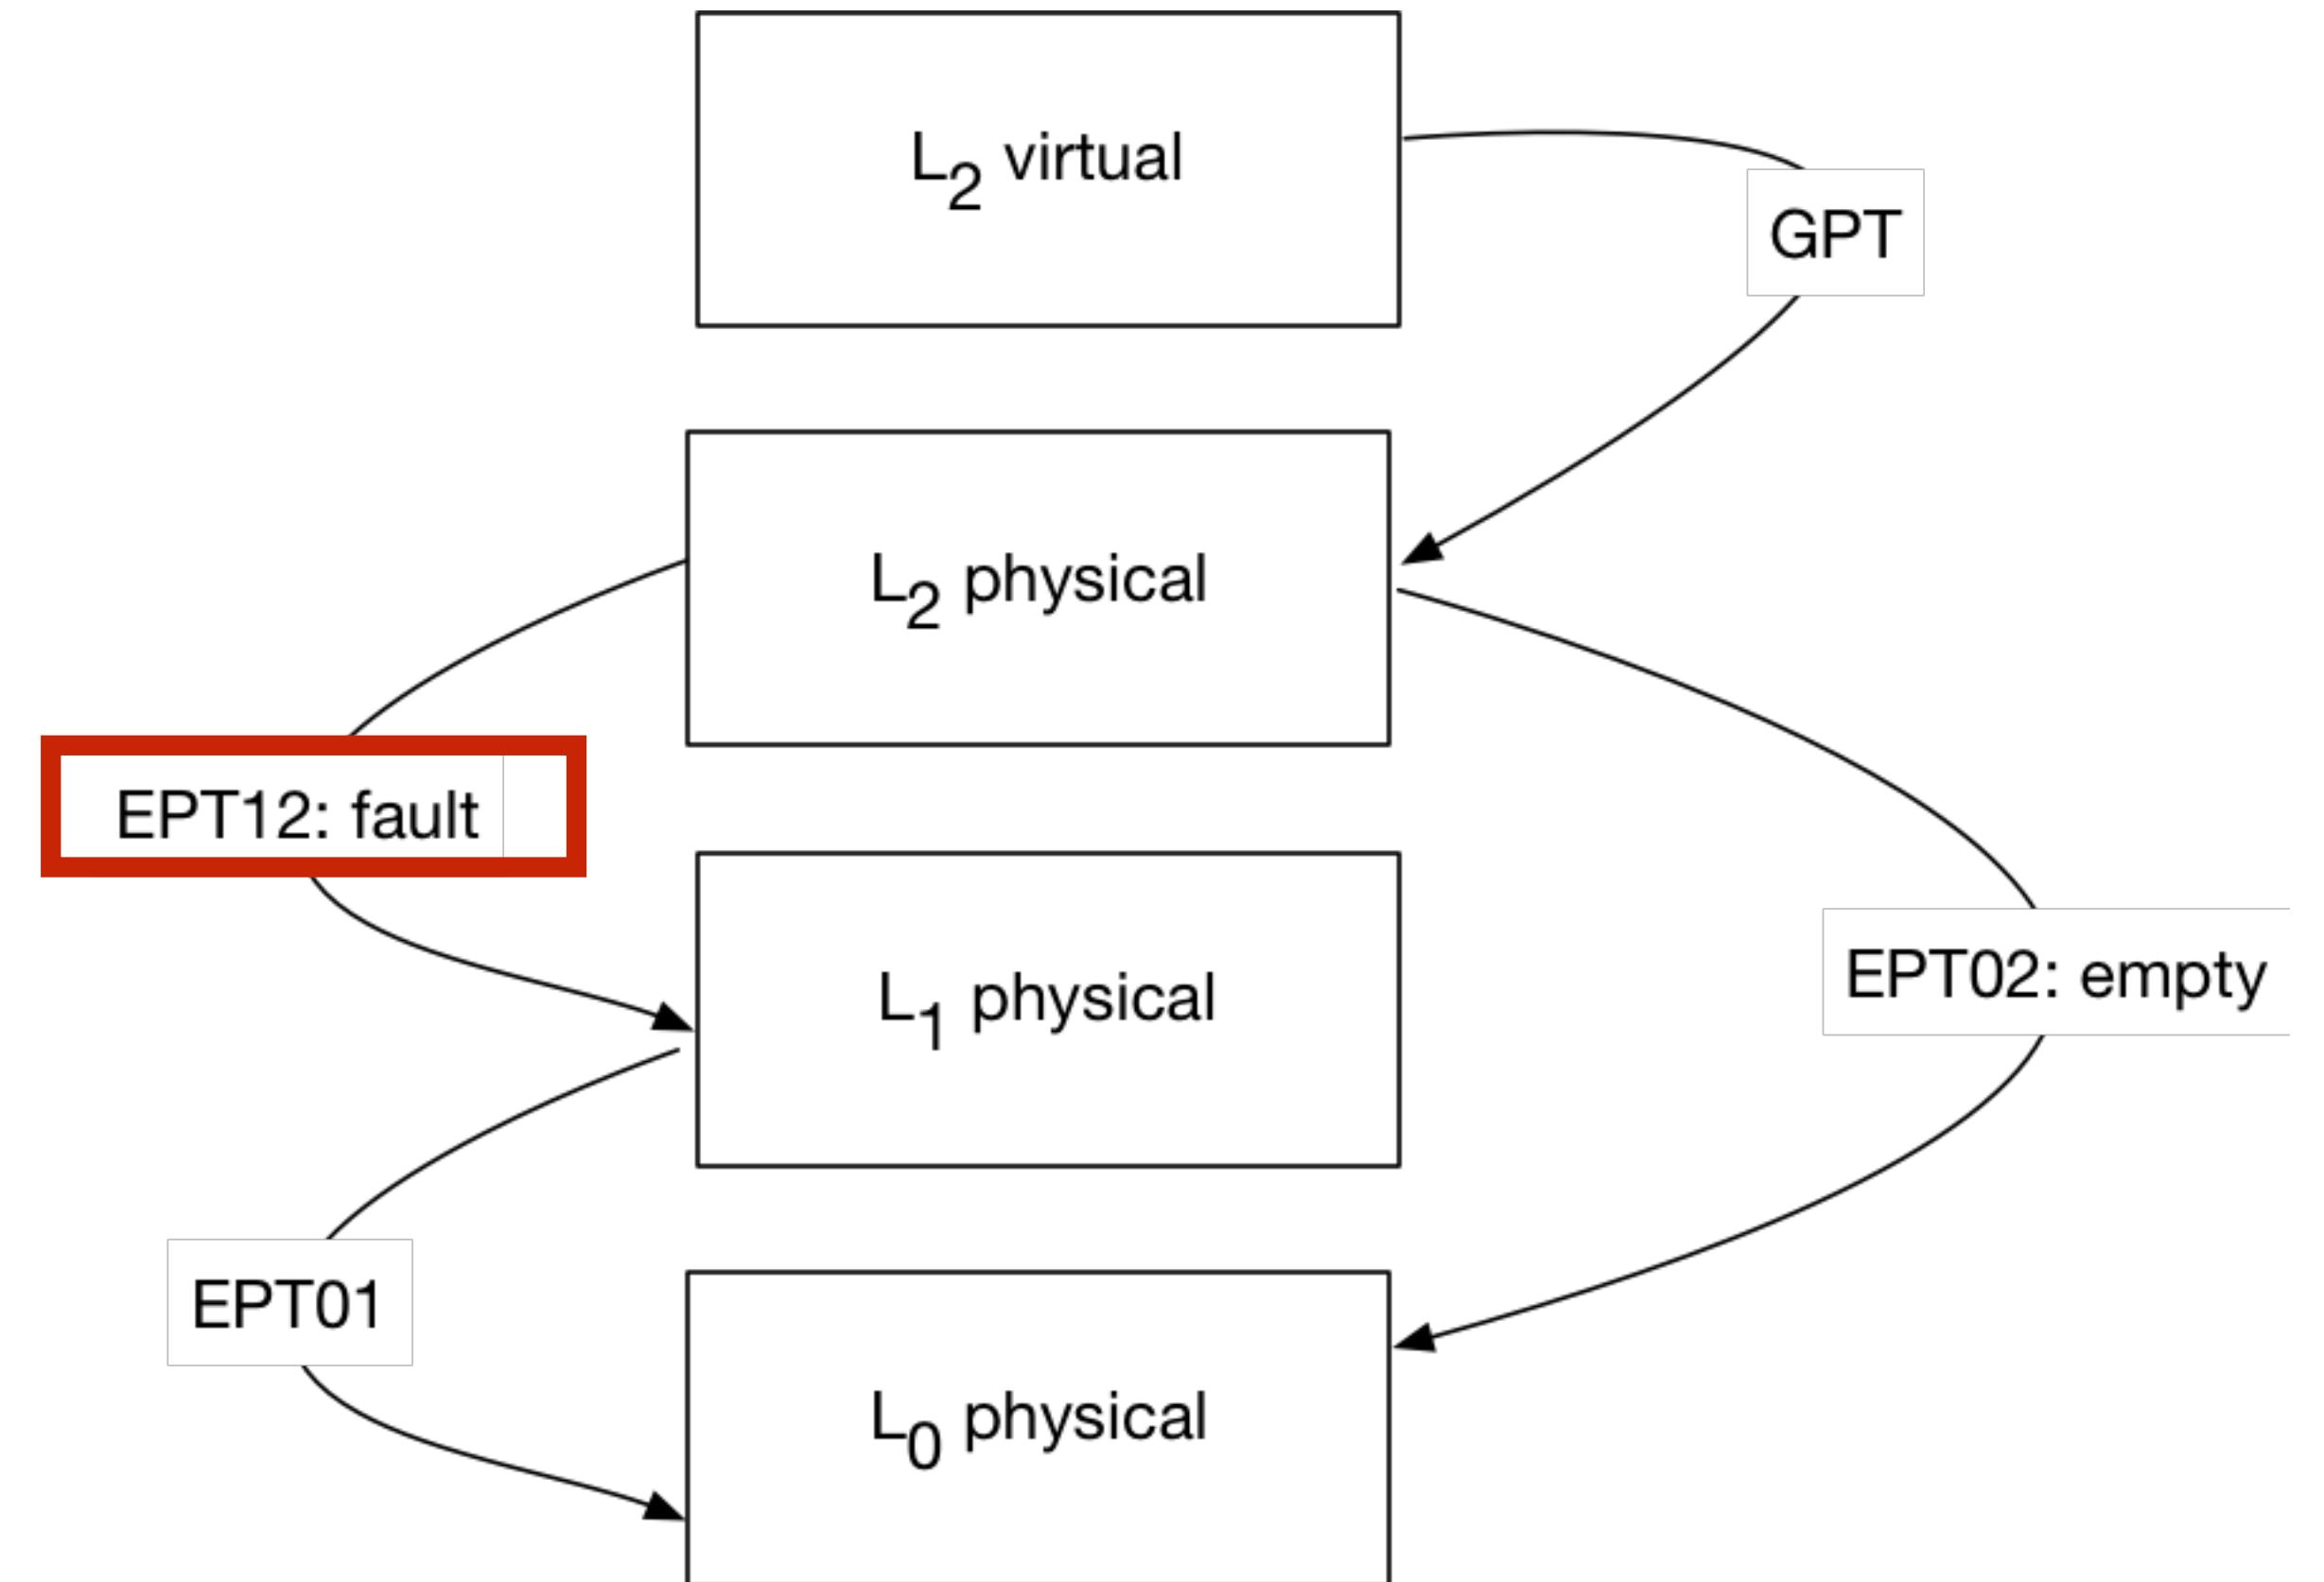
# Building the Compressed EPT

- Start with empty  $EPT_{0 \rightarrow 2}$



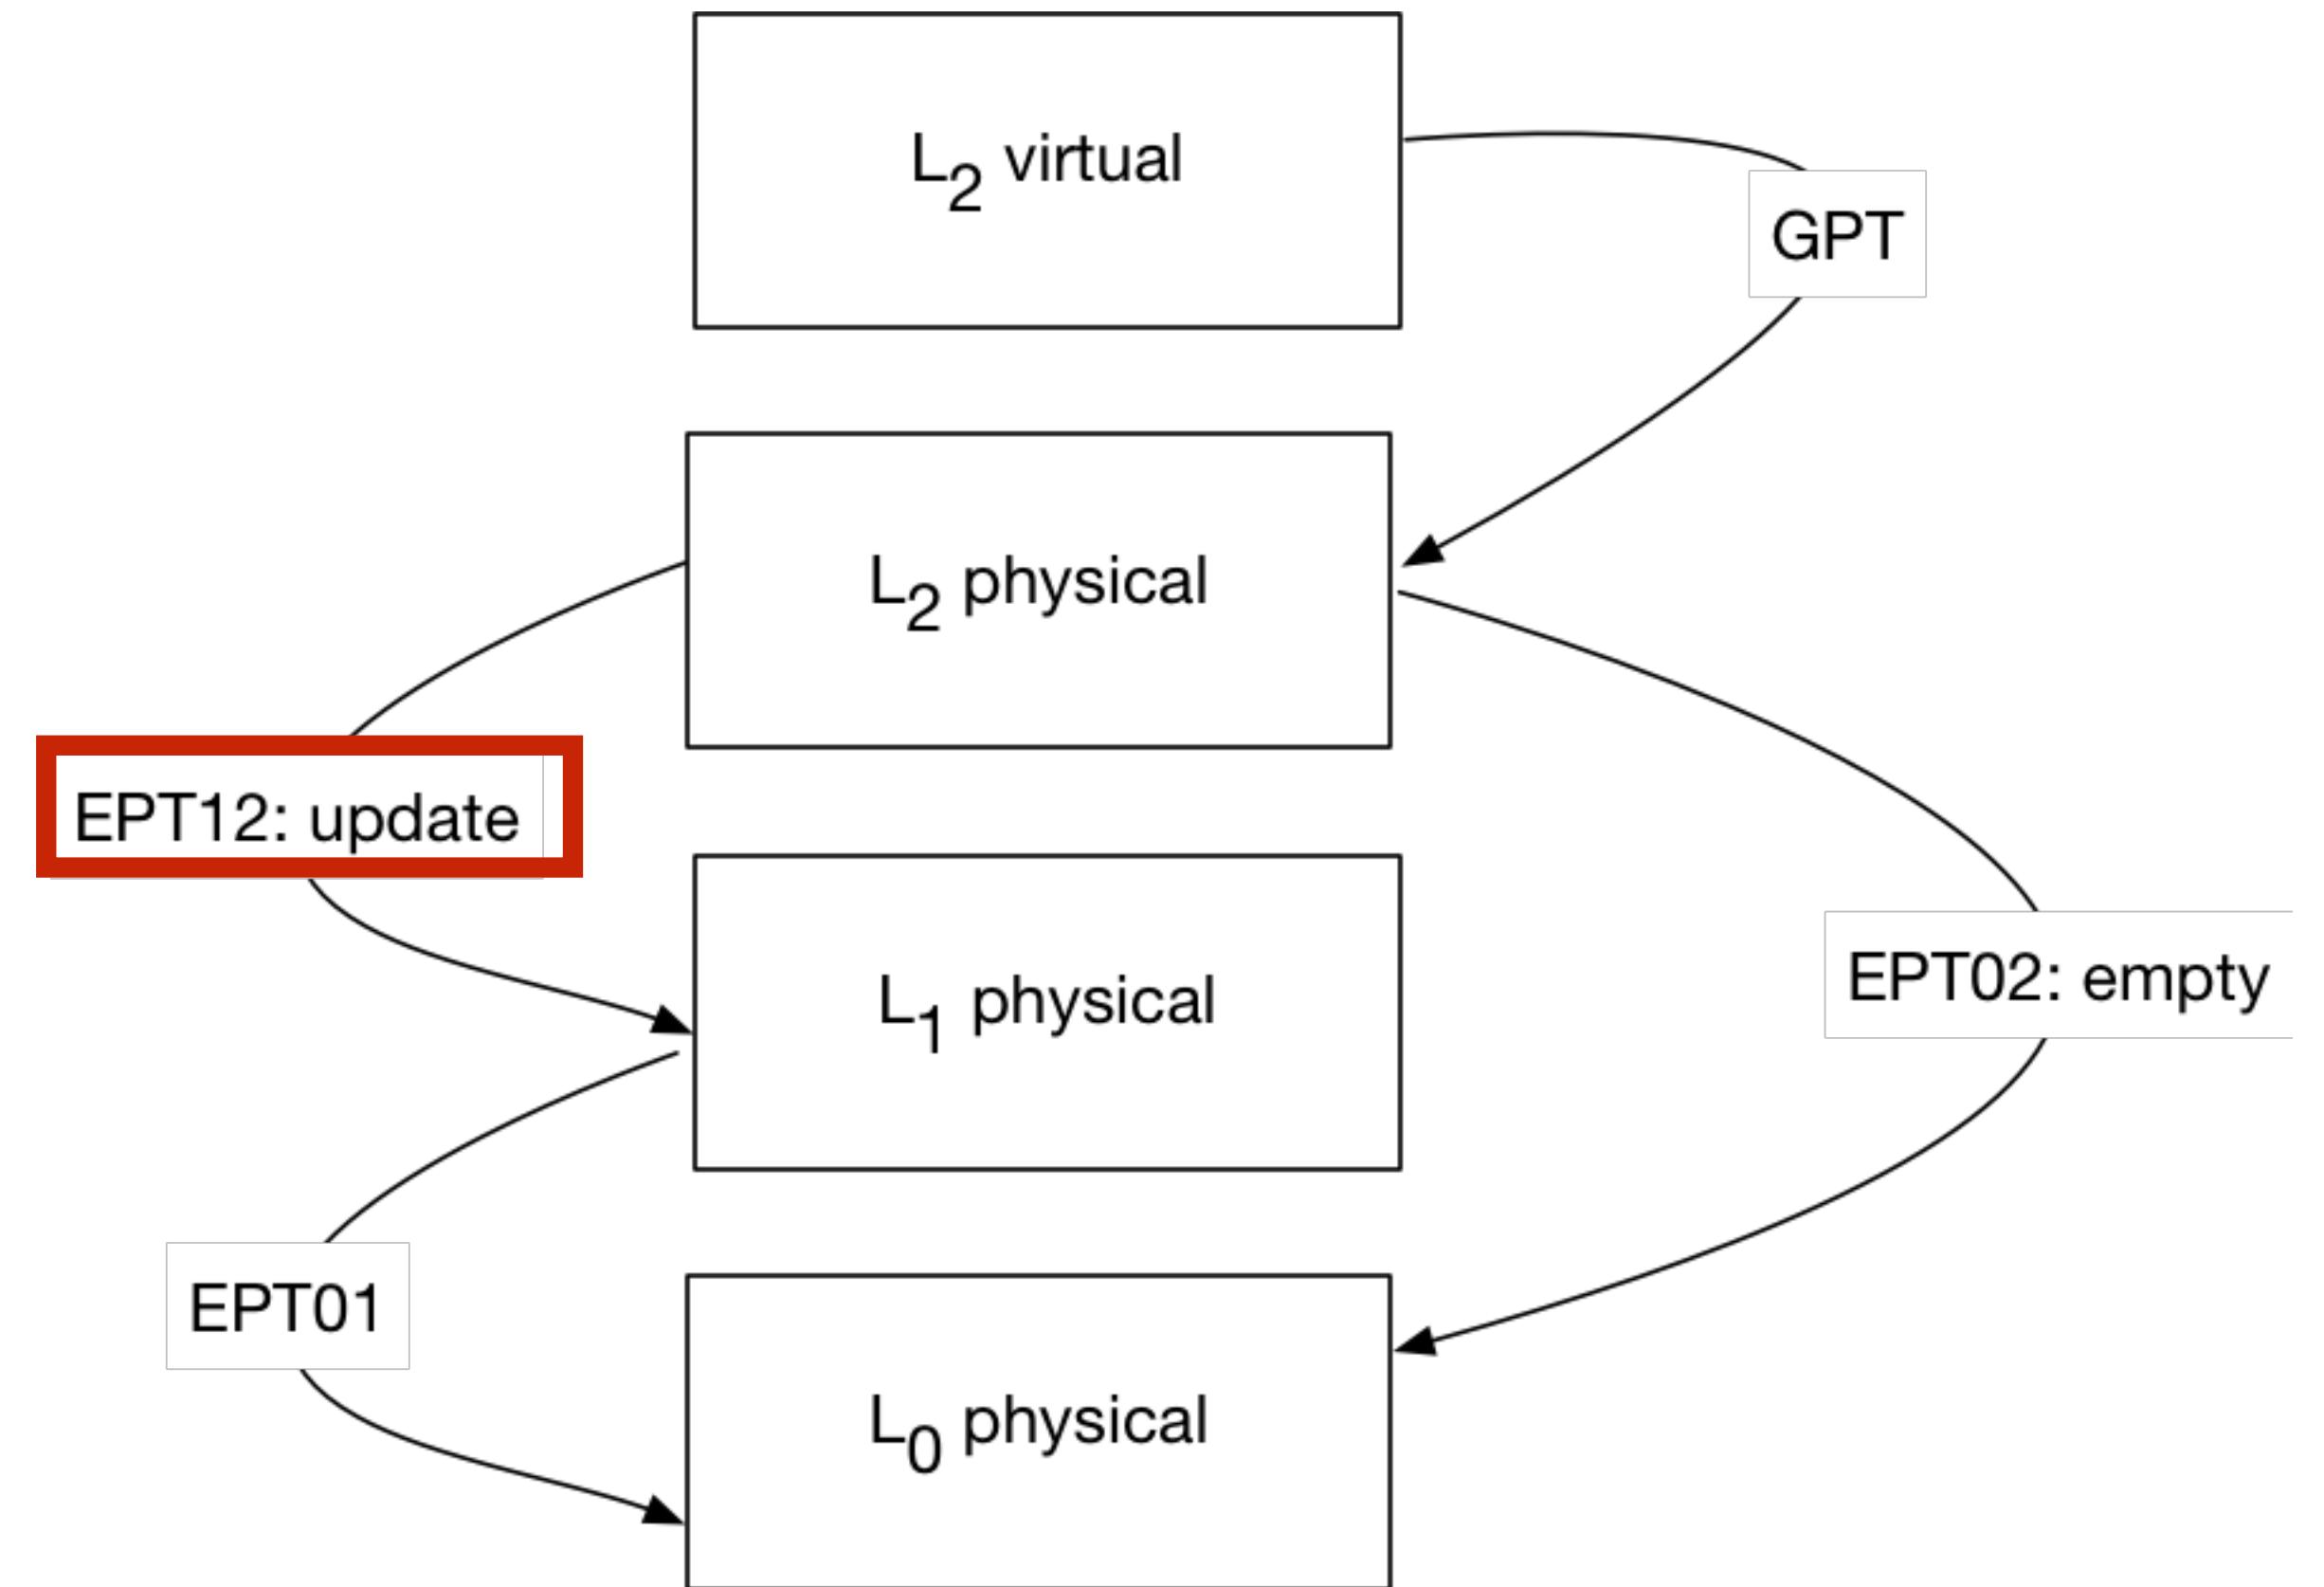
# Building the Compressed EPT

- Start with empty  $EPT_{0 \rightarrow 2}$
- When  $L_2$  causes EPT-violation in  $EPT_{1 \rightarrow 2}$



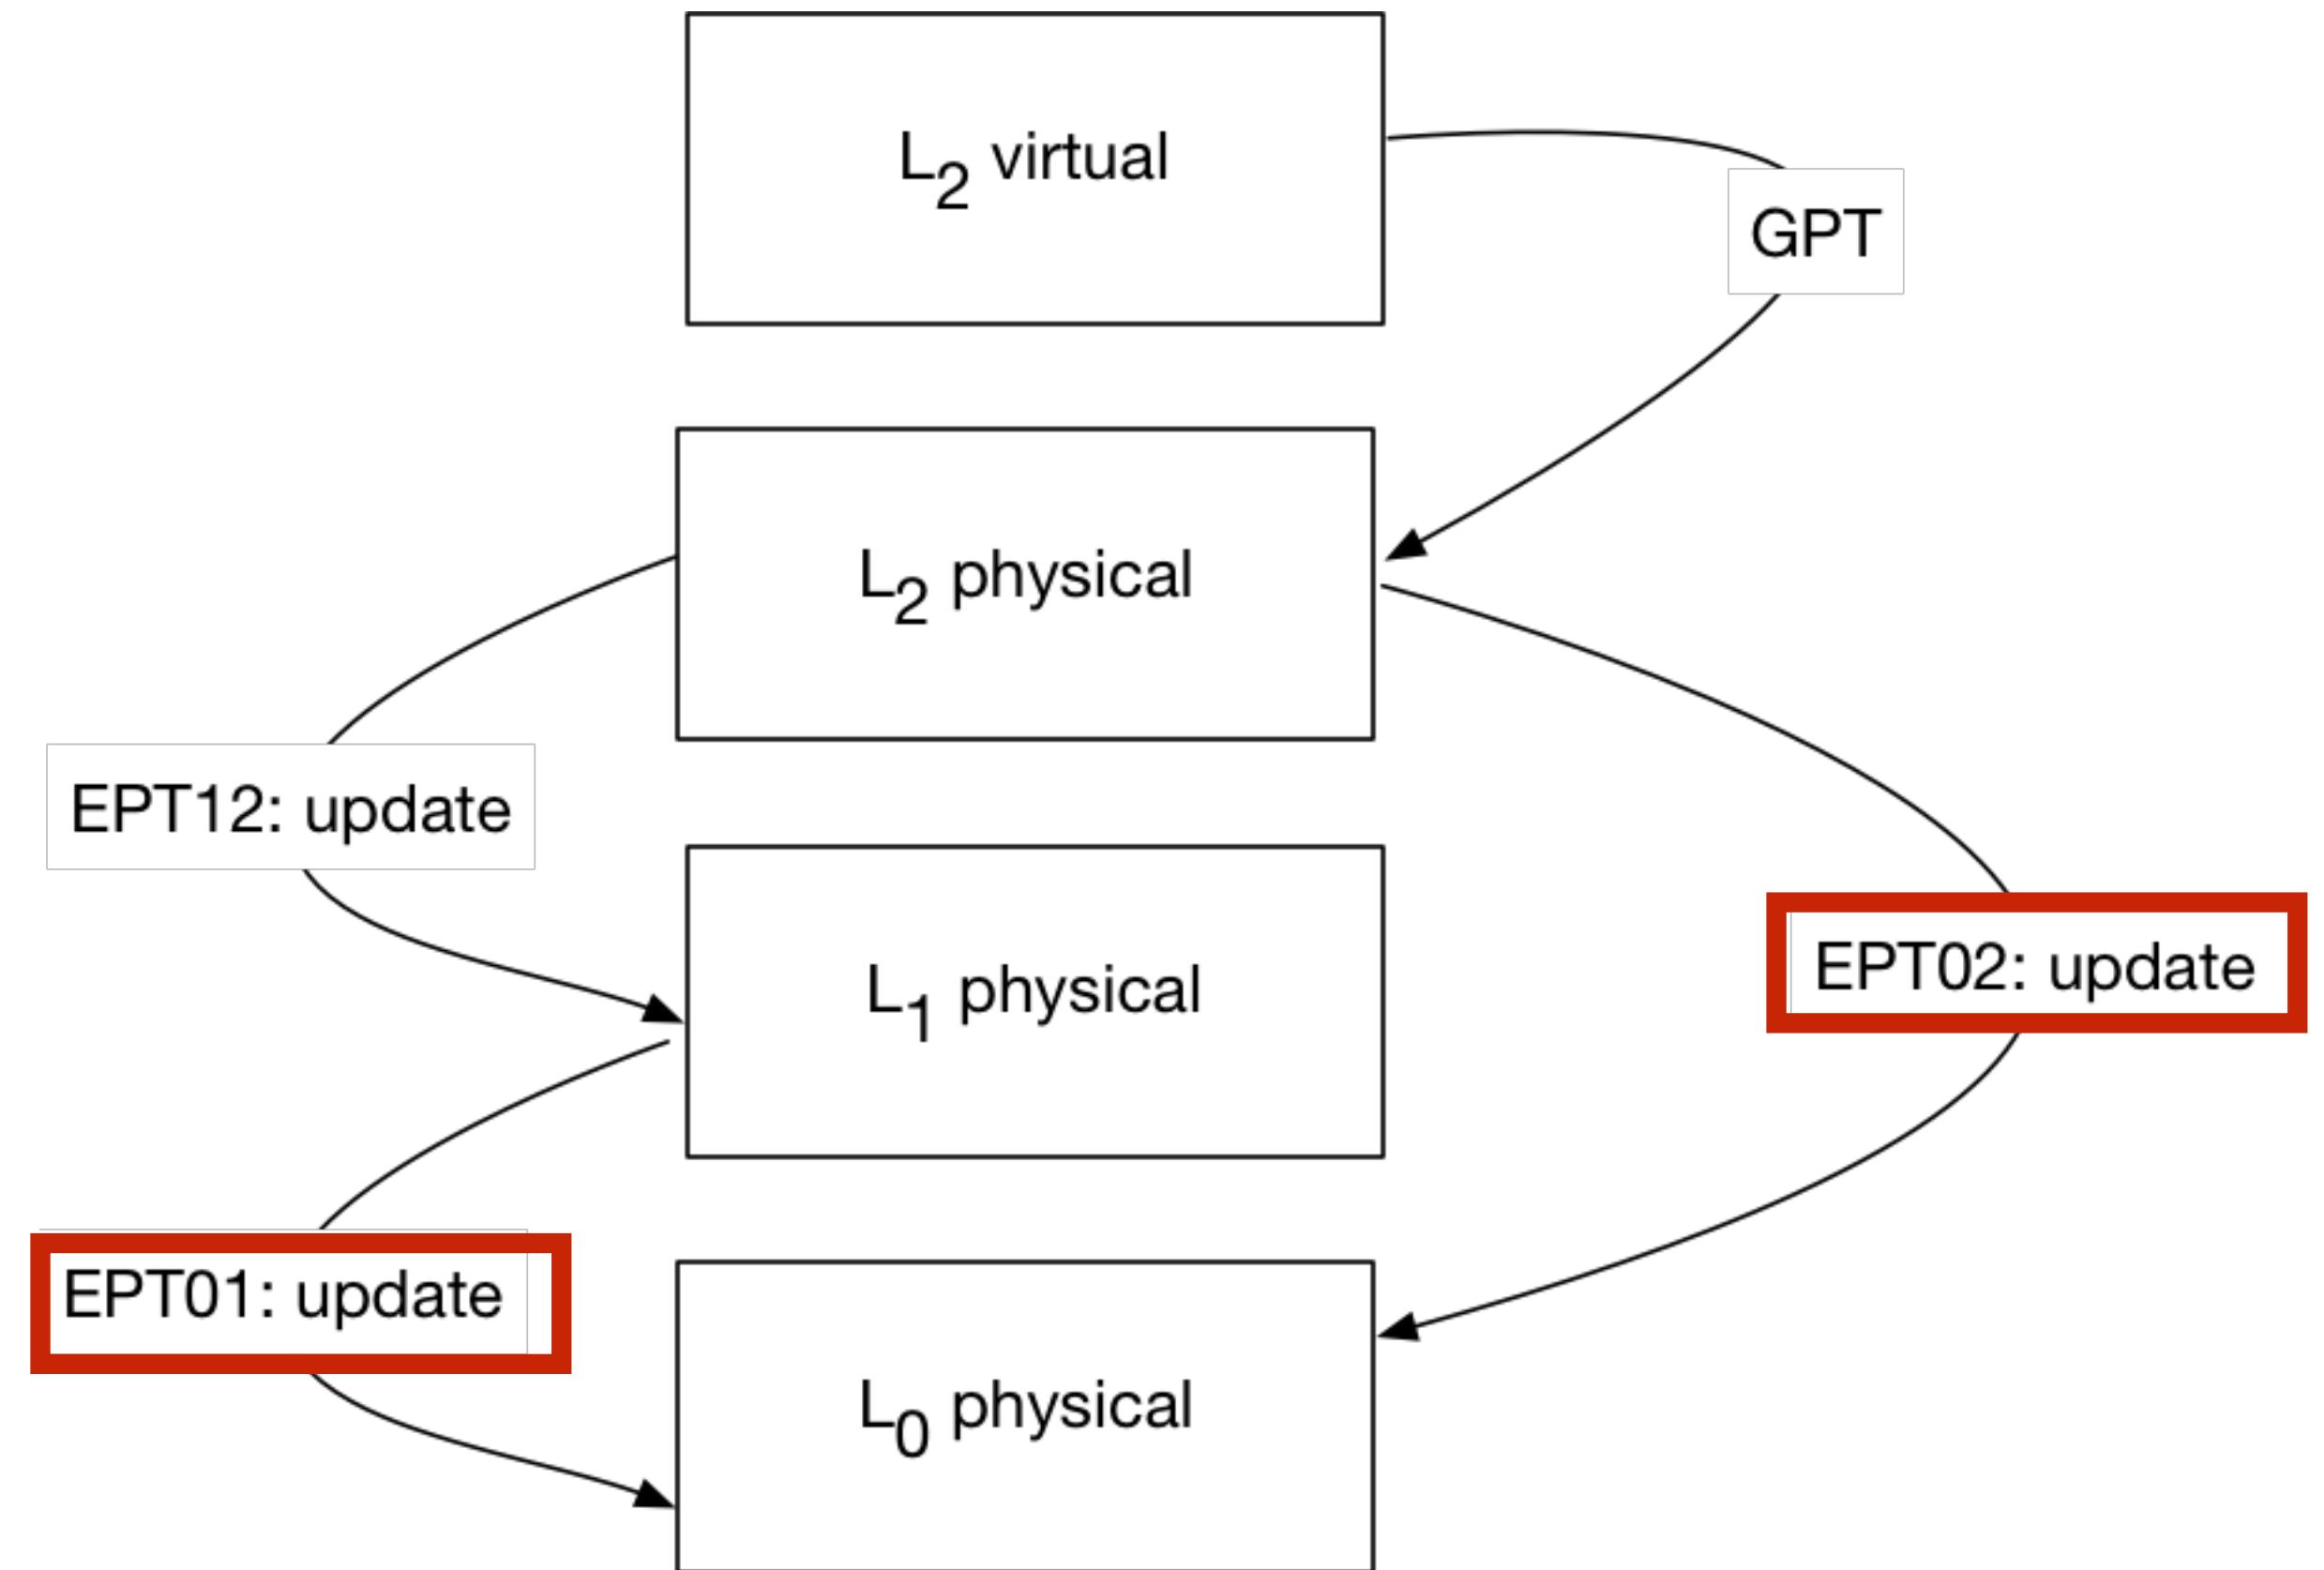
# Building the Compressed EPT

- Start with empty  $EPT_{0 \rightarrow 2}$
- When  $L_2$  causes EPT-violation in  $EPT_{1 \rightarrow 2}$ 
  - $L_1$  handles, updates  $EPT_{1 \rightarrow 2}$



# Building the Compressed EPT

- Start with empty  $EPT_{0 \rightarrow 2}$
- When  $L_2$  causes EPT-violation in  $EPT_{1 \rightarrow 2}$ 
  - $L_1$  handles, updates  $EPT_{1 \rightarrow 2}$
  - $L_0$  uses  $EPT_{0 \rightarrow 1}$  and  $EPT_{1 \rightarrow 2}$  to build  $EPT_{0 \rightarrow 2}$



# Implementing Turtles

- Nested VMX Virtualization for CPUs
- Multidimensional Paging for MMUs
- **Multi-level Device Assignment for I/O**



# Multi-level device assignment

- I/O remains an outstanding challenge



# Multi-level device assignment

- I/O remains an outstanding challenge
  - *Emulate* known devices



# Multi-level device assignment

- I/O remains an outstanding challenge
  - *Emulate* known devices
  - Para-virtual drivers (hypercalls to base hypervisor)



# Multi-level device assignment

- I/O remains an outstanding challenge
  - *Emulate* known devices
  - Para-virtual drivers (hypercalls to base hypervisor)
  - These are not performant enough!



# Multi-level device assignment

- I/O remains an outstanding challenge
  - *Emulate* known devices
  - Para-virtual drivers (hypercalls to base hypervisor)
  - These are not performant enough!
- We have hardware support for virtual devices DMA via IOMMU
  - Translates guest physical addresses for I/O



# Emulating the IOMMU

- L<sub>0</sub> emulates an IOMMU for L<sub>1</sub>



# Emulating the IOMMU

- L<sub>0</sub> emulates an IOMMU for L<sub>1</sub>
- L<sub>1</sub> sets up IOMMU<sub>1→2</sub> mapping



# Emulating the IOMMU

- L<sub>0</sub> emulates an IOMMU for L<sub>1</sub>
- L<sub>1</sub> sets up IOMMU<sub>1→2</sub> mapping
- L<sub>0</sub> builds the IOMMU<sub>0→2</sub> mapping



# Emulating the IOMMU

- L<sub>0</sub> emulates an IOMMU for L<sub>1</sub>  
**“It’s turtles all the way down!”**
- L<sub>1</sub> sets up IOMMU<sub>2→1</sub> mapping
- L<sub>0</sub> builds the IOMMU<sub>2→0</sub> mapping



# Optimizing Exit Handling

- Handling VMExits is slower in L<sub>1</sub> than in L<sub>0</sub>
  - Updates to VMCS in exit handling is privileged
  - Solution: Trap the first time privileged exit code occurs, then replay in non-trapping environment
    - L<sub>1</sub> directly write to VMCS<sub>1→2</sub>
    - Called DRW (direct read + write)



# Overall Evaluation

- Two nested layers ( $L_1$  and  $L_2$ )
- Macro and Microbenchmarks to identify general and worst cases
- KVM as base hypervisor ( $L_0$ )
  - Running KVM ( $L_1$ )
  - Running Ubuntu 9.04 ( $L_2$ )



# Overall Evaluation

**Kernbench**

	Host	Guest	Nested	Nested <sub>DRW</sub>
Runtime	324.3	355	406.3	391.5
% overhead vs. host	-	9.5	25.3	20.7
% overhead vs. guest	-	-	14.5	10.3
%CPU	93	97	99	99



# Overall Evaluation

## Kernbench

	Host	Guest	Nested	Nested <sub>DRW</sub>
Runtime	324.3	355	406.3	391.5
% overhead vs. host	-	9.5	25.3	20.7
% overhead vs. guest	-	-	14.5	10.3
%CPU	93	97	99	99



# Overall Evaluation

## SPECjbb

	Host	Guest	Nested	Nested <sub>DRW</sub>
Score	90493	83599	77065	78347
% degradation vs. host	-	7.6	14.8	13.4
% degradation vs. guest	-	-	7.8	6.3
%CPU	100	100	100	100



# Overall Evaluation

## SPECjbb

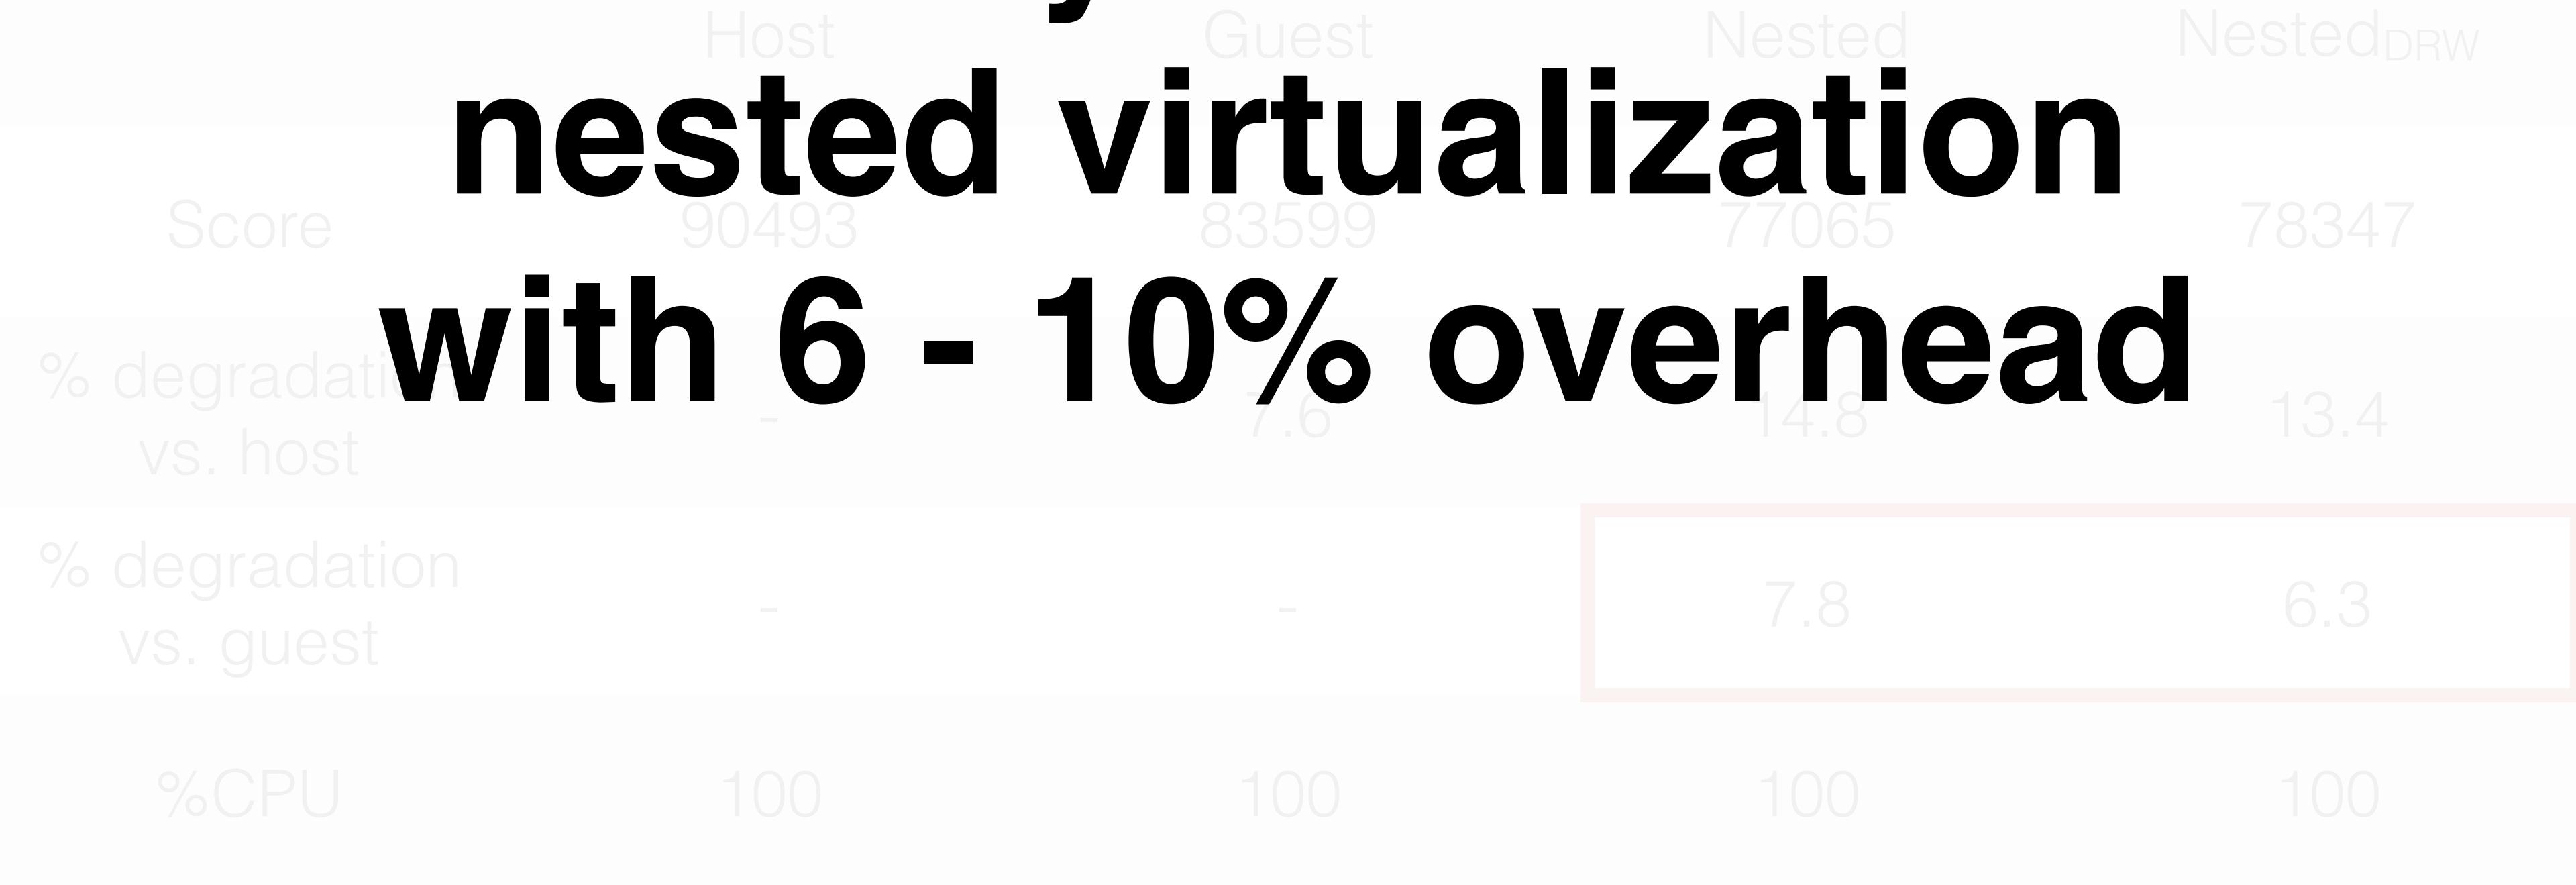
	Host	Guest	Nested	Nested <sub>DRW</sub>
Score	90493	83599	77065	78347
% degradation vs. host	-	7.6	14.8	13.4
% degradation vs. guest	-	-	7.8	6.3
%CPU	100	100	100	100



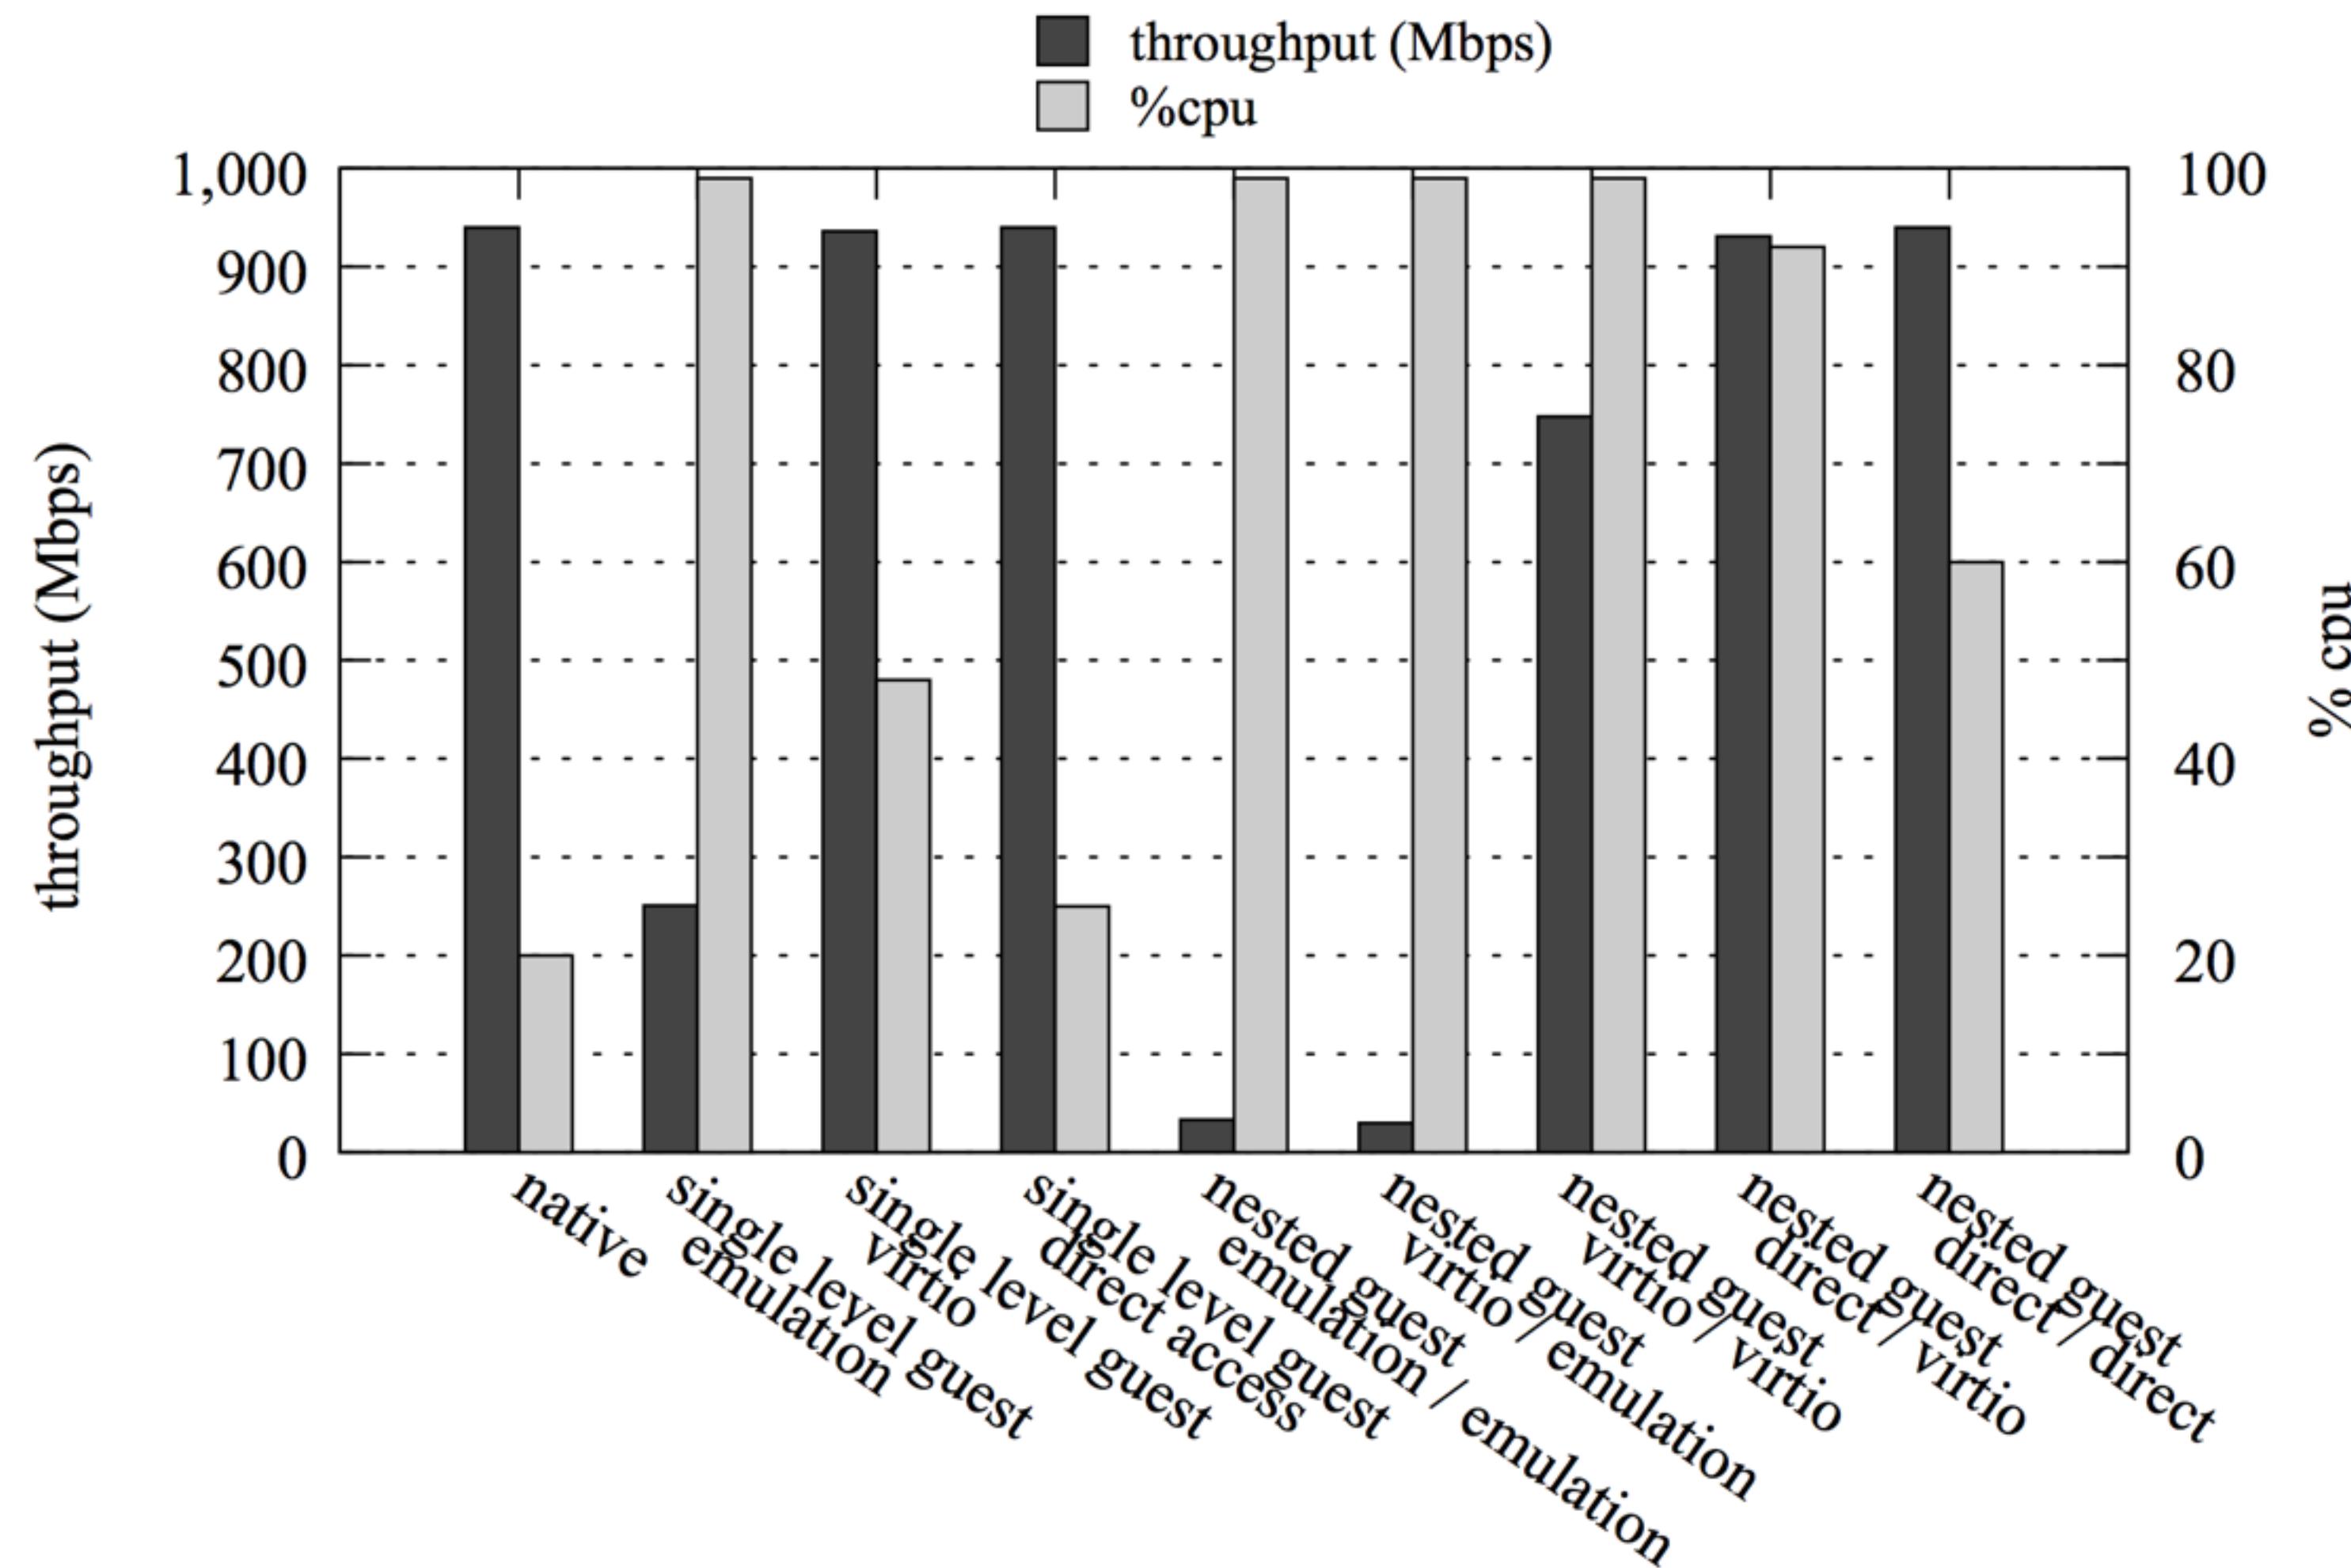
# Overall Evaluation

**Takeaway: Can achieve  
nested virtualization**

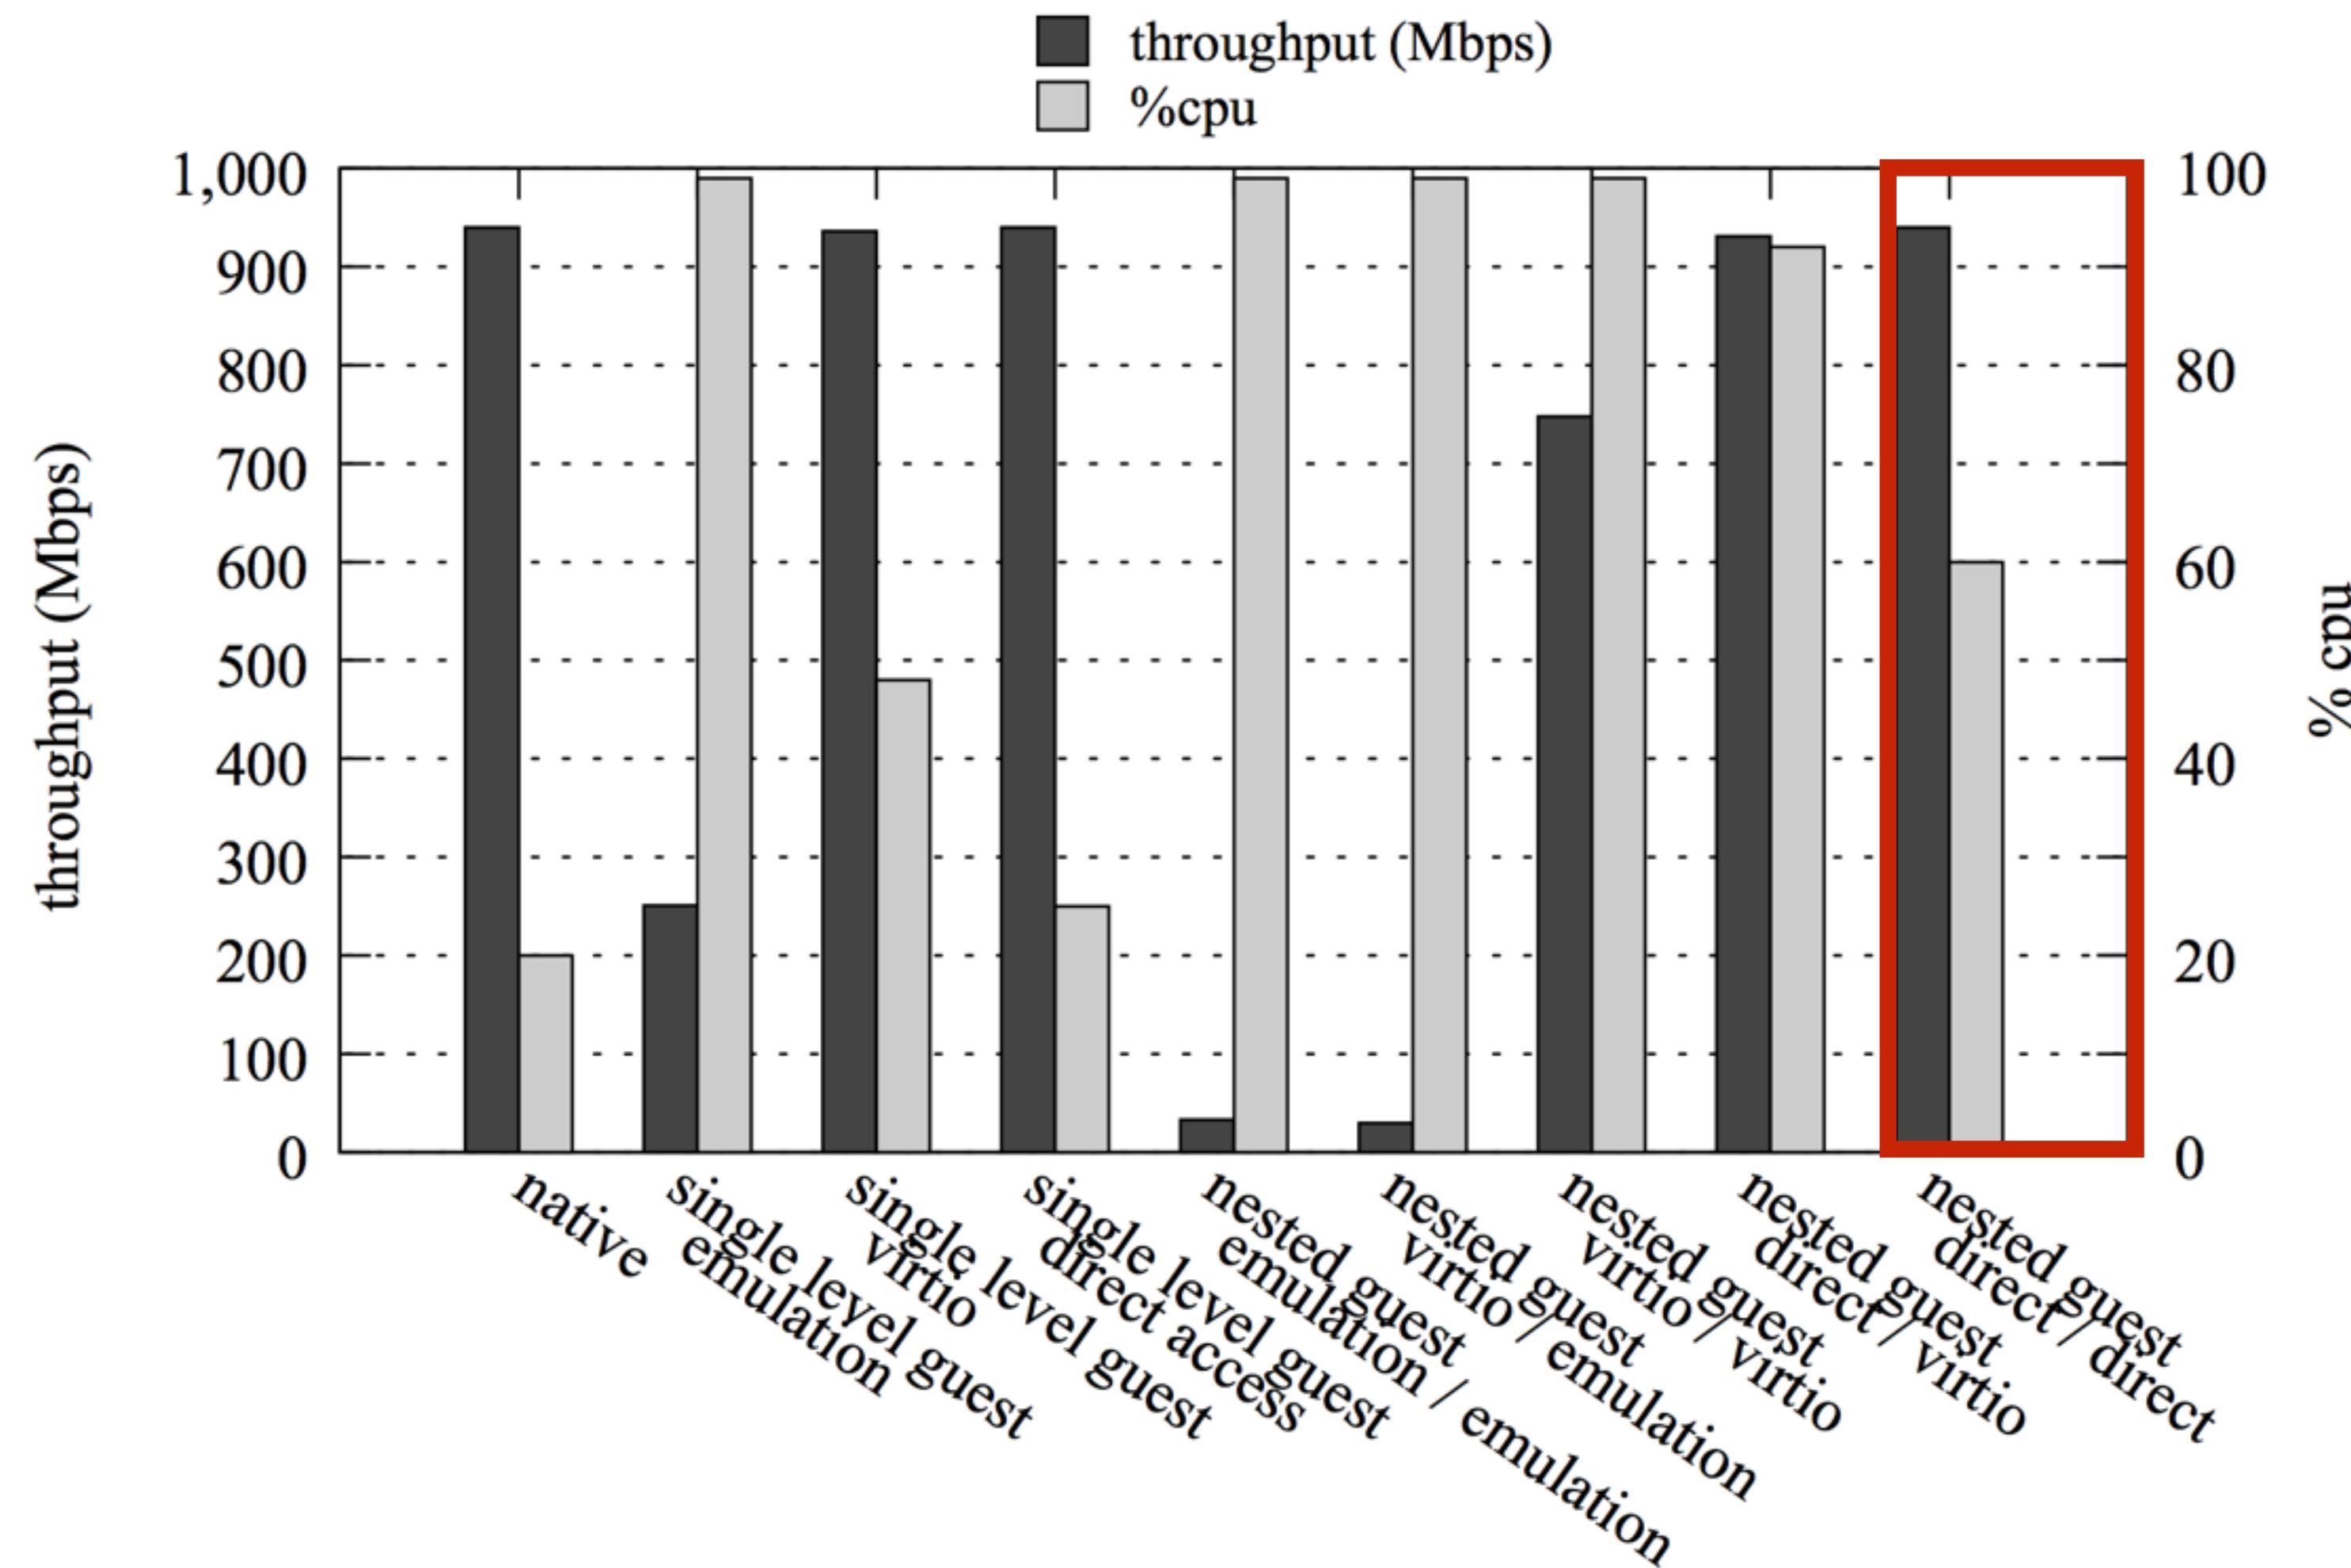
**with 6 - 10% overhead**



# I/O Intensive Workloads

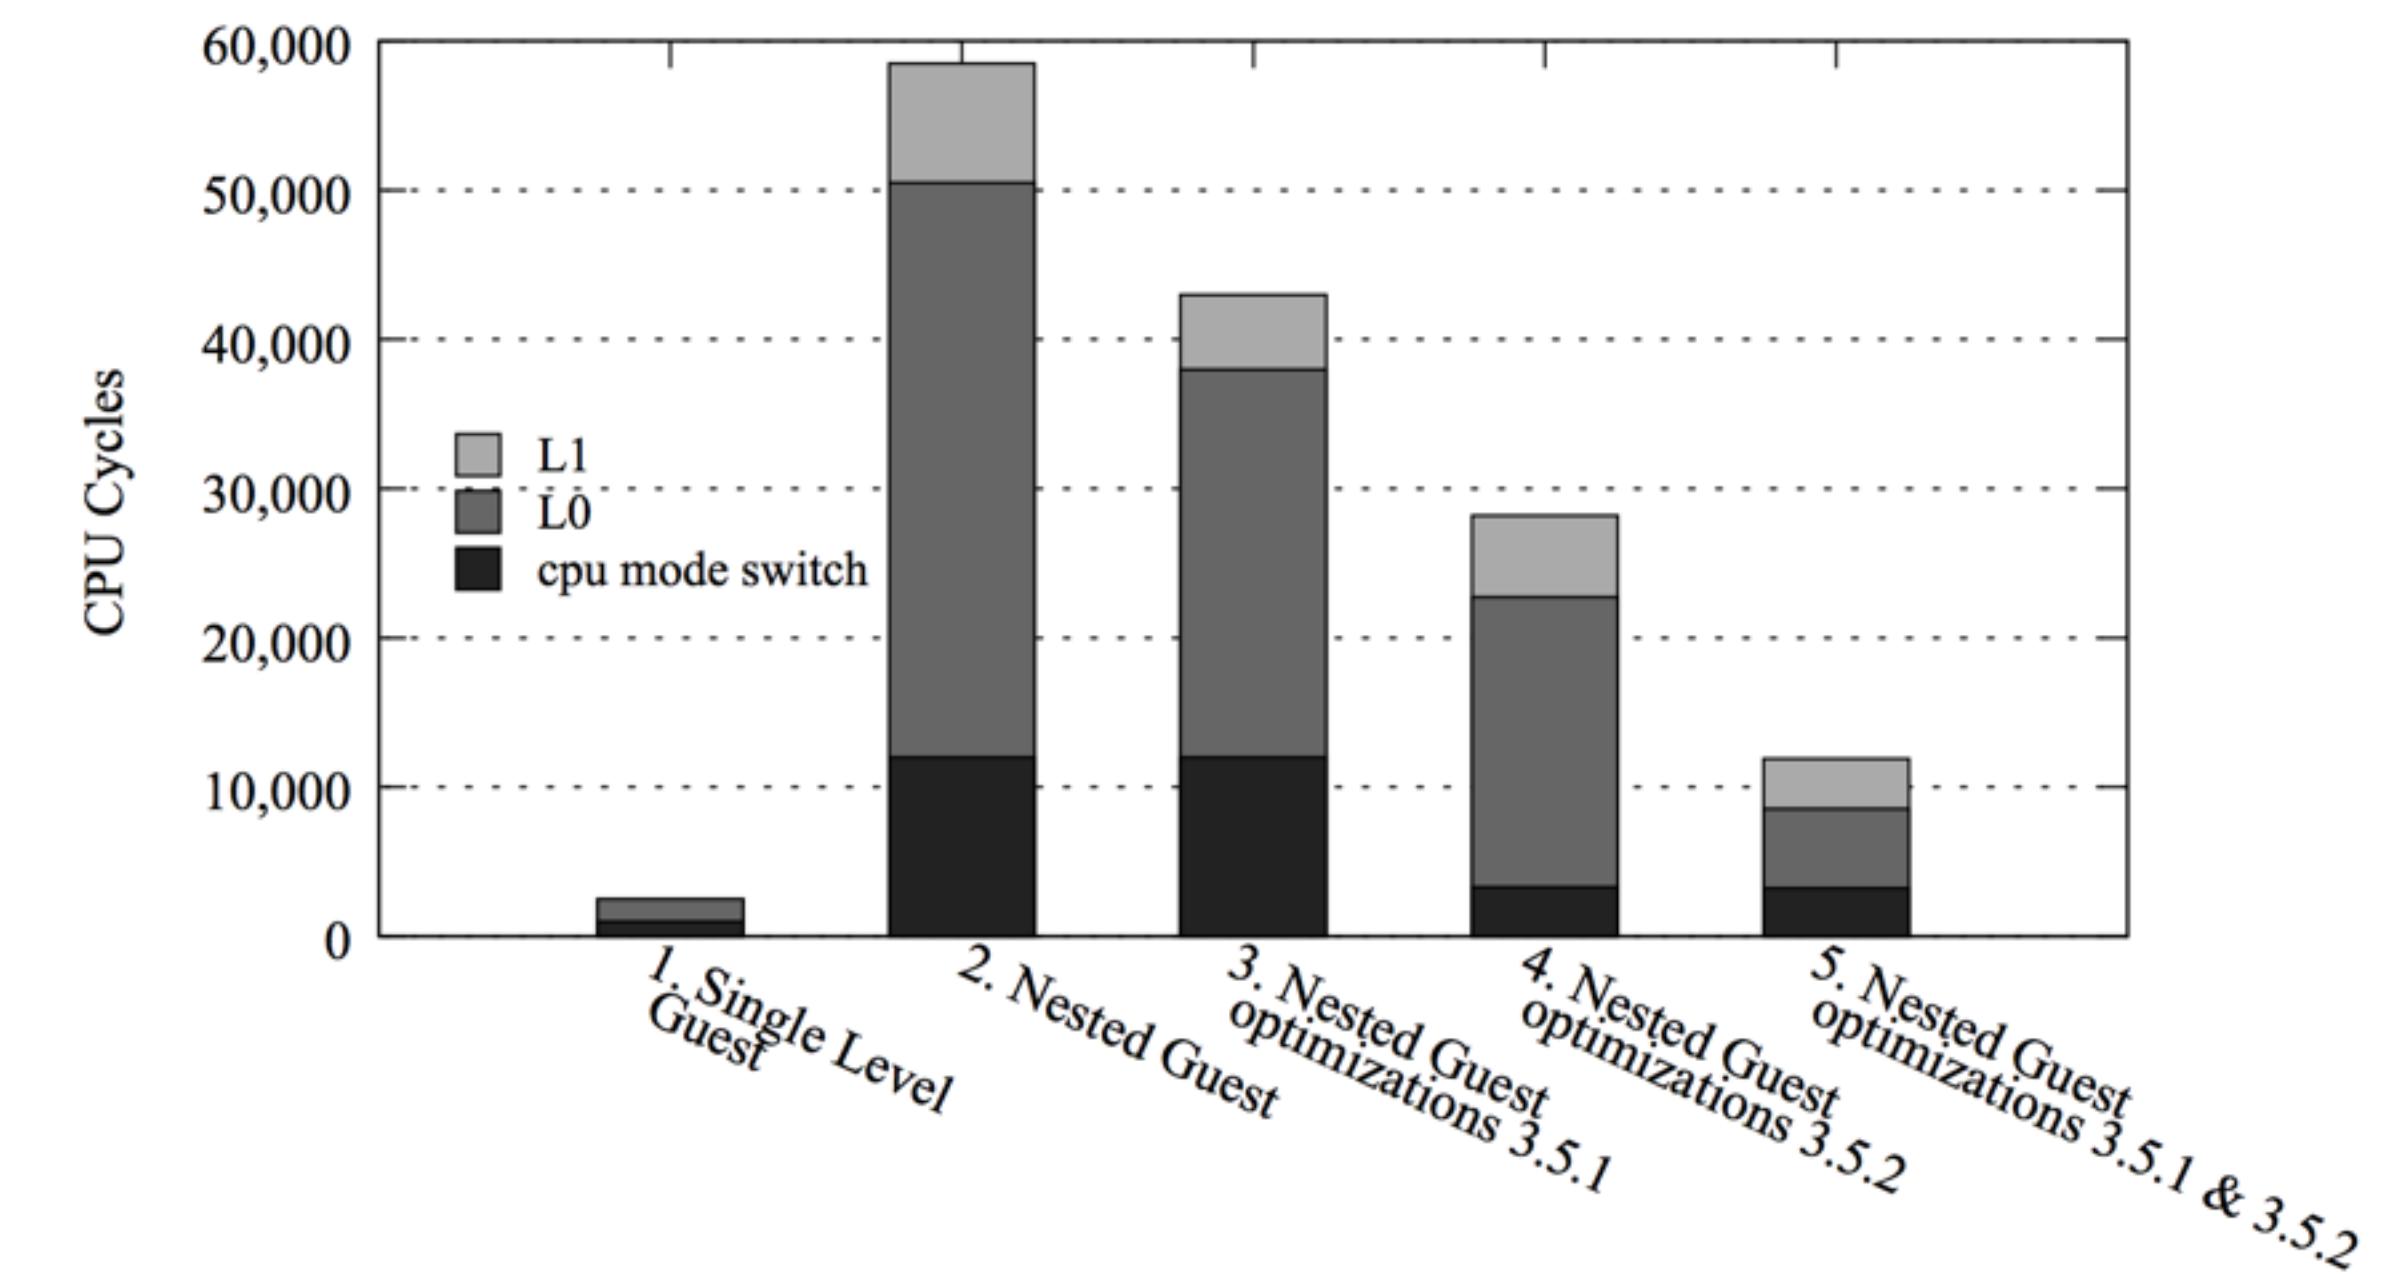


# I/O Intensive Workloads



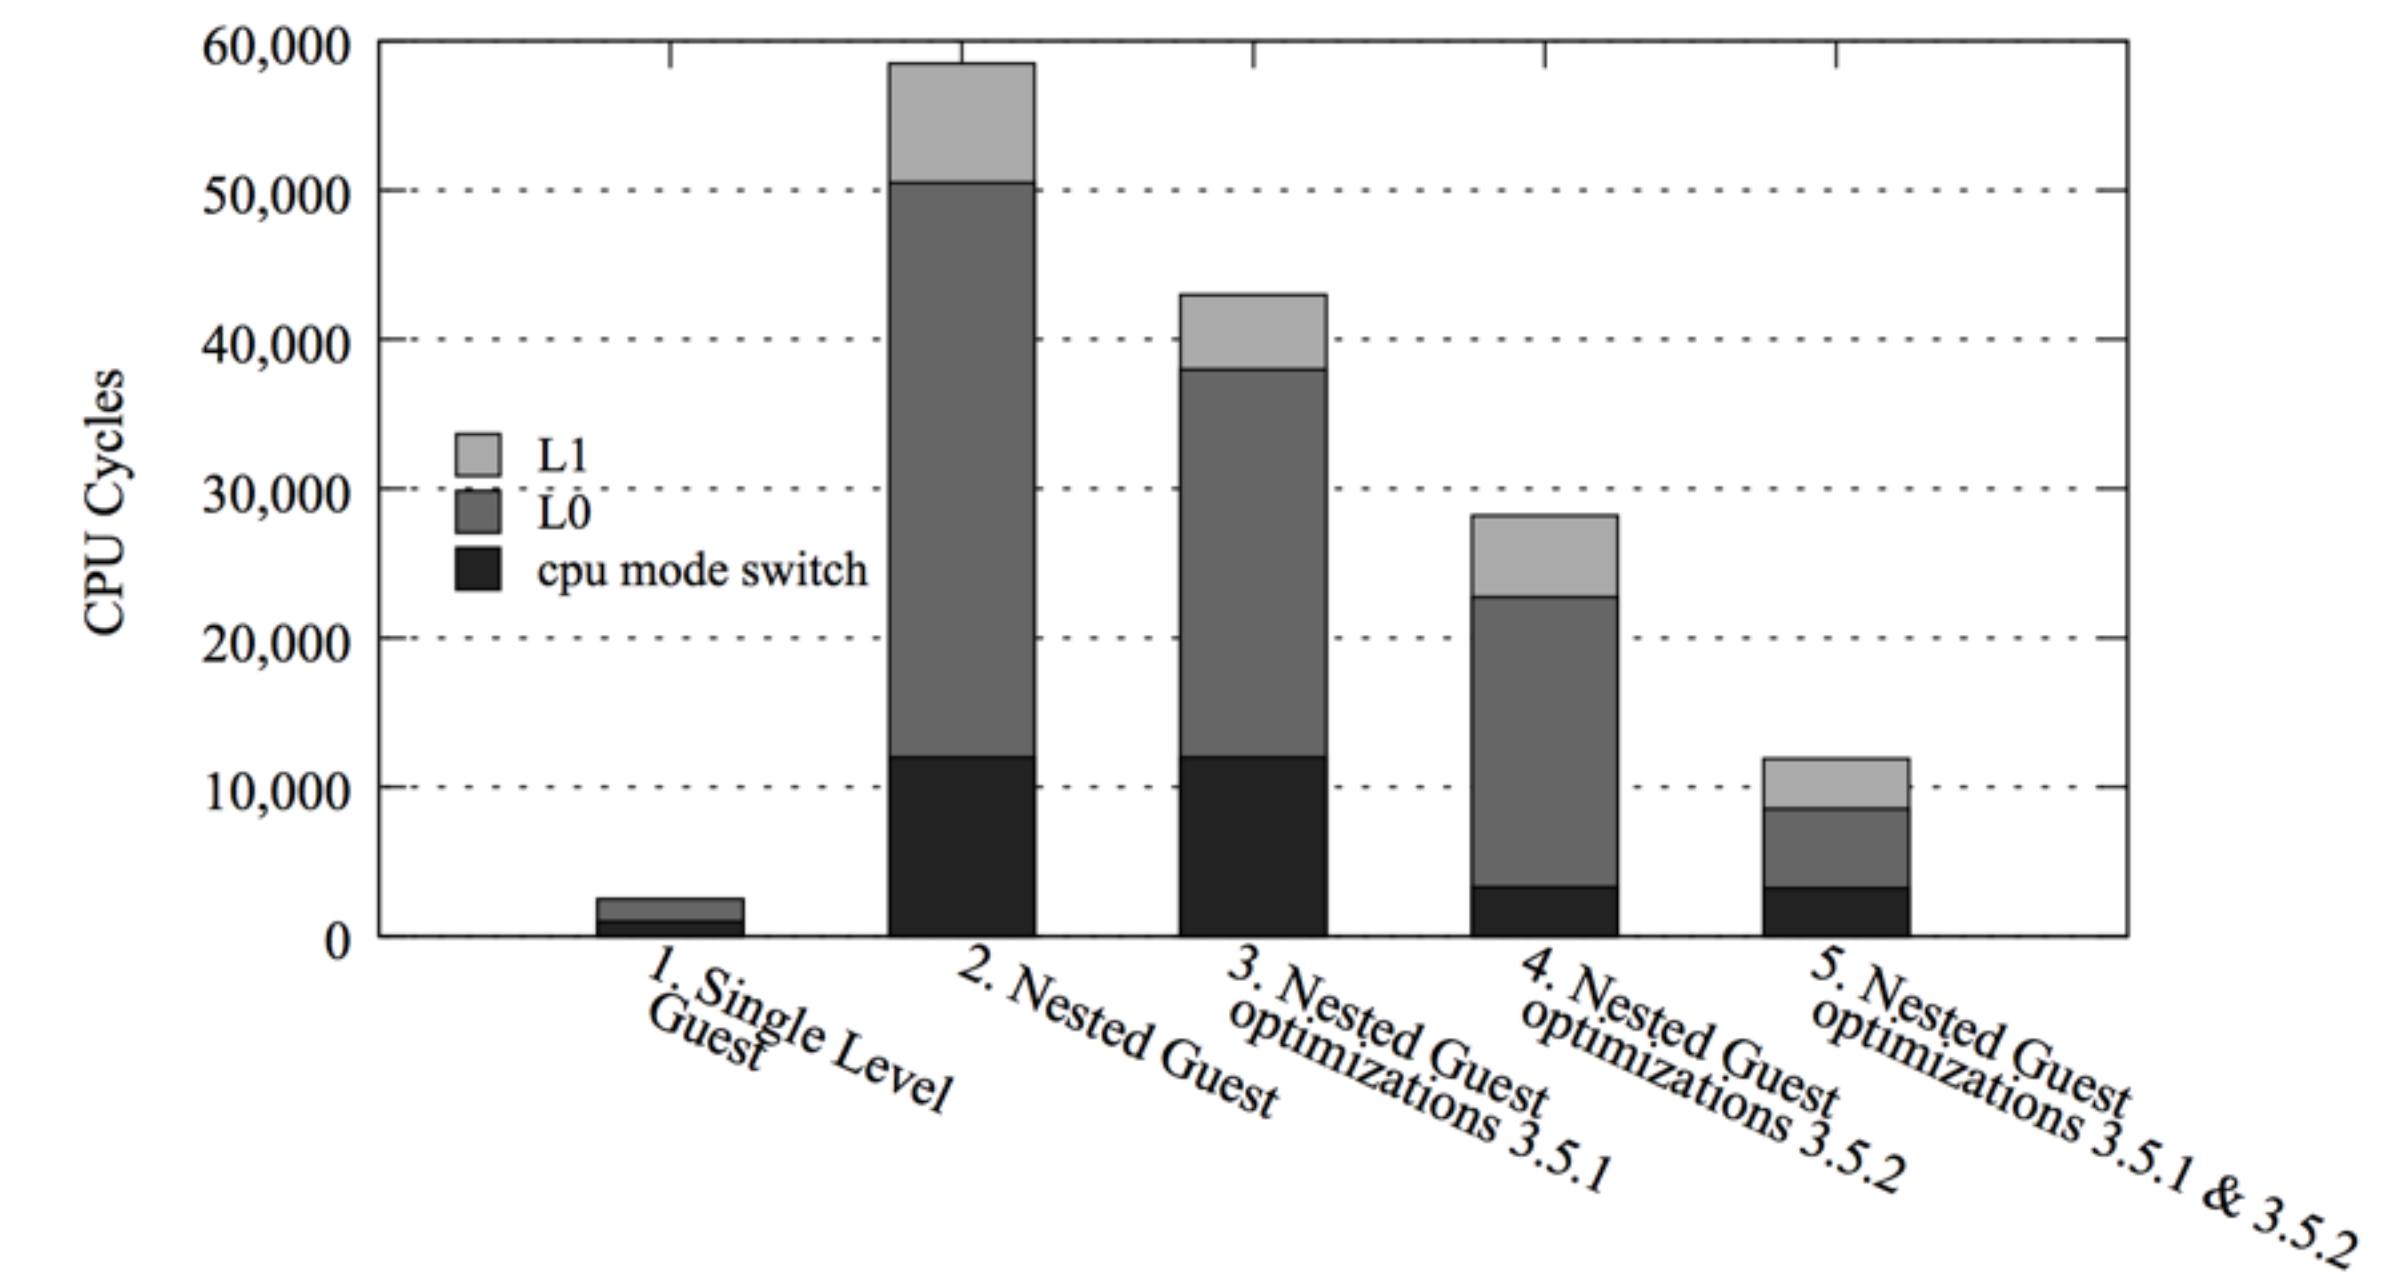
# Microbenchmark

- Run “cpuid” in a loop to generate exits
  - Run 4M times



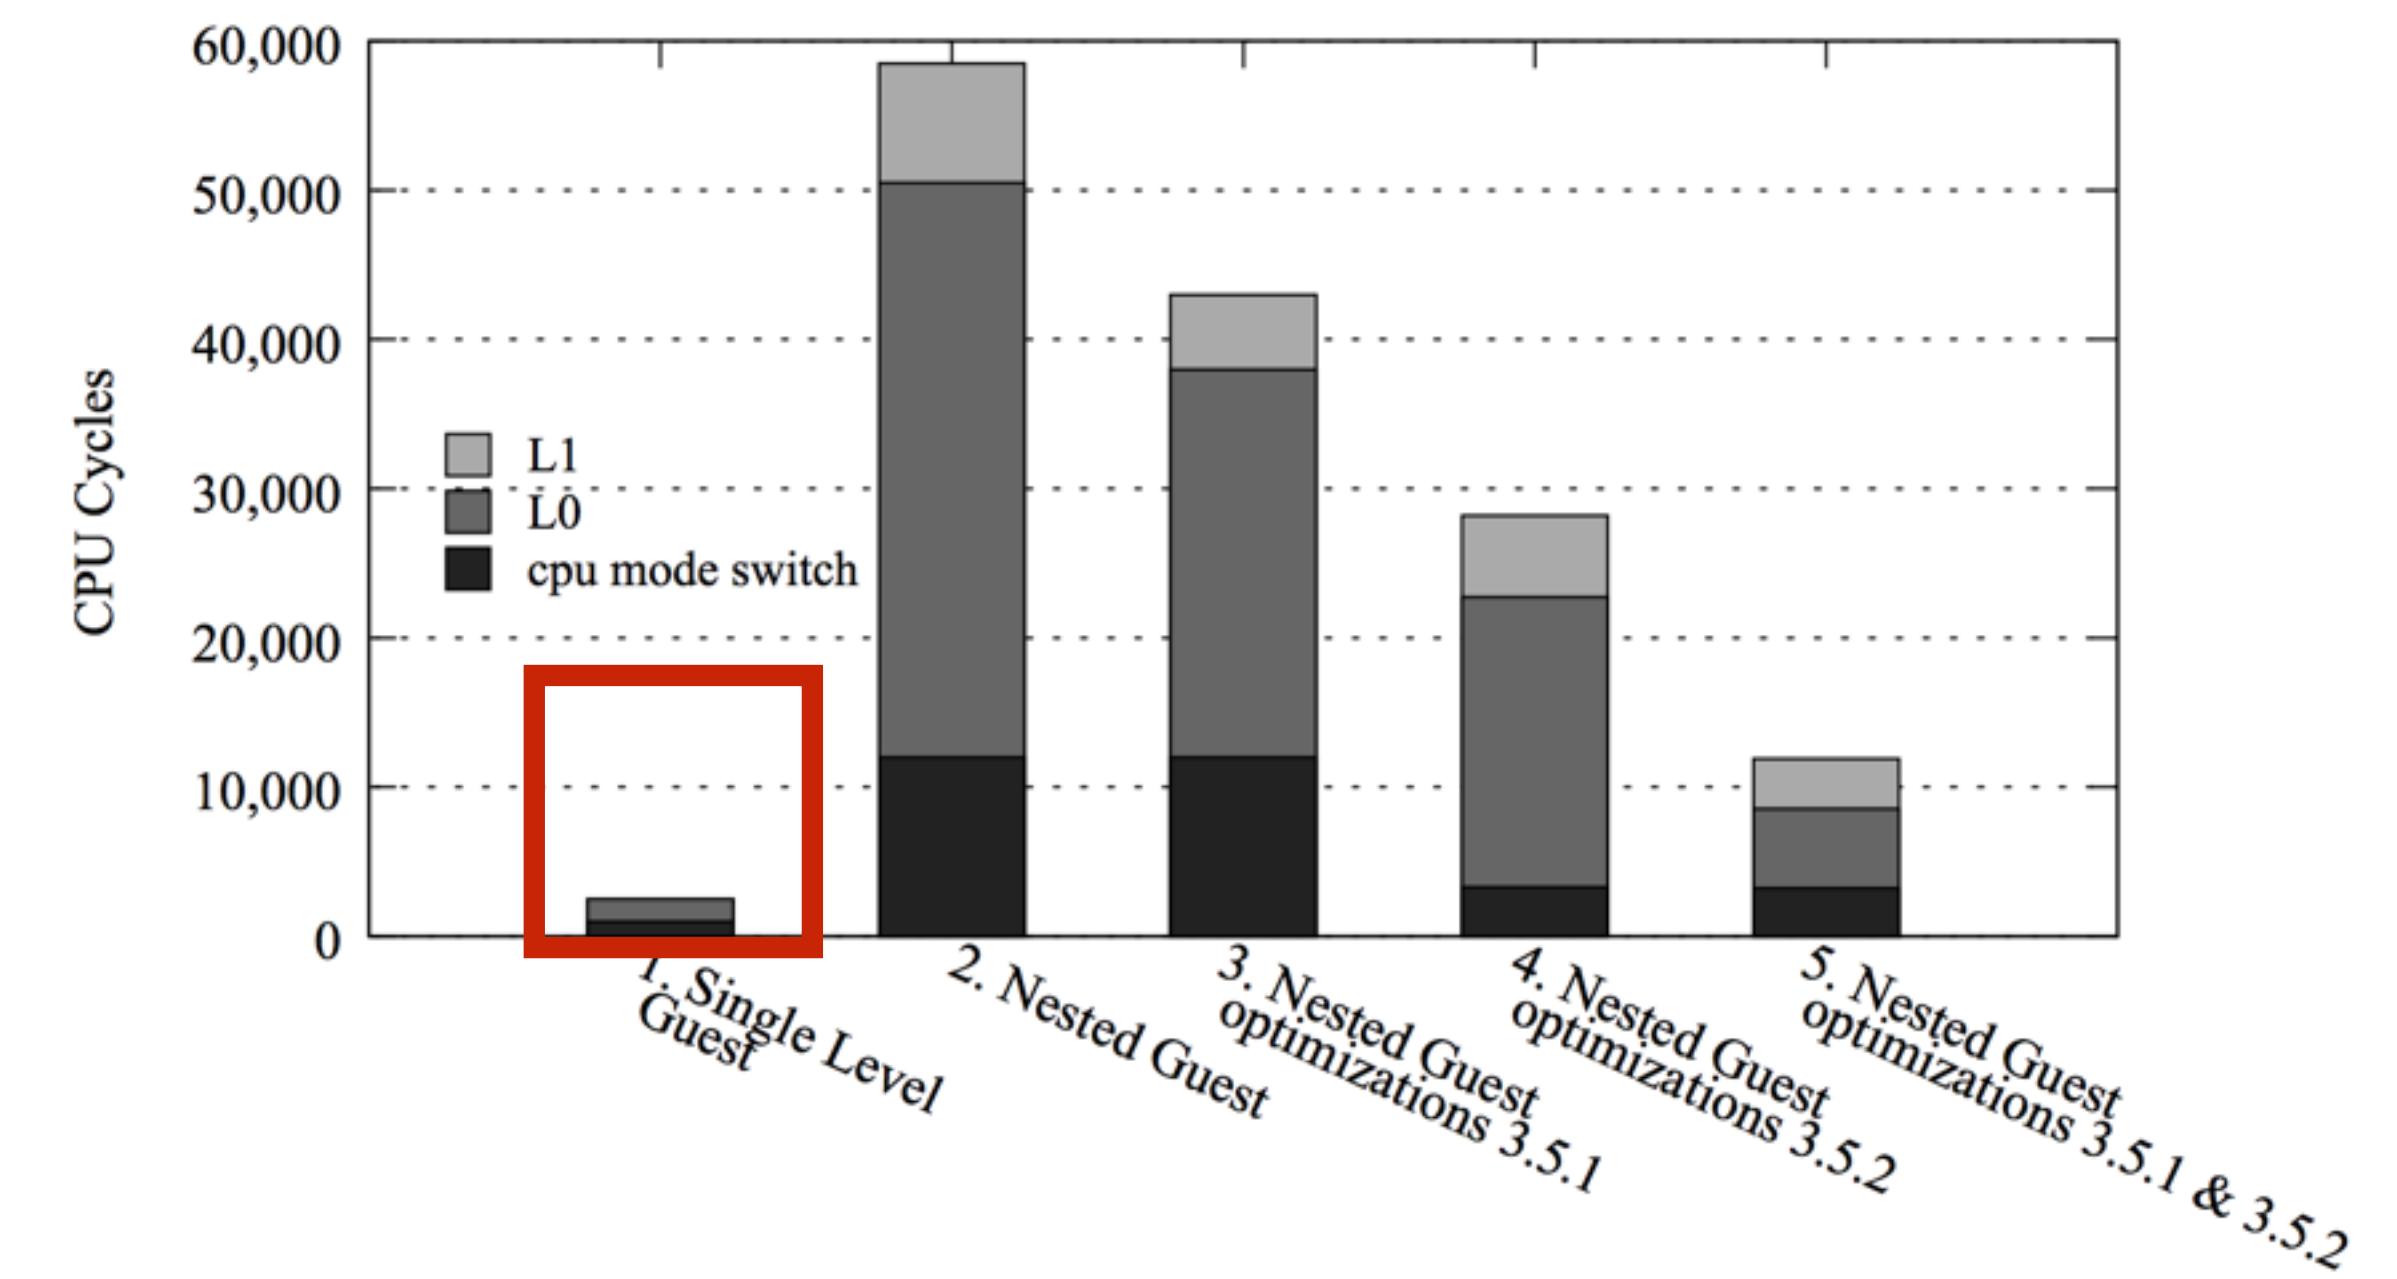
# Microbenchmark

- Run “cpuid” in a loop to generate exits
  - Run 4M times
  - Bare metal: 100 cycles



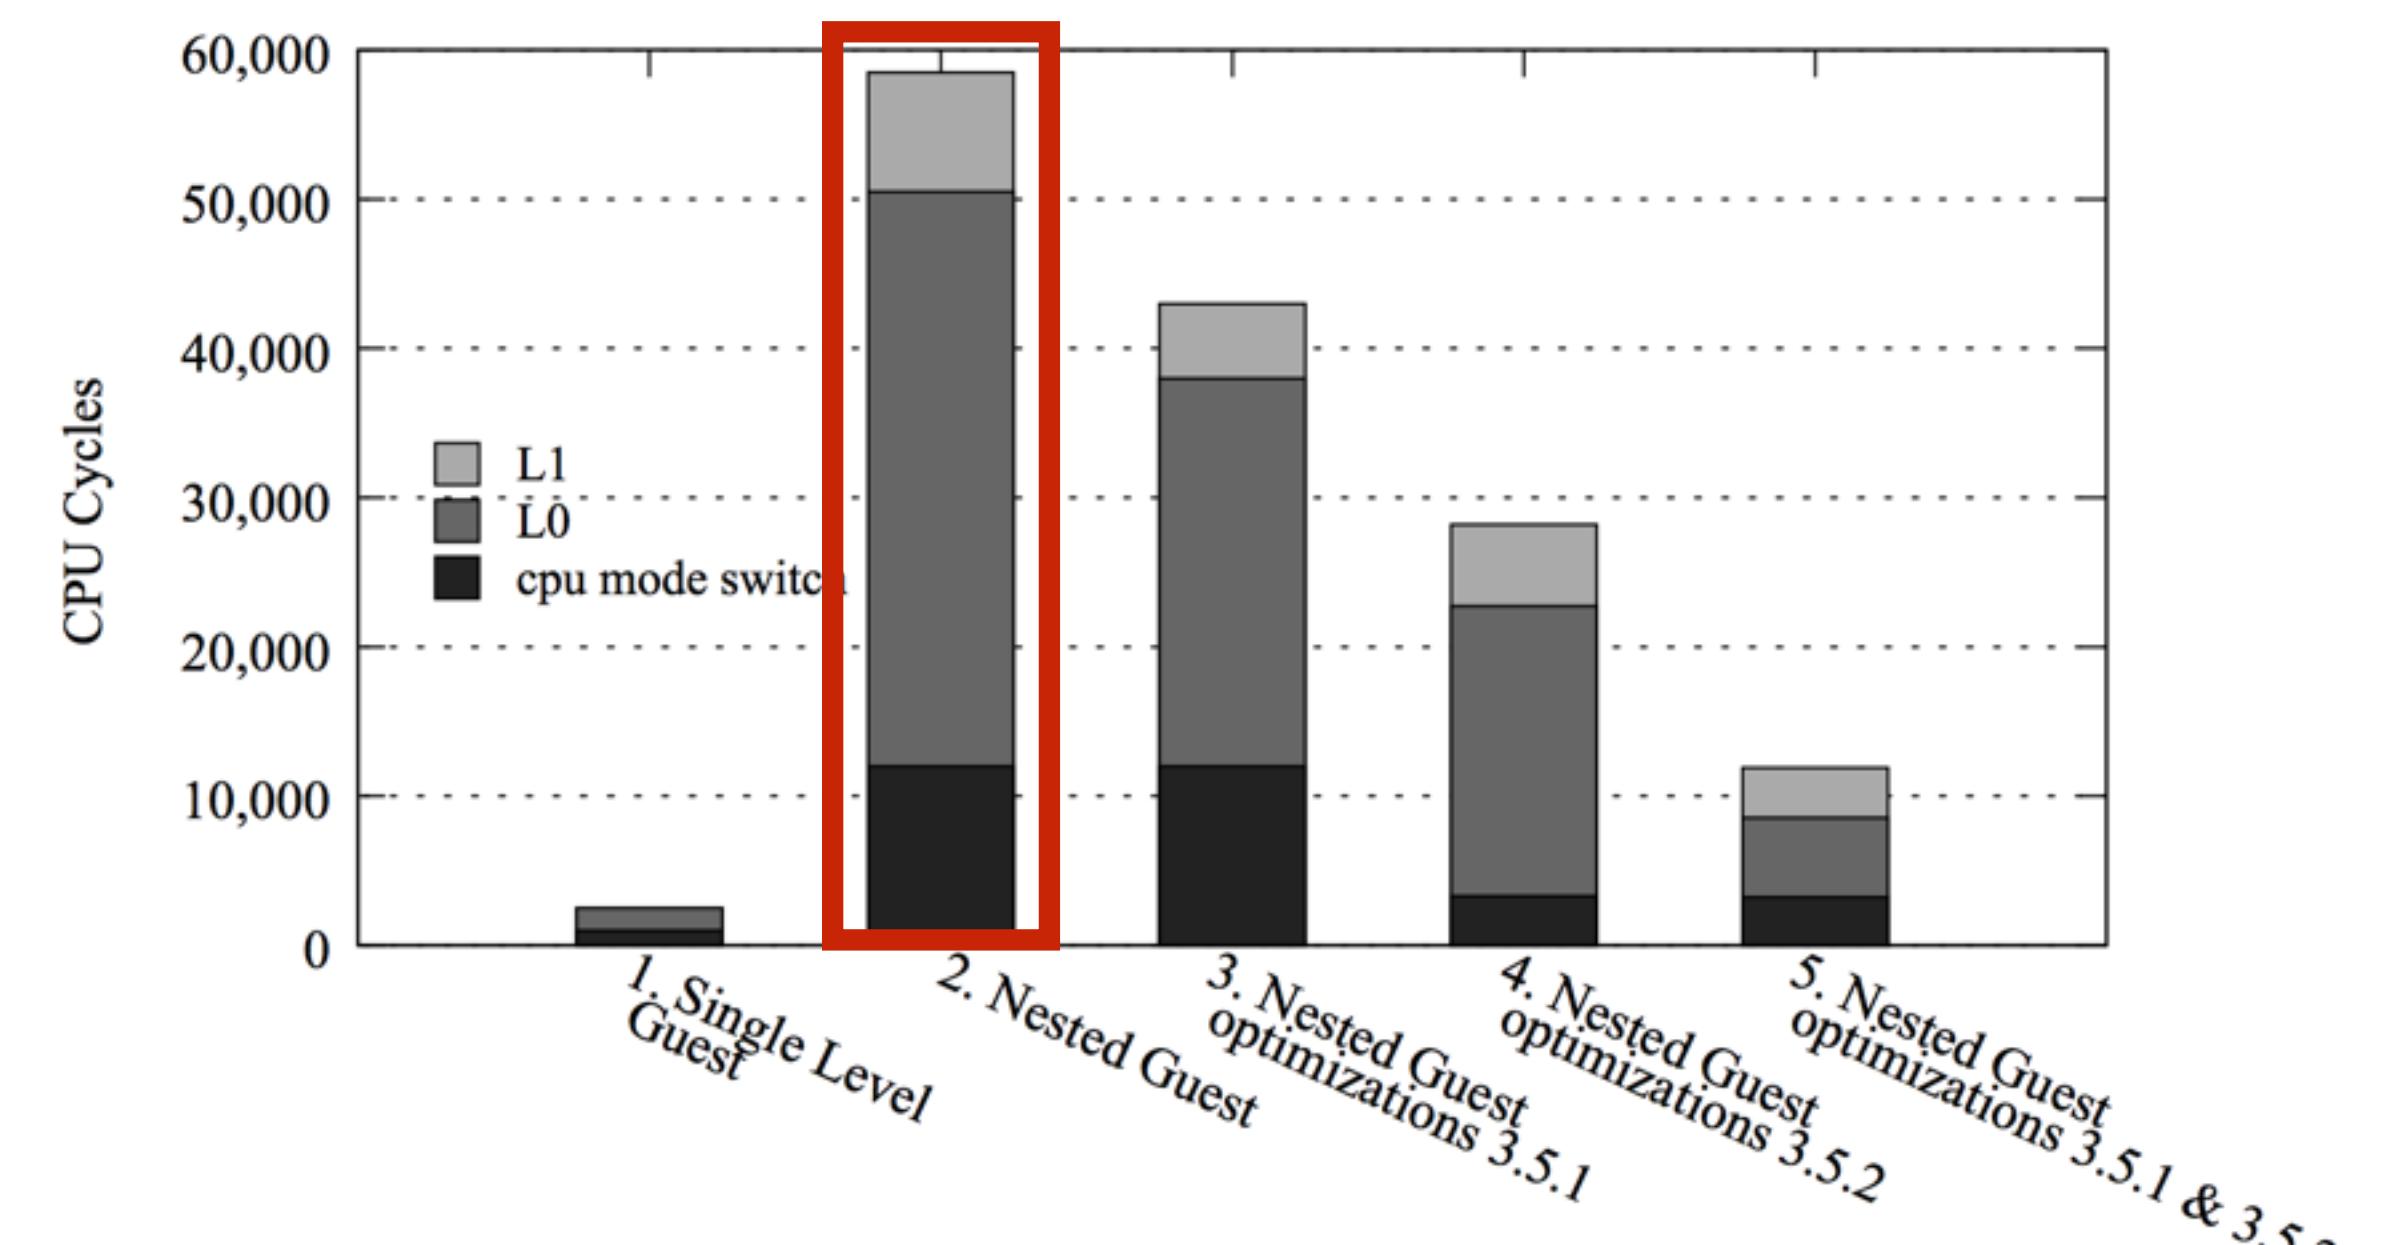
# Microbenchmark

- Run “cpuid” in a loop to generate exits
  - Run 4M times
  - Bare metal: 100 cycles
  - Single level guest: 2,600 cycles



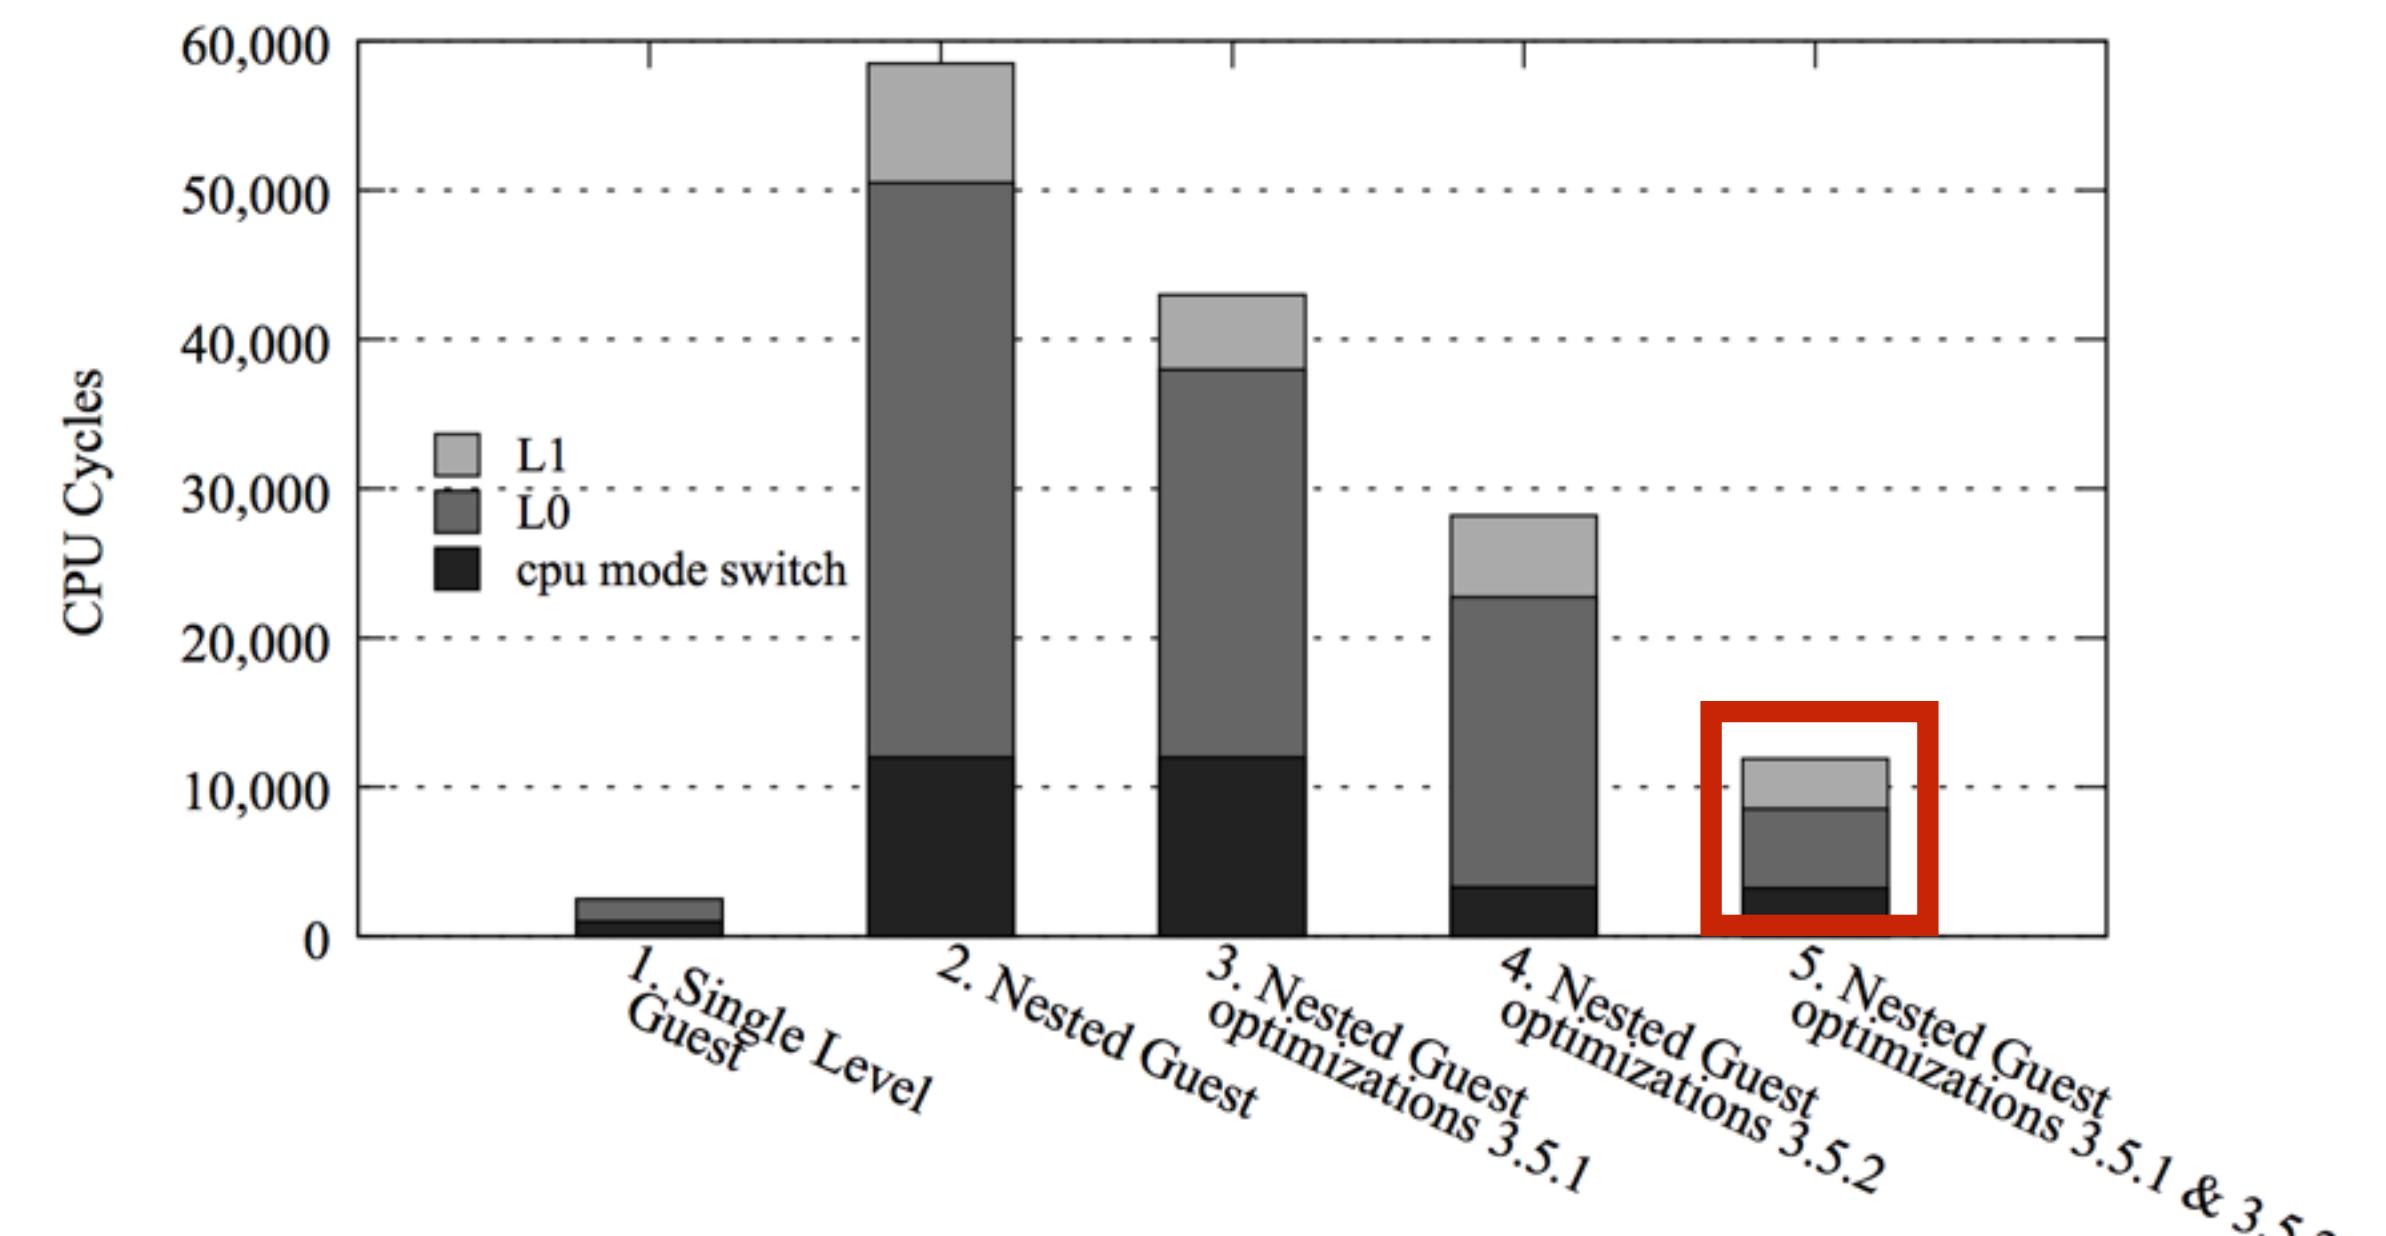
# Microbenchmark

- Run “cpuid” in a loop to generate exits
  - Run 4M times
  - Bare metal: 100 cycles
  - Single level guest: 2,600 cycles
  - Nested guest: 58,000 cycles



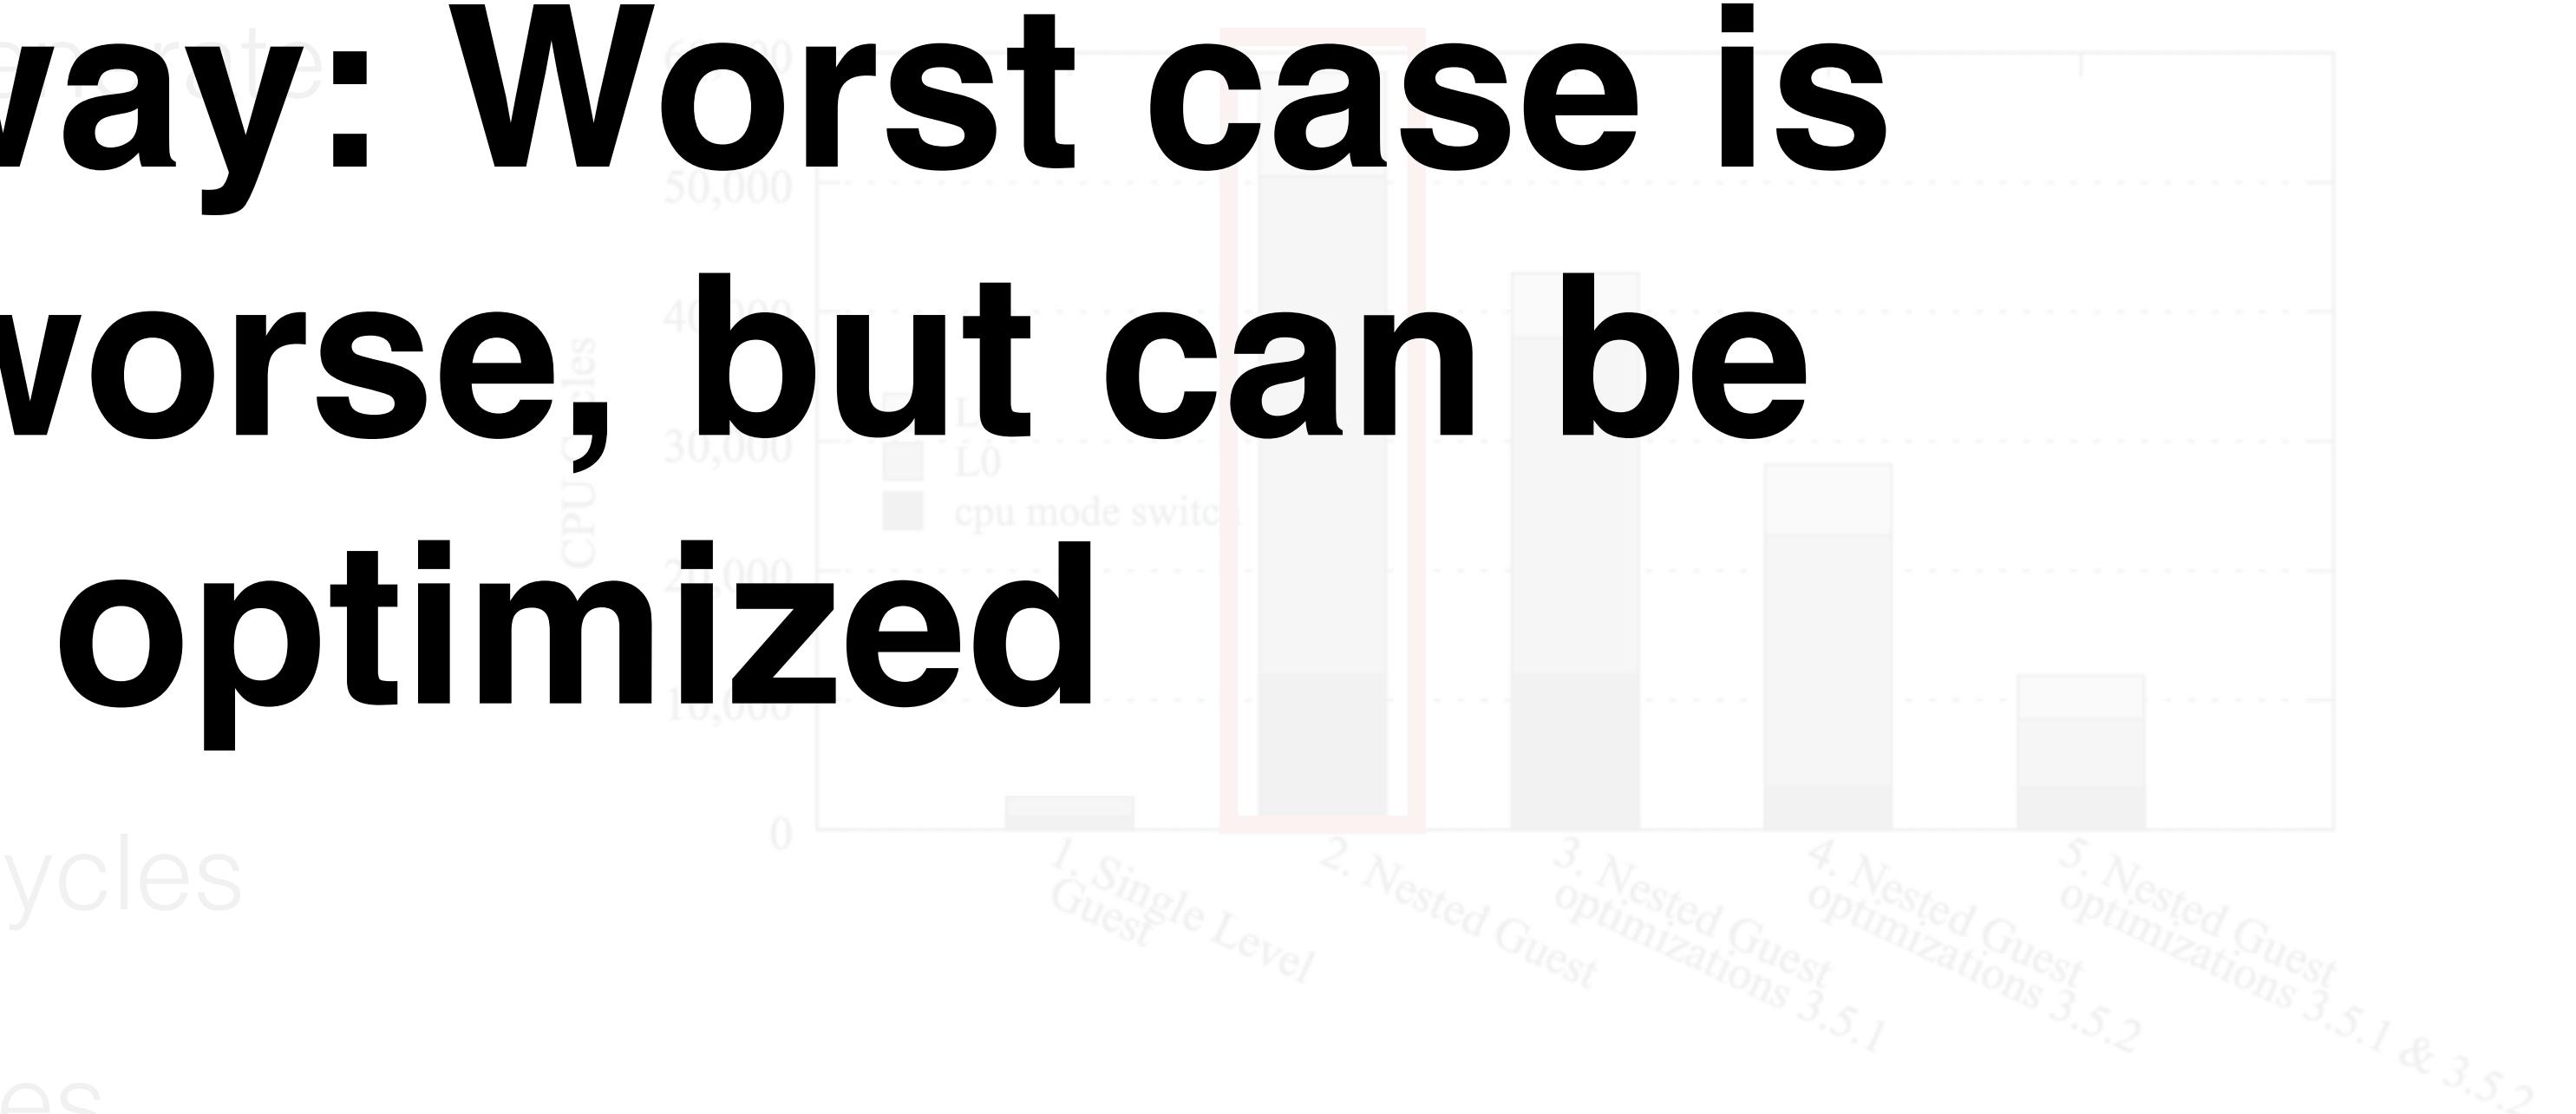
# Microbenchmark

- Run “cpuid” in a loop to generate exits
  - Run 4M times
  - Bare metal: 100 cycles
  - Single level guest: 2,600 cycles
  - Nested guest: 58,000 cycles



# Microbenchmark

- Run “cpuid” exits
- Run 4M times
- Bare metal: 100 cycles
- Single level guest: 2,600 cycles
- Nested guest: 58,000 cycles



# My Takeaways

- “Turtles” provides a clean solution to a tough problem
  - In particular, beefing up the base hypervisor to handle all traps is an intuitive solution
  - However, software will always be slower than hardware



# My Takeaways

- “Turtles” provides a clean solution to a tough problem
  - In particular, beefing up the base hypervisor to handle all traps is an intuitive solution
  - However, software will always be slower than hardware
- **Paper:** Theoretical lower bounds for virtualization



# My Takeaways

- “Turtles” provides a clean solution to a tough problem
  - In particular, beefing up the base hypervisor to handle all traps is an intuitive solution
  - However, software will always be slower than hardware
- **Paper:** Theoretical lower bounds for virtualization
- Nested virtualization is not practical for all workloads
  - I/O intensive workloads have unreasonable costs (60% vs 20% CPU overhead)
  - Largely because of trap-and-emulate, which requires VM + Host to share the same core



# My Takeaways

- “Turtles” provides a clean solution to a tough problem
  - In particular, beefing up the base hypervisor to handle all traps is an intuitive solution
  - However, software will always be slower than hardware
- **Paper:** Theoretical lower bounds for virtualization
- Nested virtualization is not practical for all workloads
  - I/O intensive workloads have unreasonable costs (60% vs 20%)
  - Largely because of trap-and-emulate, which requires VM + Host to share the same core
- **Paper:** Rebuilding trap-and-emulate for I/O intensive workloads on multicore



# References

- [1] <https://www.statista.com/statistics/282087/number-of-monthly-active-twitter-users/>

